

COL216

Computer Architecture

Input/Output – 1

14th March 2022

Diversity of Input/Output devices

Communicate with

- Human
- Computer
- Physical systems
- Storage

Variations in terms of

- direction
- speed
- data type . . .



Communication with humans

- Key board, mouse, joy stick, tablet, touch, gesture
 - very slow, 10^1 - 10^2 bytes per sec
- Documents: Scanners, printers, plotters
 - $8\text{-}16 \text{ ppm}$, $300\text{-}1200 \text{ dpi}$, $10^2\text{-}10^3 \text{ KB/sec}$
- Sound (voice/music): mikes, speakers
 - $10^1\text{-}10^3 \text{ KB/sec}$
- Graphic displays
 - mega pixels, refresh rate, $10^1\text{-}10^3 \text{ MB/sec}$

Communication with machines

- Wired network controllers

$10^1\text{-}10^3$ MB/sec

- Wireless network controllers

 - ethernet, blue tooth, zigbee etc

$10^0\text{-}10^1$ MB/sec

- Modems

$10^1\text{-}10^3$ KB/sec

Communication with physical systems

■ Sensors

- pressure, temperature, humidity, flow, level, position, speed, proximity, gas
- audio, visual

■ Actuators

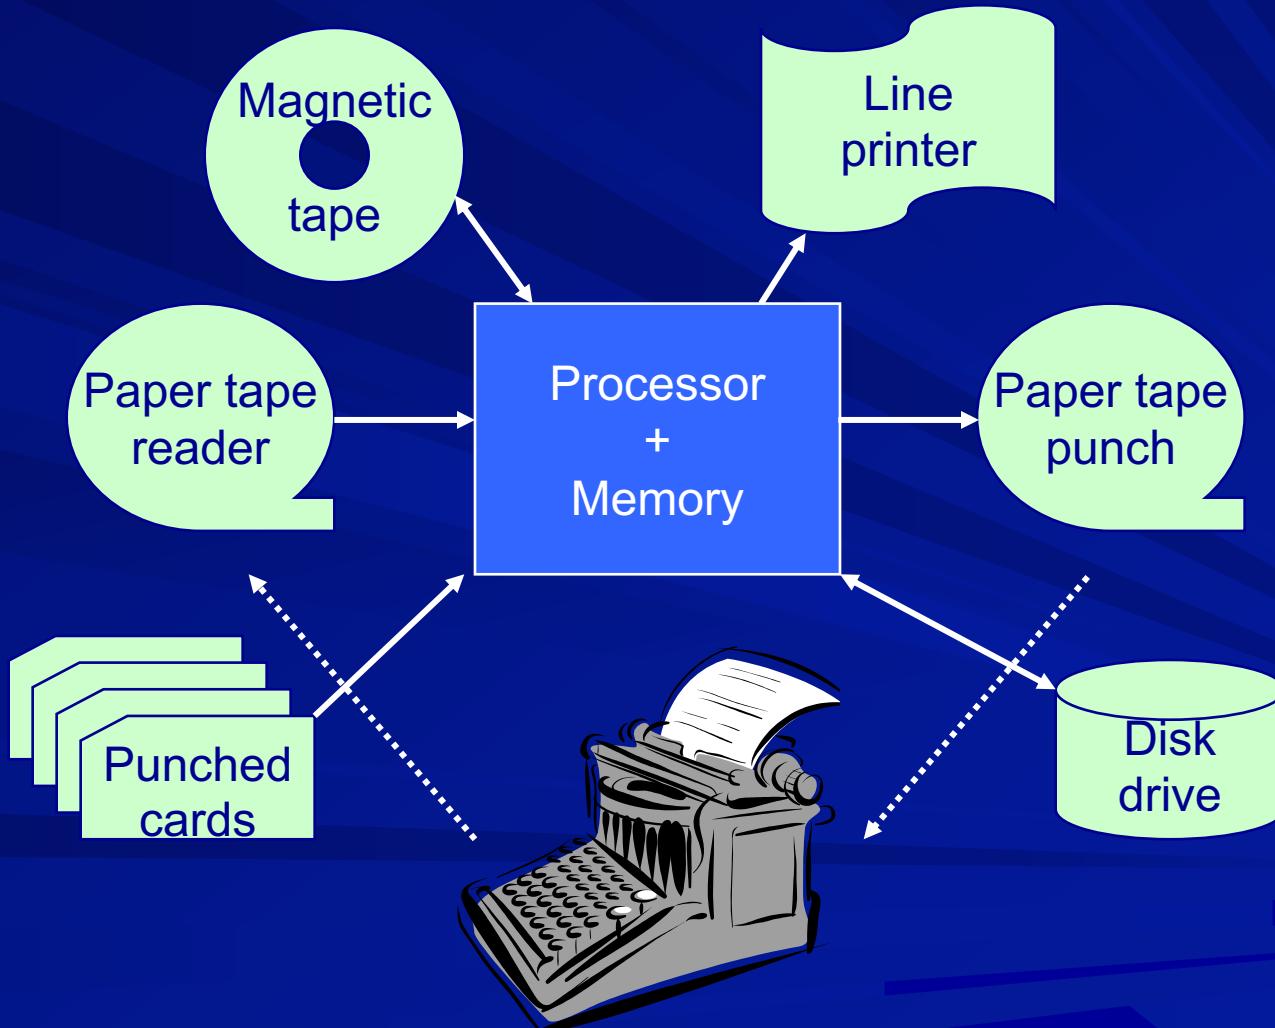
- switches, valves, solenoid, motors

Signal forms: on/off, analog, pulse train etc

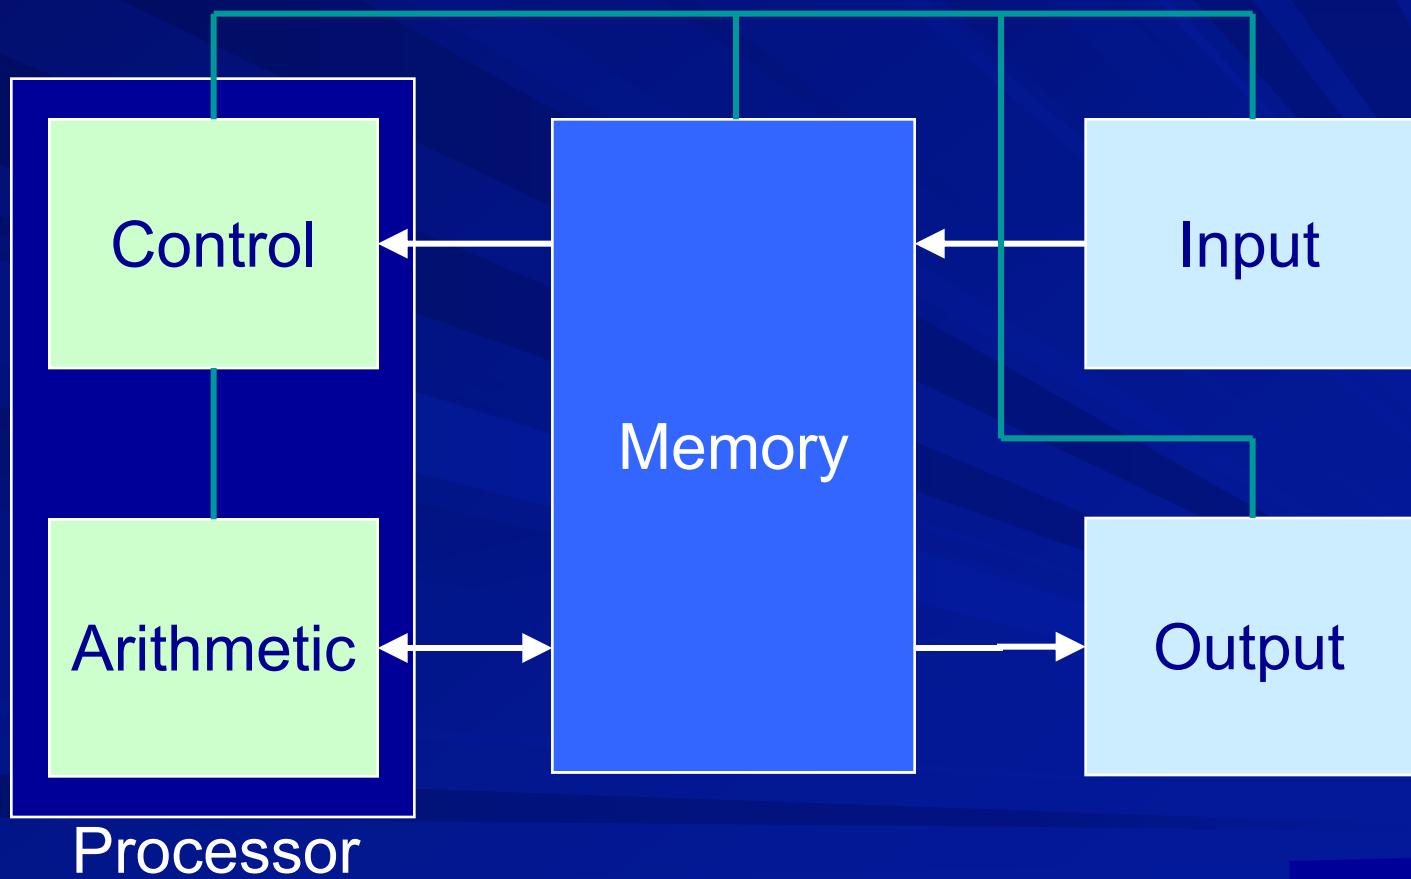
Storage devices

- Magnetic
 - disks, tapes
- Optical
 - CDROM, DVD, Blue Ray
- Semi-conductor
 - Nor Flash, Nand Flash

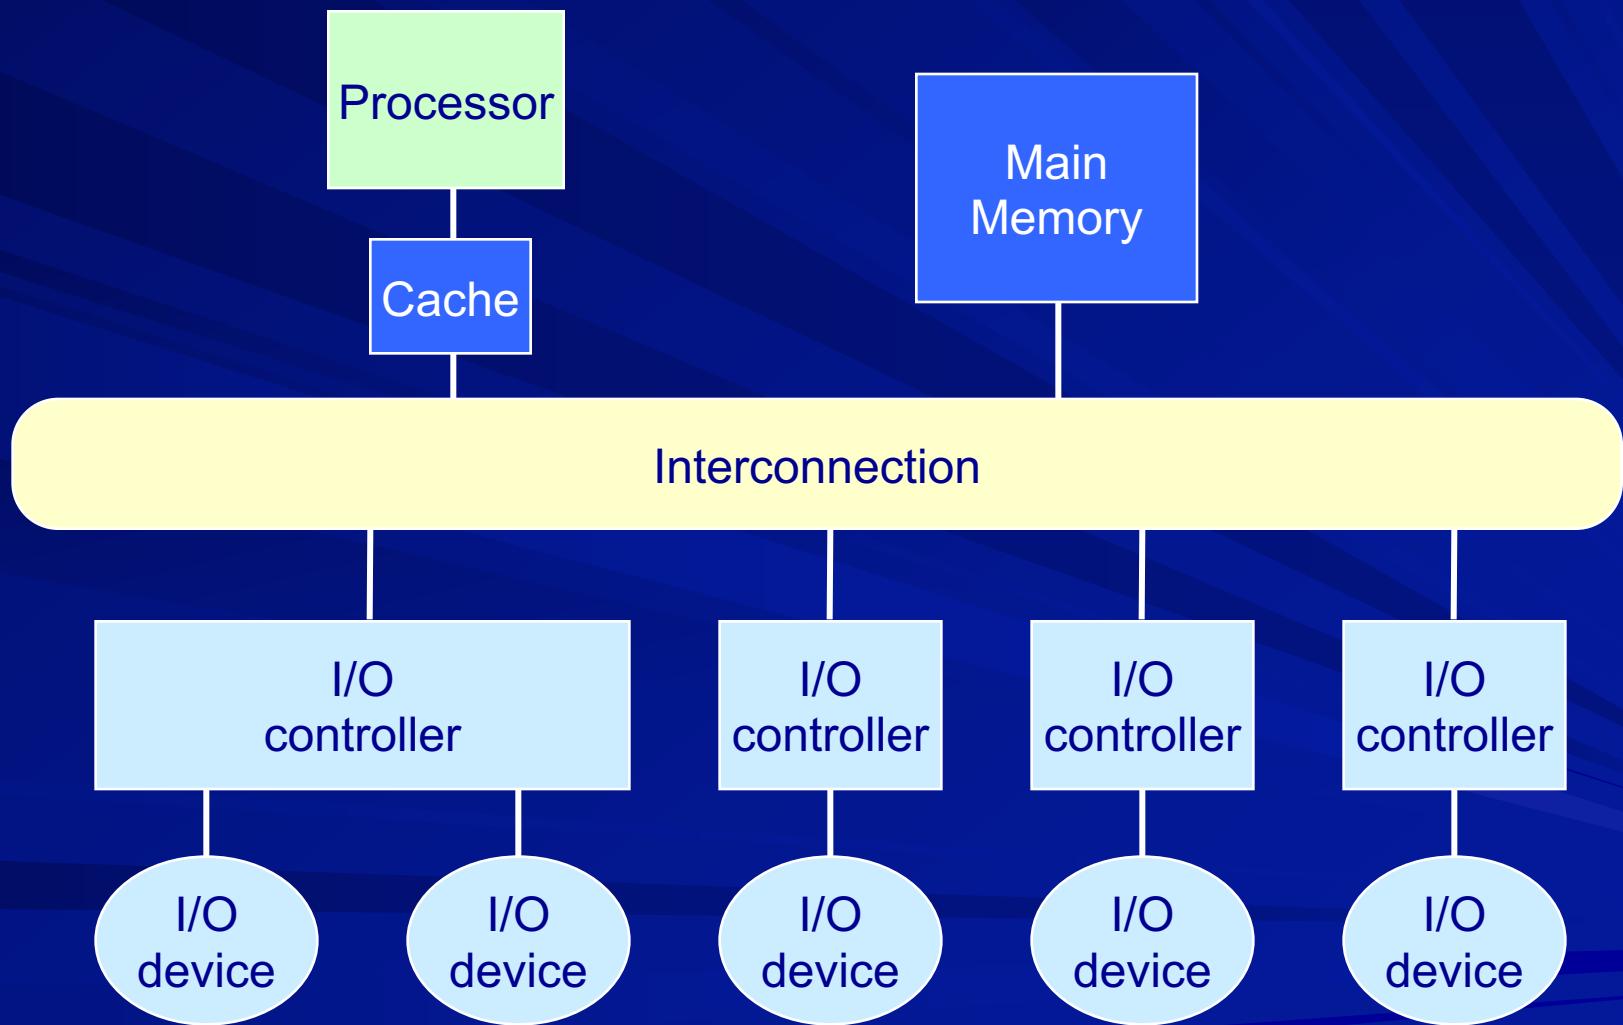
Peripherals circa 1970 : ICL 1909



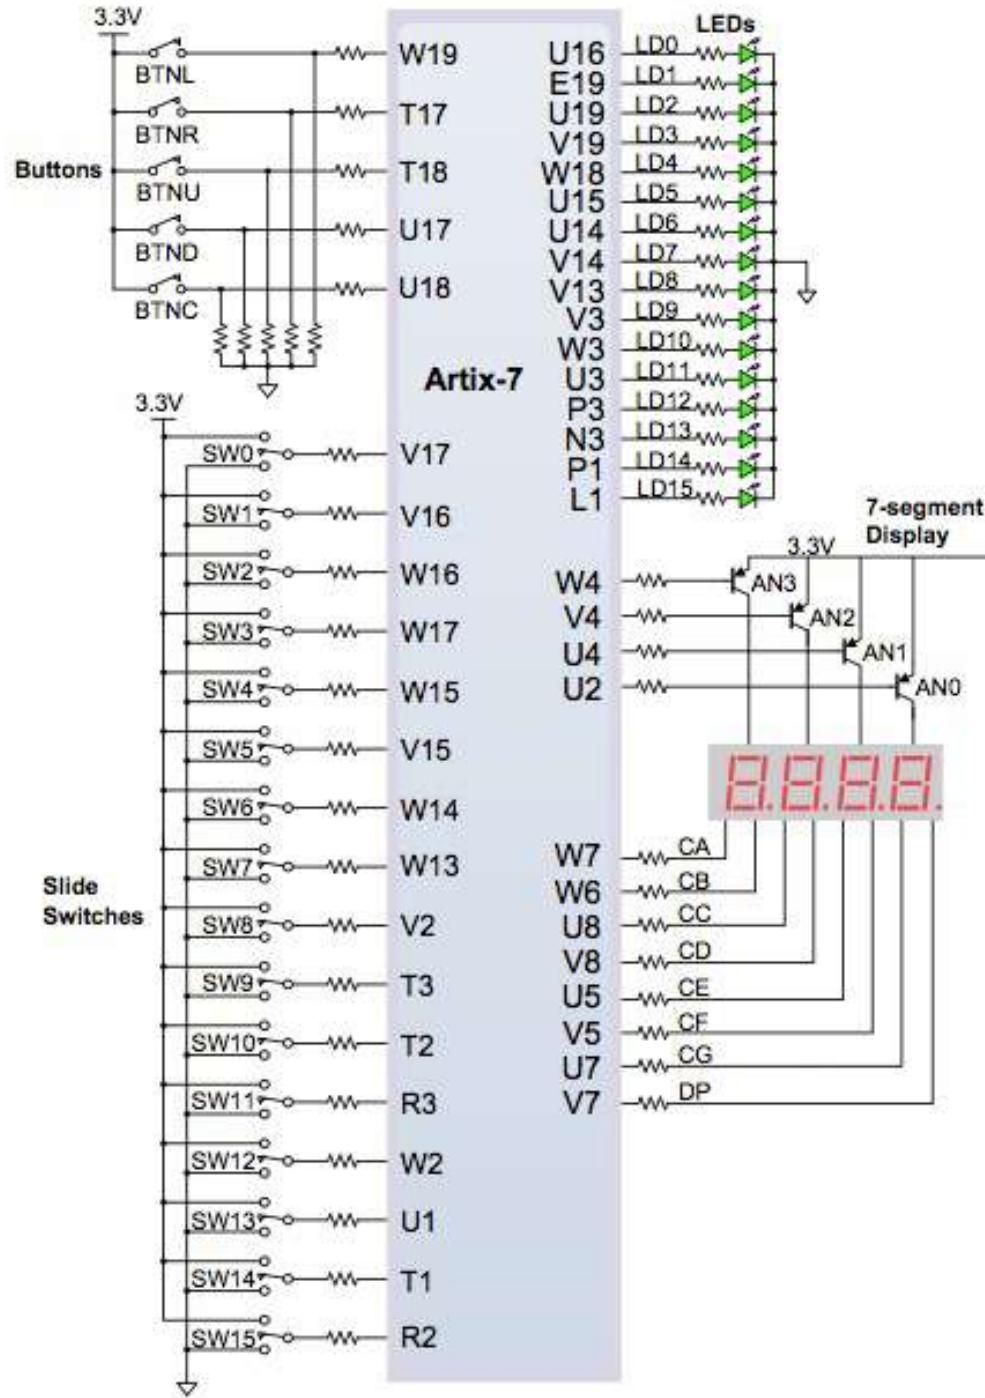
Computer System: Simplified Block Diagram



More realistic block diagram



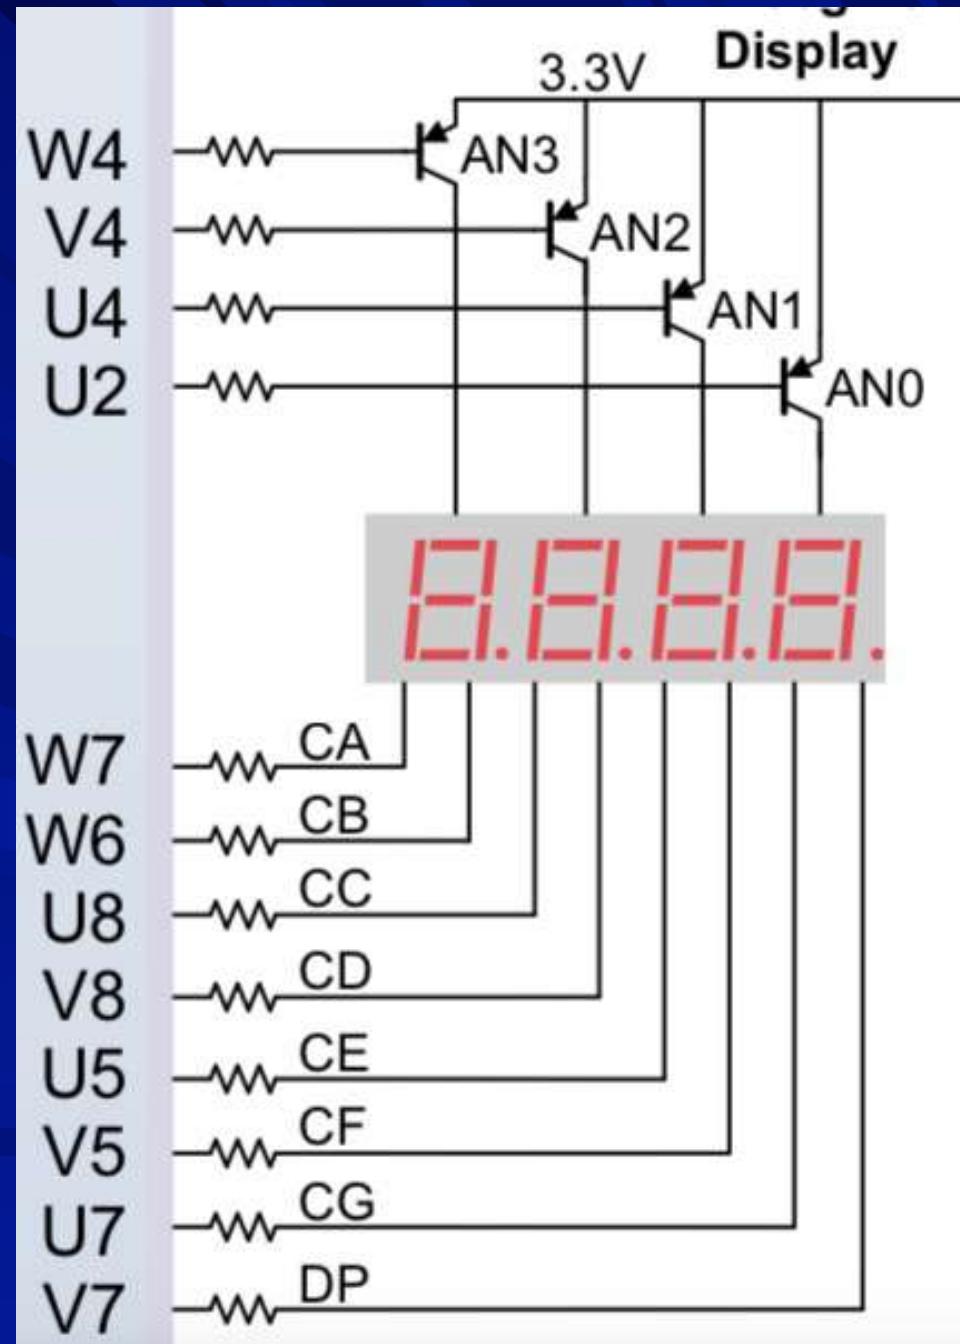
Switches and displays on BASYS3 board



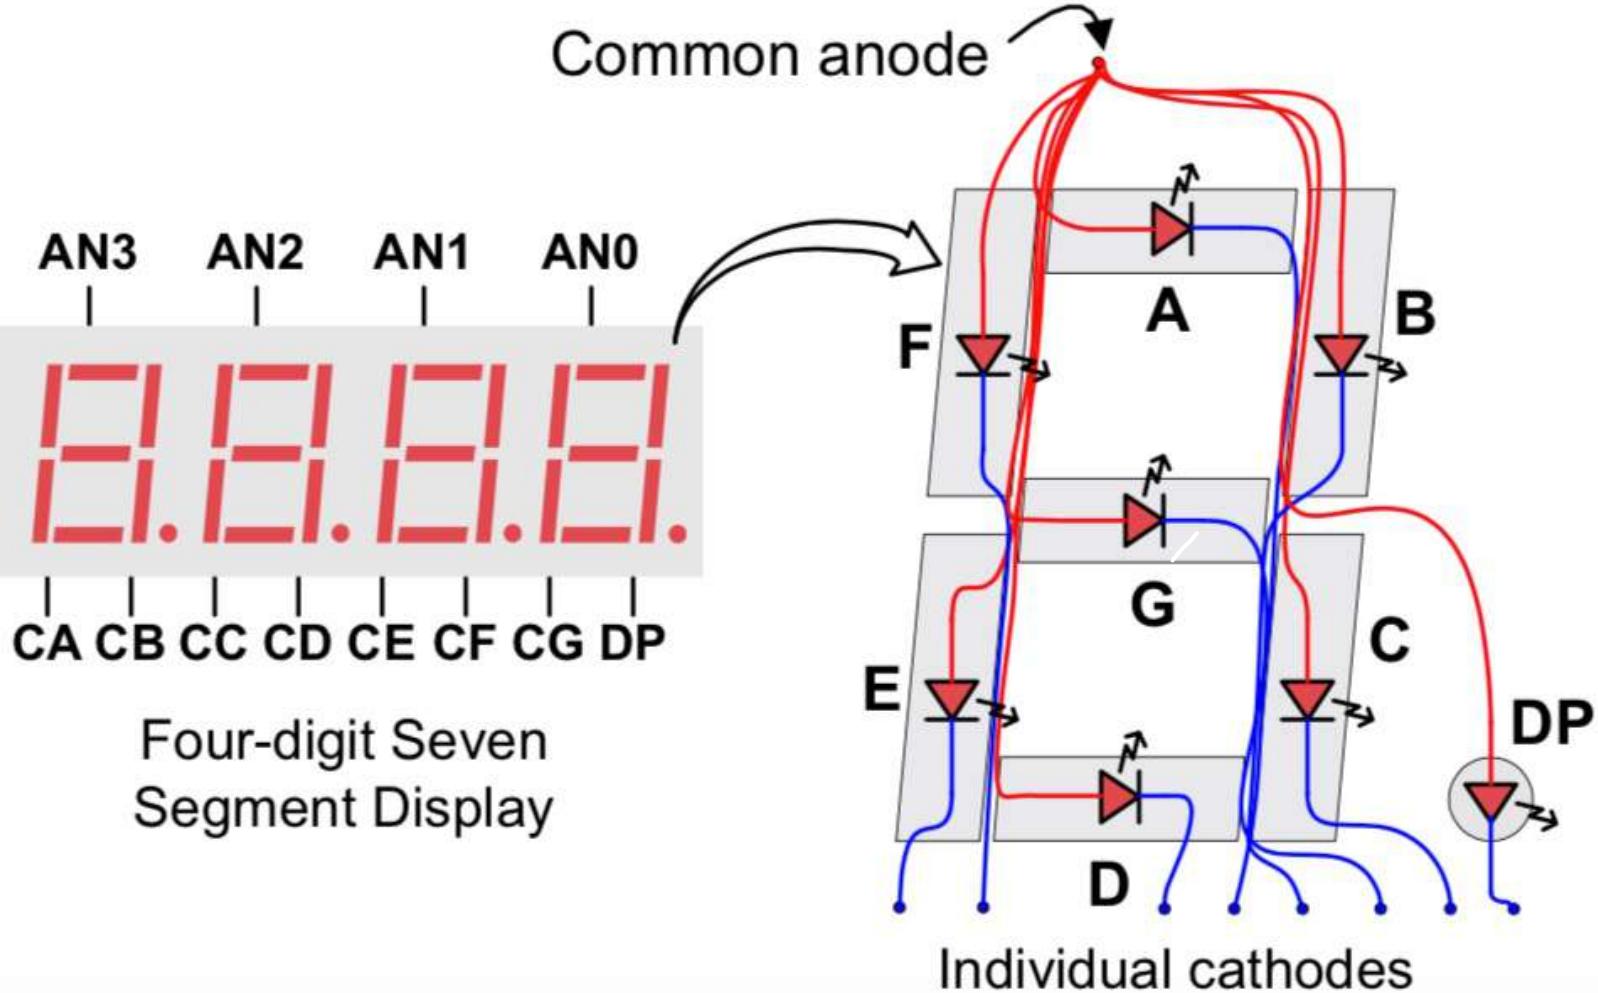
16 Button Key-pad



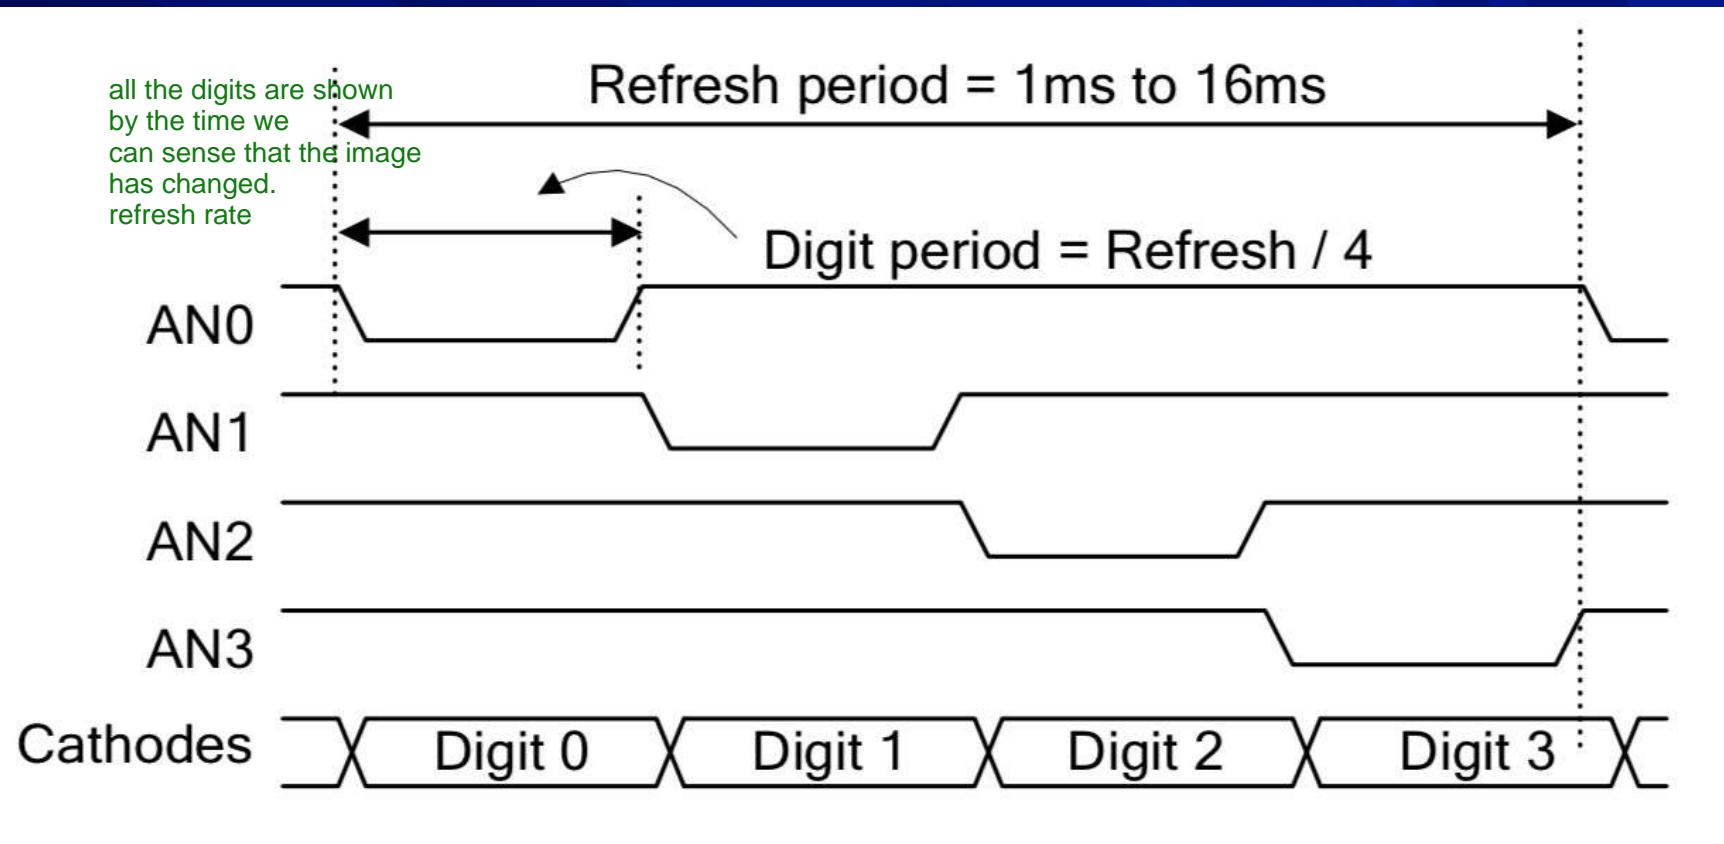
4-digit Display on BASYS 3 Board



7-segment Display

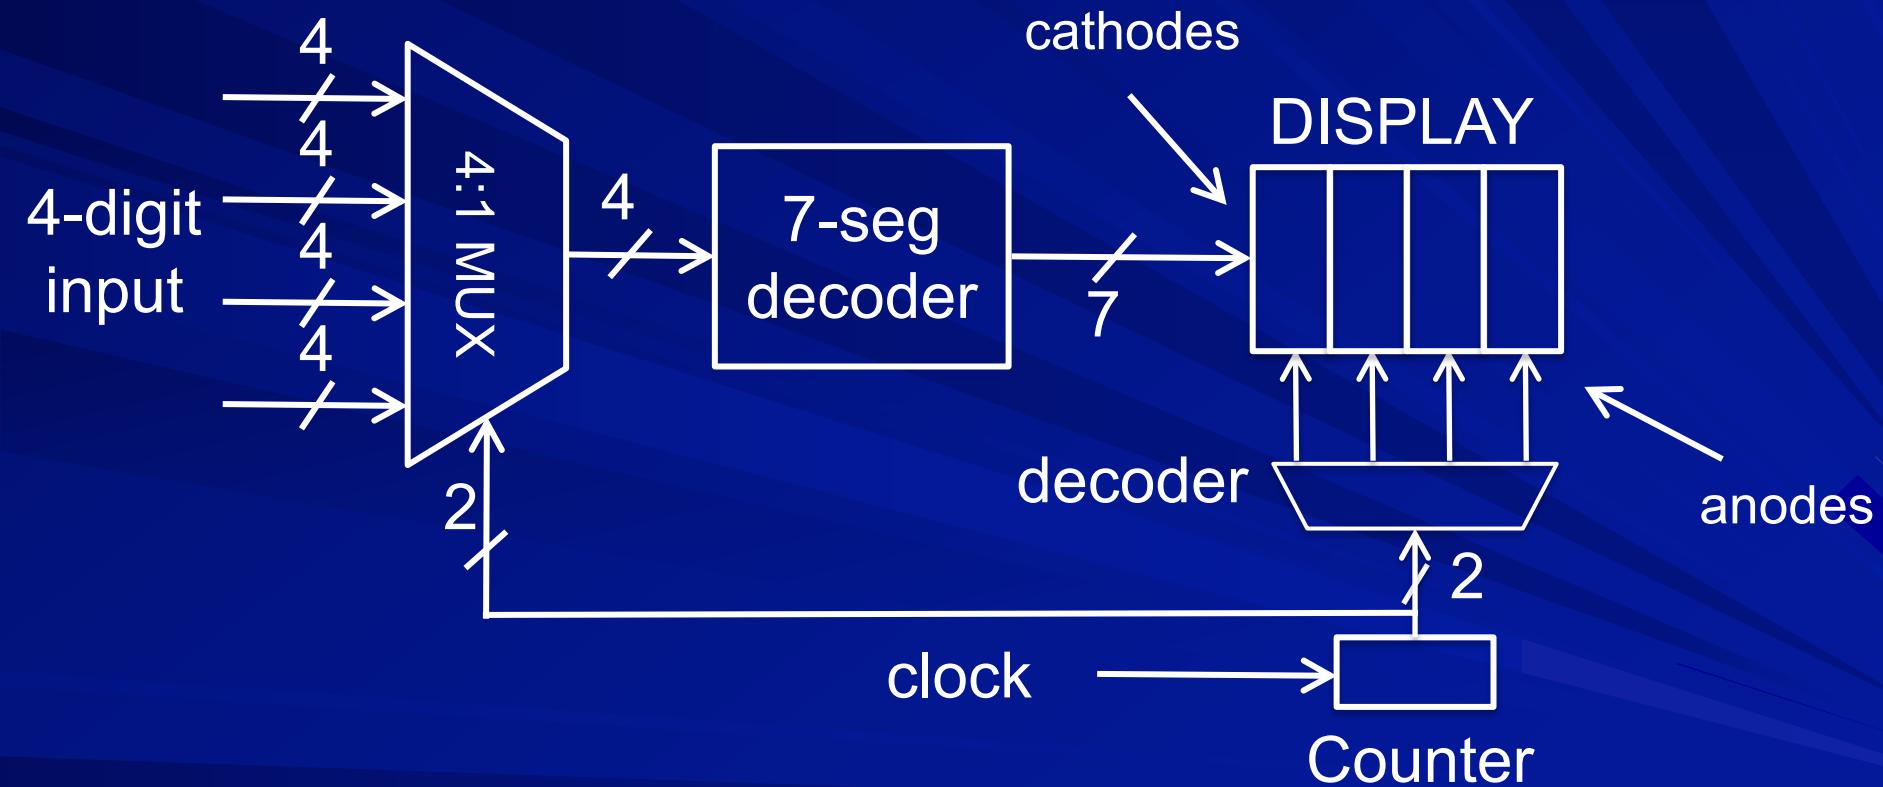


Refresh timings



Processor simply gives the digits, display controller displays the digit

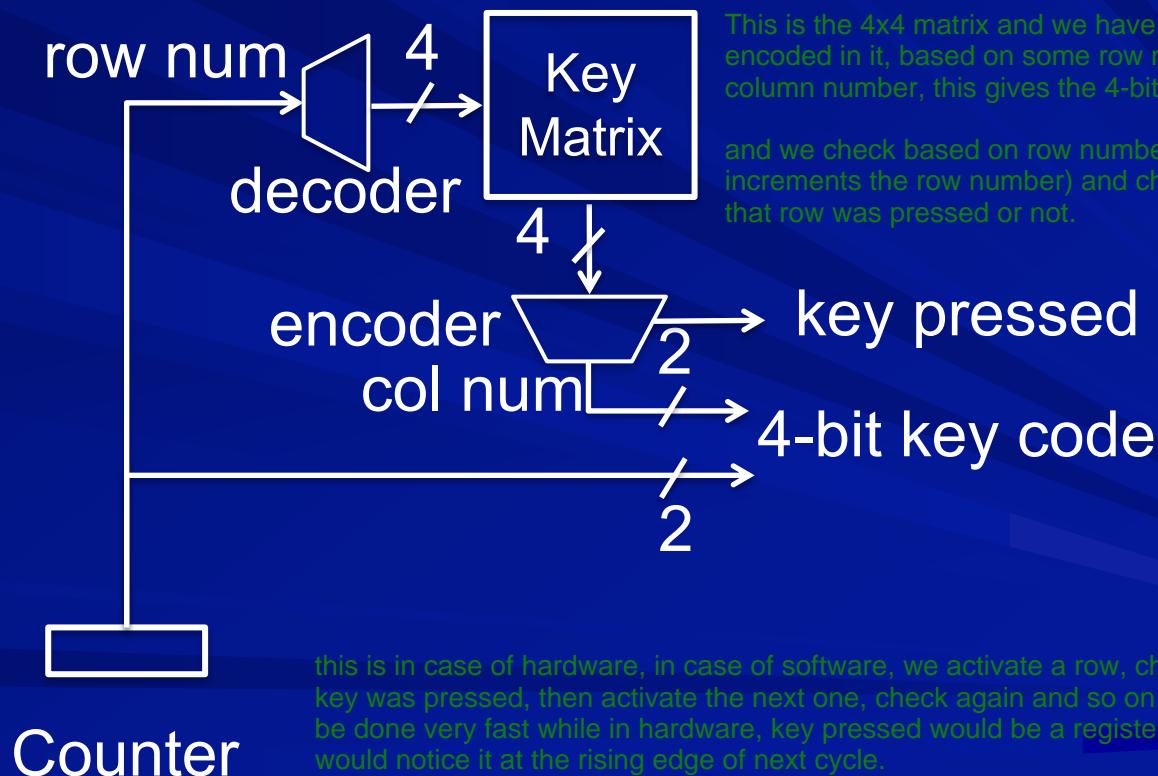
7-segment Display Controller



digits are being displayed in sequence, at a rate faster than our persistence of vision so we are able to see all the lit up pixels simultaneously when actually they are happening in sequence

4 x 4 Keyboard Encoder

We have many keys on the keyboard, so we can't have so many wires moving from the keyboard to the screen, hence we have a decoder.

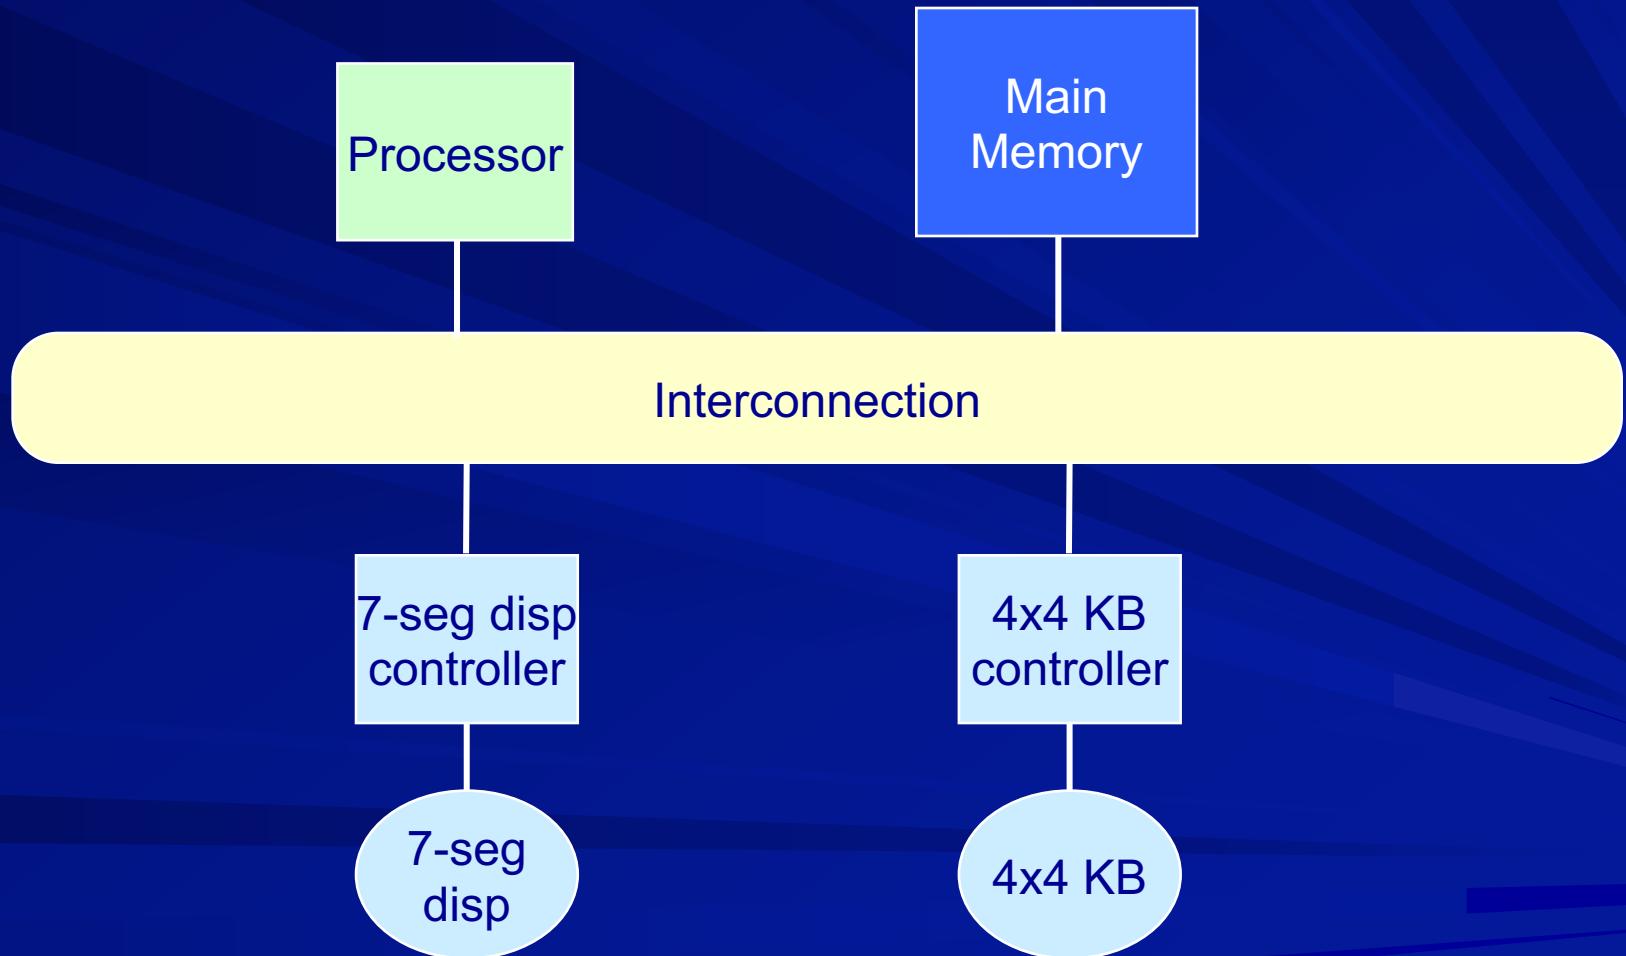


This is the 4x4 matrix and we have each key encoded in it, based on some row number and column number, this gives the 4-bit key code.

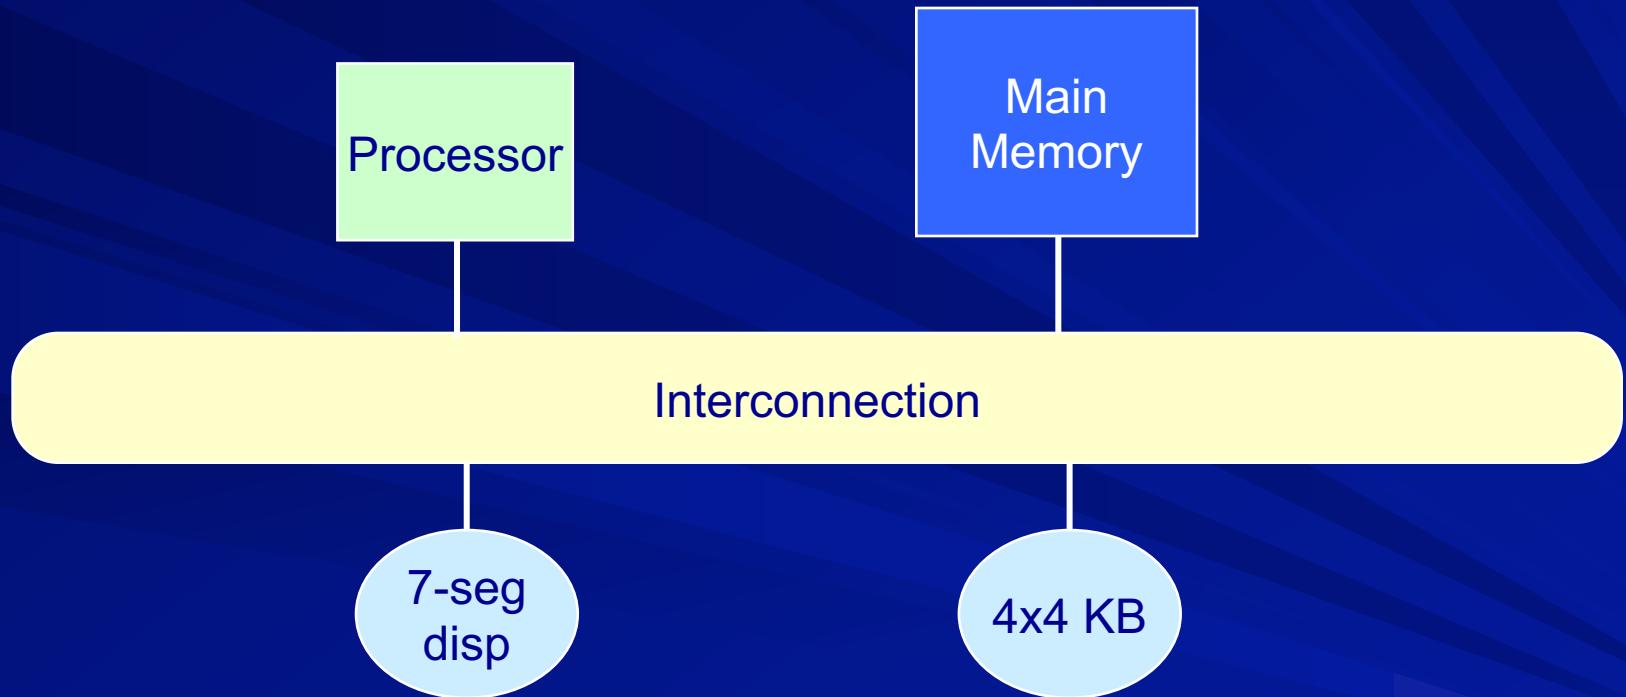
and we check based on row number (counter increments the row number) and check if any key in that row was pressed or not.

this is in case of hardware, in case of software, we activate a row, check if any column key was pressed, then activate the next one, check again and so on and this needs to be done very fast while in hardware, key pressed would be a register so processor would notice it at the rising edge of next cycle.

Block diagram



Block diagram



Low level control by software

Instructions for I/O

Use SWI instruction/system call

Require instruction for I/O inside handler

- Special input/output instructions

- specify device
 - specify register/memory location

=> Dedicated I/O

- Use LDR, STR instructions

- use devices as memory locations

=> Memory mapped I/O

I/O address space

- Memory mapped I/O
 - part of memory address space reserved for I/O
 - usual load/store instructions are used for I/O
 - memory ignores these addresses
- Dedicated I/O
 - separate address space for I/O - much smaller than the memory space
 - control signal from processor indicates which space is being addressed
 - separate instructions are required

Addressing I/O devices

Each I/O device (controller) viewed as a set of registers or ports by CPU

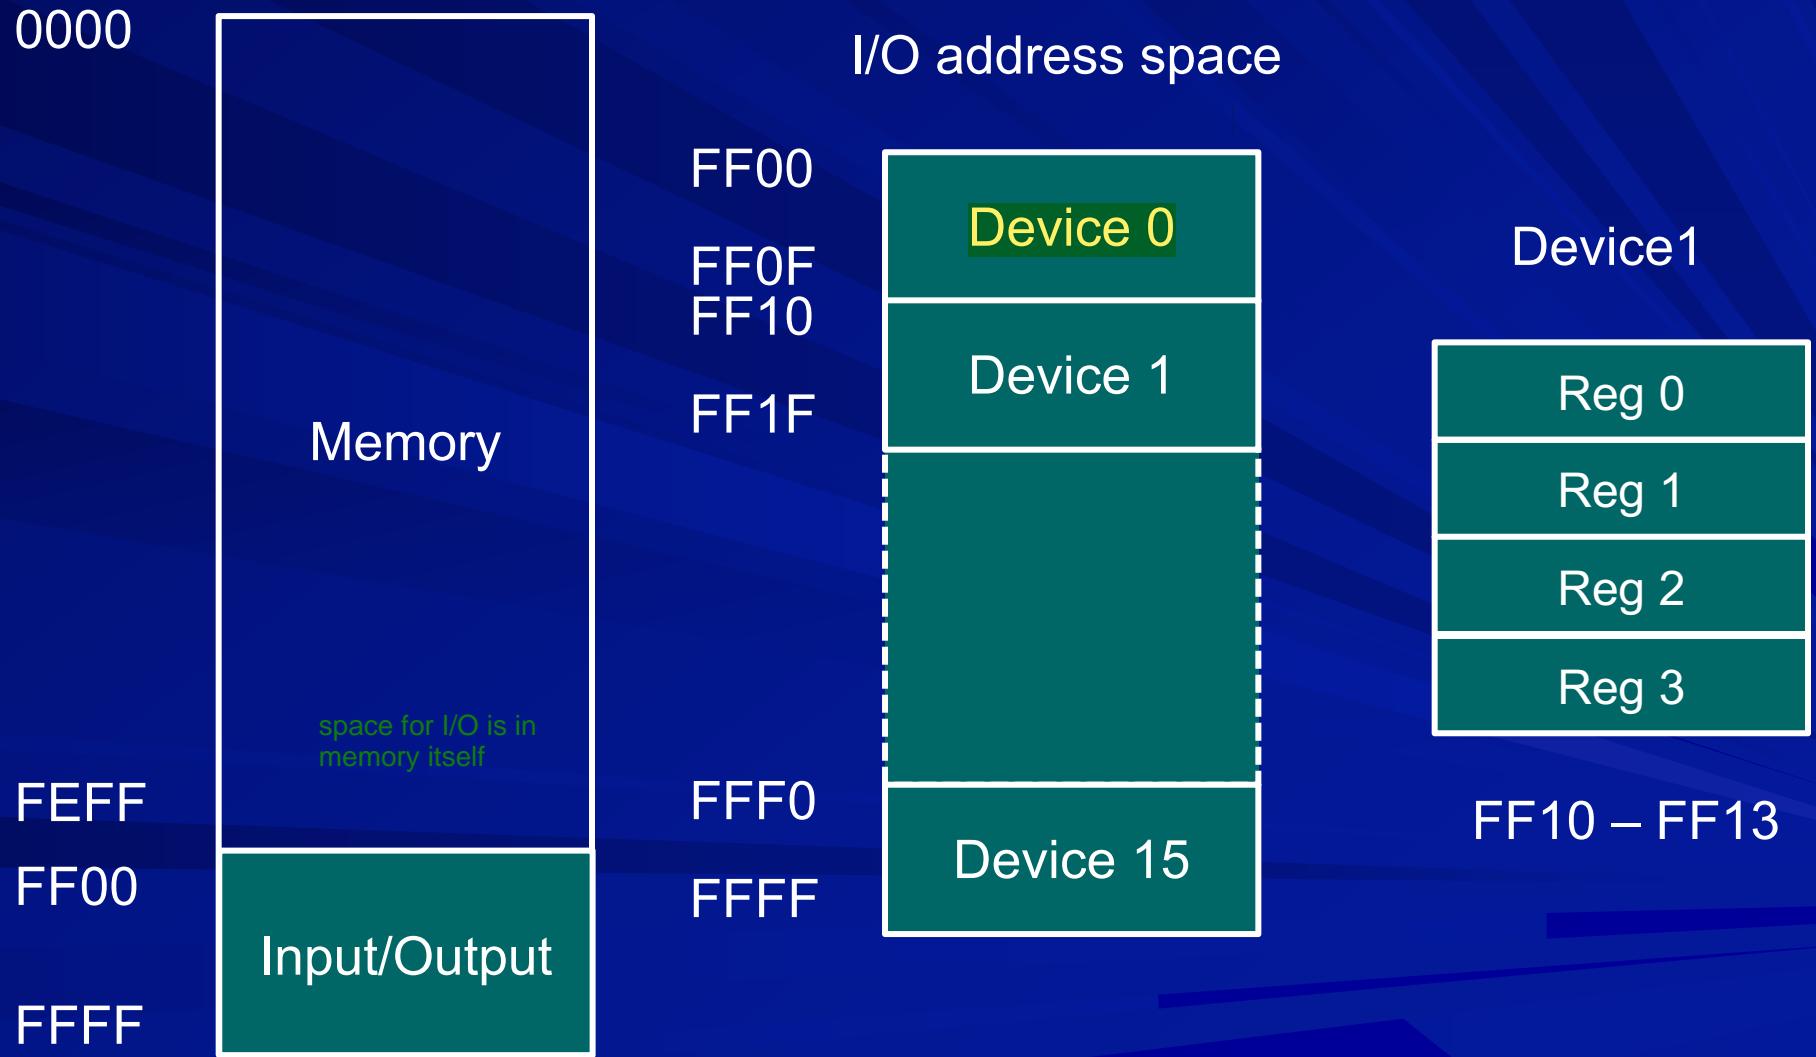
CPU can use these registers to

- write commands and parameters
- read status
- read/write data

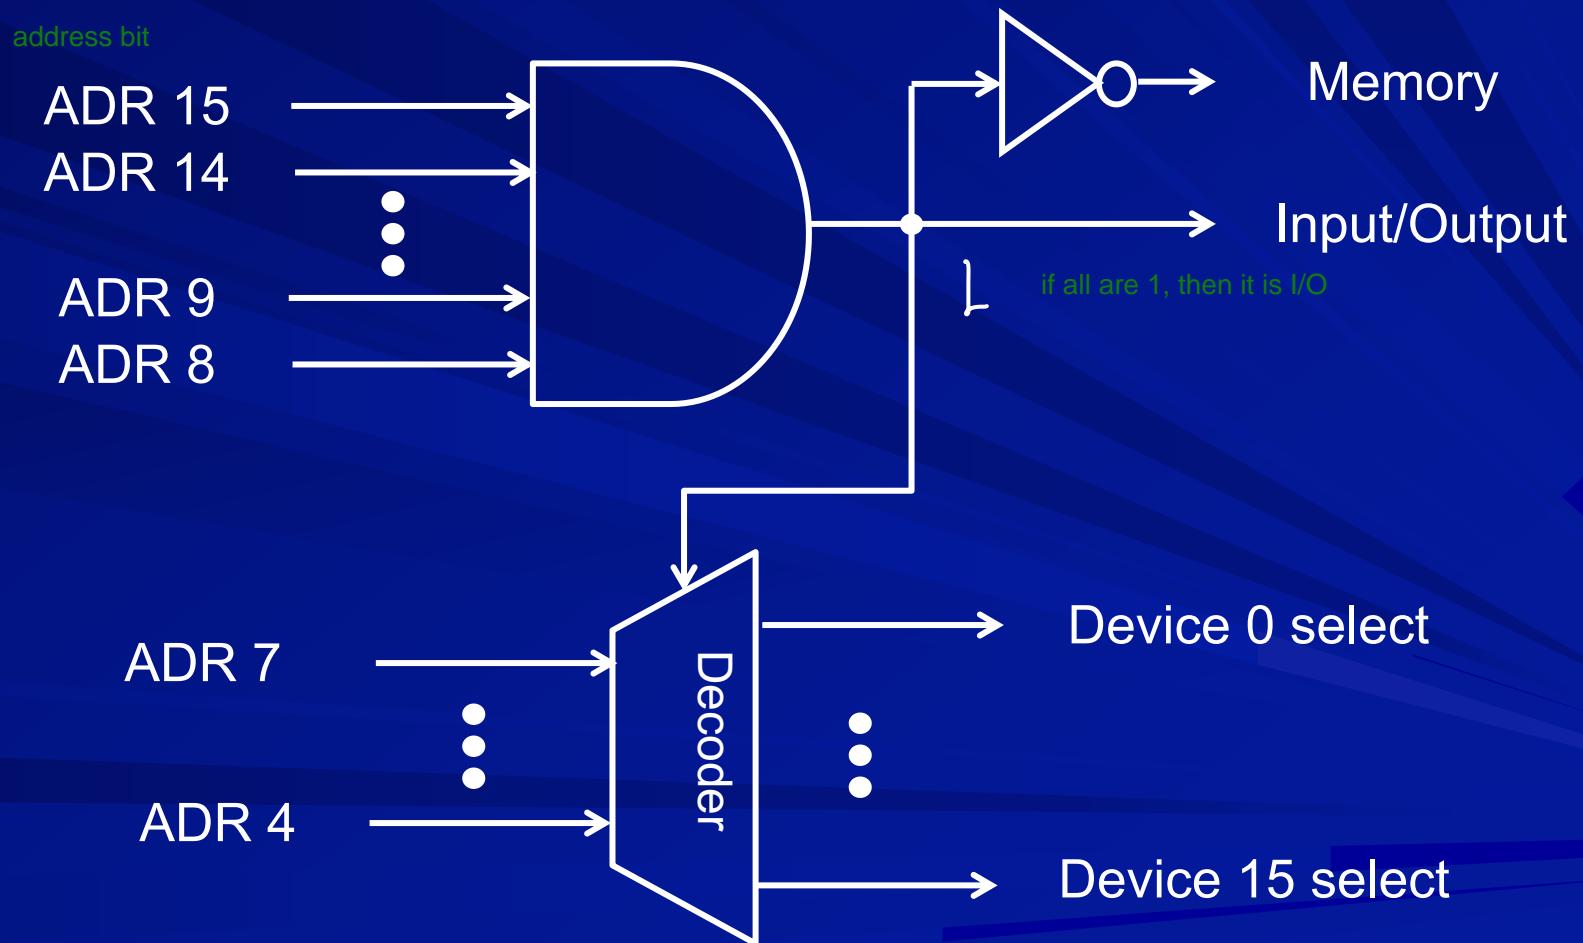
These registers are addressable

- Memory mapped I/O
- Dedicated I/O

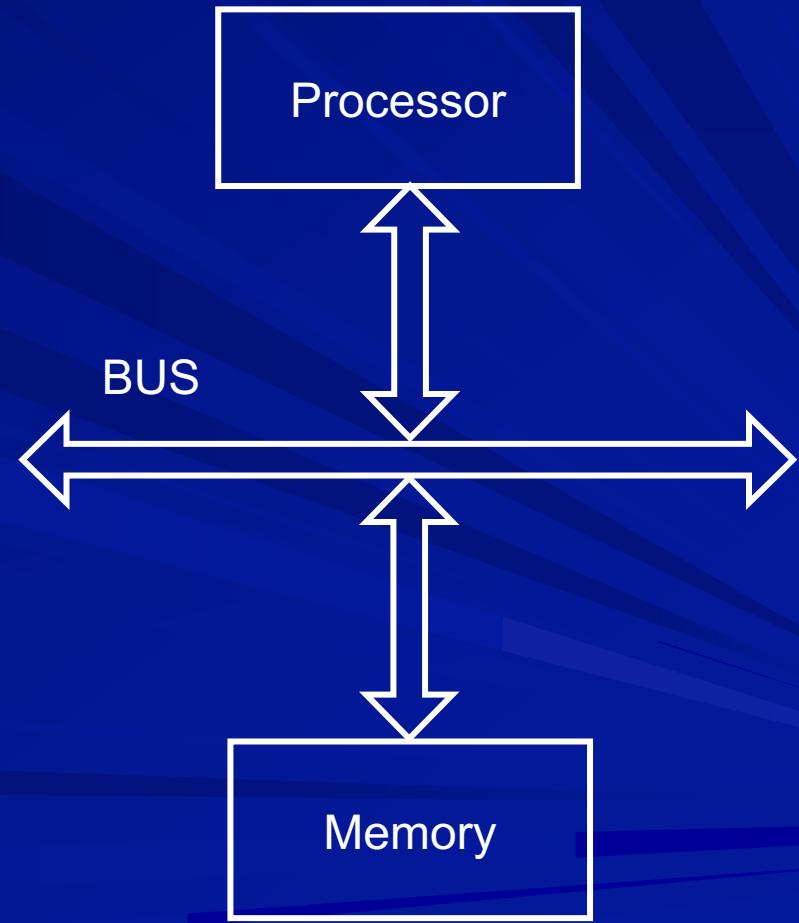
Address Map Example



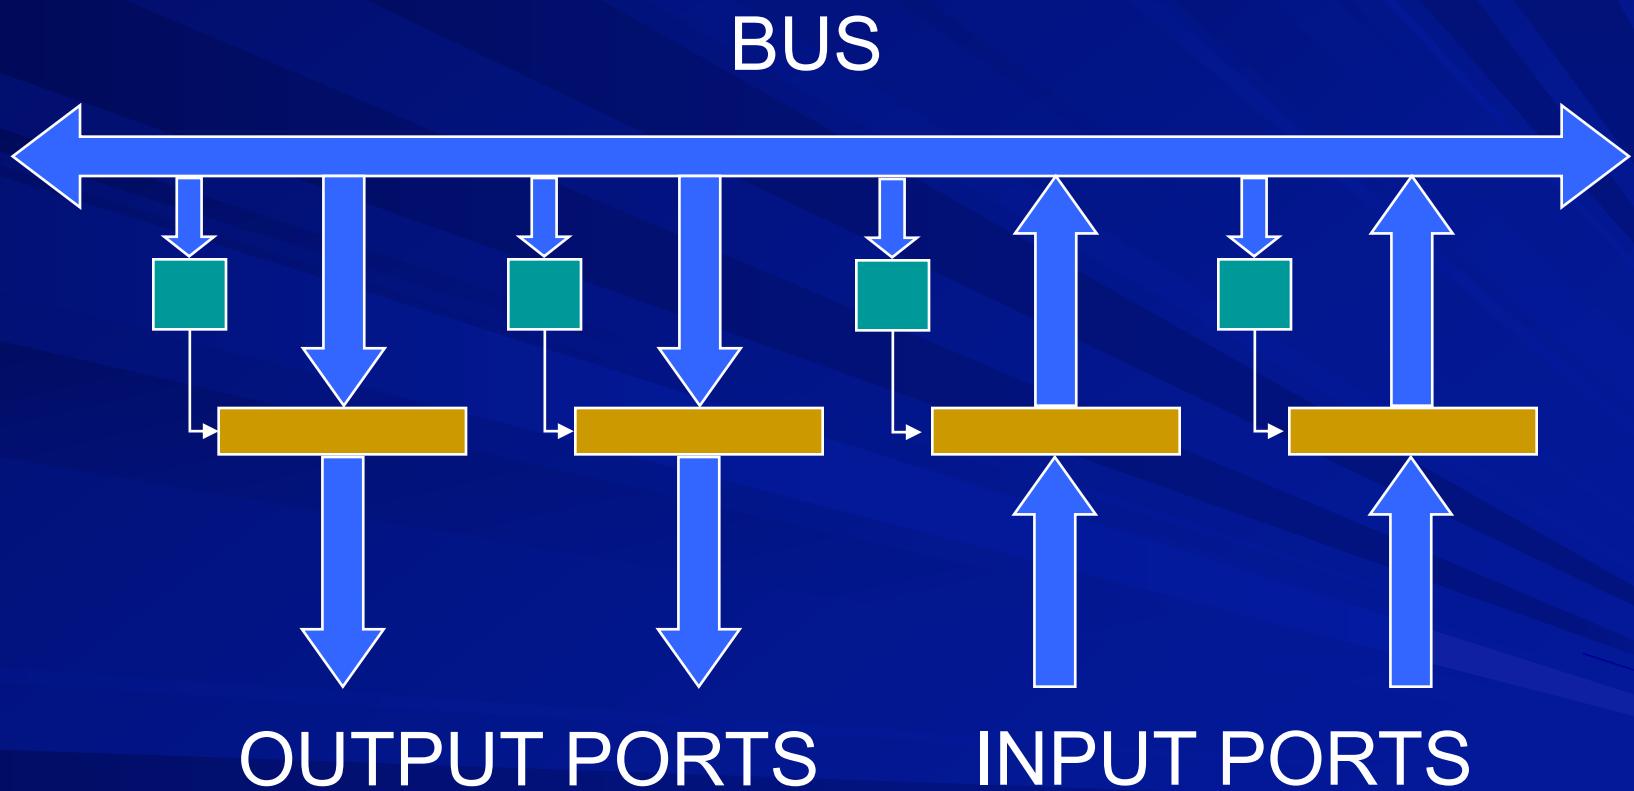
Decoding Address



Interconnection using a bus



INPUT / OUTPUT PORTS



Ports for Disp, KB controllers

7 segment display controller

Display

out	Digit 0
out	Digit 1
out	Digit 2
out	Digit 3

input

Key Board

key has been pressed or not

in	✓ status
in	✓ key code

controlled by software

Ports for raw Disp, KB

raw means low level control

Display

out
out

anode pattern
cathode pattern

4 bit anode
7 bit cathode

since we do this for 1 digit, so we need to
refresh every 10 ms

refresh required

Key Board

out
in

row pattern
column pattern

then we won't deal
with code, we shall be
dealing with row
pattern and column
pattern directly.

repeated scan required

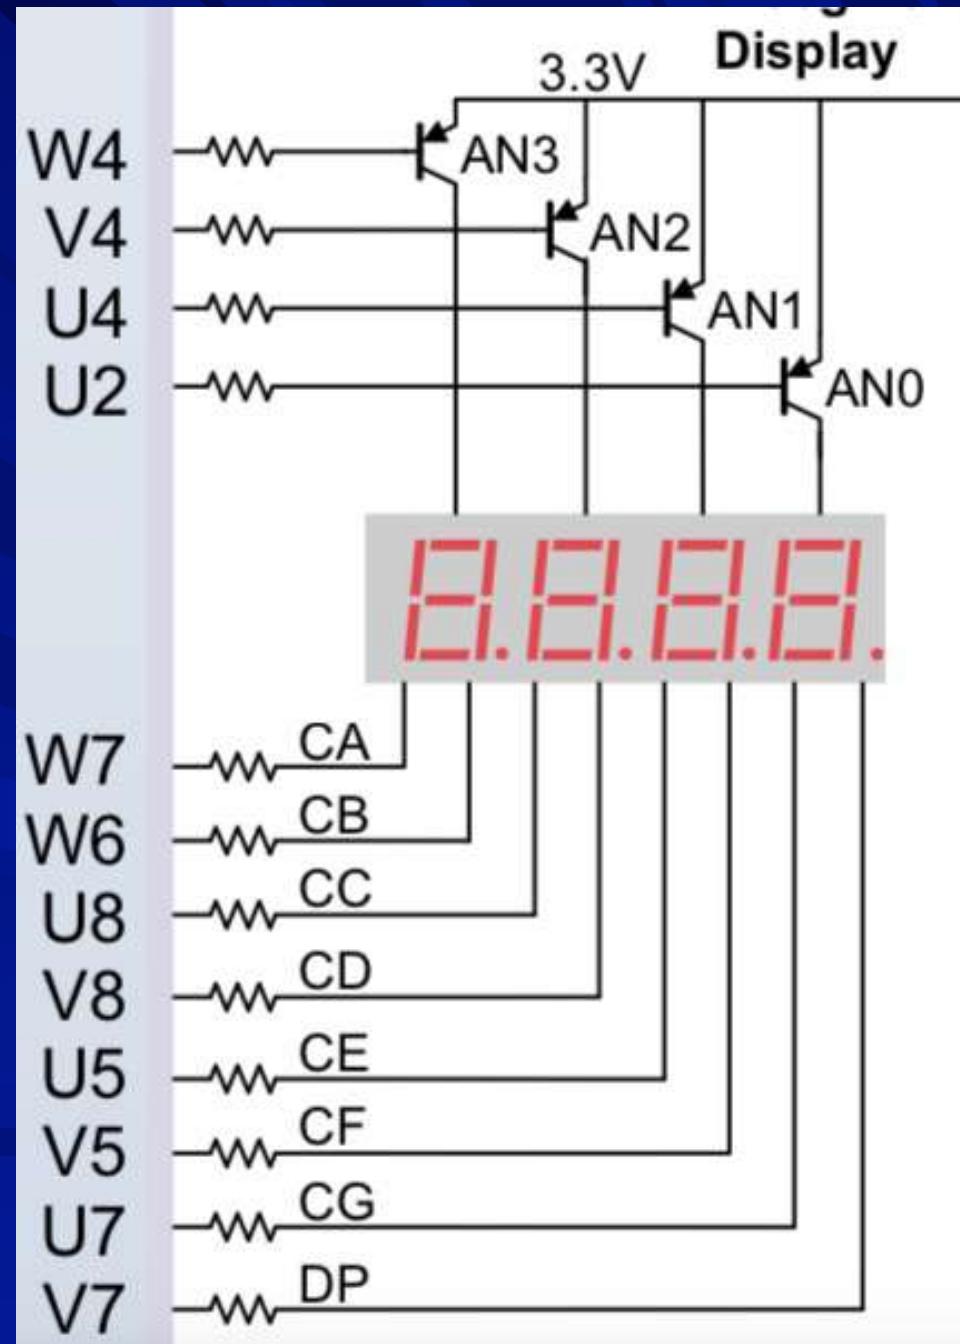
COL216

Computer Architecture

Input/Output – 2

17th March 2022

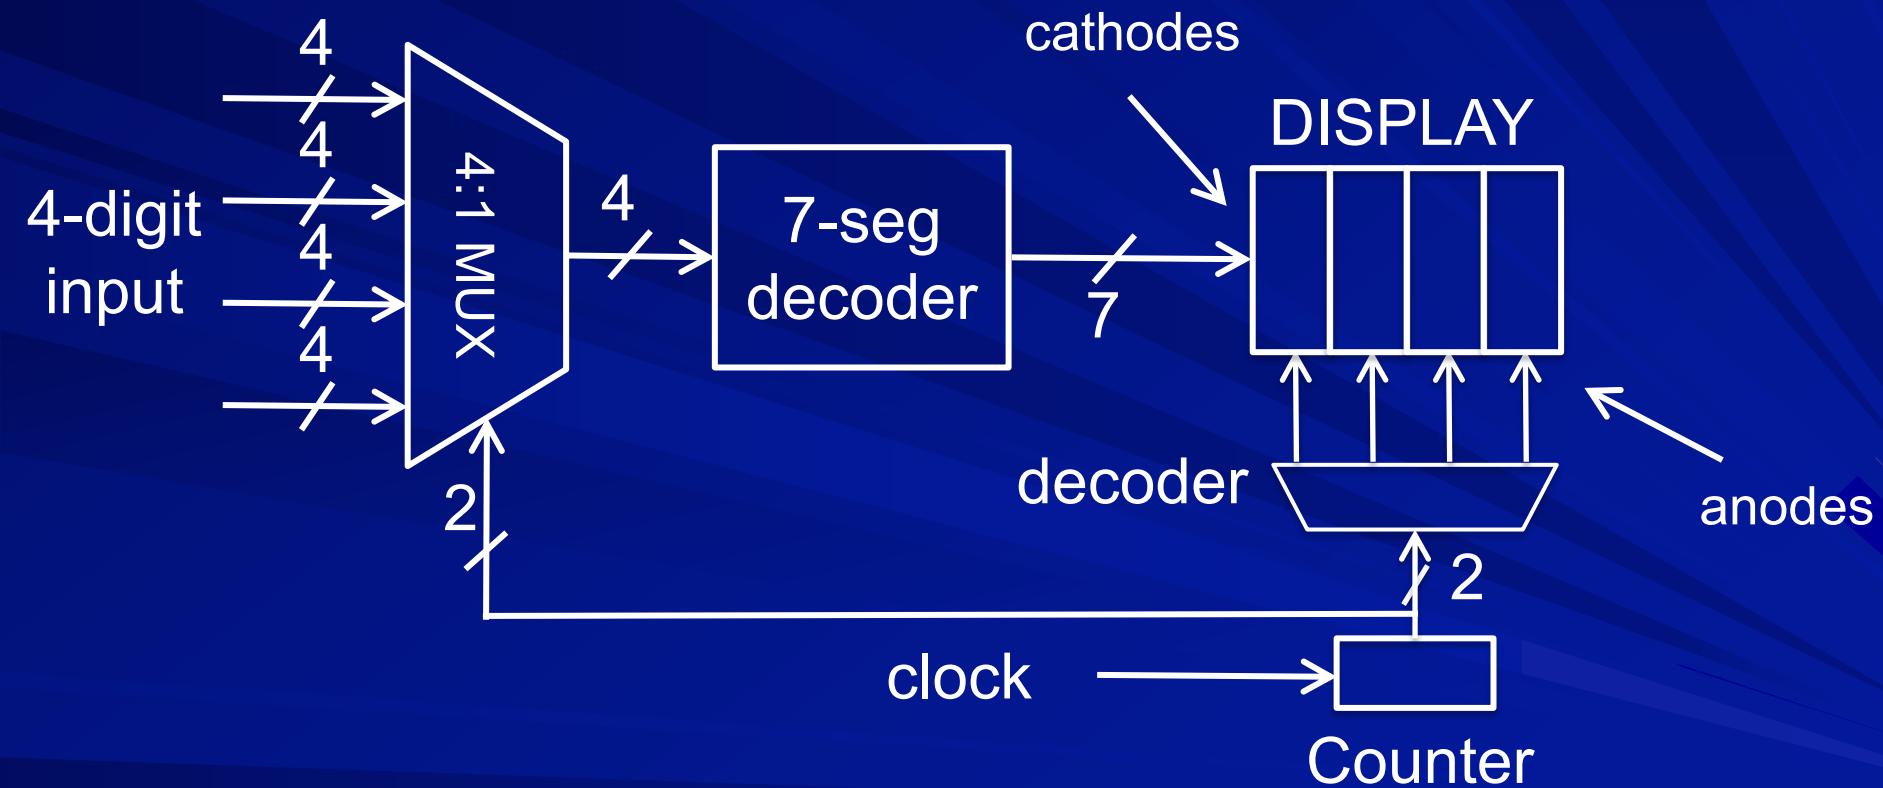
4-digit Display on BASYS 3 Board



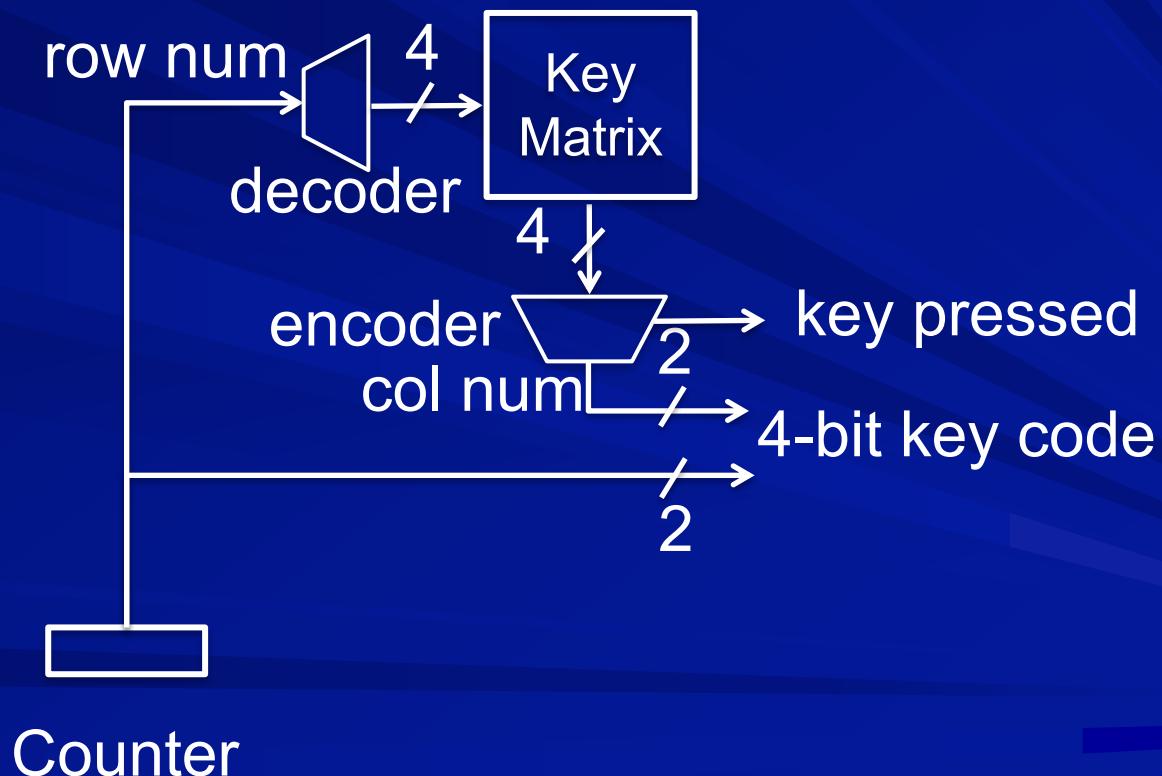
16 Button Key-pad



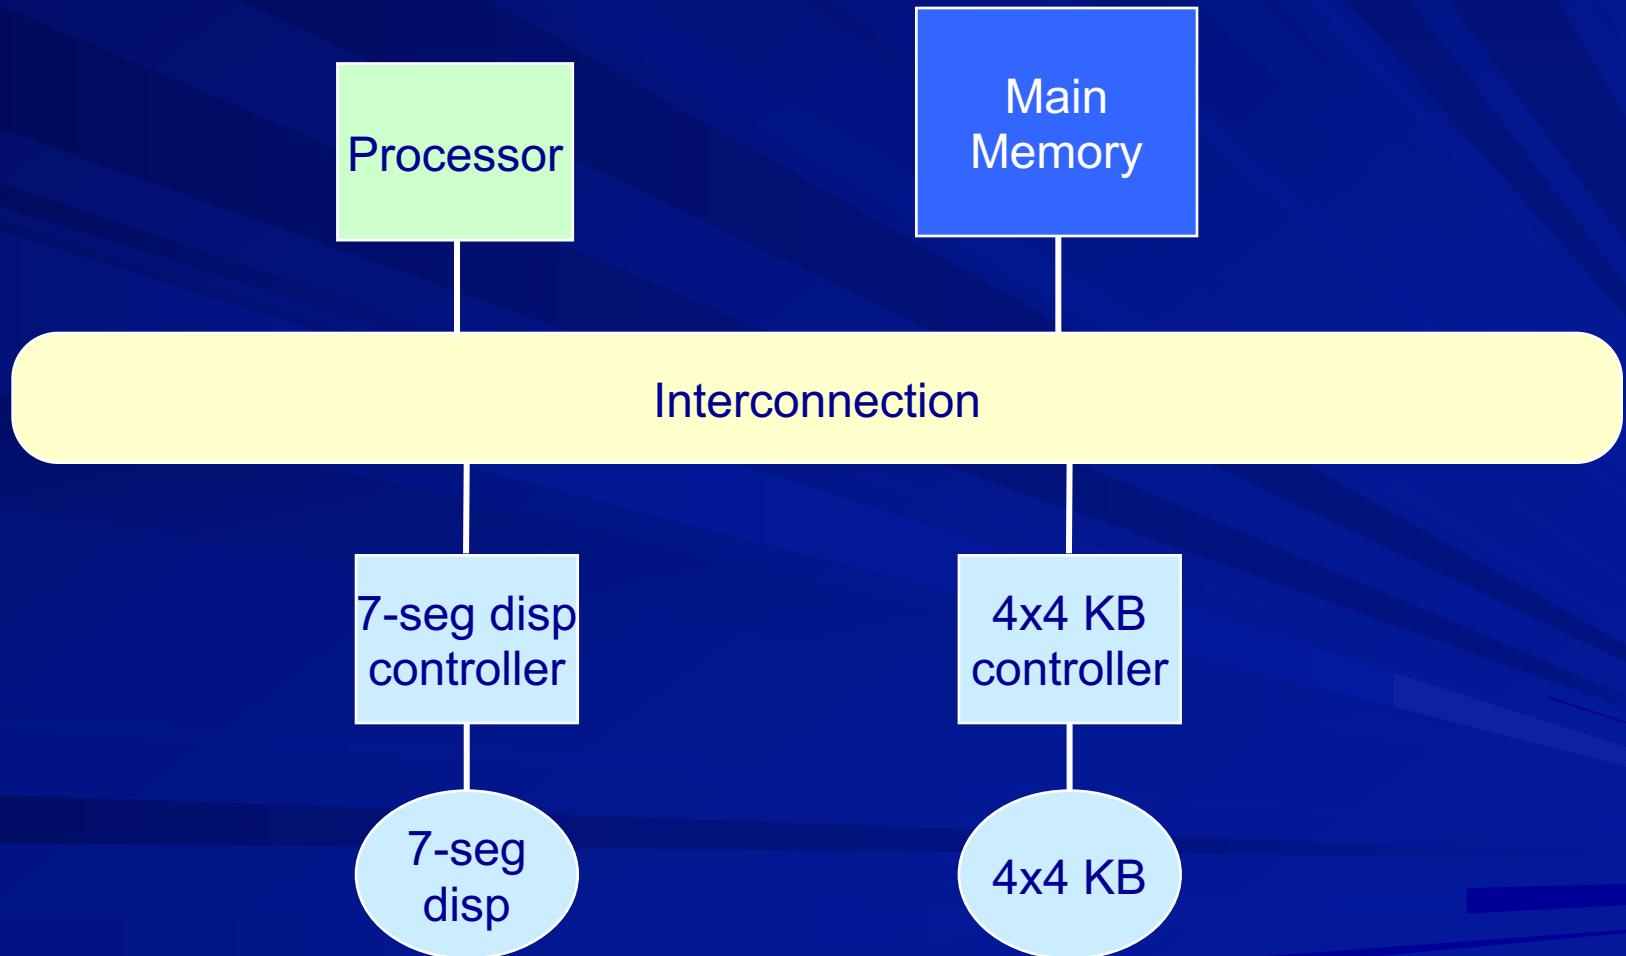
7-segment Display Controller



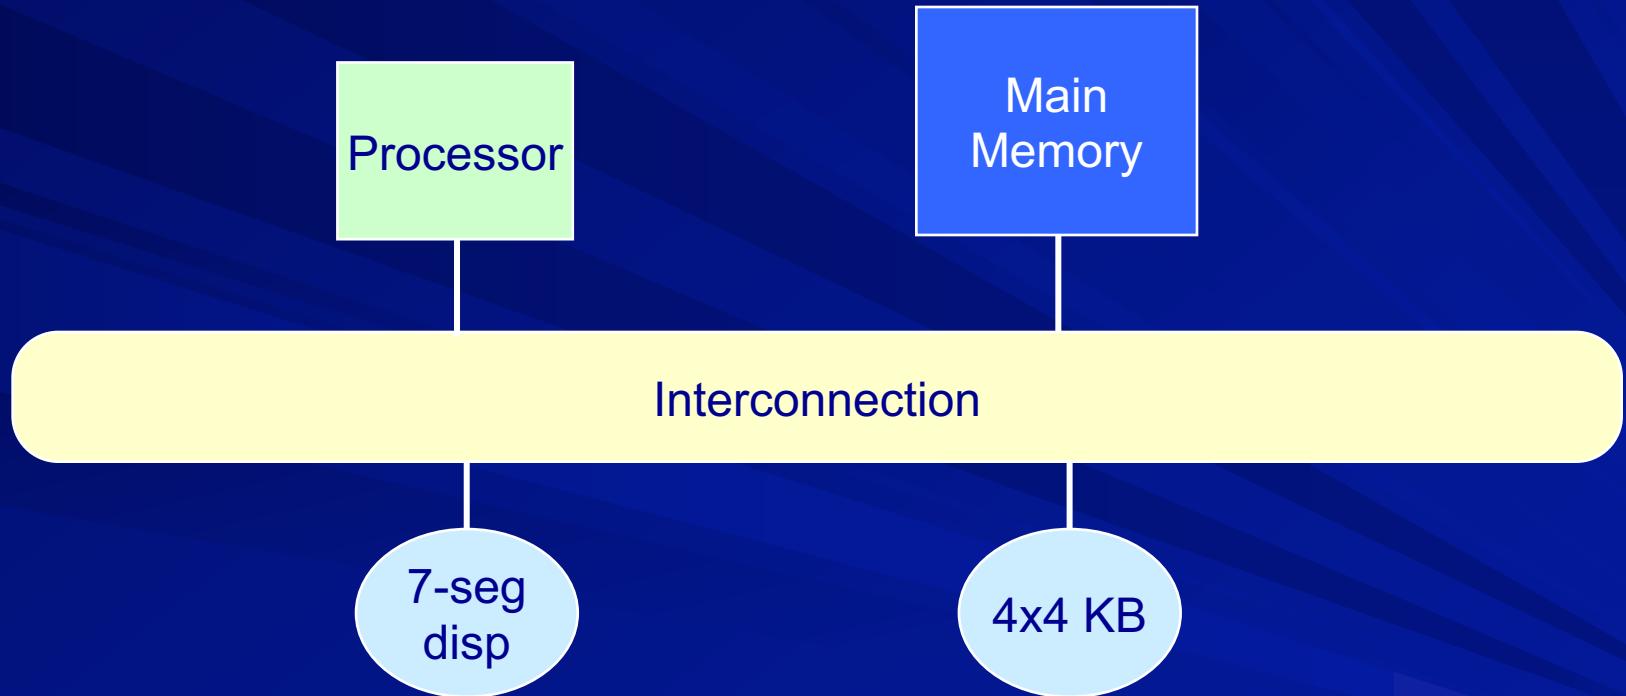
4 x 4 Keyboard Encoder



Block diagram

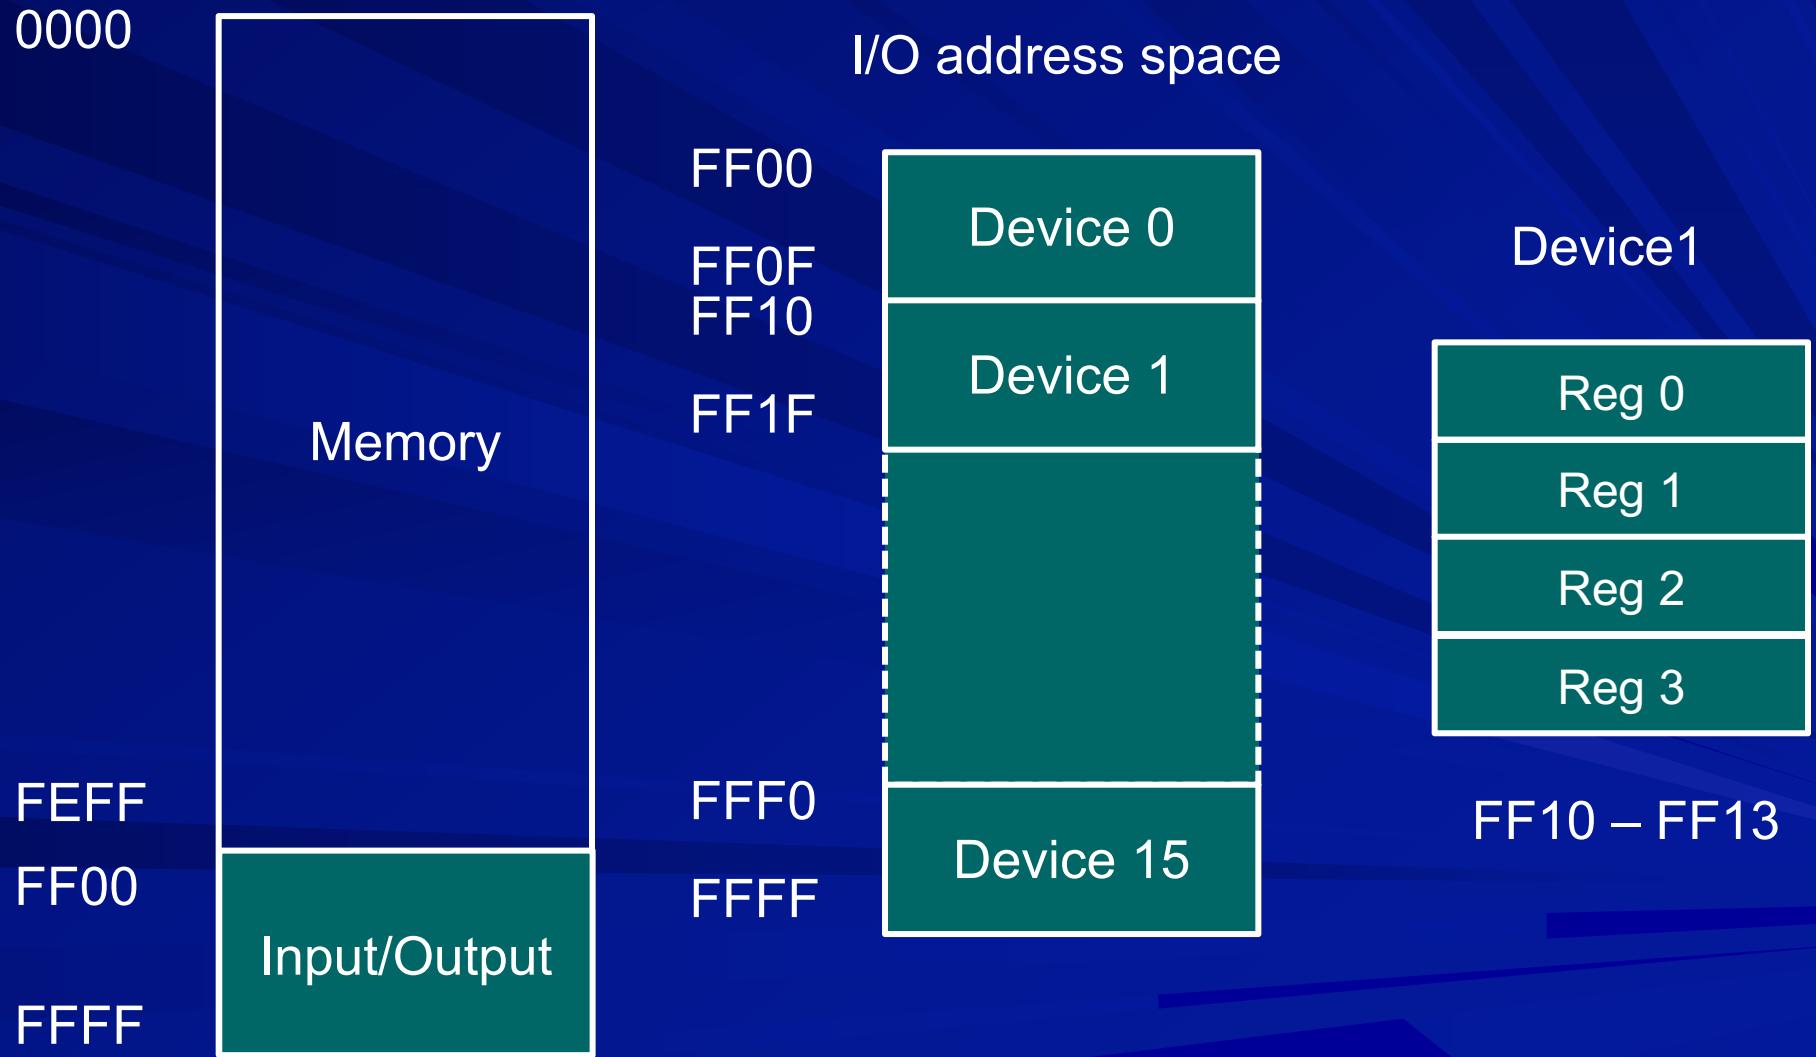


Block diagram

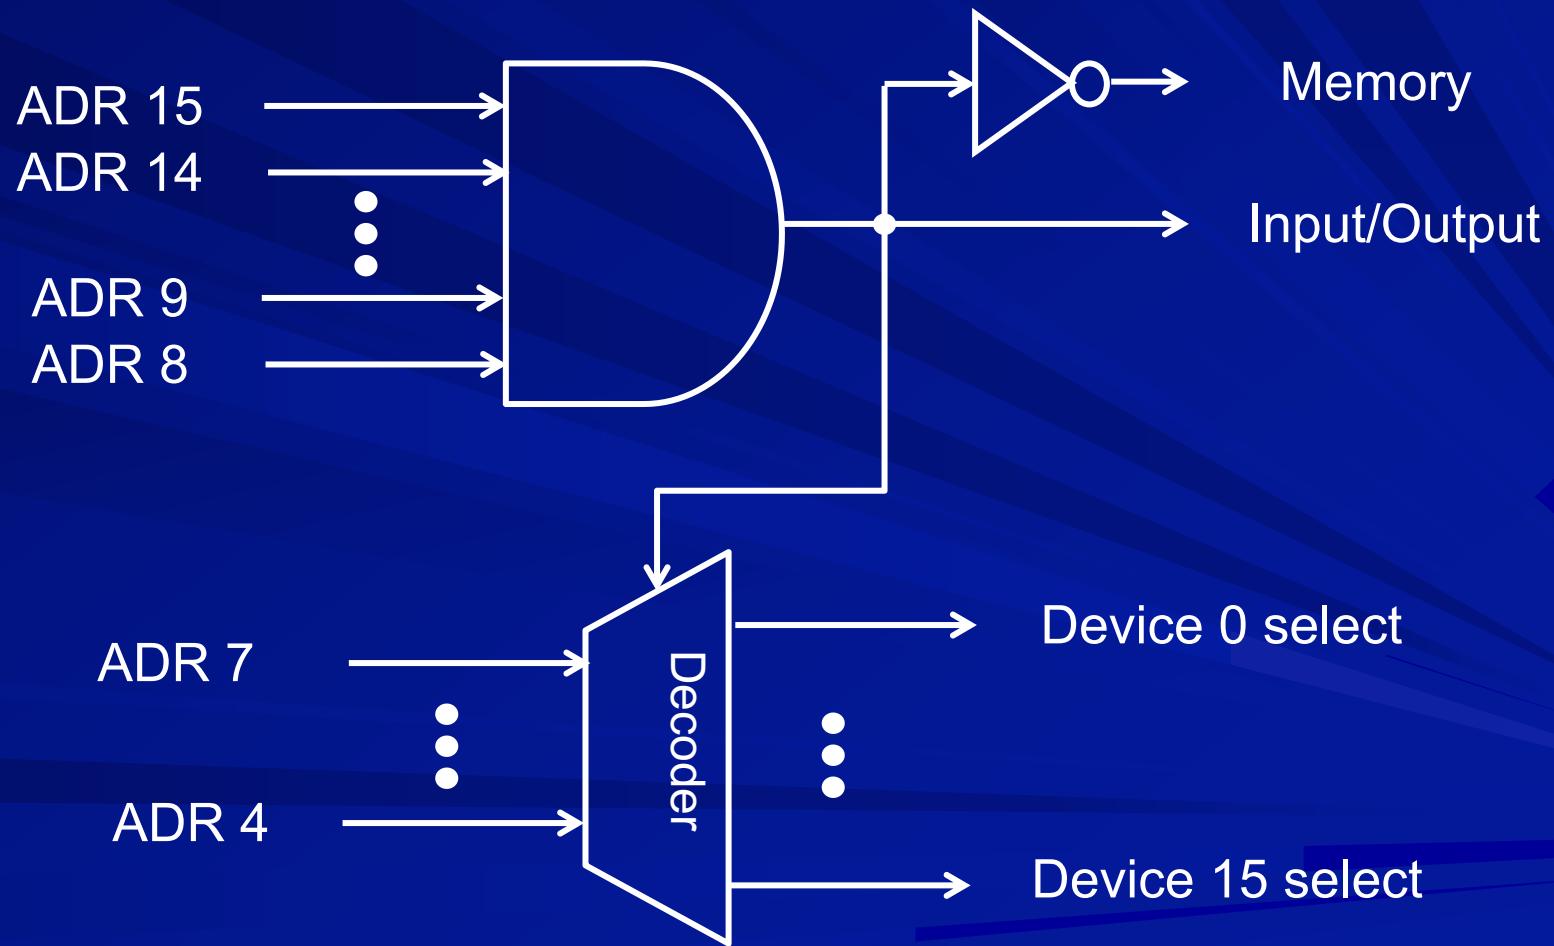


Low level control by software

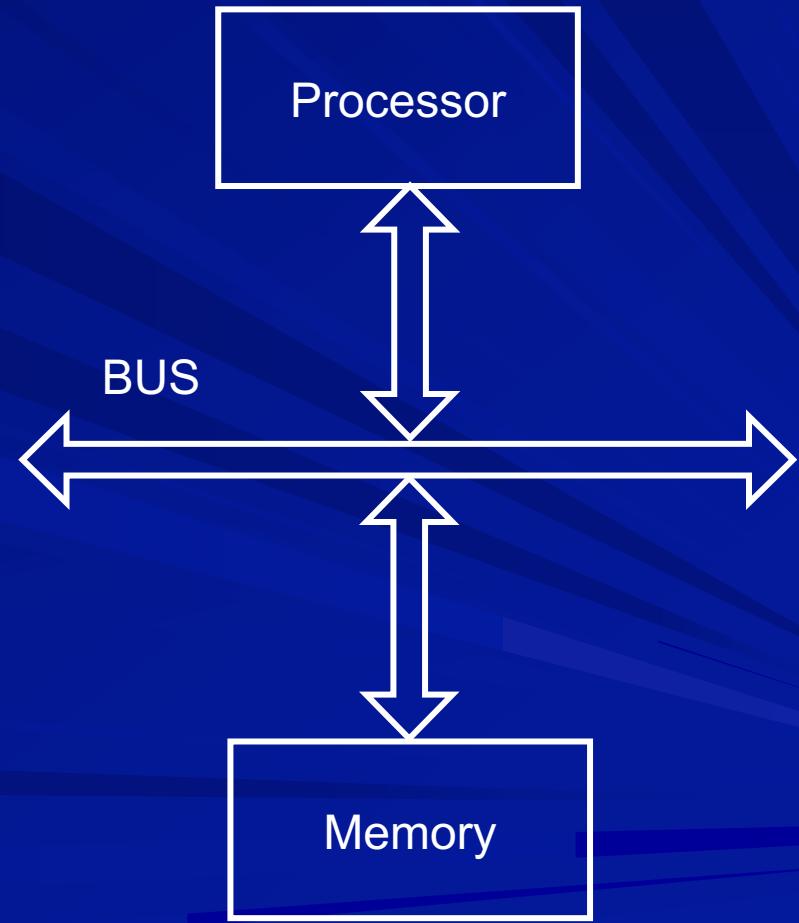
Address Map Example



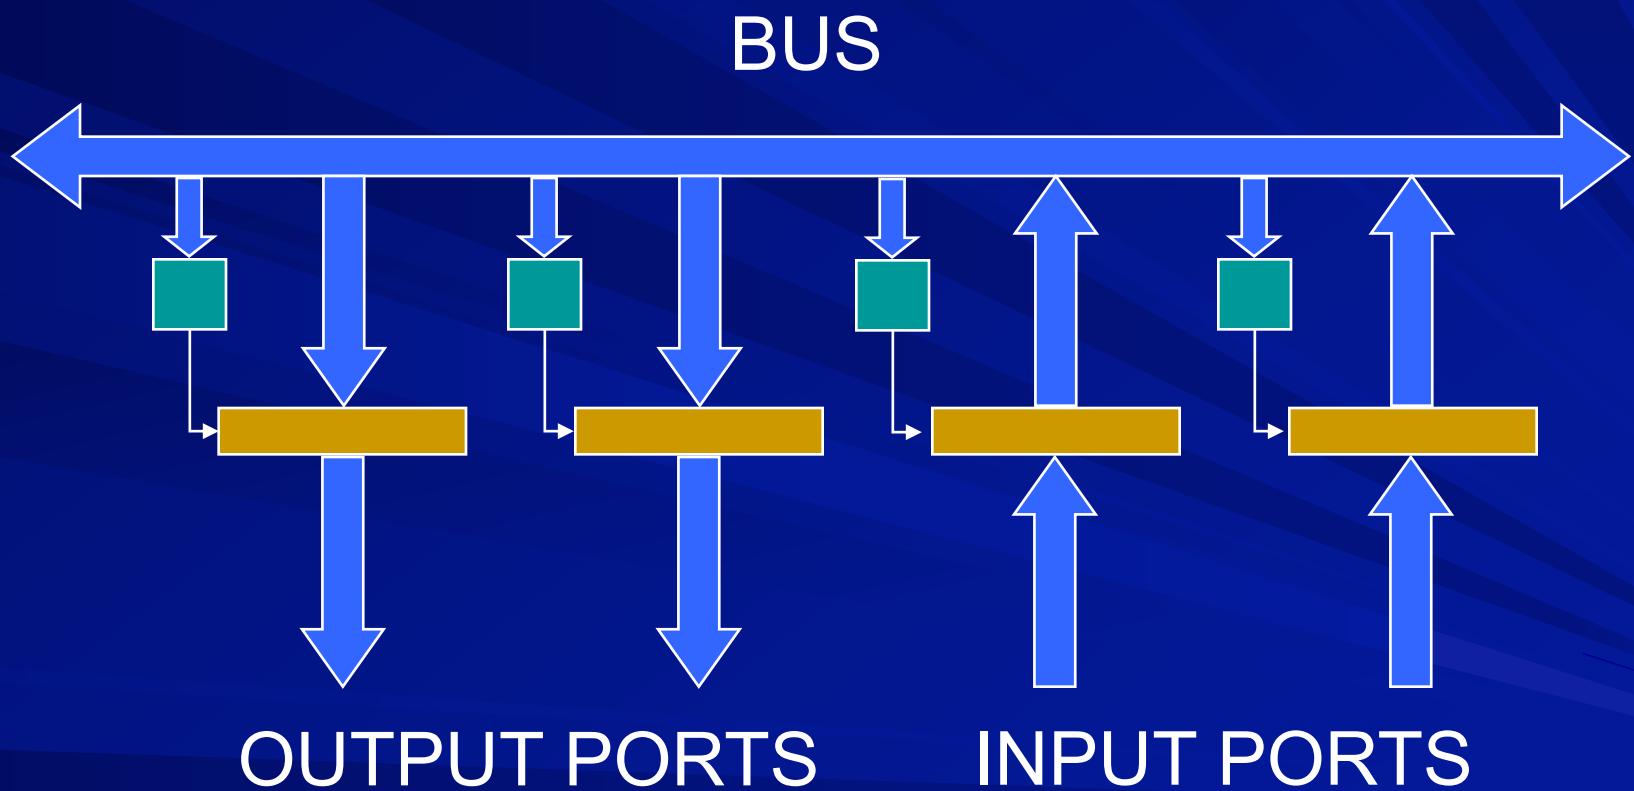
Decoding Address



Interconnection using a bus



INPUT / OUTPUT PORTS



Ports for Disp, KB controllers



Ports for raw Disp, KB

Display

out	anode pattern
out	cathode pattern

refresh required

Key Board

out	row pattern
in	column pattern

repeated scan required

Taking care of device timings

Devices may have latency of hundreds or thousands of cycles.

users are much slower than processor clock

Can the instruction execution be stretched for so long?

How CPU checks the device status

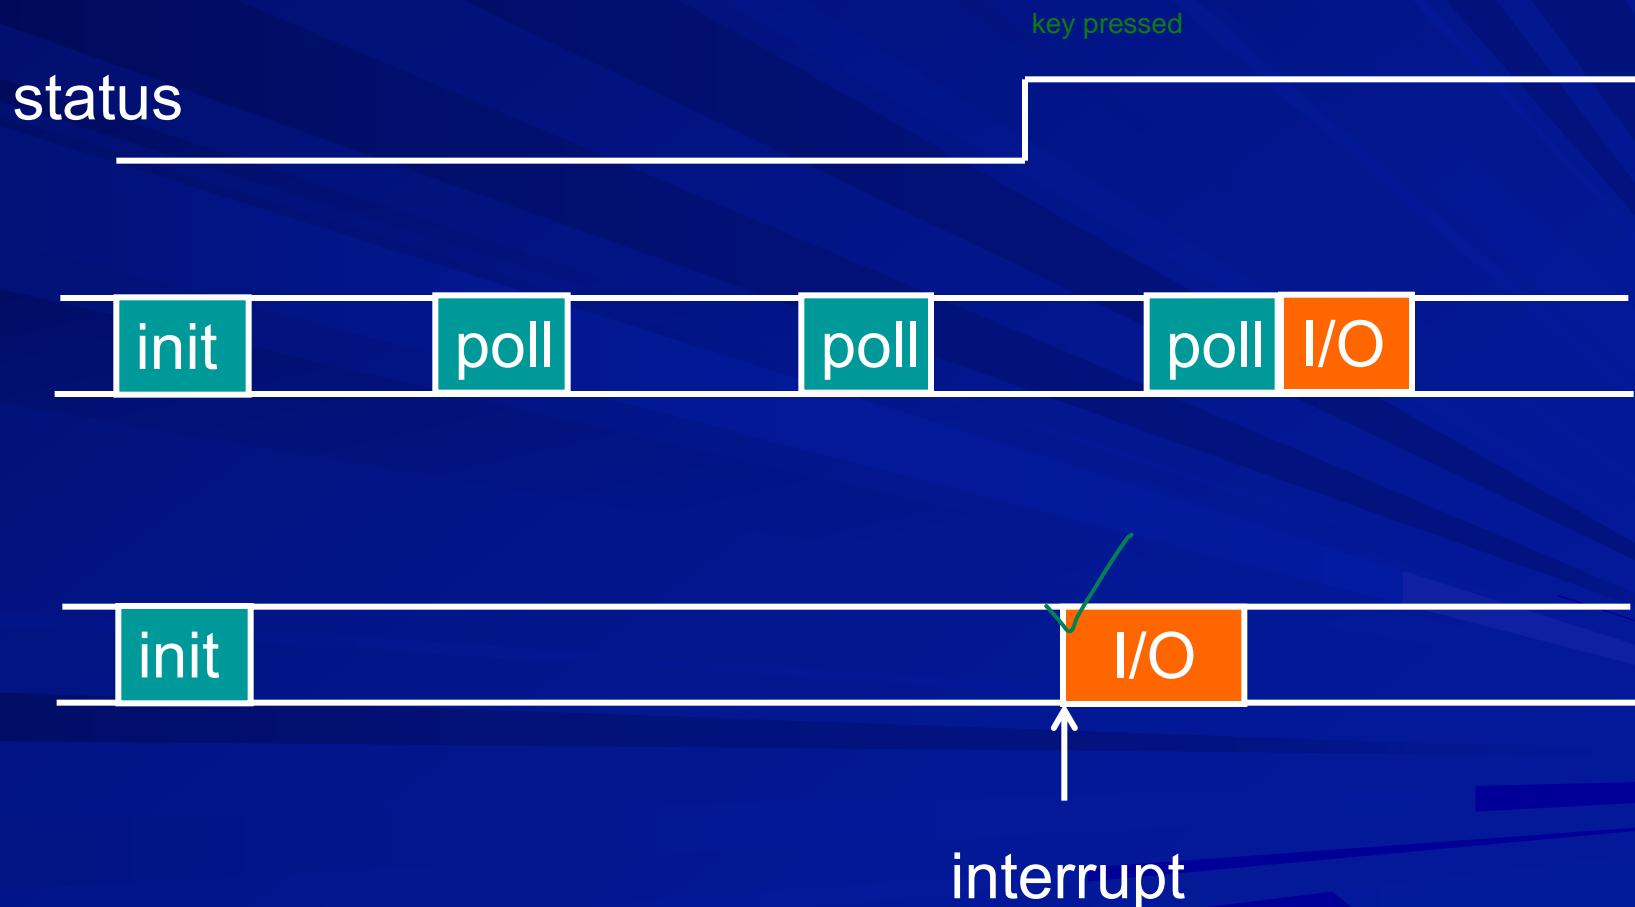
■ Polling

- periodically read the status register and figure out whether the device is ready or not
- should be frequent enough so that no events are missed out

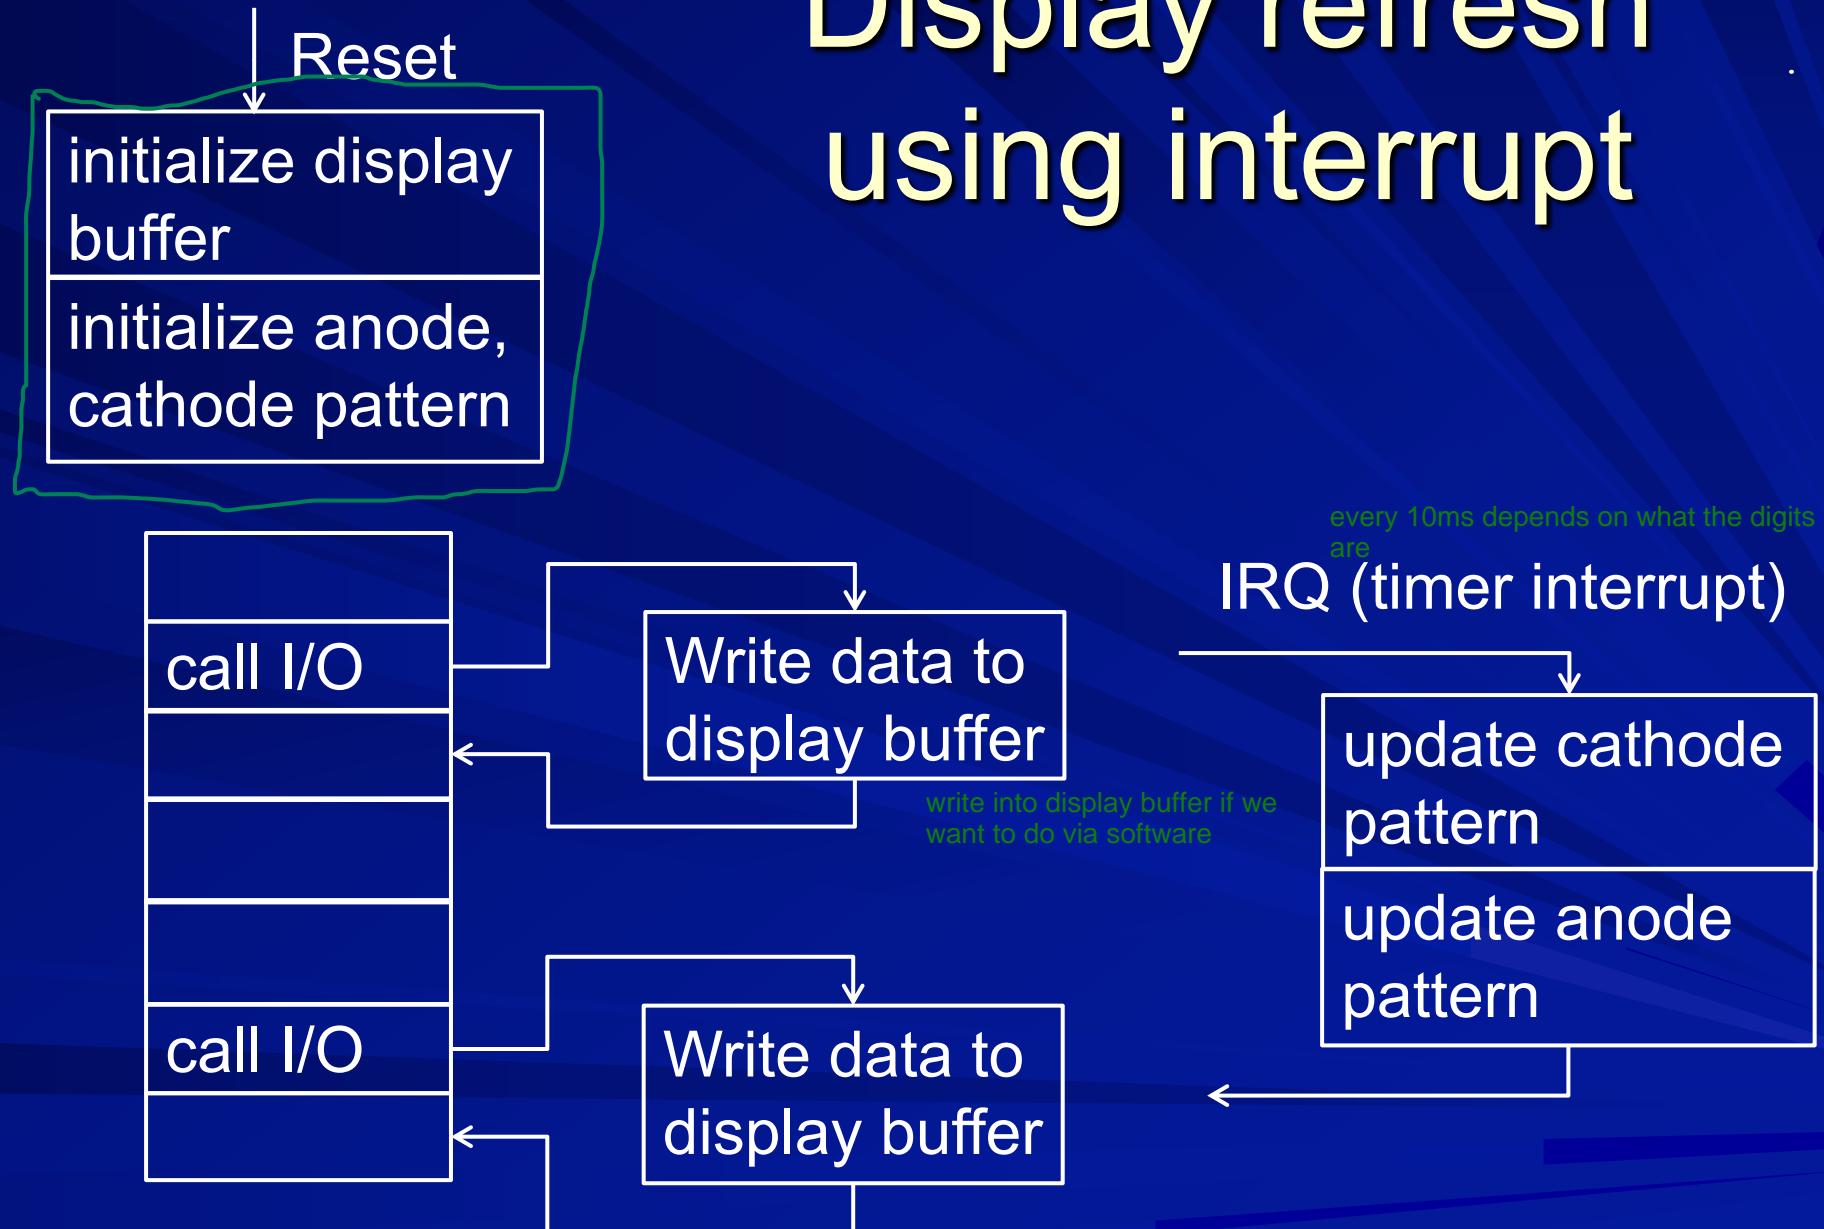
■ Interrupt

- keep busy with some other task and allow the device to intimate its readiness by sending a signal

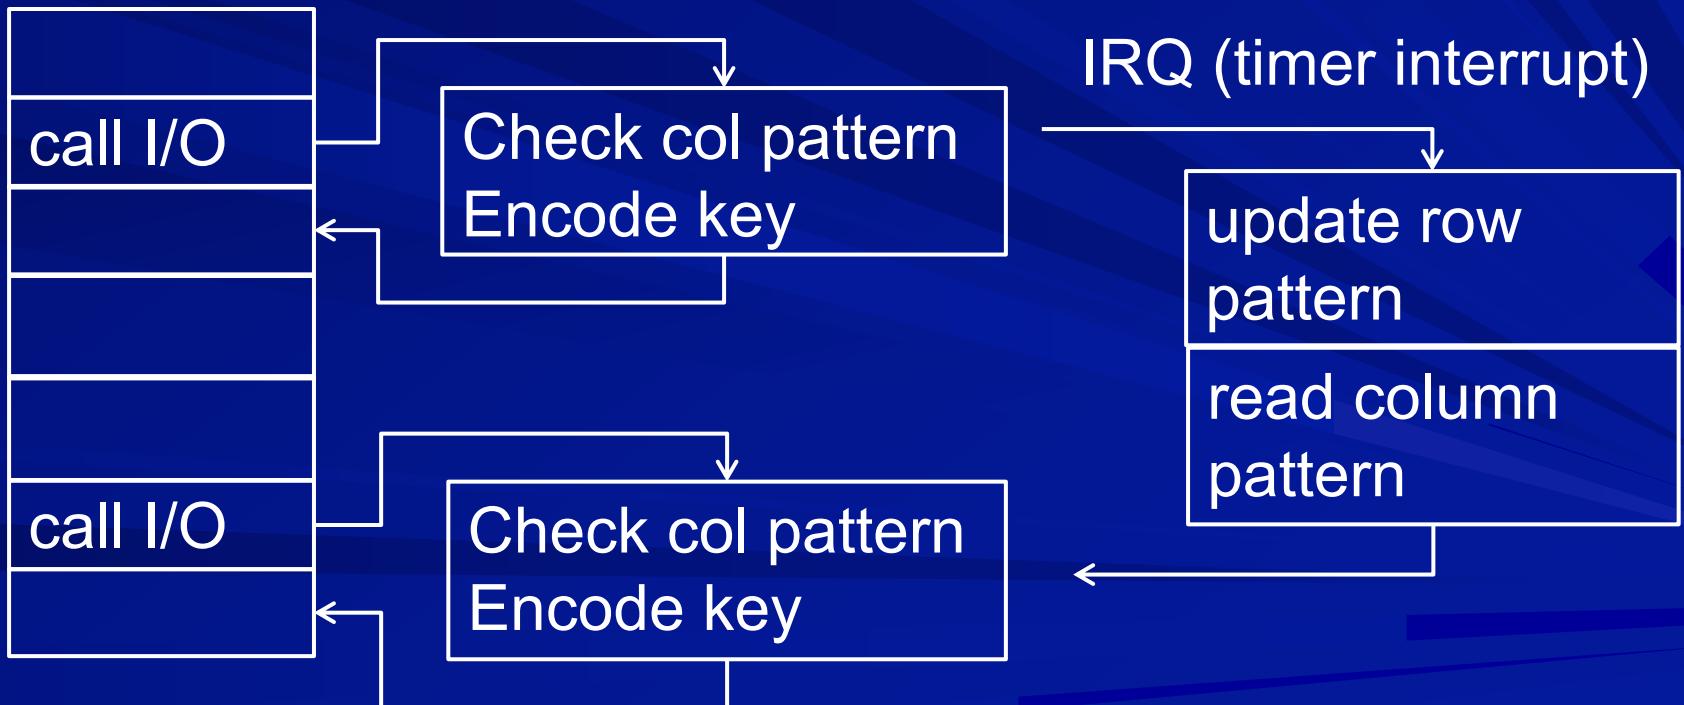
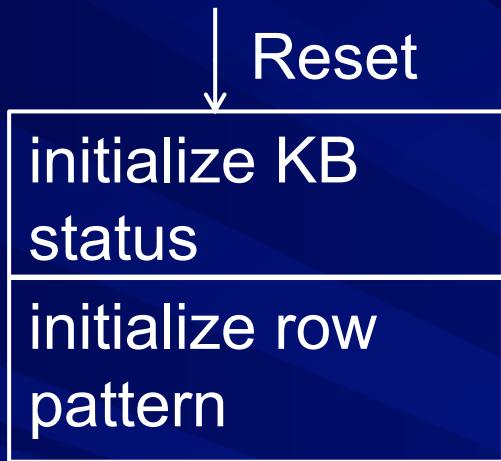
Polling vs interrupt



Display refresh using interrupt



Keyboard scan using interrupt



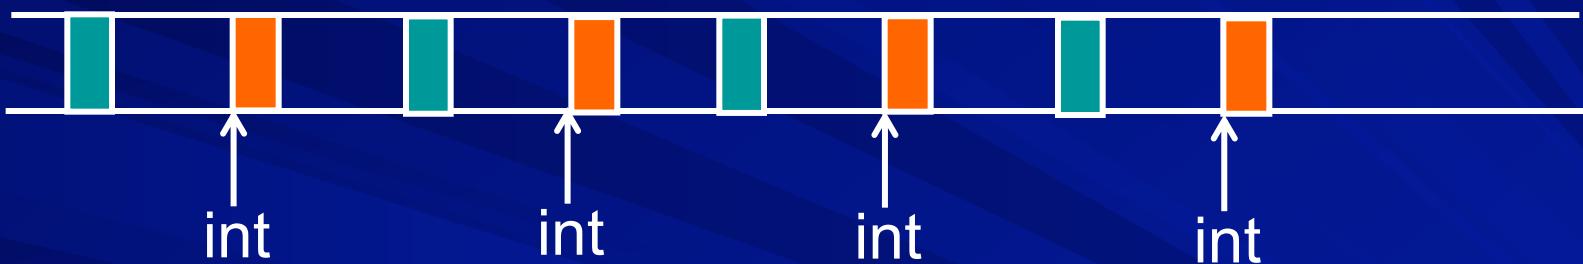
Direct memory access

- Allow direct data transfer between the device and the memory transfer big chunk of data
- The process is initiated by the processor
- When the entire job is done, the device informs the processor by an interrupt
- A specialized controller called DMA controller is used

processor has no role here, it is an independent controller to communicate between device and memory

Benefit of DMA

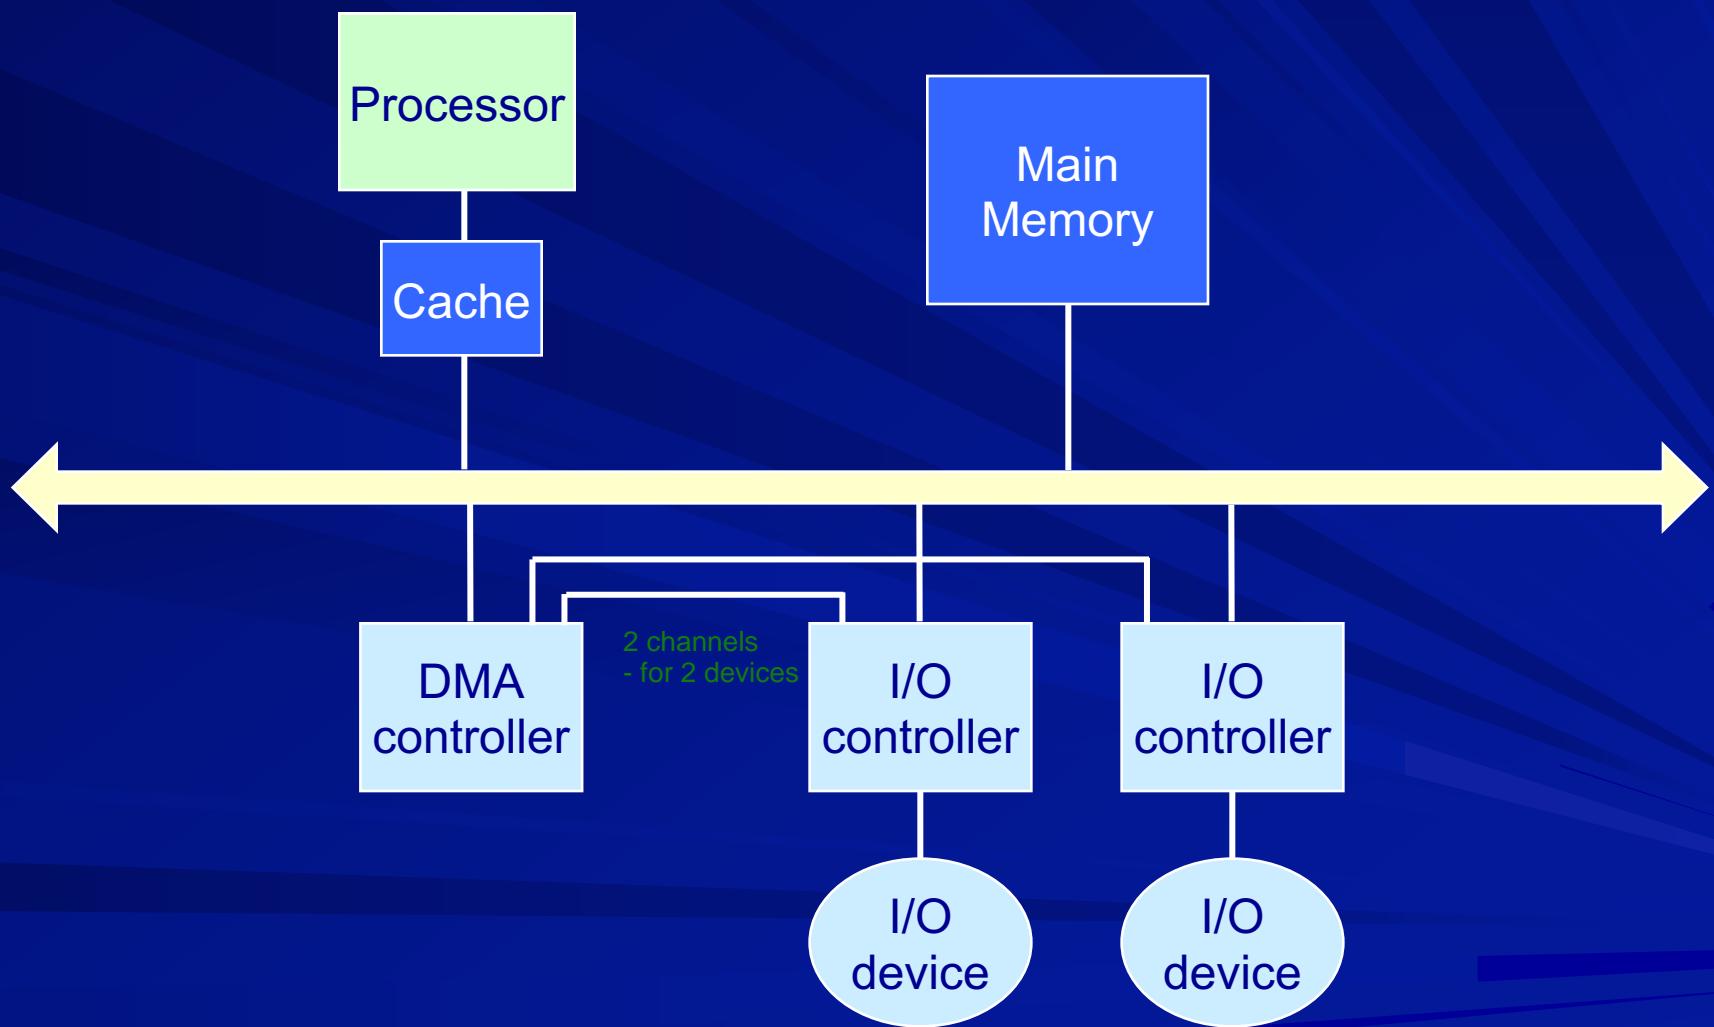
Without DMA



With DMA



Direct memory access



COL216

Computer Architecture

Input/Output – 3

21st March 2022

Interrupts/Exceptions

What are exceptions

- Unusual events/conditions
 - Source: internal or external
 - Synchronous or asynchronous
 - Intentional or unintentional
- Require change in flow of control
 - Mechanism to alter the control flow
 - Event identification and response to it
 - Halt or resume original control flow

Hardware and Software Interrupts

- Software interrupts
 - caused by specific instructions to get some system service
 - alternative terms : system calls, traps, exceptions
- Hardware interrupts are caused by events/conditions detected by hardware
 - intentional : e.g., I/O event
 - unintentional : e.g., hardware fault, power outage, arithmetic overflow

Terminology

Exceptions vs. Interrupts

- Use both the terms interchangeably
- MIPS, ARM: exception - internal as well as external, interrupt - external
- Intel 80x86: interrupt - internal as well as external
- PowerPC: exception - unusual event, interrupt - change in flow of control

Examples of exceptions

◆ synchronous

◆ asynchronous

	Intentional	Unintentional
Internal	Invoke OS function, Trace/debug	Access to privileged inst, Overflow/underflow, Undefined instruction, Hardware malfunction
External	I/O device request	Mem access exception, Alignment error, Timeouts, Power down, Hardware malfunction

Interrupt identification and response

- S/w interrupt identified during instruction decoding
- H/w interrupt identified by checking specific signals
- Response mechanism in both cases (h/w and s/w interrupts) is same
 - execution of some code : interrupt handler or interrupt service routine (**ISR**)
 - after serving interrupt, terminate or resume

Checking for exceptions



undefined instruction
checked here

overflow
checked here



hardware interrupts can be checked any time

why can't we have a normal subroutine? Like when we face an exception, we transfer the control to some other flow so the part of the program where we transfer our function call to is actually an ISR

ISR vs normal subroutine

- Processors have two or more modes
 - normal mode / user mode
 - privileged mode / kernel mode / supervisor mode
- Application program executes in user mode
 - i/o is done in kernel mode
 - virtual memory is done in supervisor mode, OS sets up page table and it cannot be modified.
- To do certain privileged tasks, execution of some code in OS kernel is required
- ISR executes in privileged mode, provides controlled access to kernel functions

How exceptions are handled

uninterrupted execution



interrupt and halt



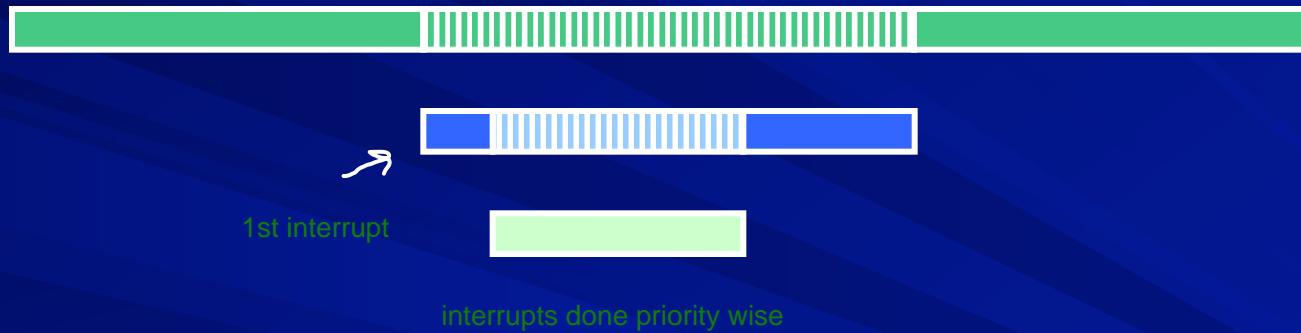
interrupt and resume
(vectored interrupt)



interrupt and resume
(non-vectored)



Nested interrupts



give a message to the user
and terminate the program

Exceptions in ARM

Exception type	Mode	Address of the ISR
Reset <small>pc = 0 and execution starts from address 0 again</small>	Supervisor	0x00000000
Undef instr	Undefined	0x00000004
S/W interrupt	Supervisor	0x00000008
Prefetch Abort	Abort	0x0000000C
Data Abort	Abort	0x00000010
Interrupt	IRQ	0x00000018
Fast interrupt	FIQ	0x0000001C

depending on type, the hardware transfers control to the address specified.

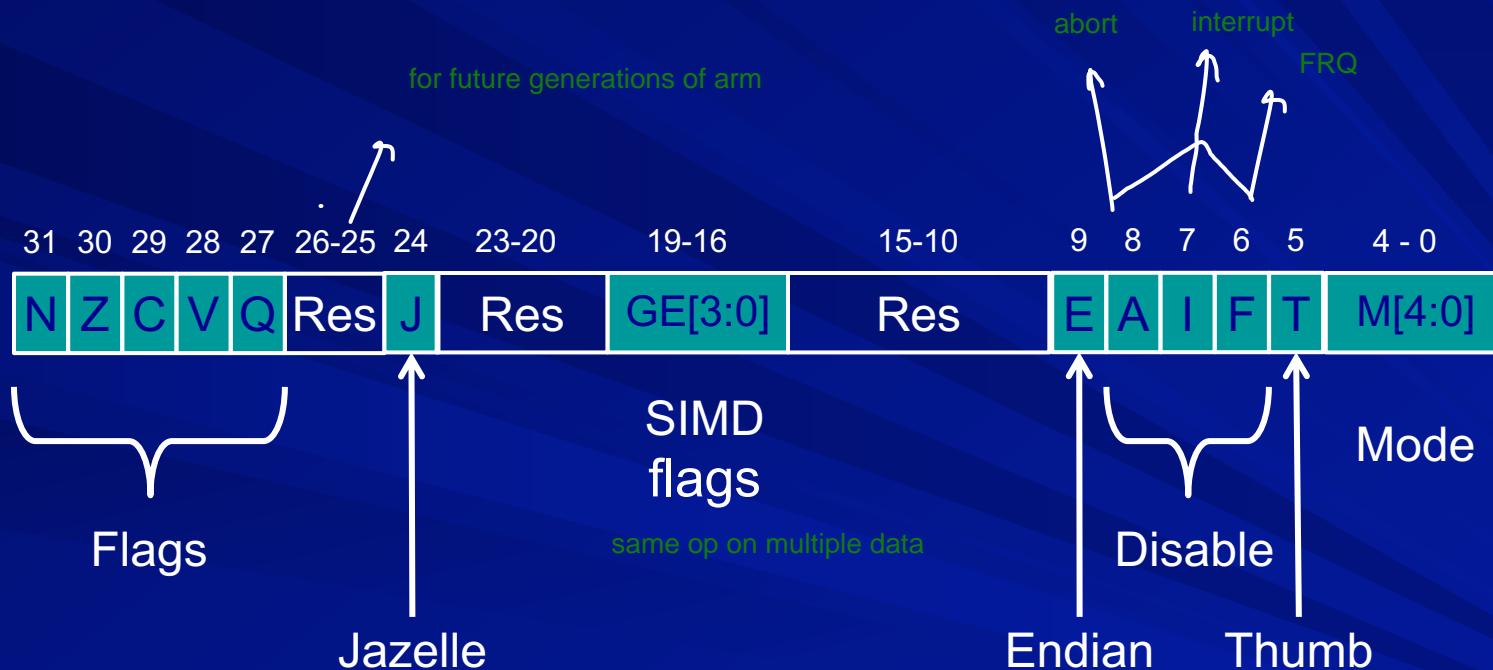
external

SWI software interrupt

cond	1111	comment
4	4	24

- Control transferred to address 0x00000008
- Mode changes to 'supervisor' mode
- CPSR (Current Processor Status Register) is saved in SPSR_{svc} (Saved Processor Status Register)
- PC is saved in R14_{svc}
- Comment field may specify a particular service/function

Processor status register



ARM register set

Mode	Registers			17th register
	SP	LR	PC	
User:	R0-R7	R8-R12	R13	R14 R15 CPSR
System:				
Supervisor:		R13 R14	<small>-> in supervisor mode</small>	SPSR
Abort:		R13 R14		SPSR
Undefined:		R13 R14		SPSR
Interrupt:		R13 R14		SPSR
Fast int:	R8-R12	R13 R14		SPSR

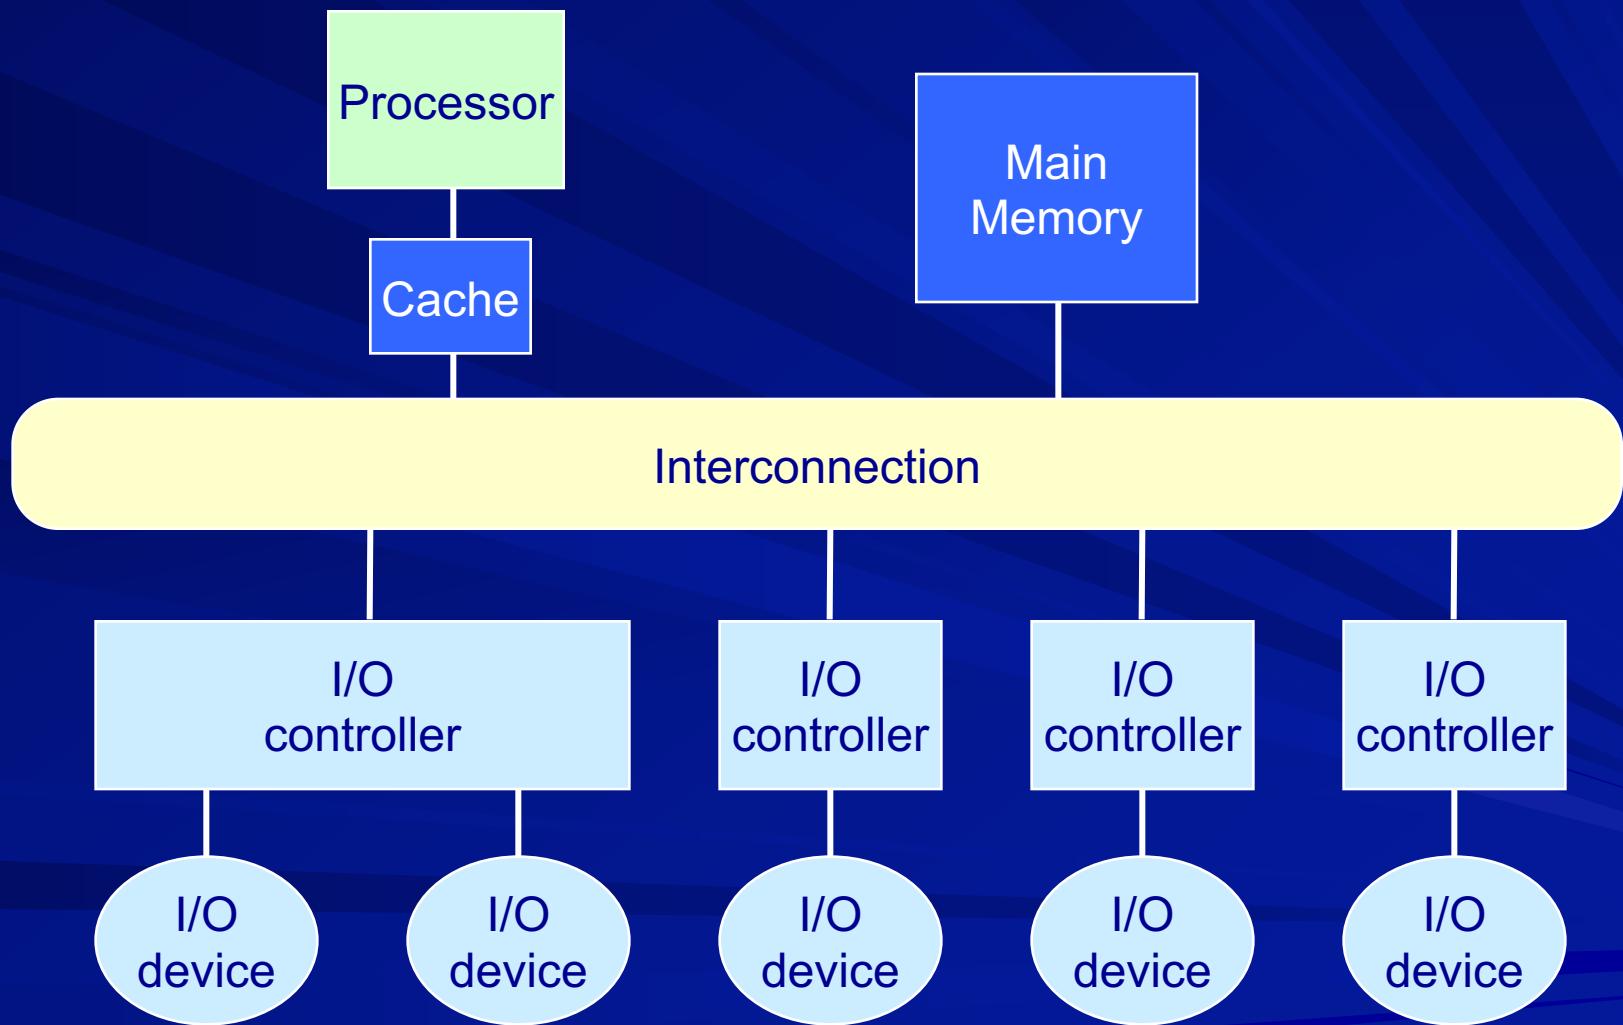
Total 37

Exception handling in pipelined design

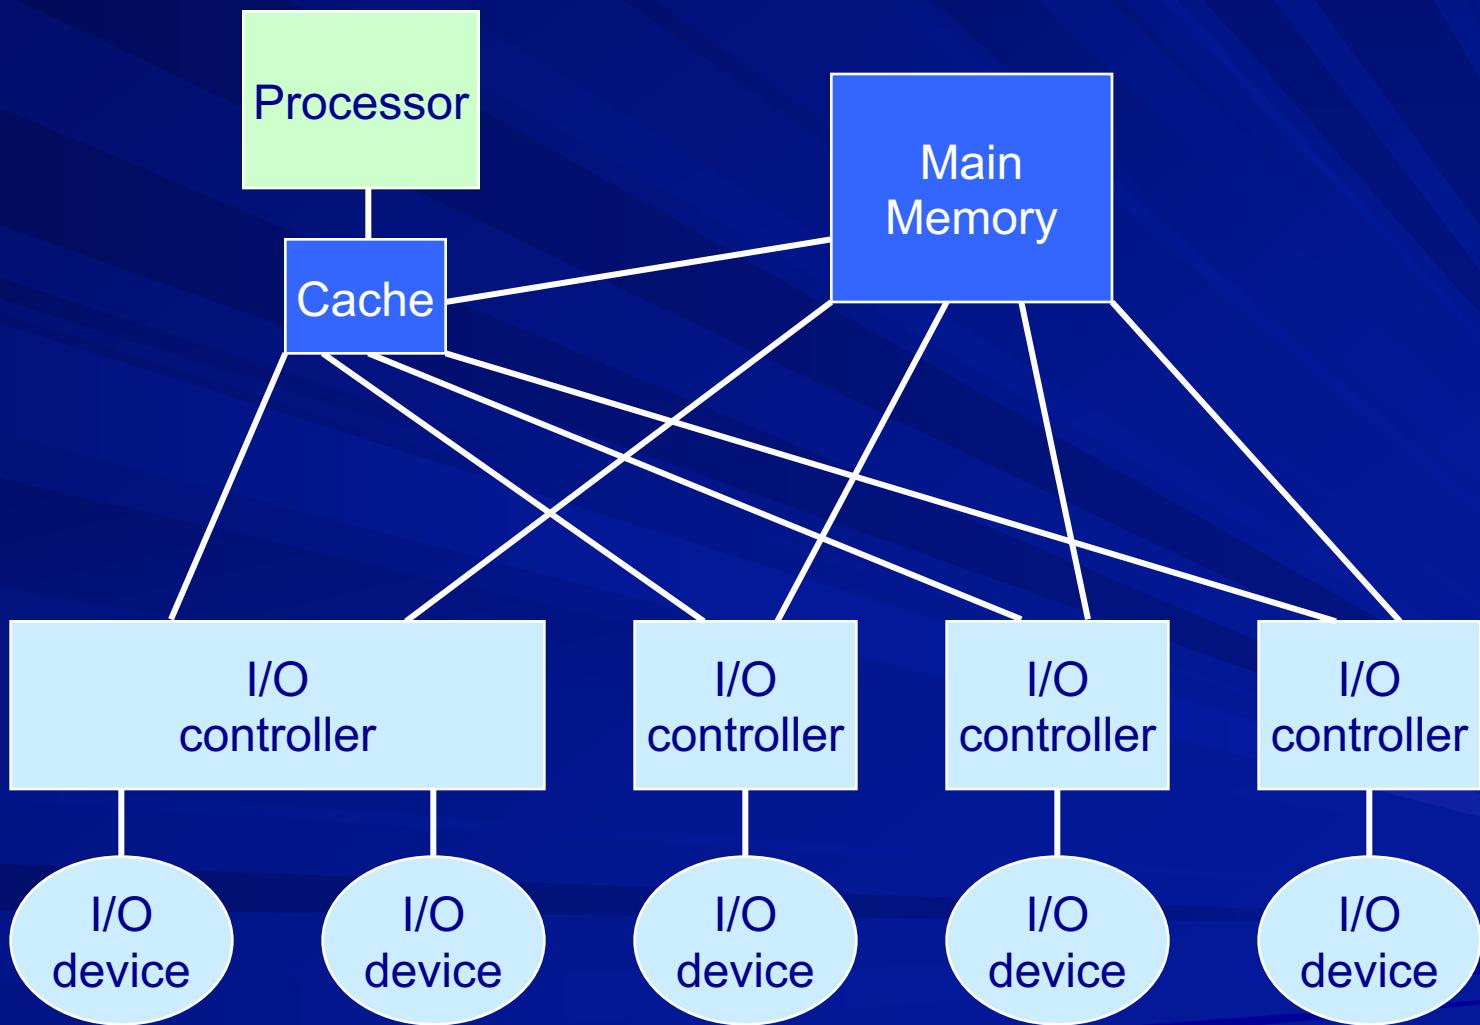
- Difficulty in status saving
 - multiple instructions under execution
- Possibility of
 - multiple exceptions in same cycle
 - change in order of exceptions
- Handling approaches
 - precise
 - imprecise

Interconnecting Subsystems

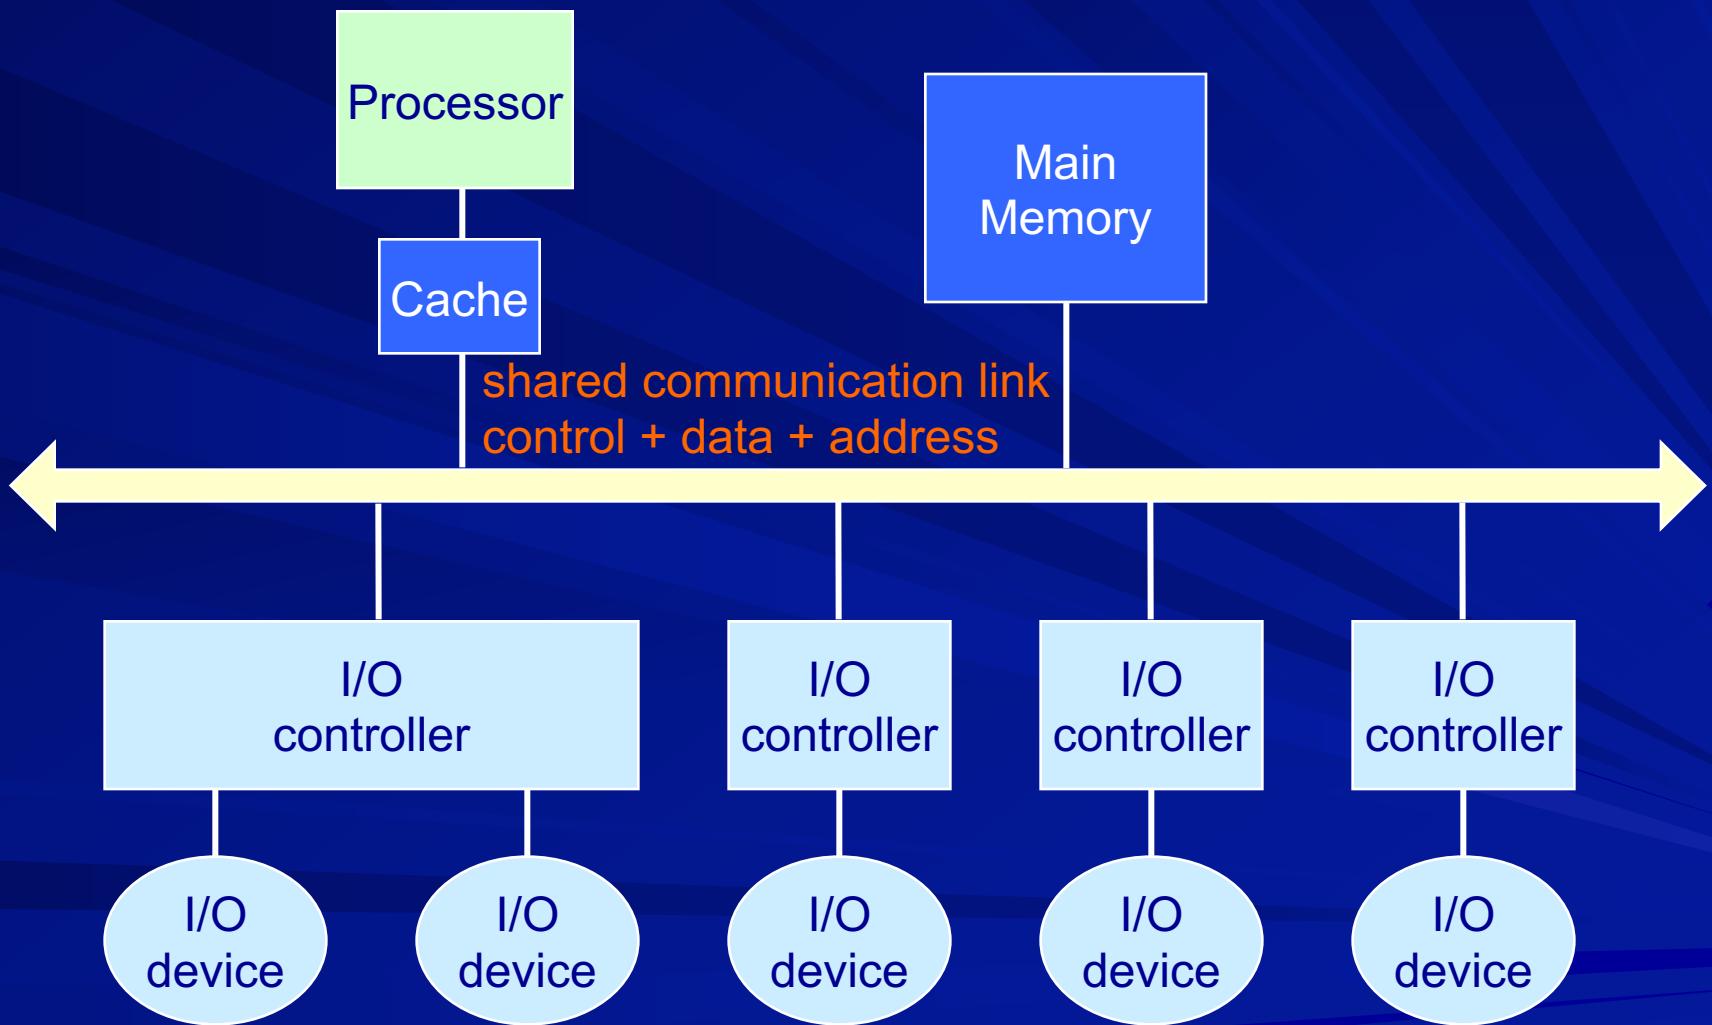
Interconnecting Subsystems



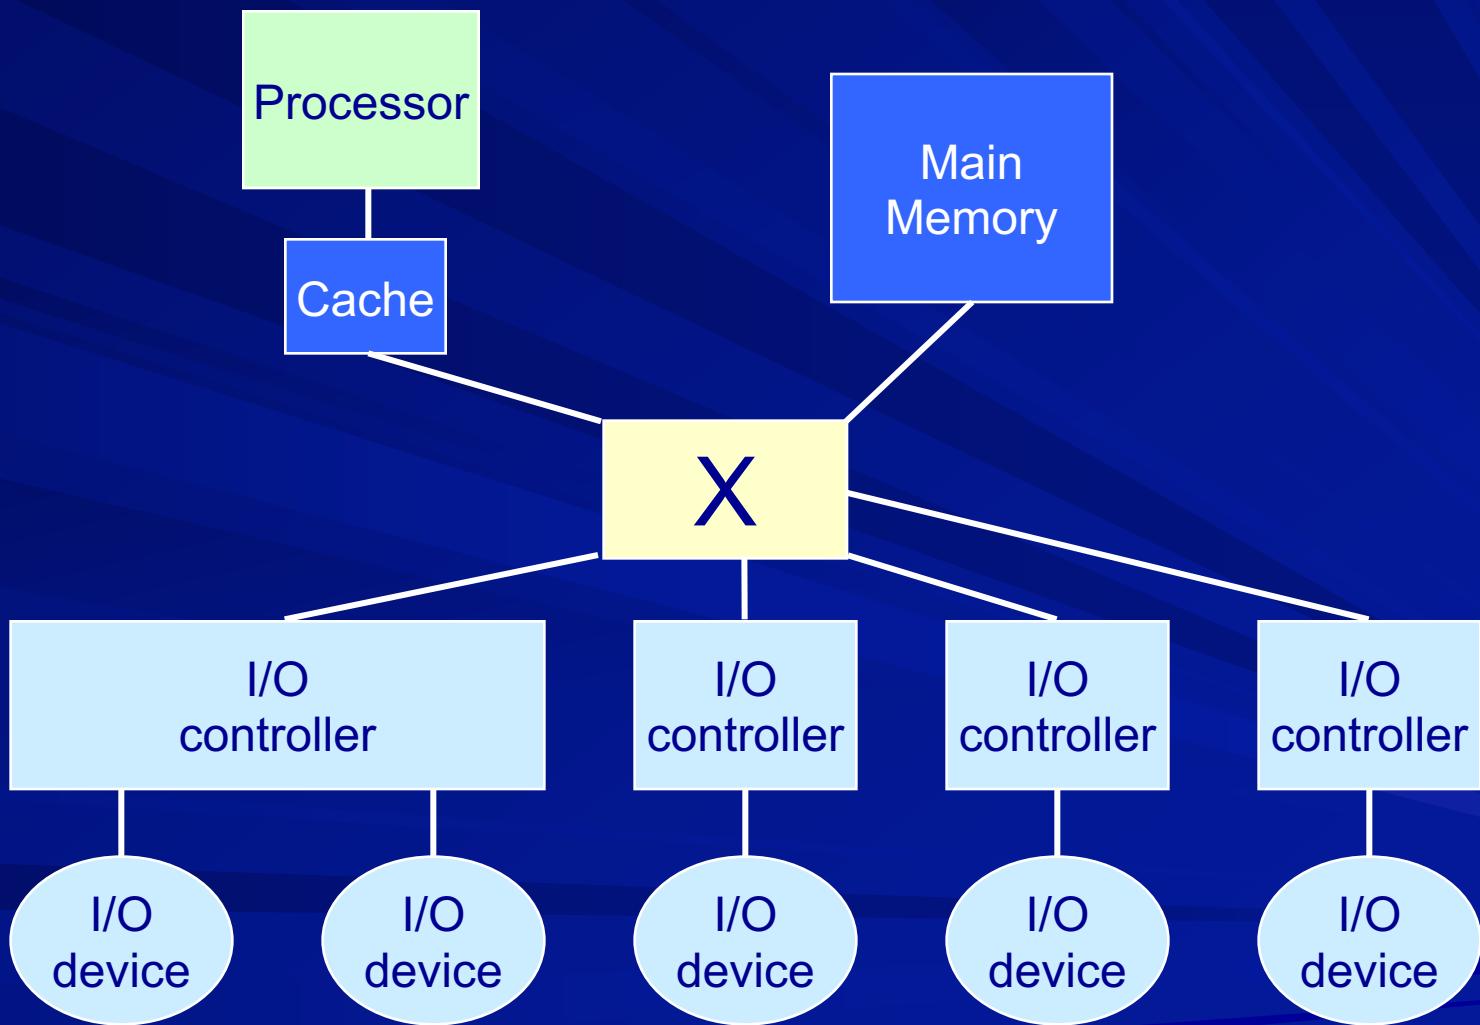
Point to point connections



Single bus connecting all



Cross bar switch



Comparison

	<u>P to P</u>	<u>Bus</u>	<u>X-bar</u>
Cost	high	low	high
Throughput	high	low <small>only one at a particular time</small>	high
Ports	multi	single	single
Expansion	difficult	easy	difficult

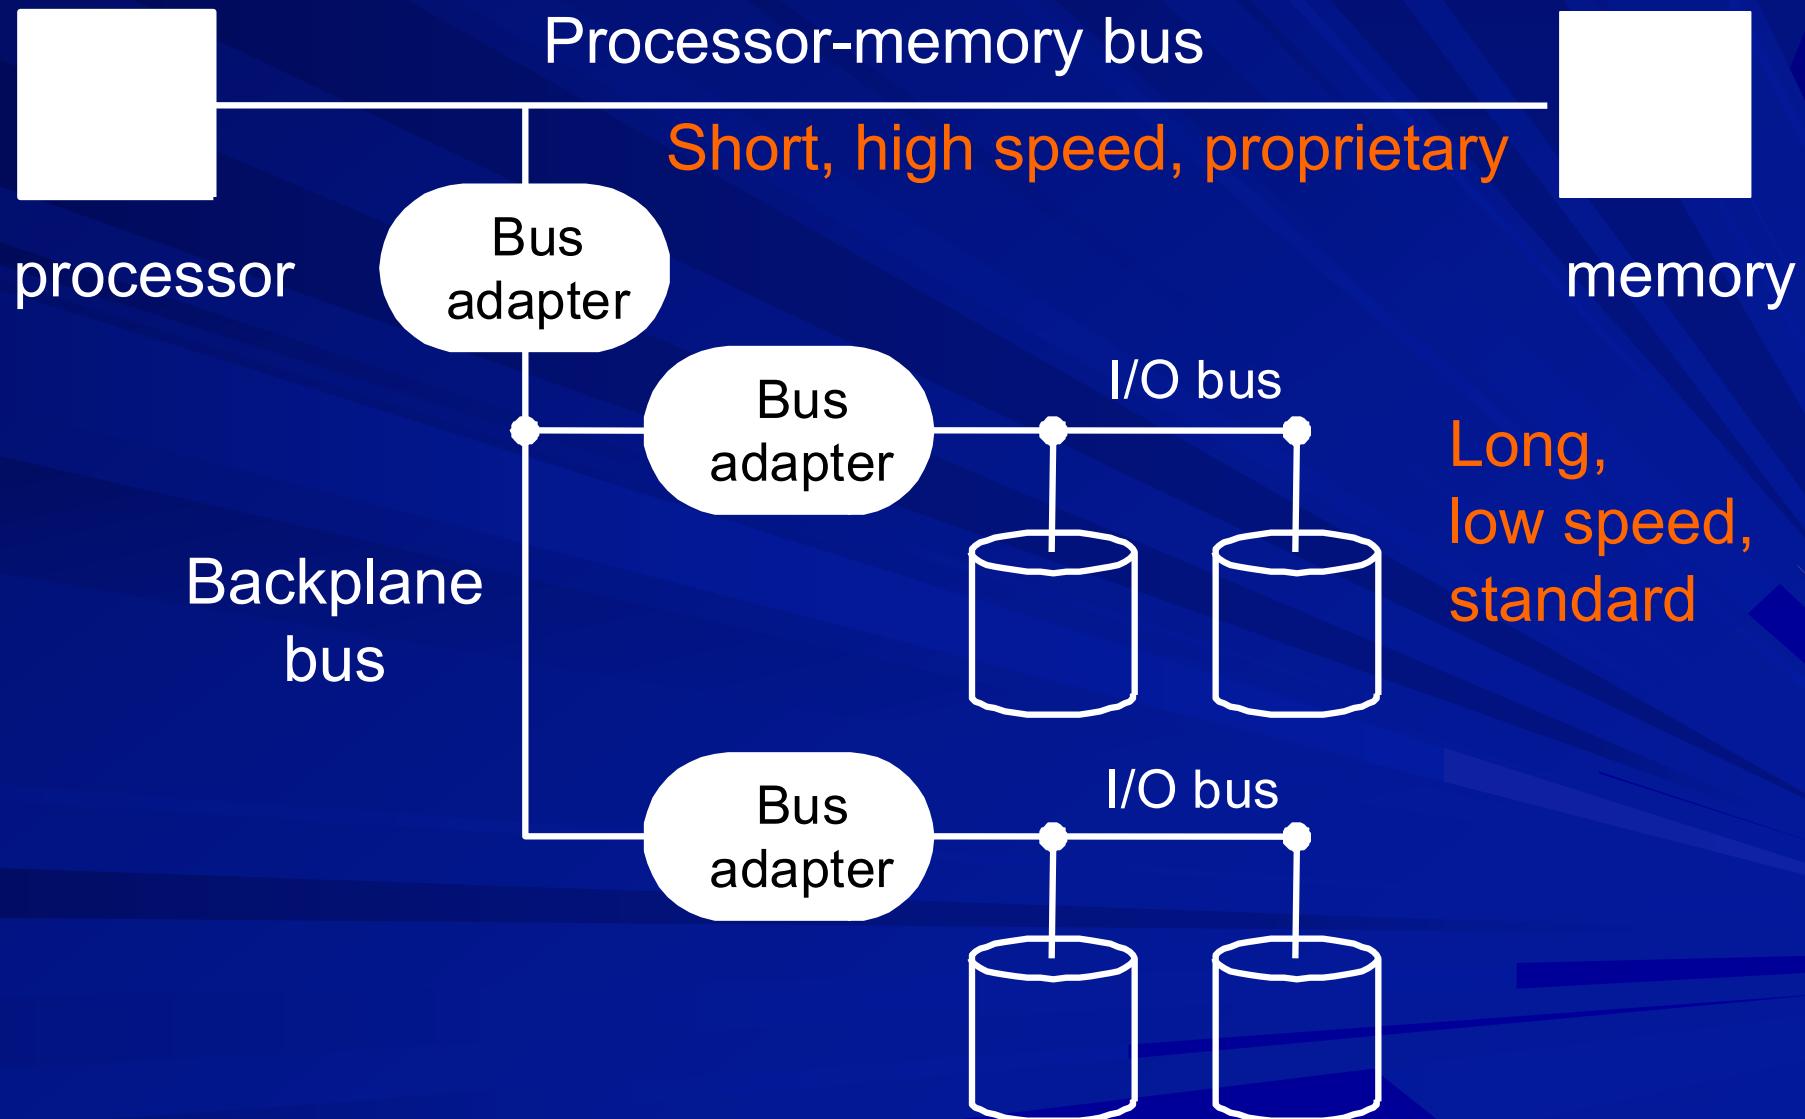
Multi-buses: More performance, more cost

Multi-switches: Lower cost, same performance

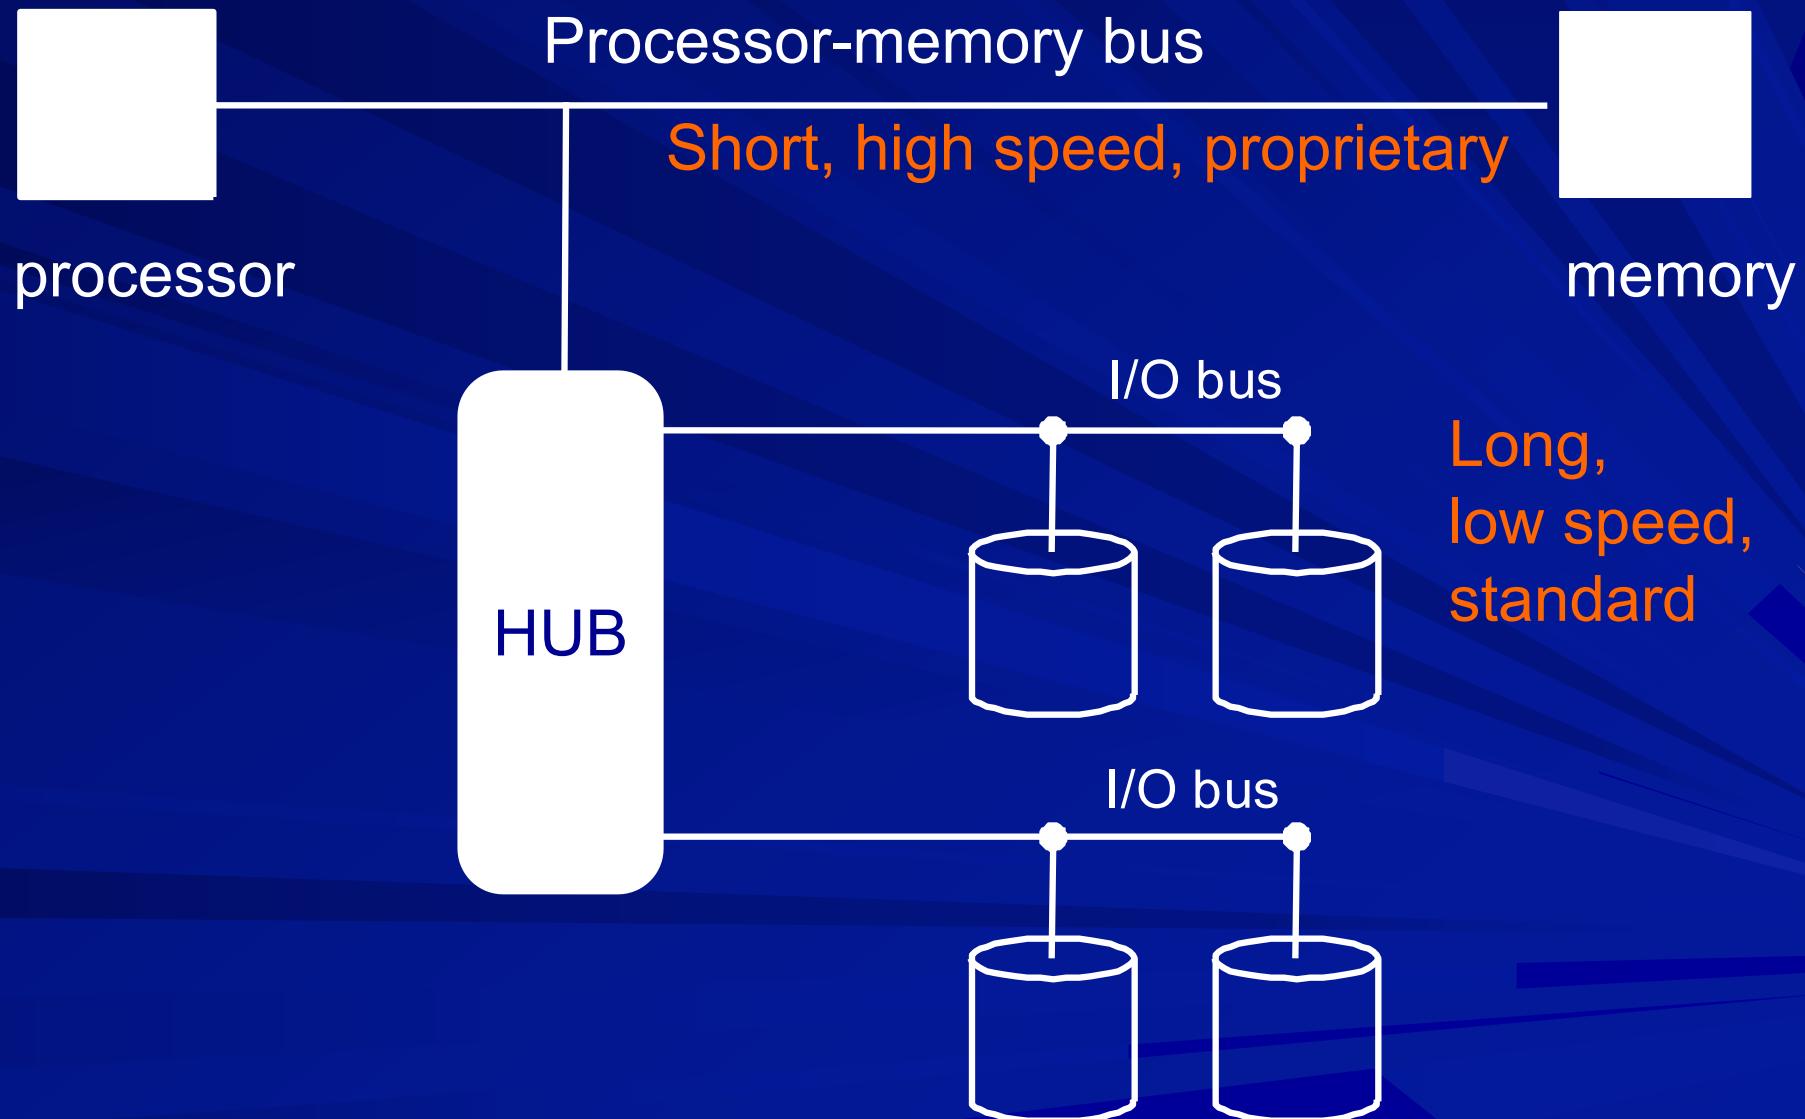
Backplane bus



System with multiple buses



System with multiple buses



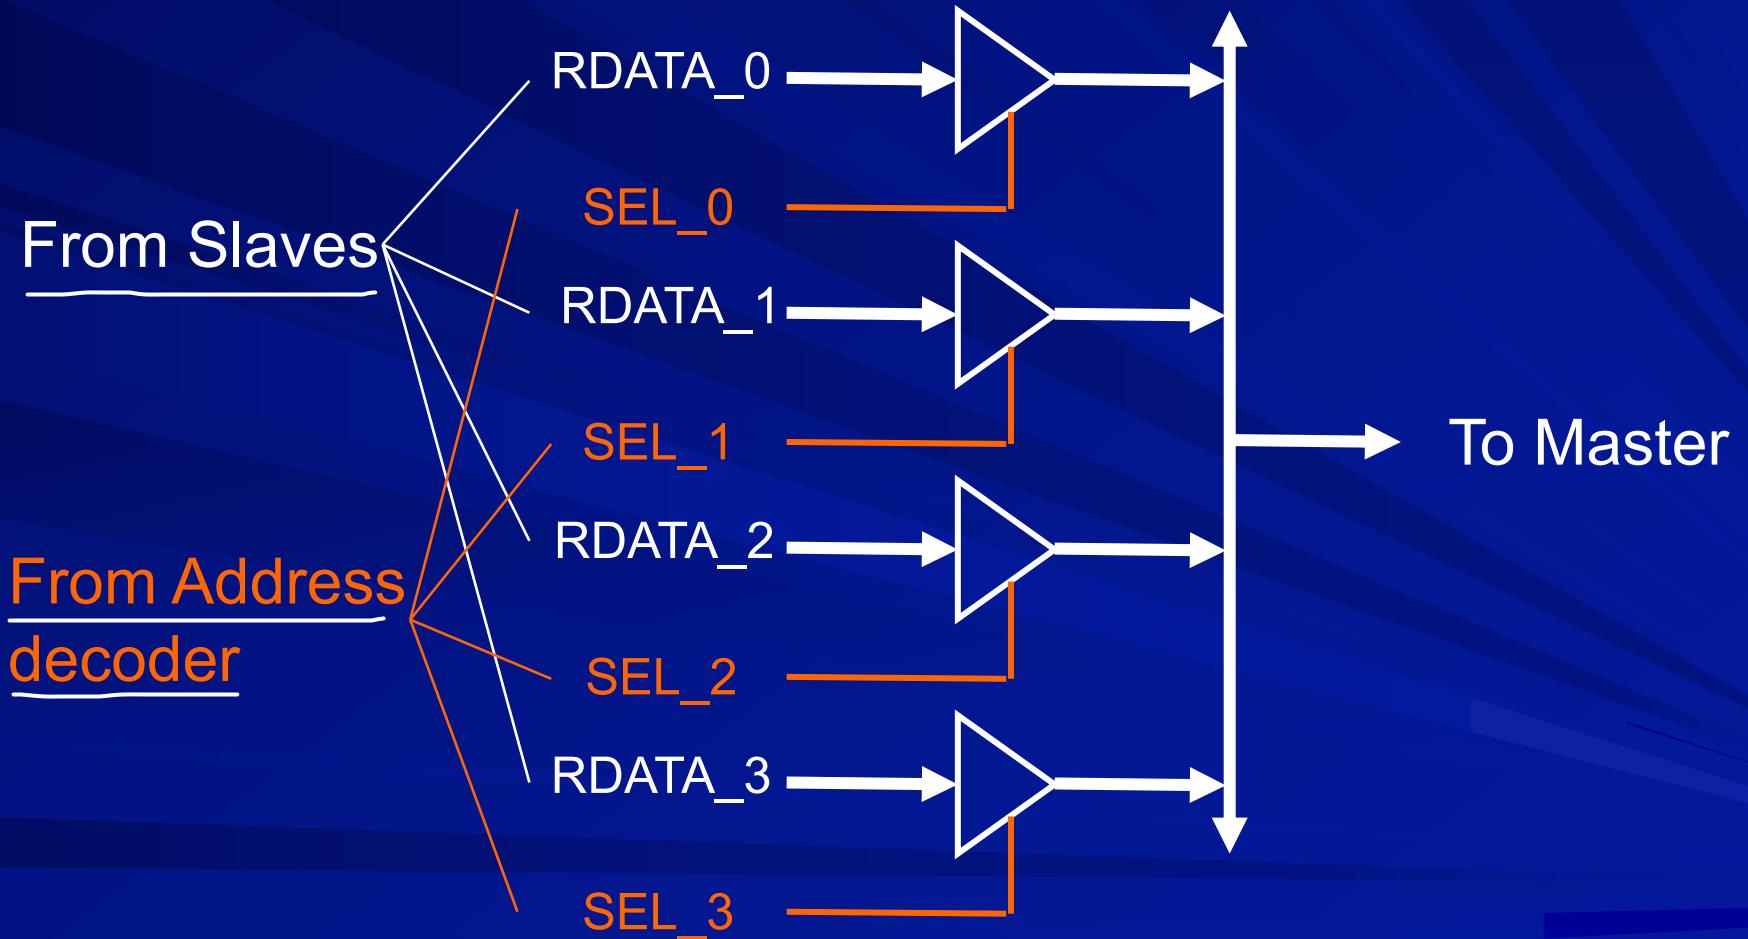
Bus Standards

- Physical / mechanical
 - Pins, connectors, cables
- Electrical
 - Voltage / current levels, impedances
- Logical
 - Definition of signals
 - Timings and protocols

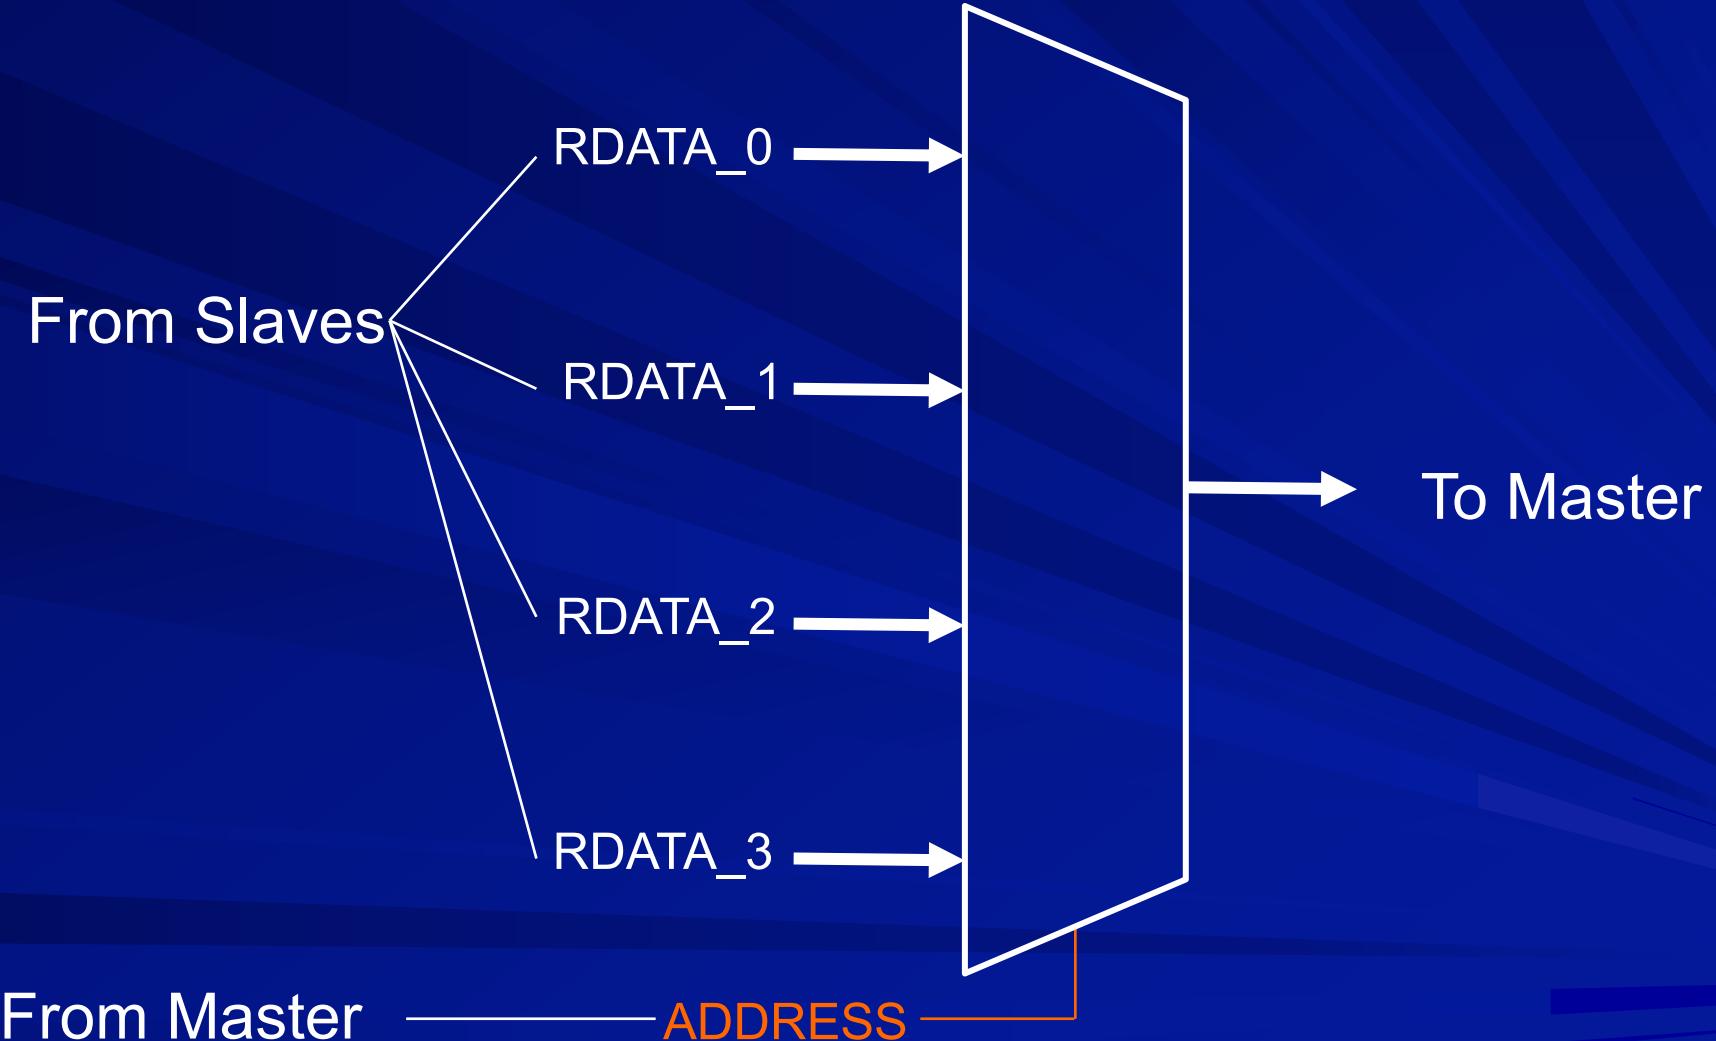
Subsystems on the buses

- Master
 - Initiates transfer
- Slave
 - Responds to master
- ✓ ■ Bridge /
bus adapter
 - Connects two buses
- Hub
 - Connects many buses
- Arbiter
 - Resolves master requests
 - Selects slaves
- Decoder

Tri-state Bus



Multiplexed Bus



AMBA : Advanced Microcontroller Bus Architecture

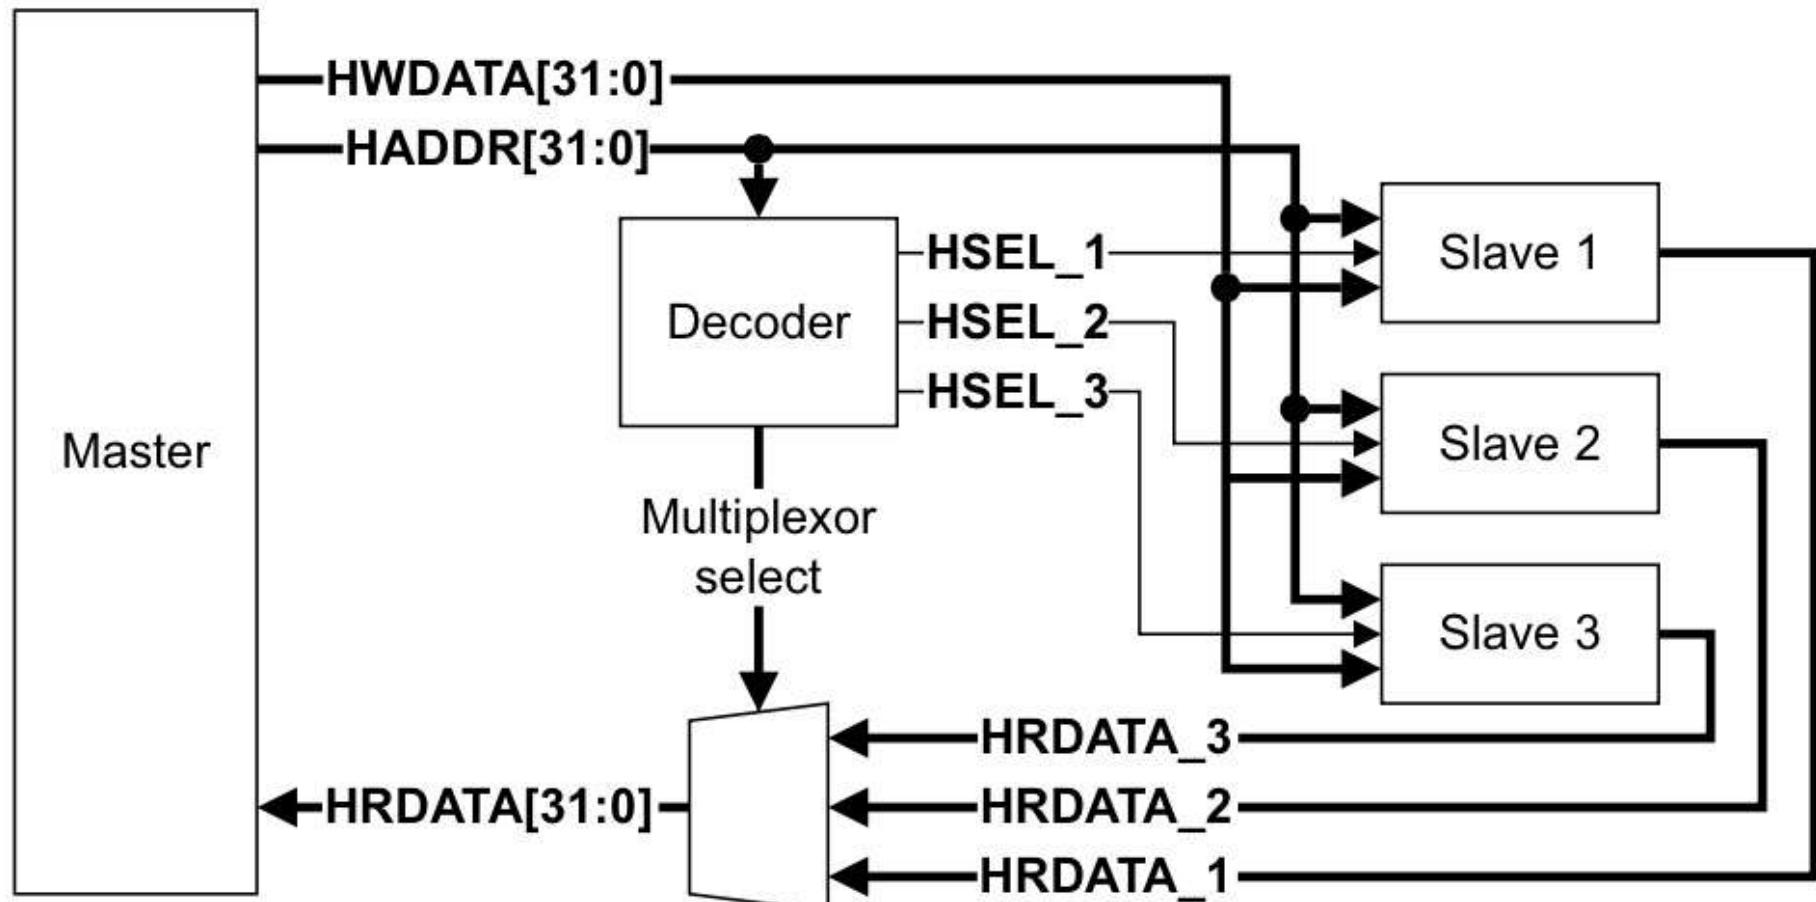
- ARM open standard for on-chip buses
- For System-on-Chip (SoC)
- Variants
 - AHB : Advanced High speed Bus
 - APB : Advanced Peripheral Bus
 - AXI : Advanced eXtensible Interface
 - ✓ AHB-Lite : subset of AHB

AHB-Lite

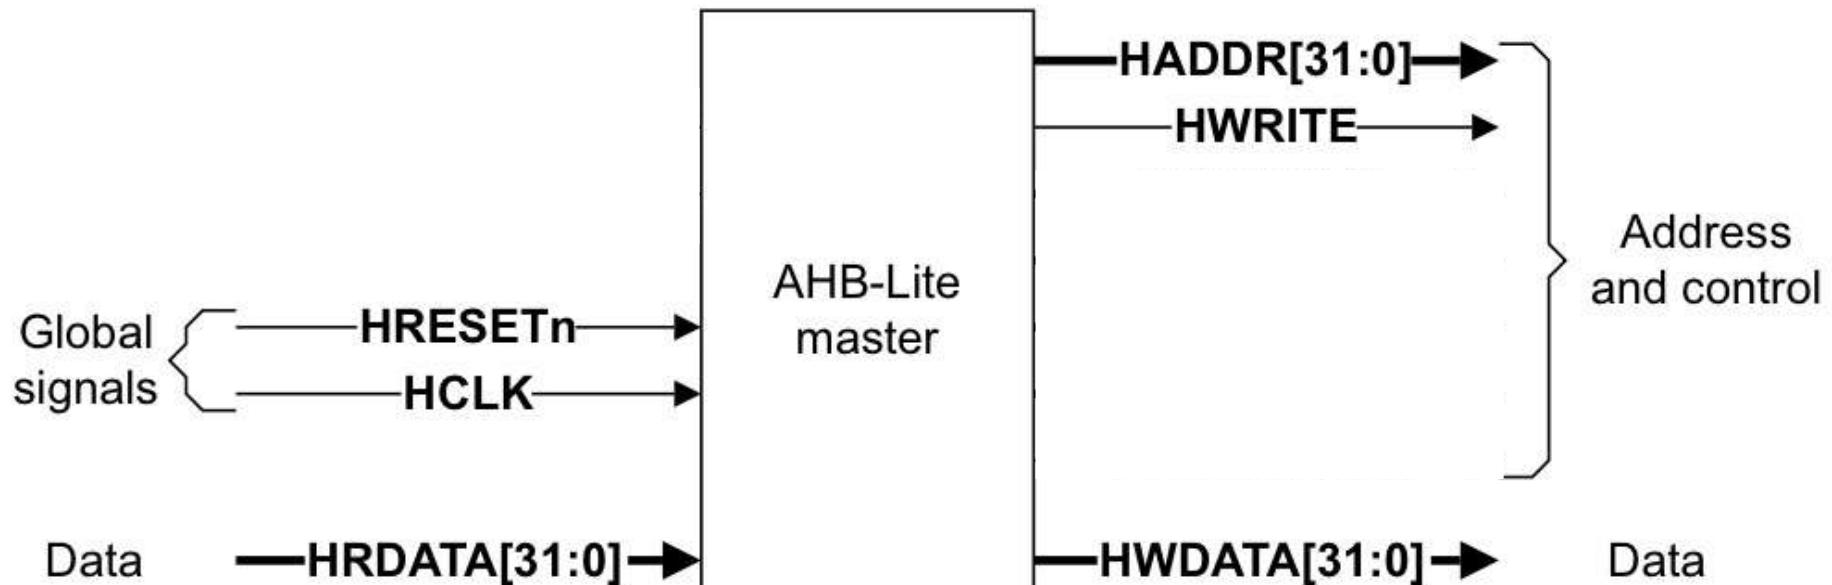


- Single master, multiple slaves
- Parallel
 - usb is serial bus- one bit at a time
- Multiplexed (not tri-state)
- Separate data, address, control
- Synchronous
 - all transitions happen synchronised with the clock
- Pipelined
- Burst transfer supported
 - each transaction involves burst of data, large data
- Split transaction not supported
 - transaction between initial and final transaction, in between we could not start another transaction

AHB-Lite Block Diagram

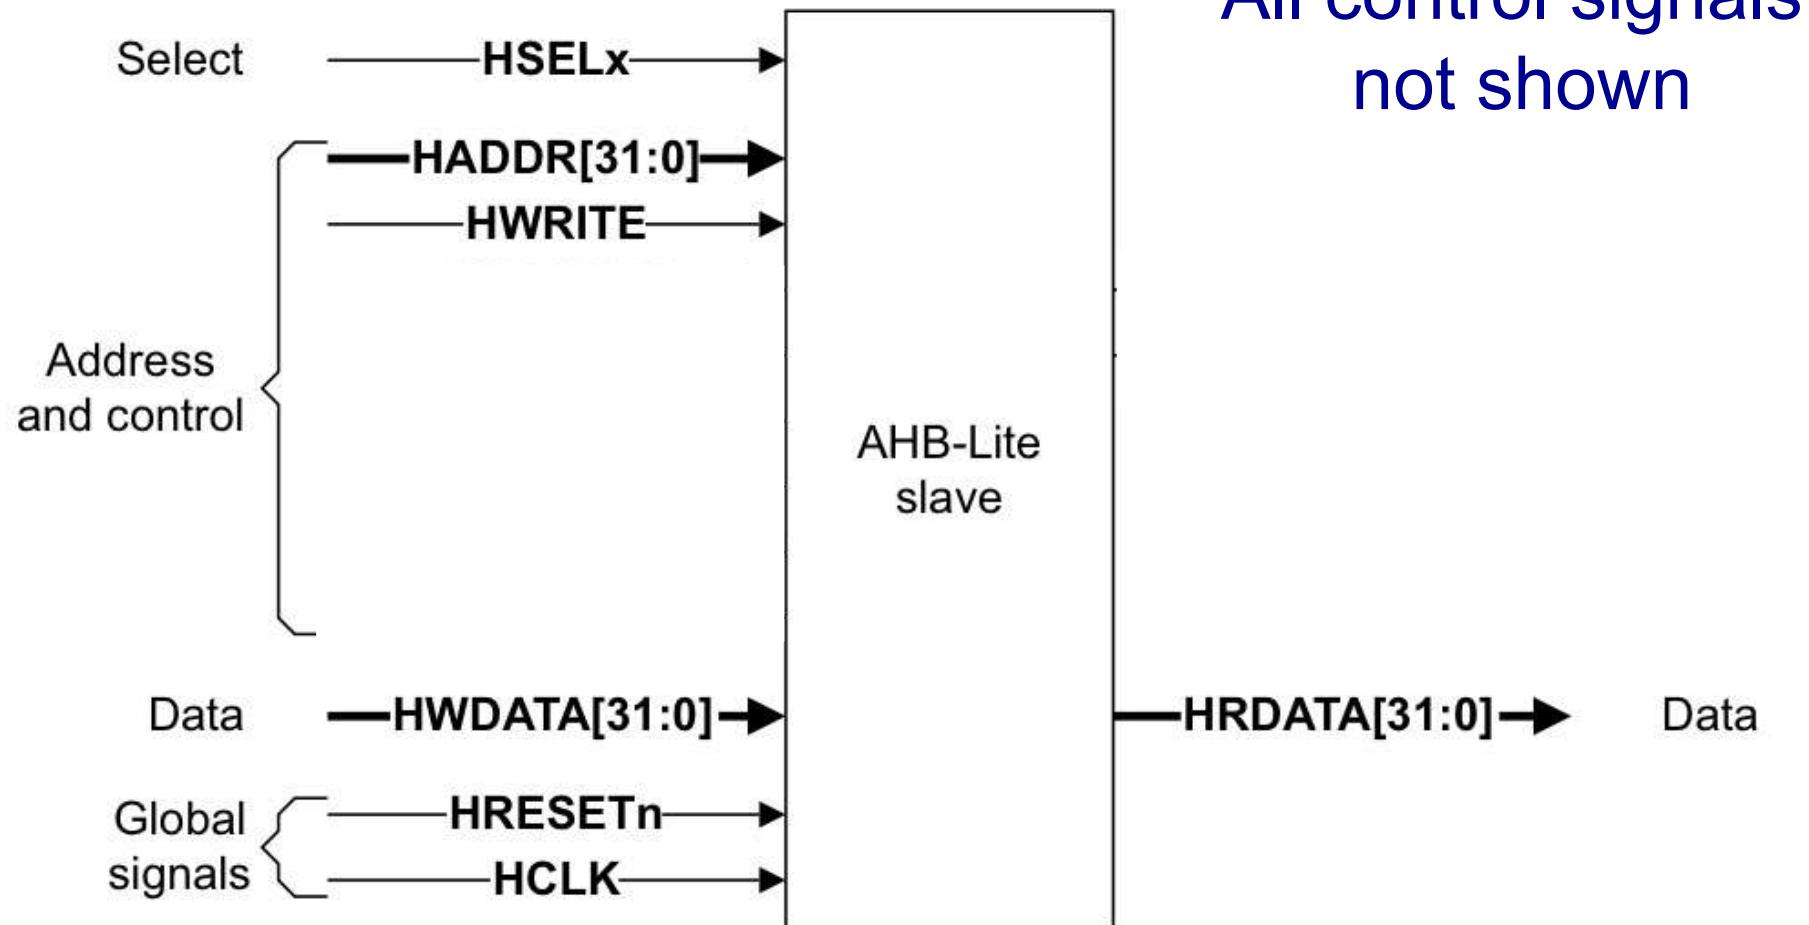


Master Interface



All control signals not shown

Slave Interface



THANKS

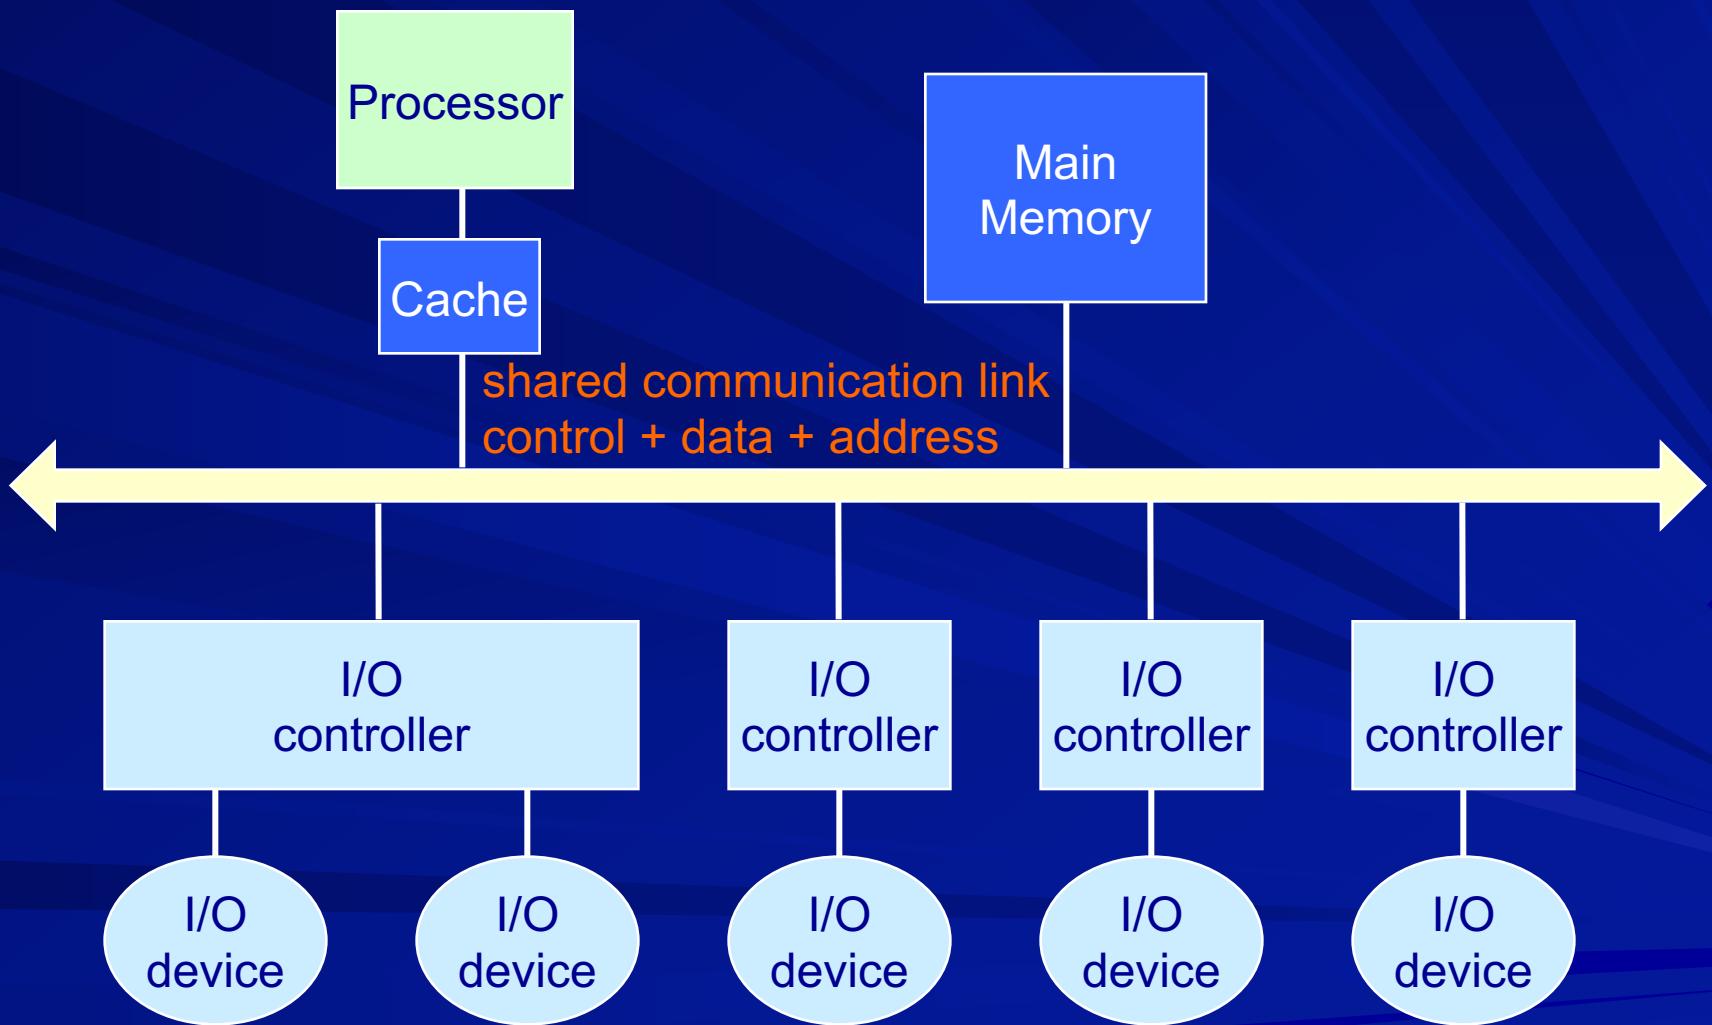
COL216

Computer Architecture

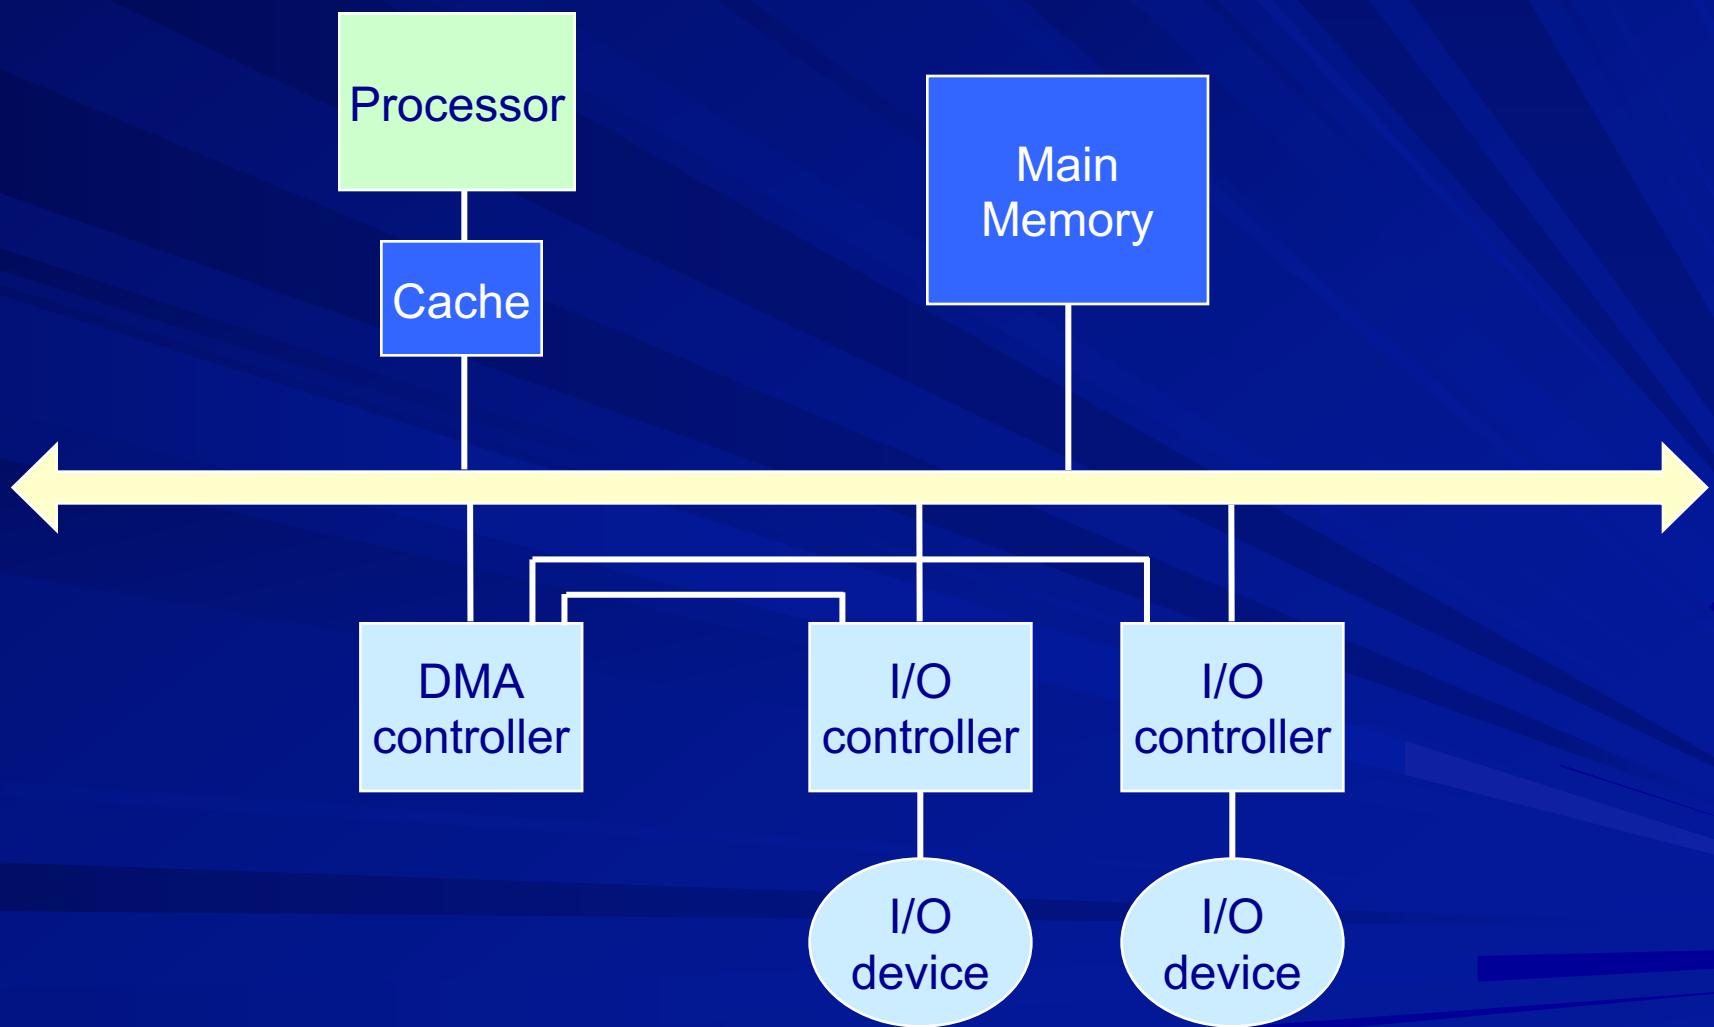
Input/Output – 4

24th March 2022

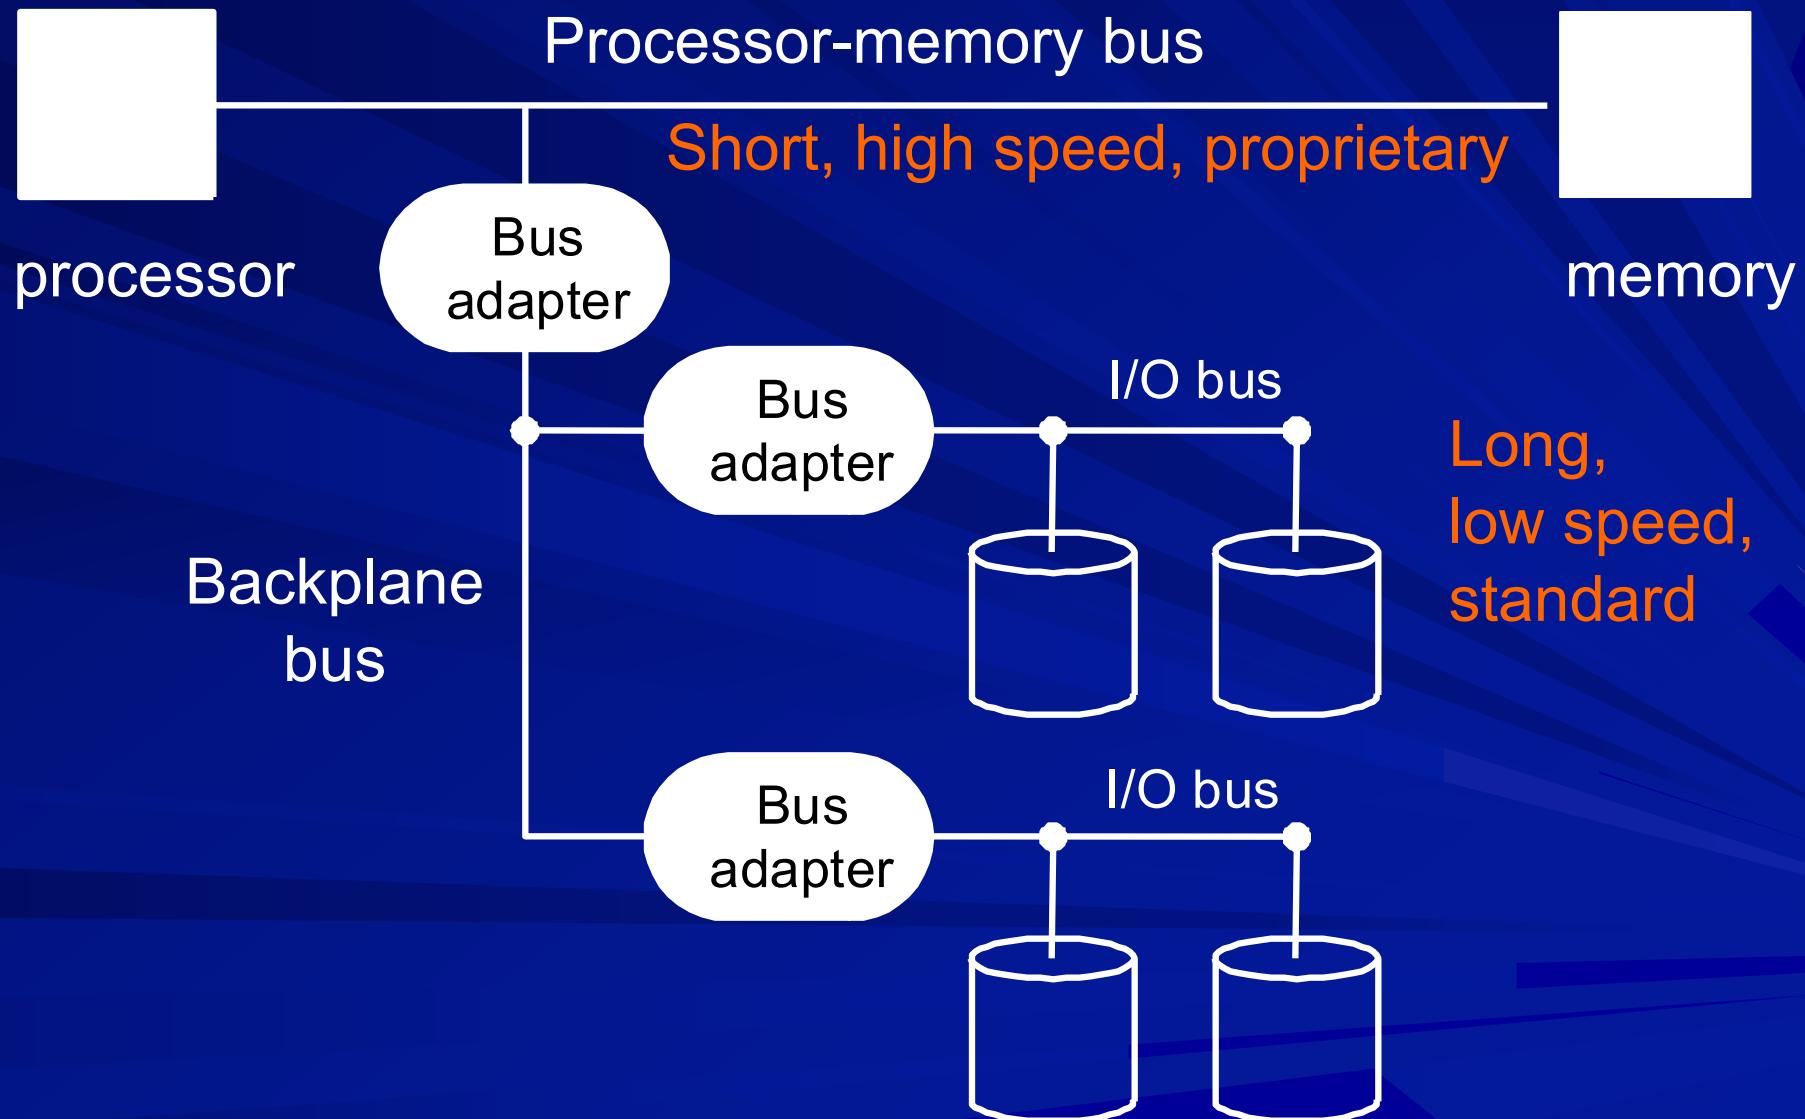
Single bus connecting all



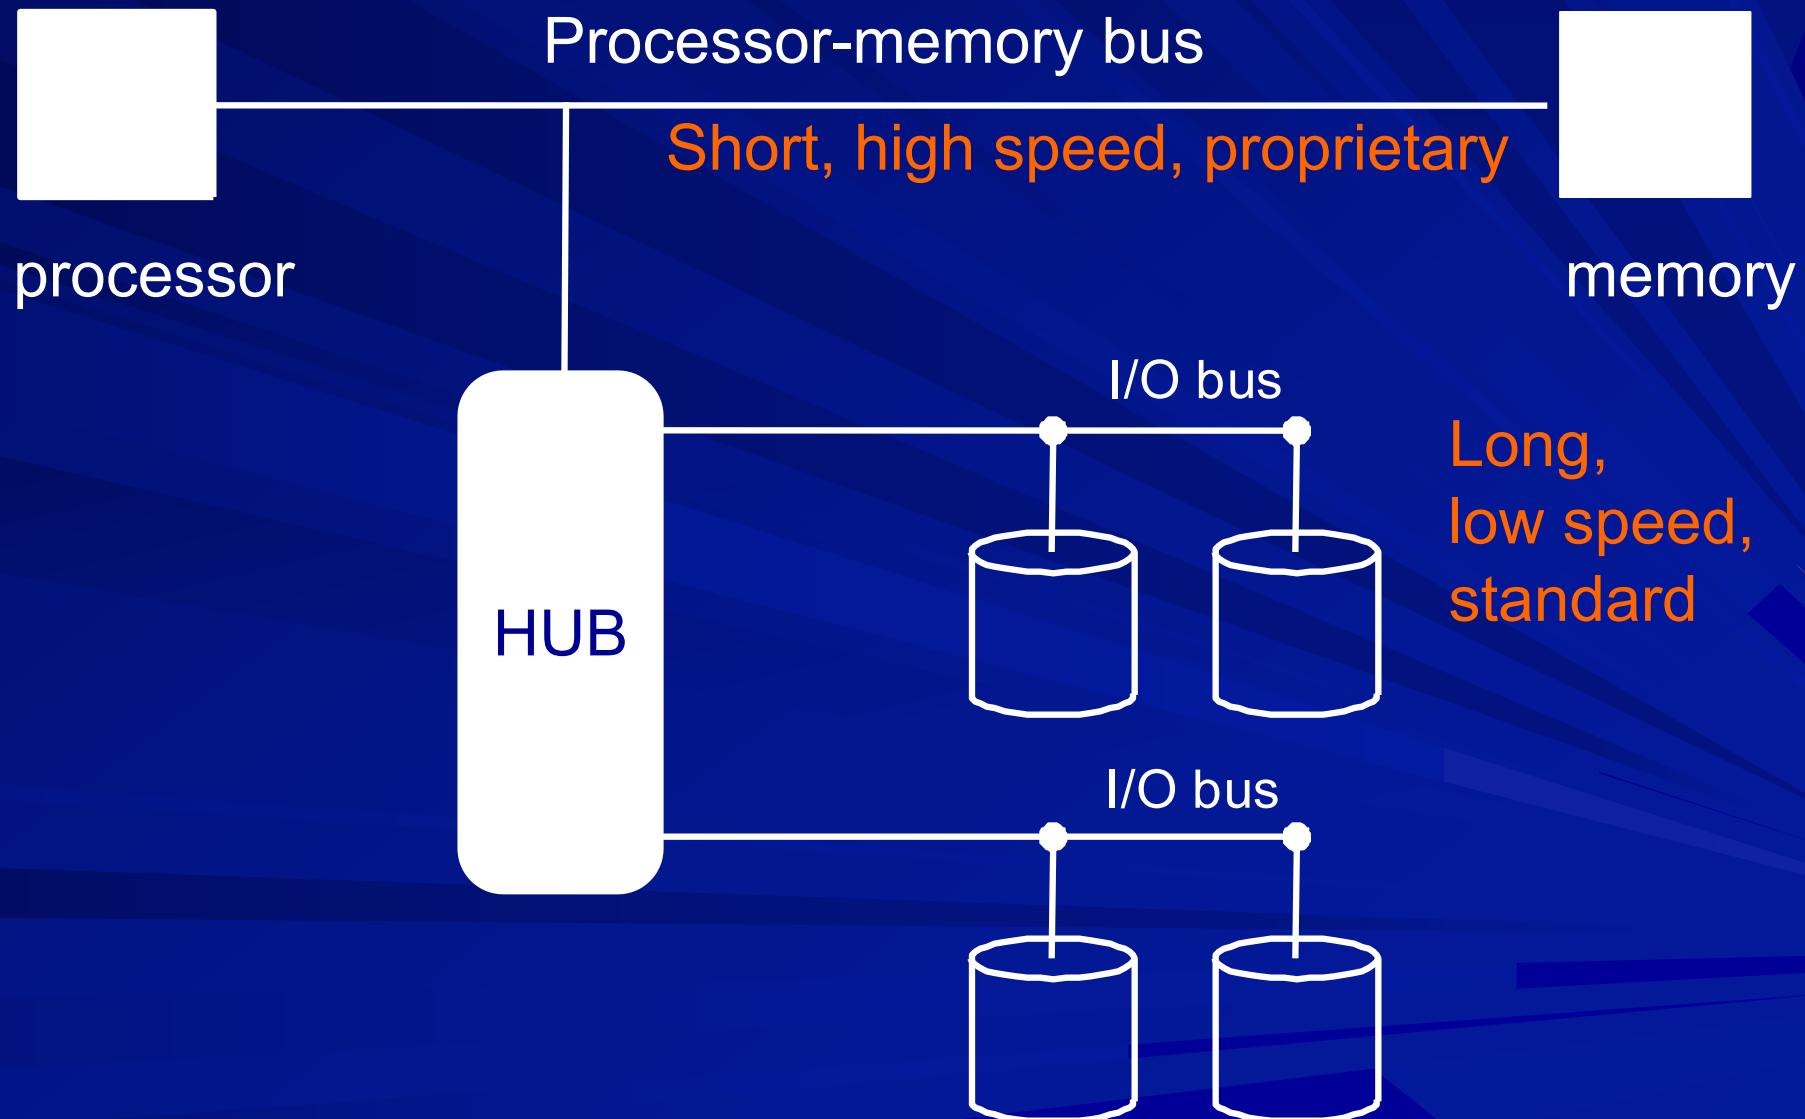
Direct memory access



System with multiple buses



System with multiple buses



Bus Standards

- Physical / mechanical
 - Pins, connectors, cables
- Electrical
 - Voltage / current levels, impedances
- Logical
 - Definition of signals
 - Timings and protocols

Standard buses/ports/interfaces

- Standardization required for subsystems from multiple sources to work together
- Technological developments lead to continuous improvement of specs, standardization implies freezing the specs
- Standards need to be revised and improved periodically
- Standardization is done by consortium of several organizations or professional bodies

AMBA : Advanced Microcontroller Bus Architecture

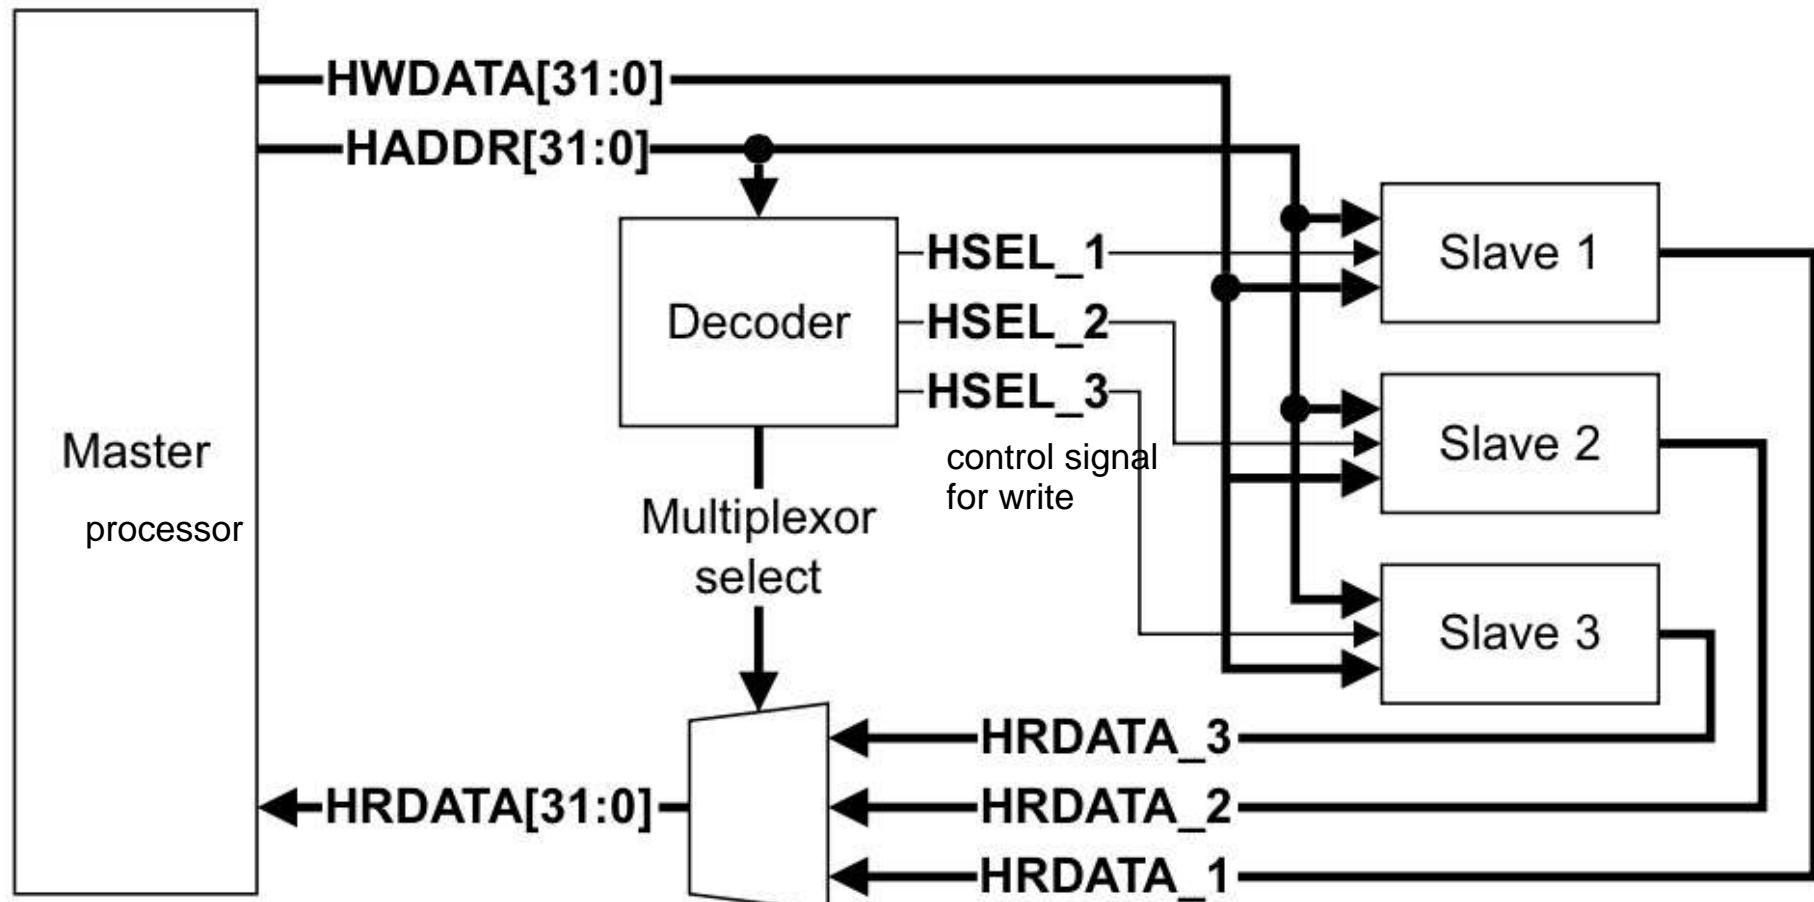
for ARM

- ARM open standard for on-chip buses
- For System-on-Chip (SoC)
- Variants
 - AHB : Advanced High speed Bus
 - APB : Advanced Peripheral Bus
 - AXI : Advanced eXtensible Interface
 - AHB-Lite : subset of AHB

AHB-Lite

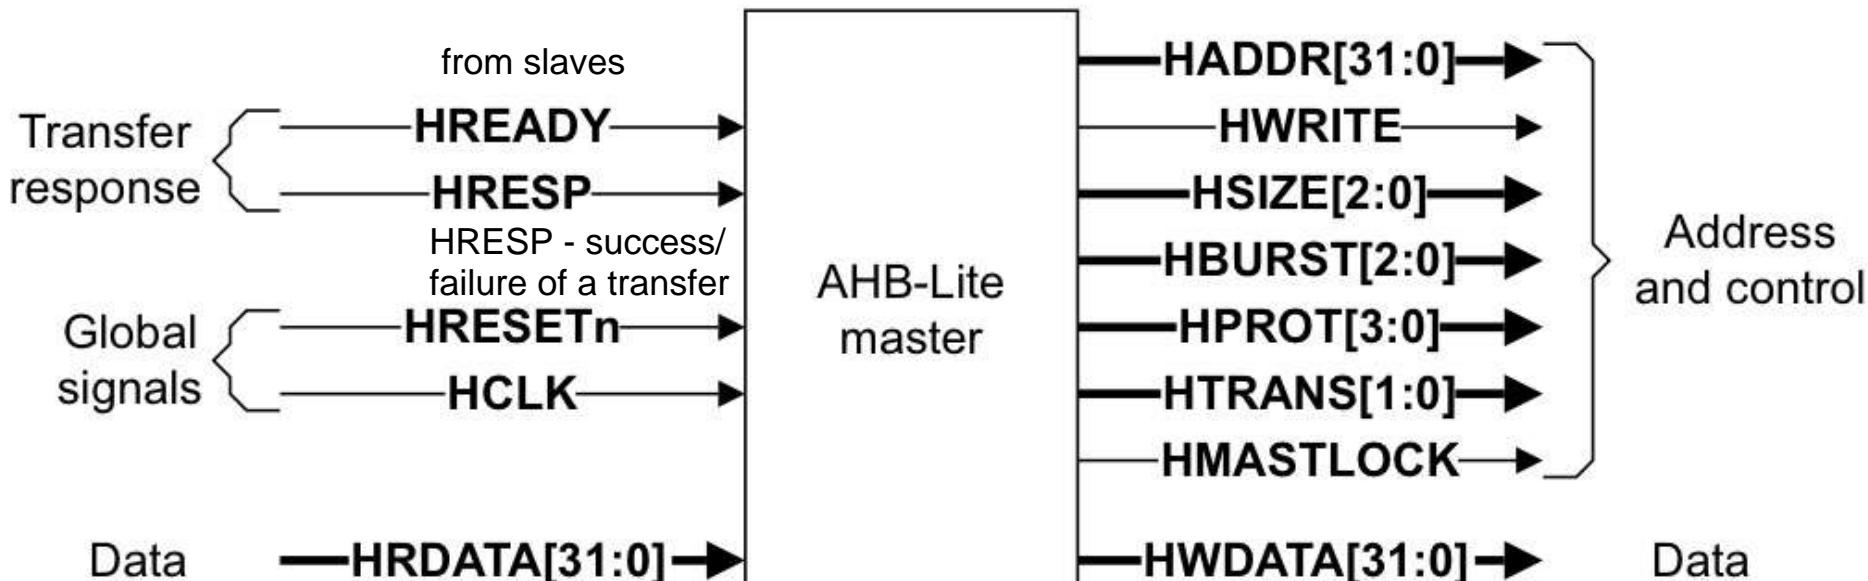
- Single master, multiple slaves
- Parallel
- Multiplexed (not tri-state)
- Separate data, address, control
- Synchronous
- Pipelined
- Burst transfer supported
- Split transaction not supported

AHB-Lite Block Diagram



Master Interface

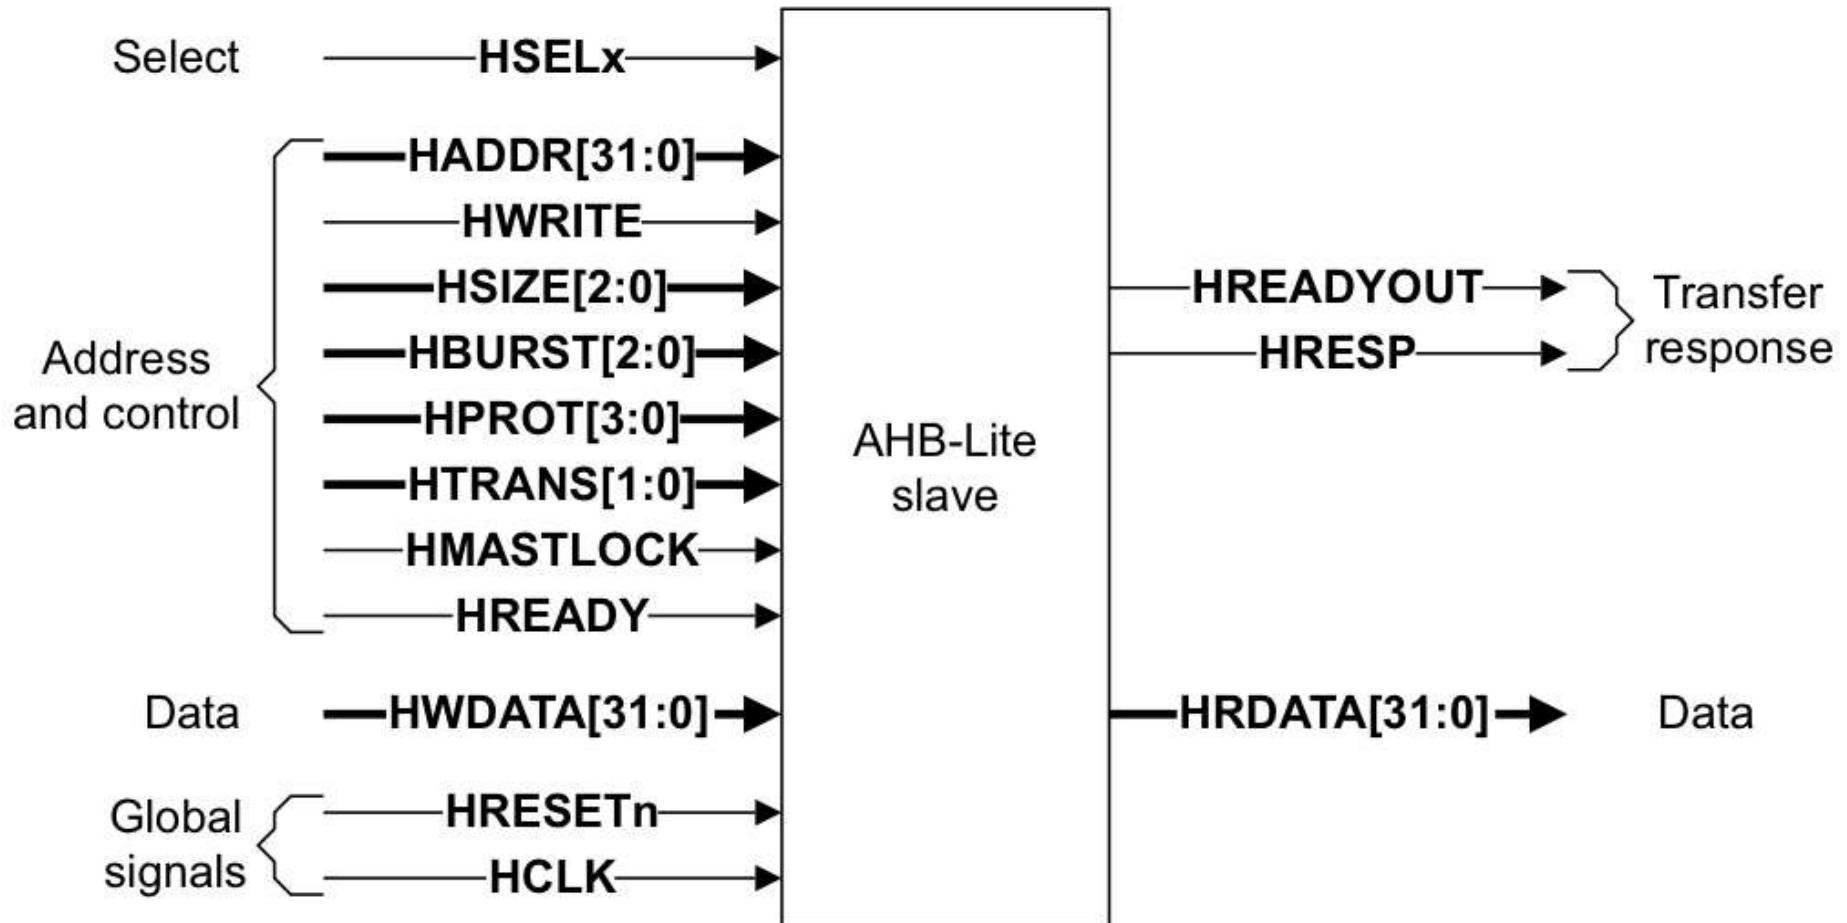
HREADY - slave asks the master
for extension of data phase by asking master to introduce wait cycles



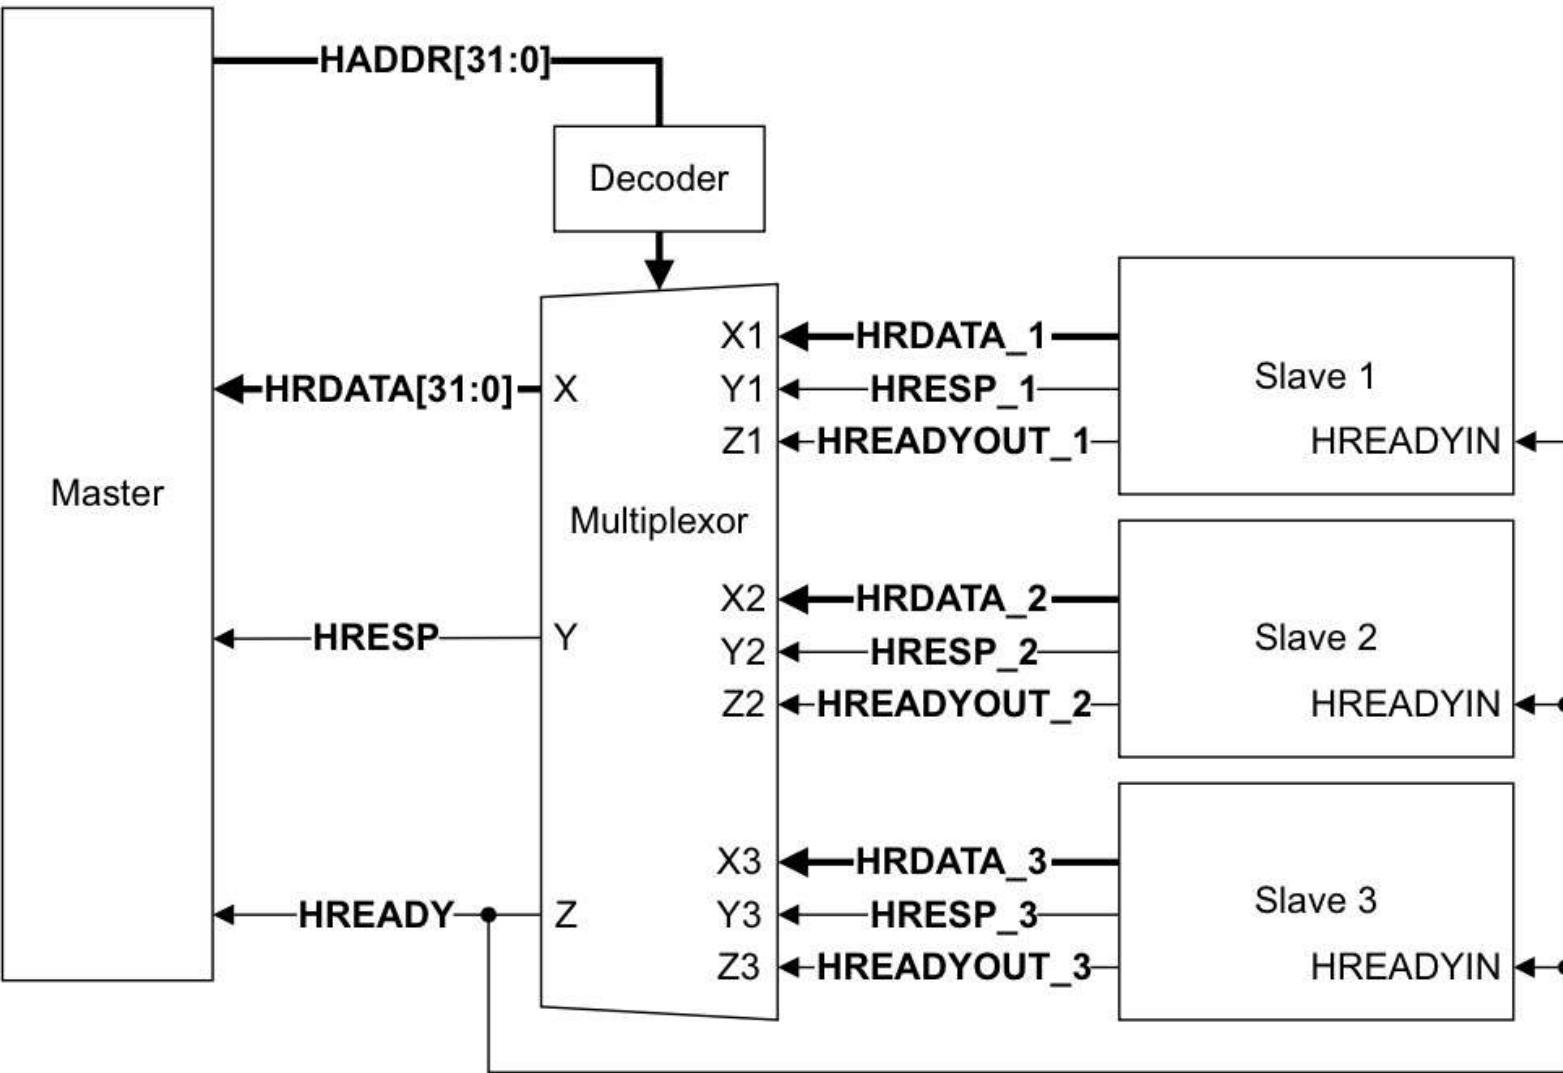
write data bus transfers data to be written
from master to slave while read data bus
moves from slave to master

Slave Interface

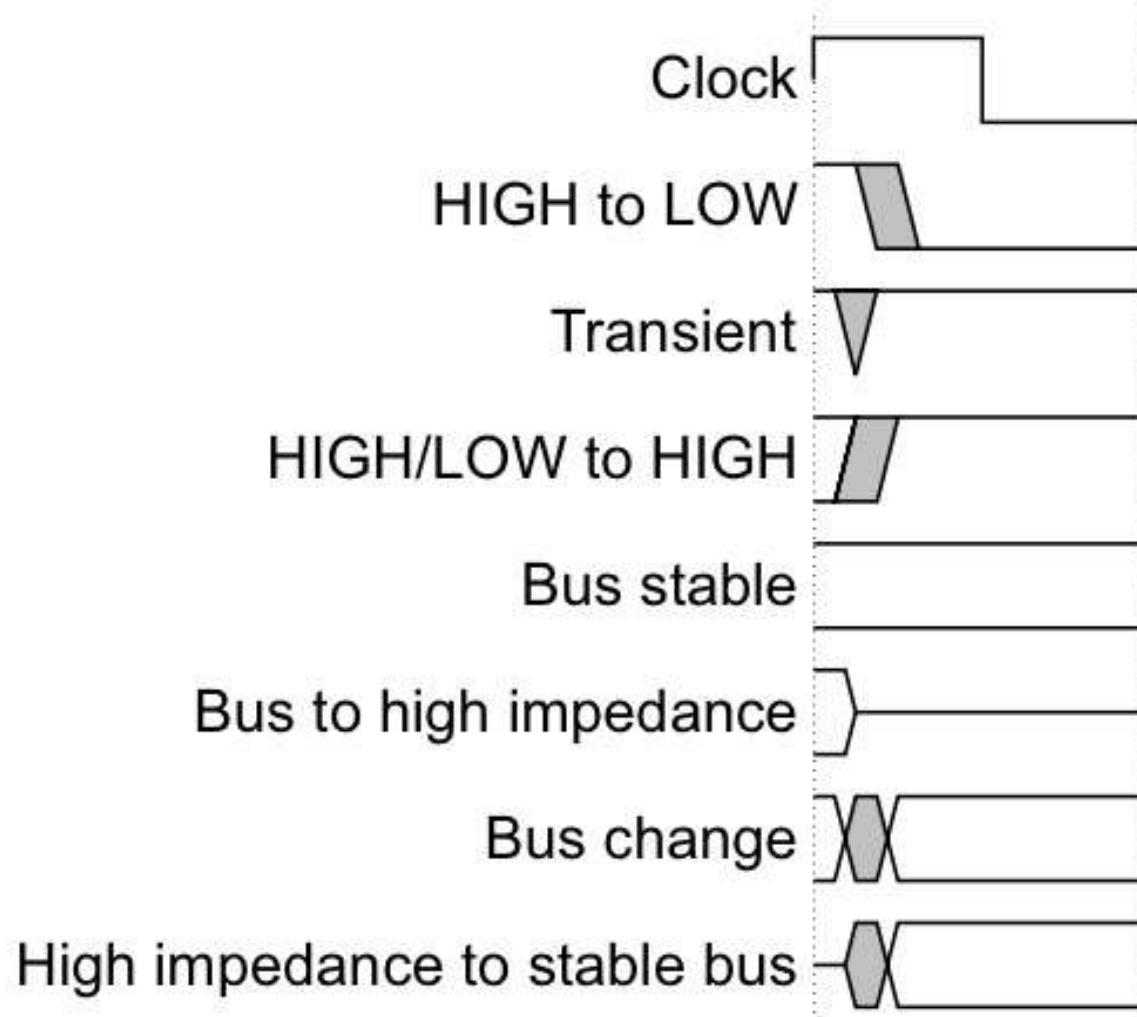
HREADY signal indicates to master and all slaves that the previous transfer is complete



Multiplexor Interconnection

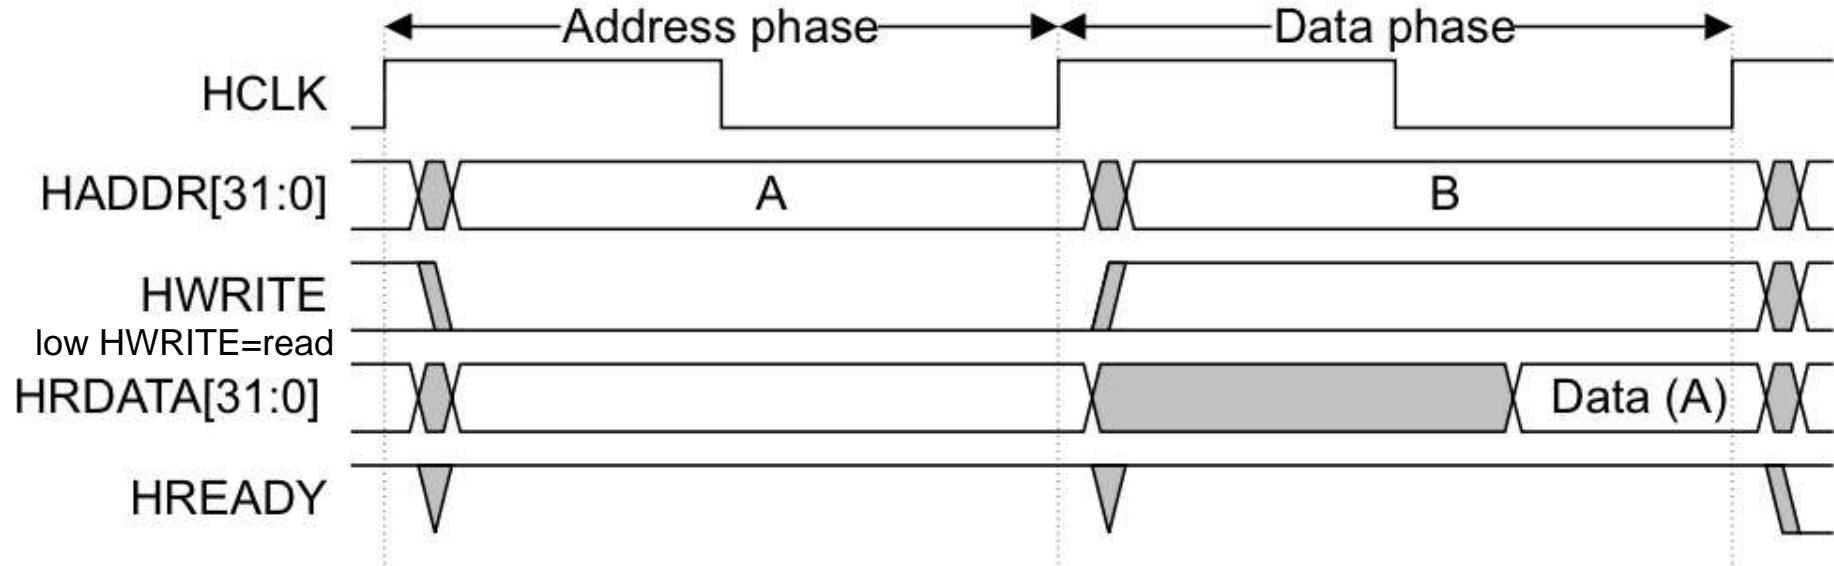


Timing Diagram Conventions



Read Transfer

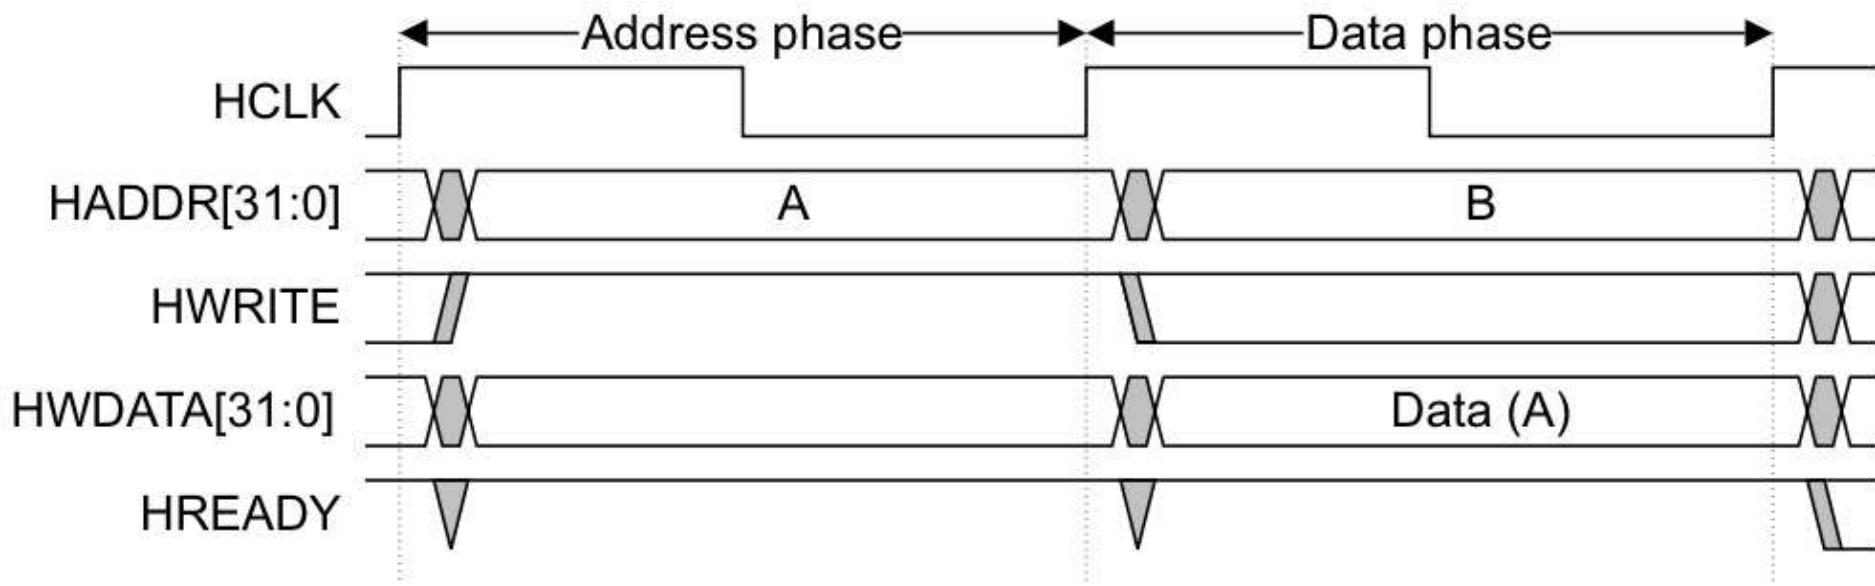
Advantage of this is pipelining, we can have overlap



Here, HREADY = 1 indicates that the master has sent the address and next HREADY = 1 indicates that the slave is ready with the data. If not, it introduces wait cycles.

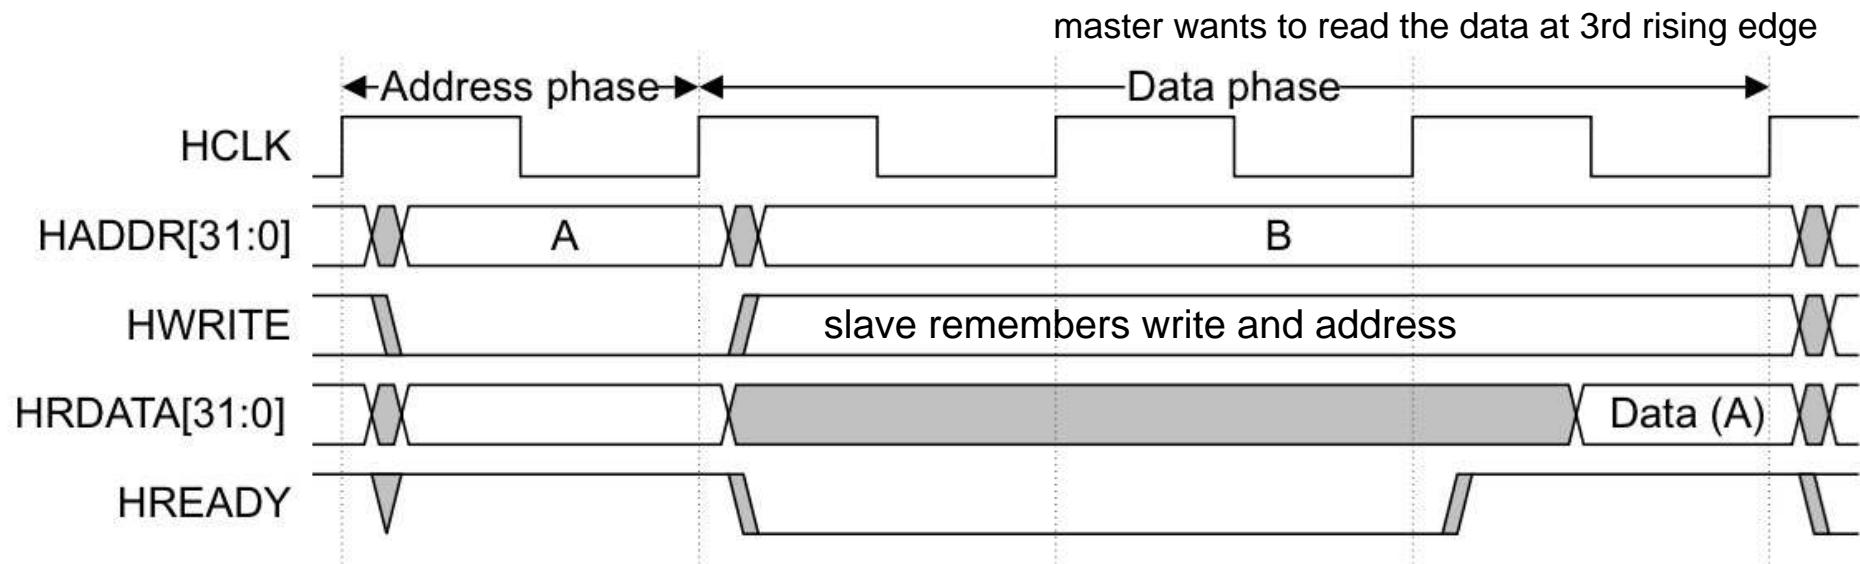
Write Transfer

note: in read transfer, data is stable a long time after the start, while in write transfer, we see data is stable just after the start.



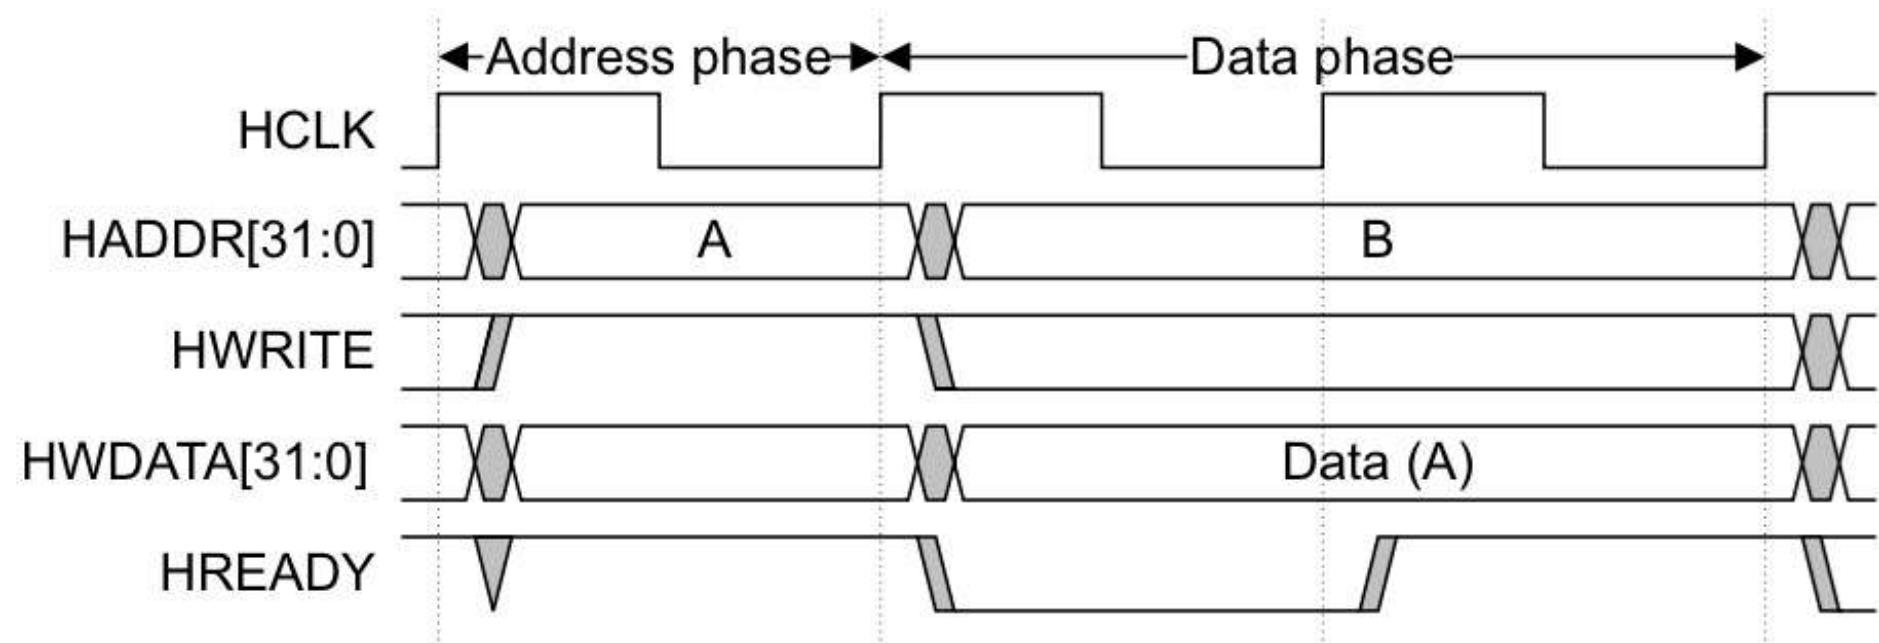
Data gets written at the rising edge of the 3rd clock cycle, after slave has seen the data to be written by keeping HREADY as 1 during the data phase itself. HWRITE is 1 to indicate the data needs to be written

Read Transfer with Wait States

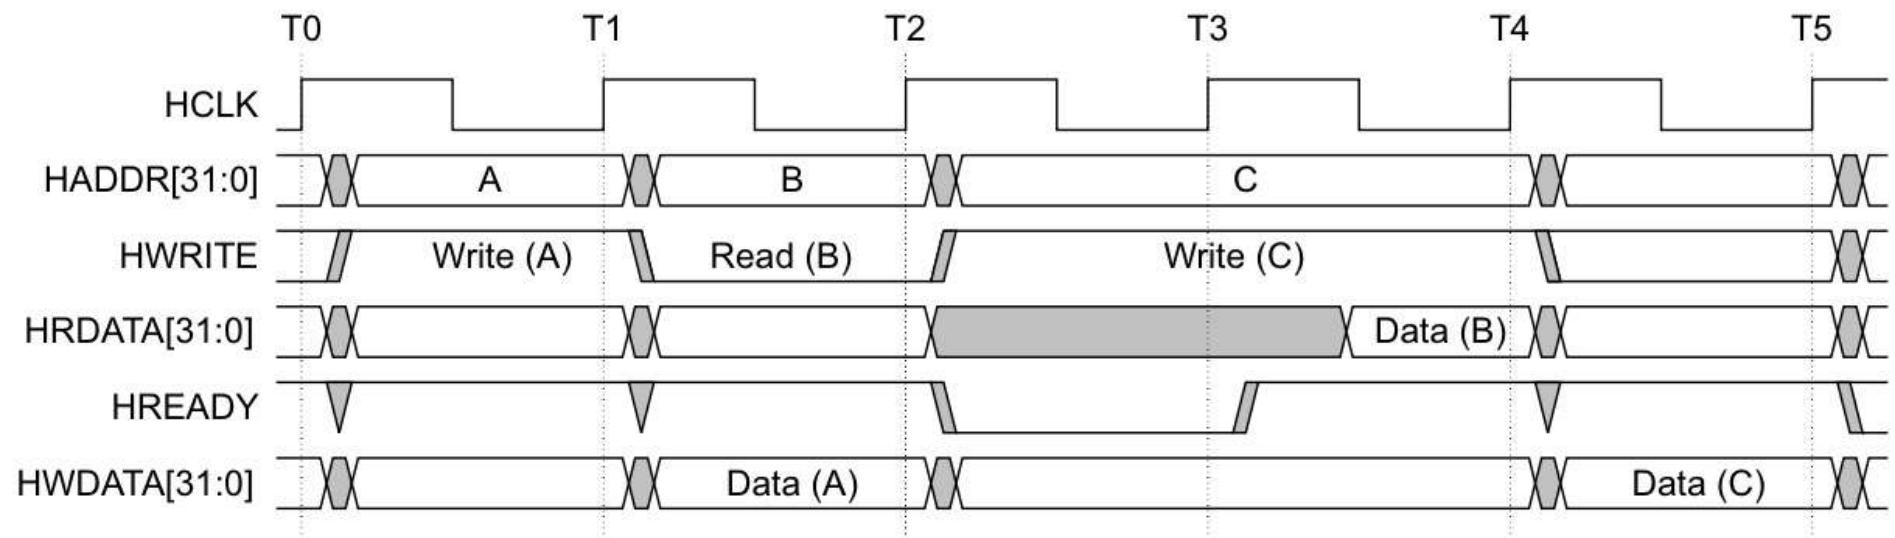


Control comes to slave when we are talking about data while control remains with master when we are talking about address

Write Transfer with Wait States



Multiple Transfers



Transfer Types

HTRANS[1:0] Type

00 IDLE

Master does not want transfer

01 BUSY

data on data bus or address bus can be ignored

Master inserts wait cycle

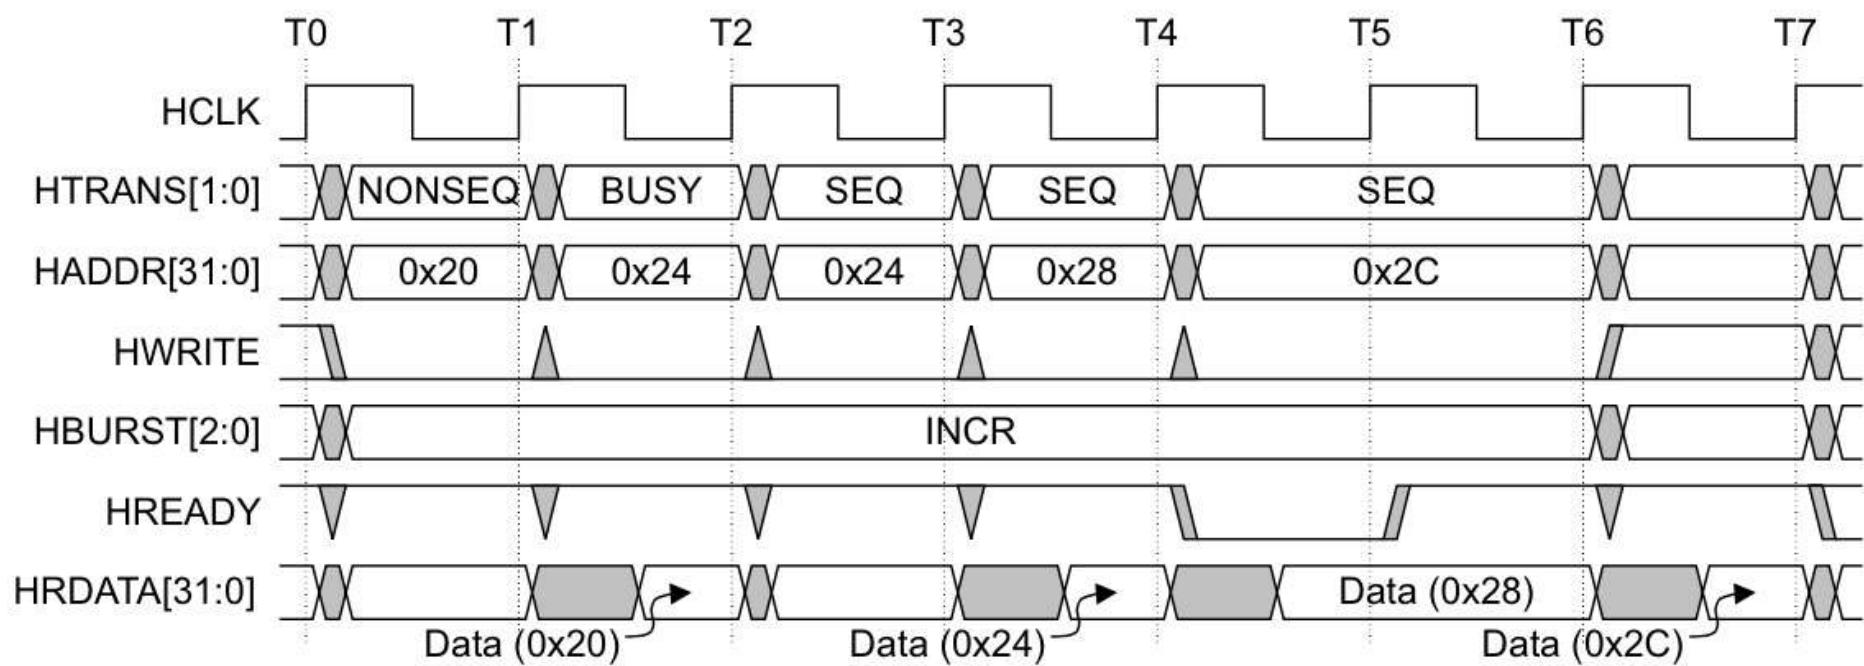
NONSEQ

Single transfer or first transfer of
a burst

SEQ

Remaining transfers of a burst

Transfer Type Examples



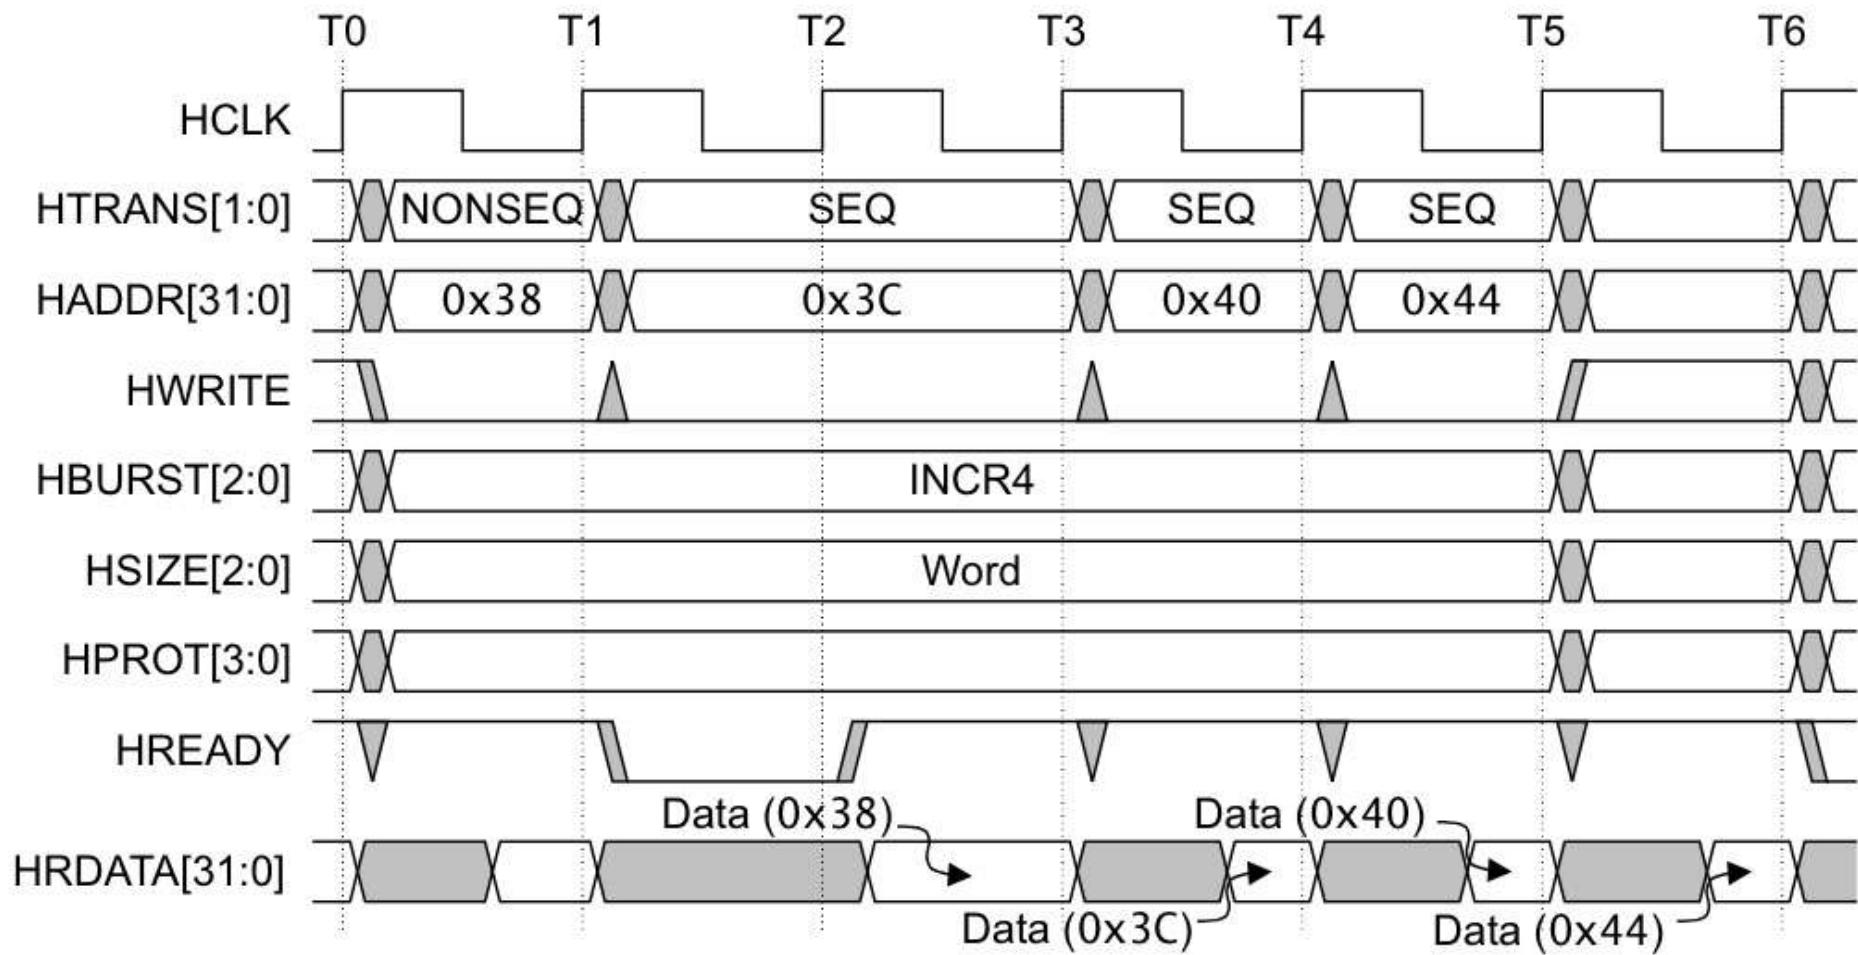
Transfer Size

HSIZE[2:0]	Size (bits)	Size (bytes)
000	8	1
001	16	2
010	32	4
011	64	8
100	128	16
101	256	32
110	512	64
111	1024	128

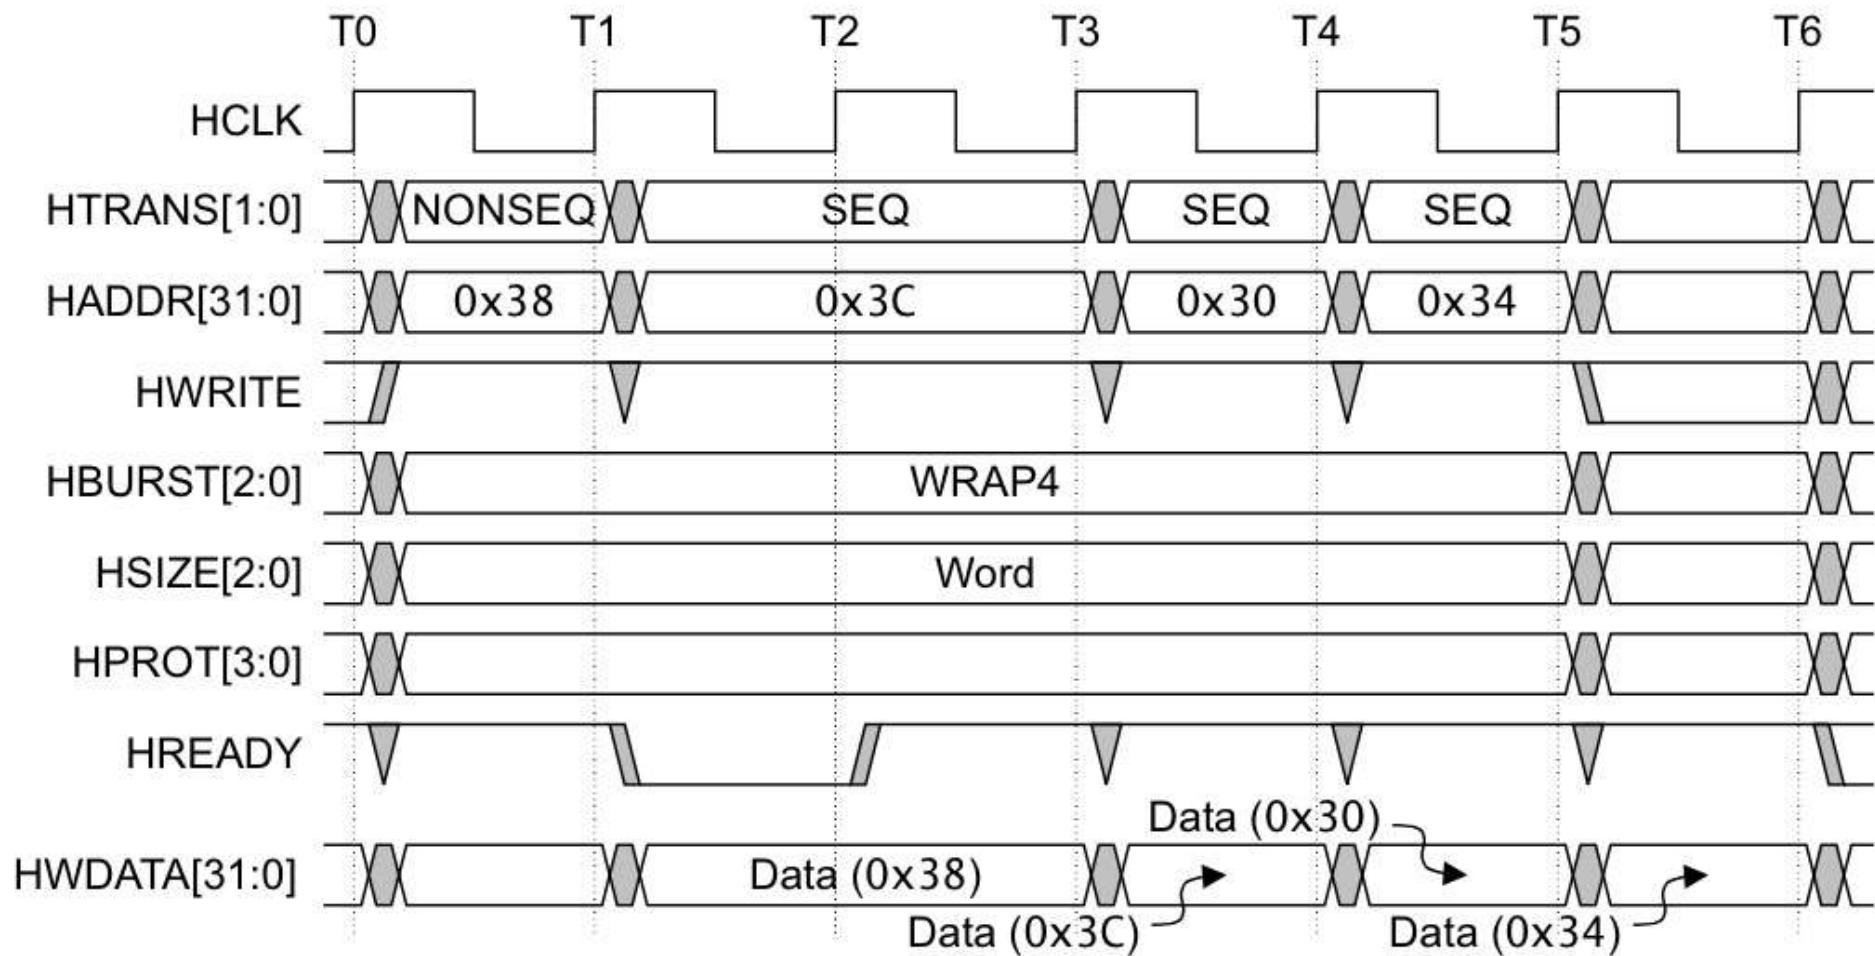
Burst type and size

HBURST[2:0]	Type and size
000	SINGLE
001	INCR
010	WRAP4
011	INCR4
100	WRAP8
101	INCR8
110	WRAP16
111	INCR16

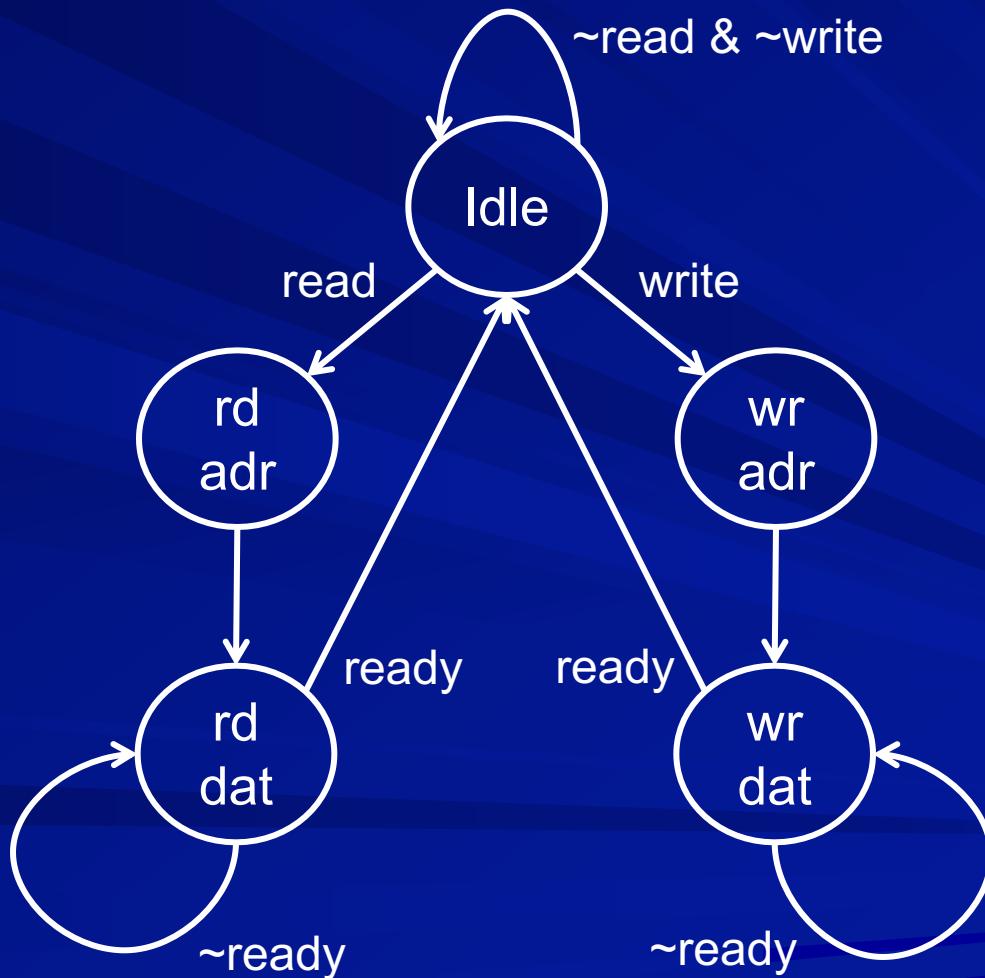
Four-Beat Incrementing Burst



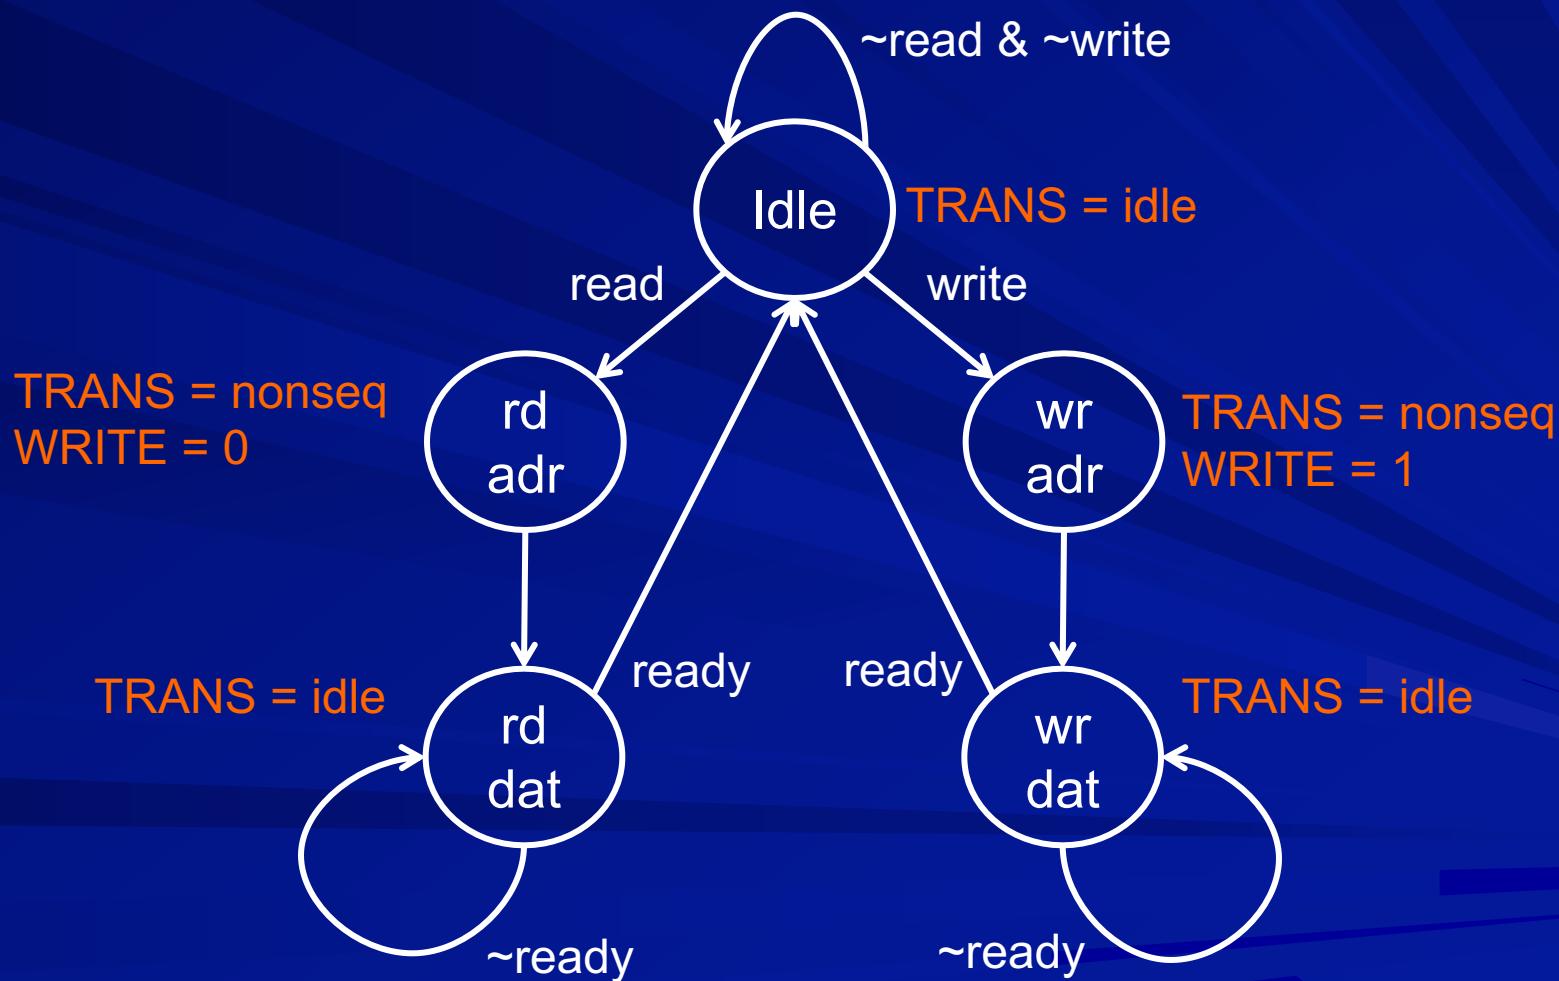
Four-Beat Wrapping Burst



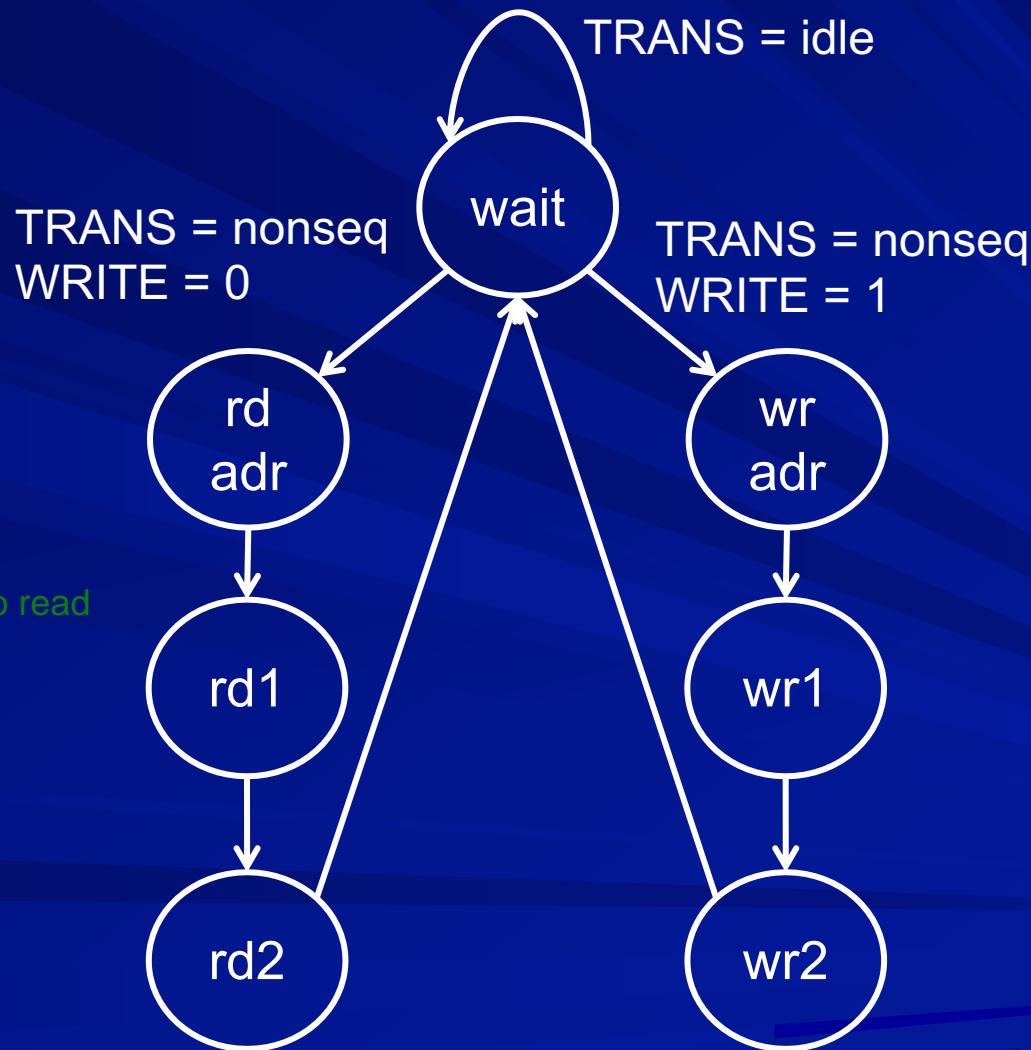
Master interface



Master interface



Slave interface



Obtaining access to a bus

Masters and slaves on a bus

I/O device are peripherals could act as master via DMA controller

masters (processors/peripherals) initiate transactions
slaves (peripherals/memories) respond to the masters

How does a master get control of a bus?

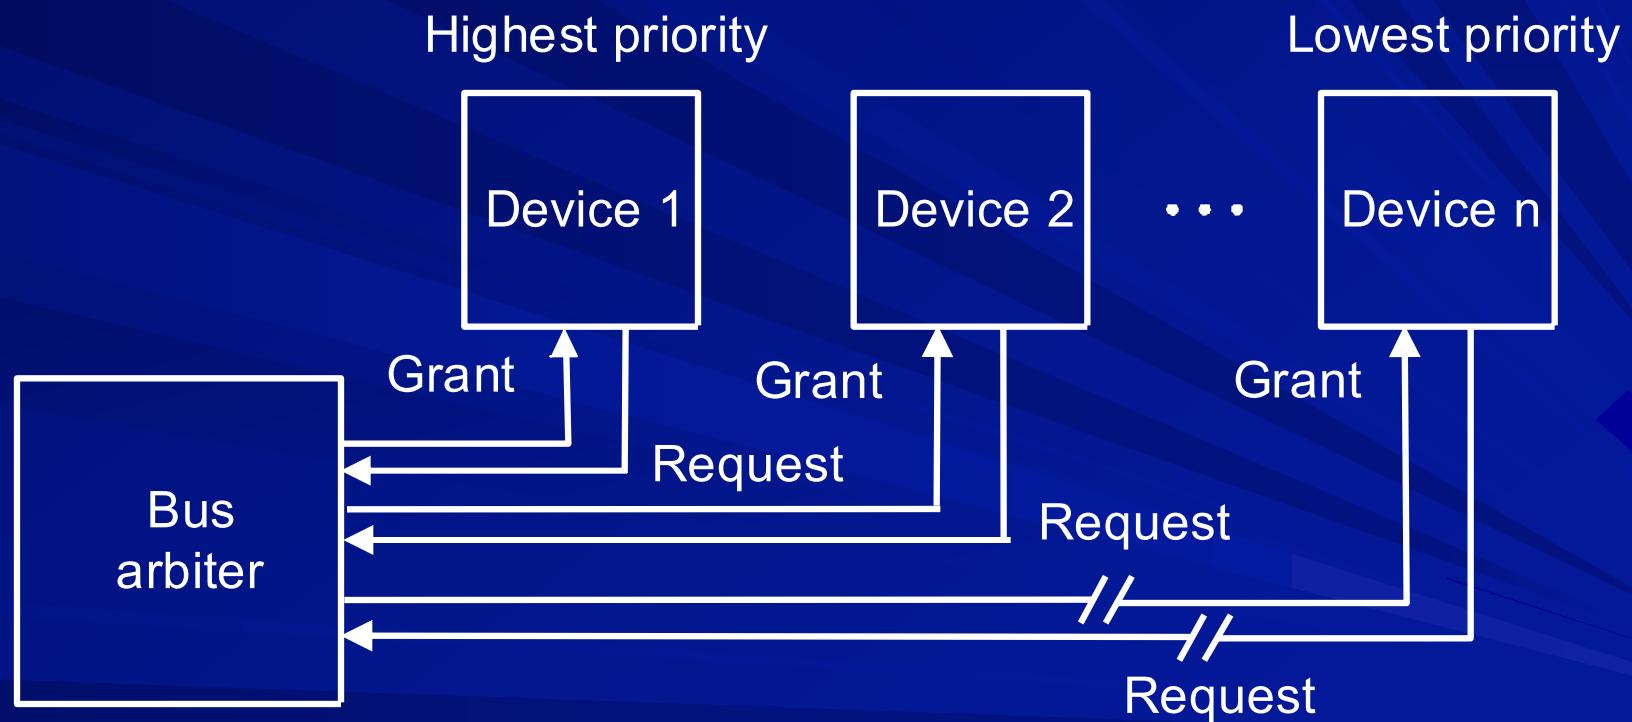
Priorities may be assigned. Fairness is essential.

- Daisy chaining
- Centralized parallel arbitration
- Distributed arbitration by self selection
- Distributed arbitration by collision detection

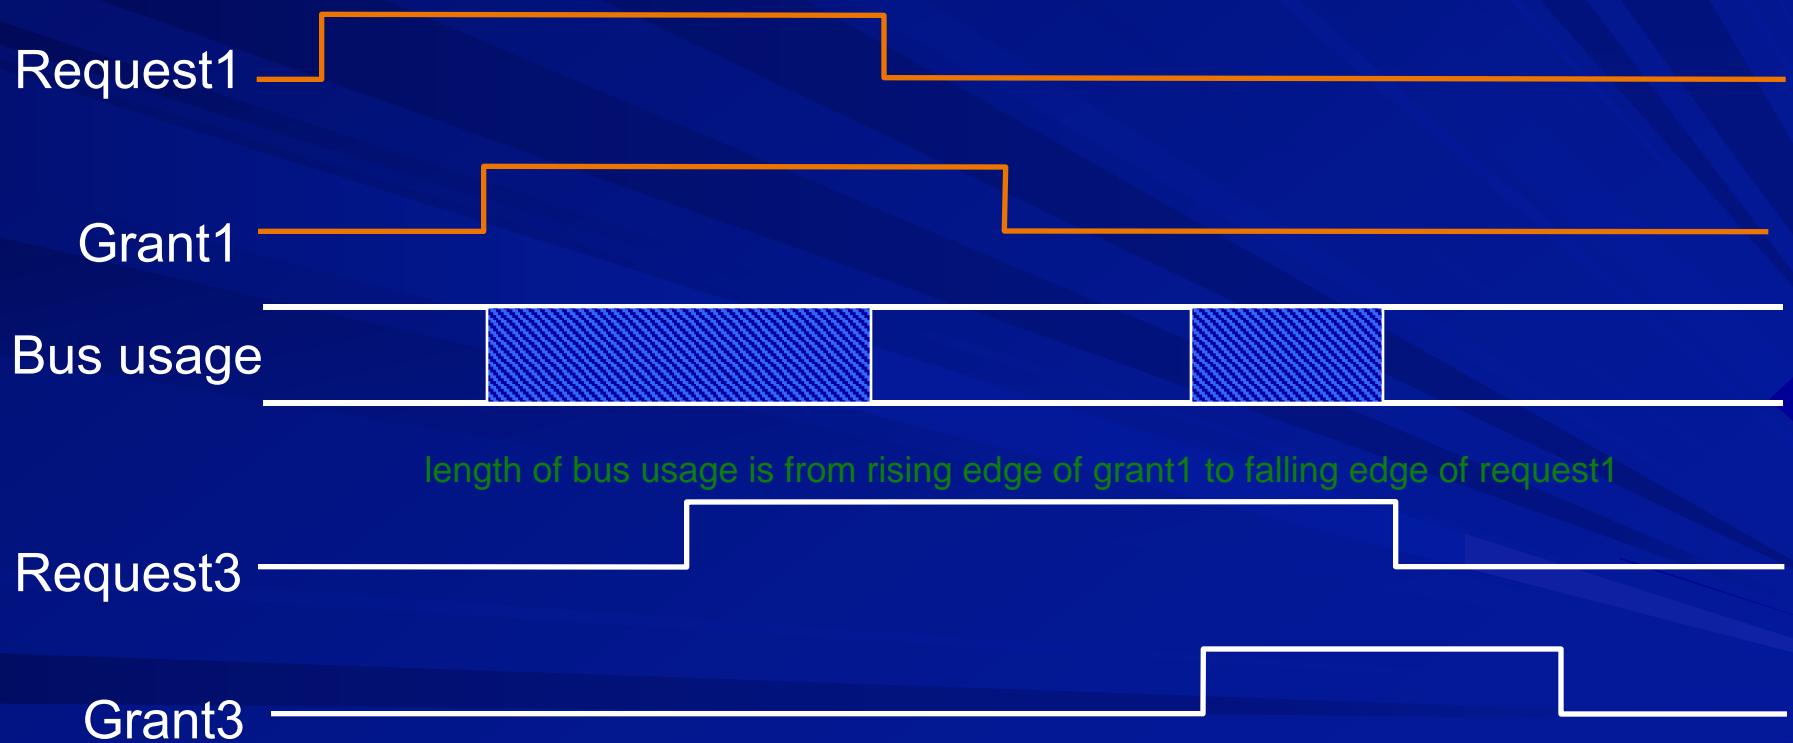
multiple masters - these devices are multiple masters

Centralized parallel arbitration

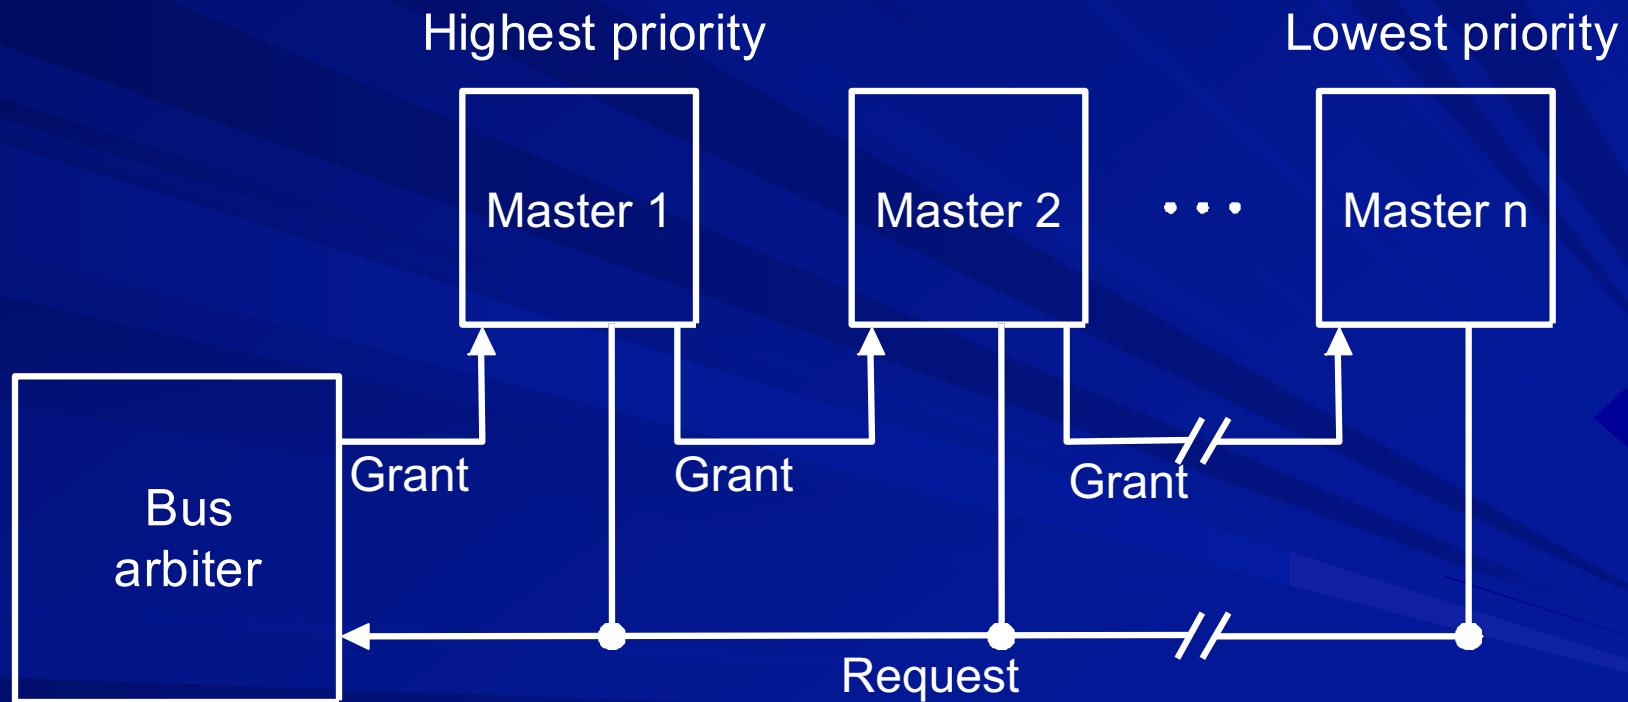
Example: PCI bus (a backplane bus)



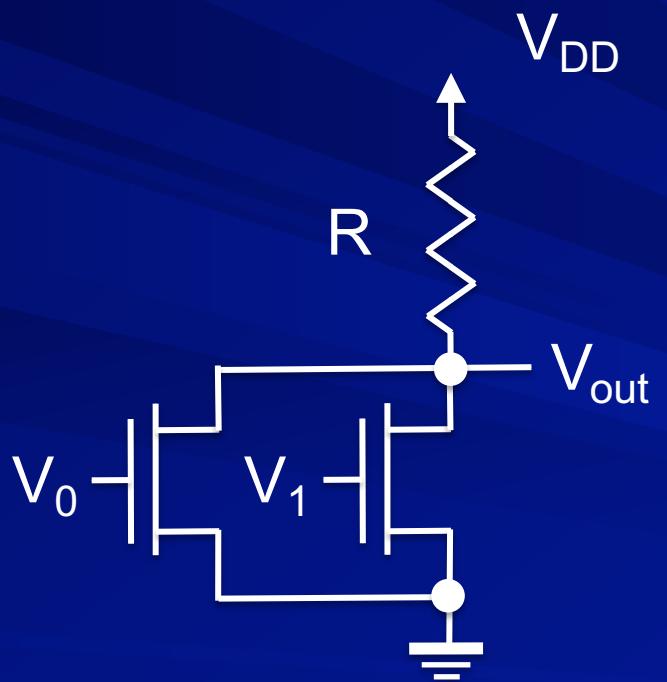
Request and Grant signals



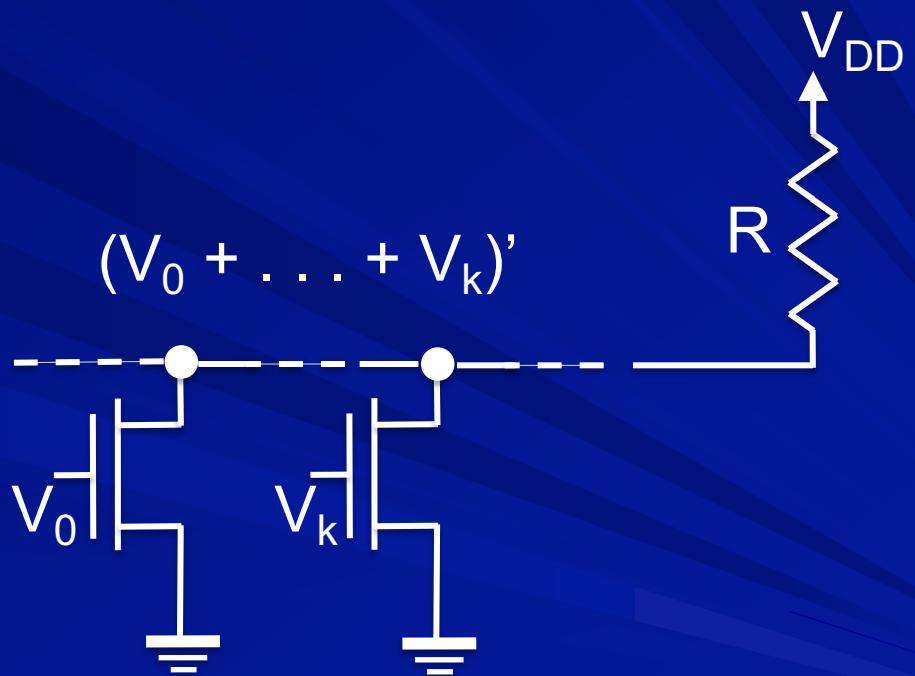
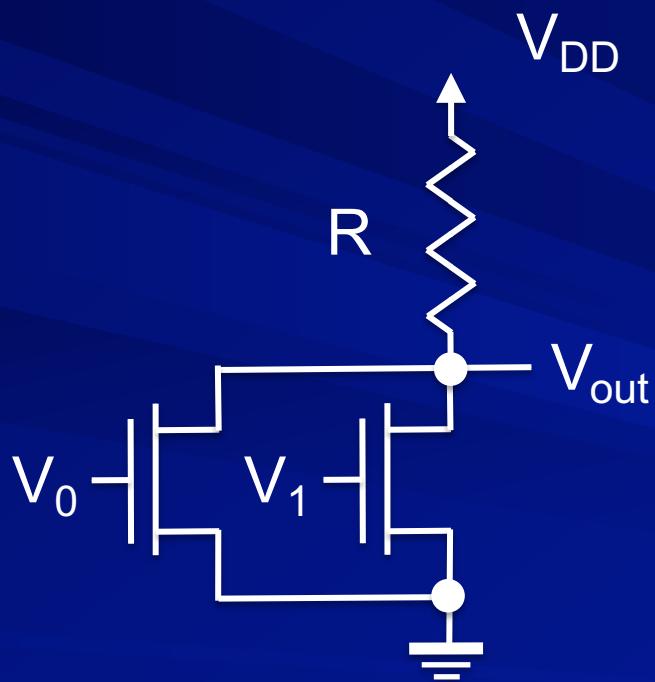
Combining/chaining signals



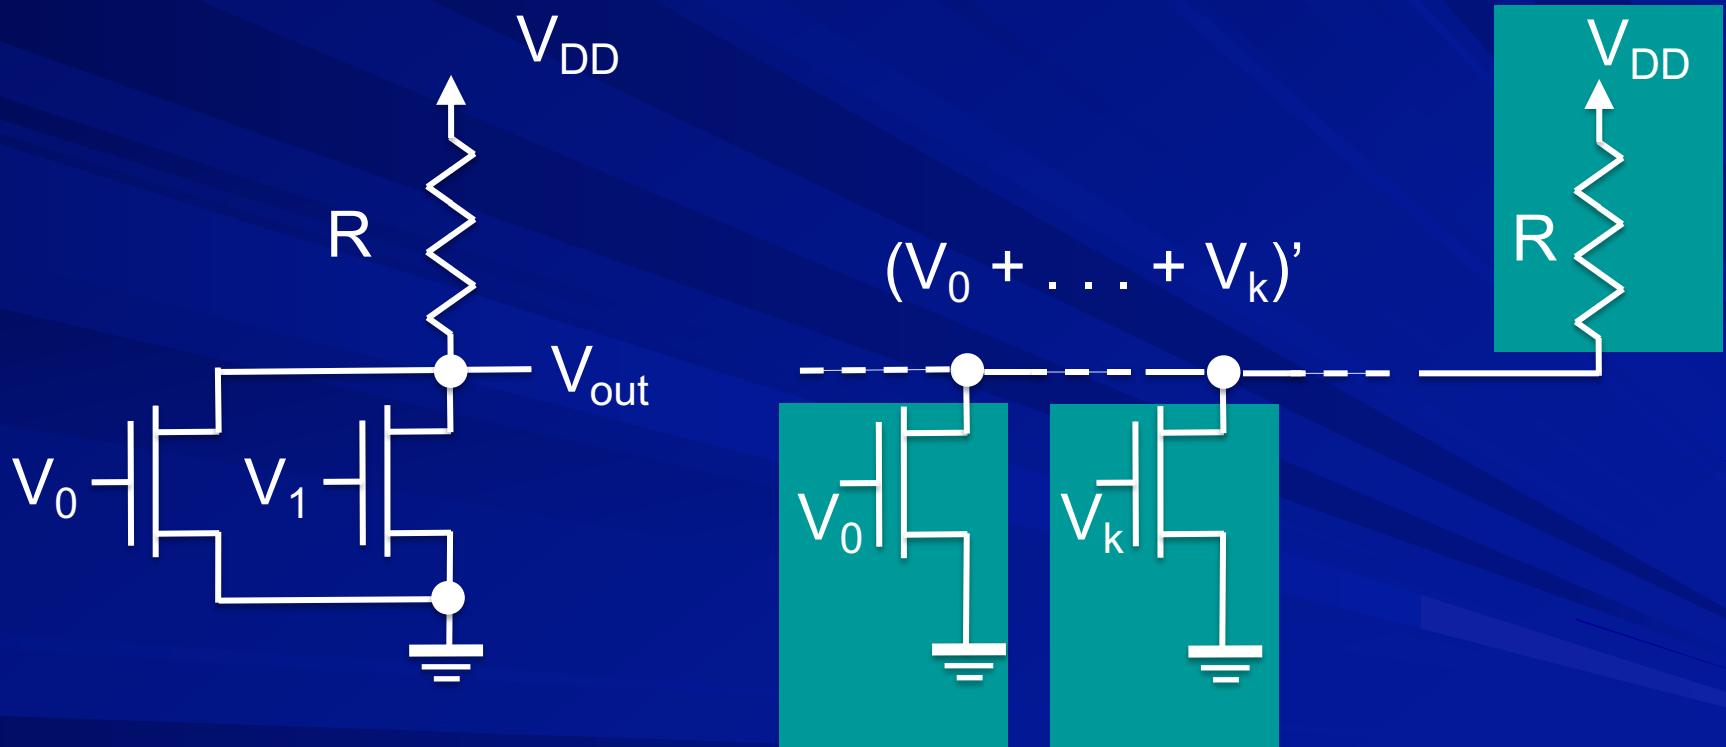
Wired OR



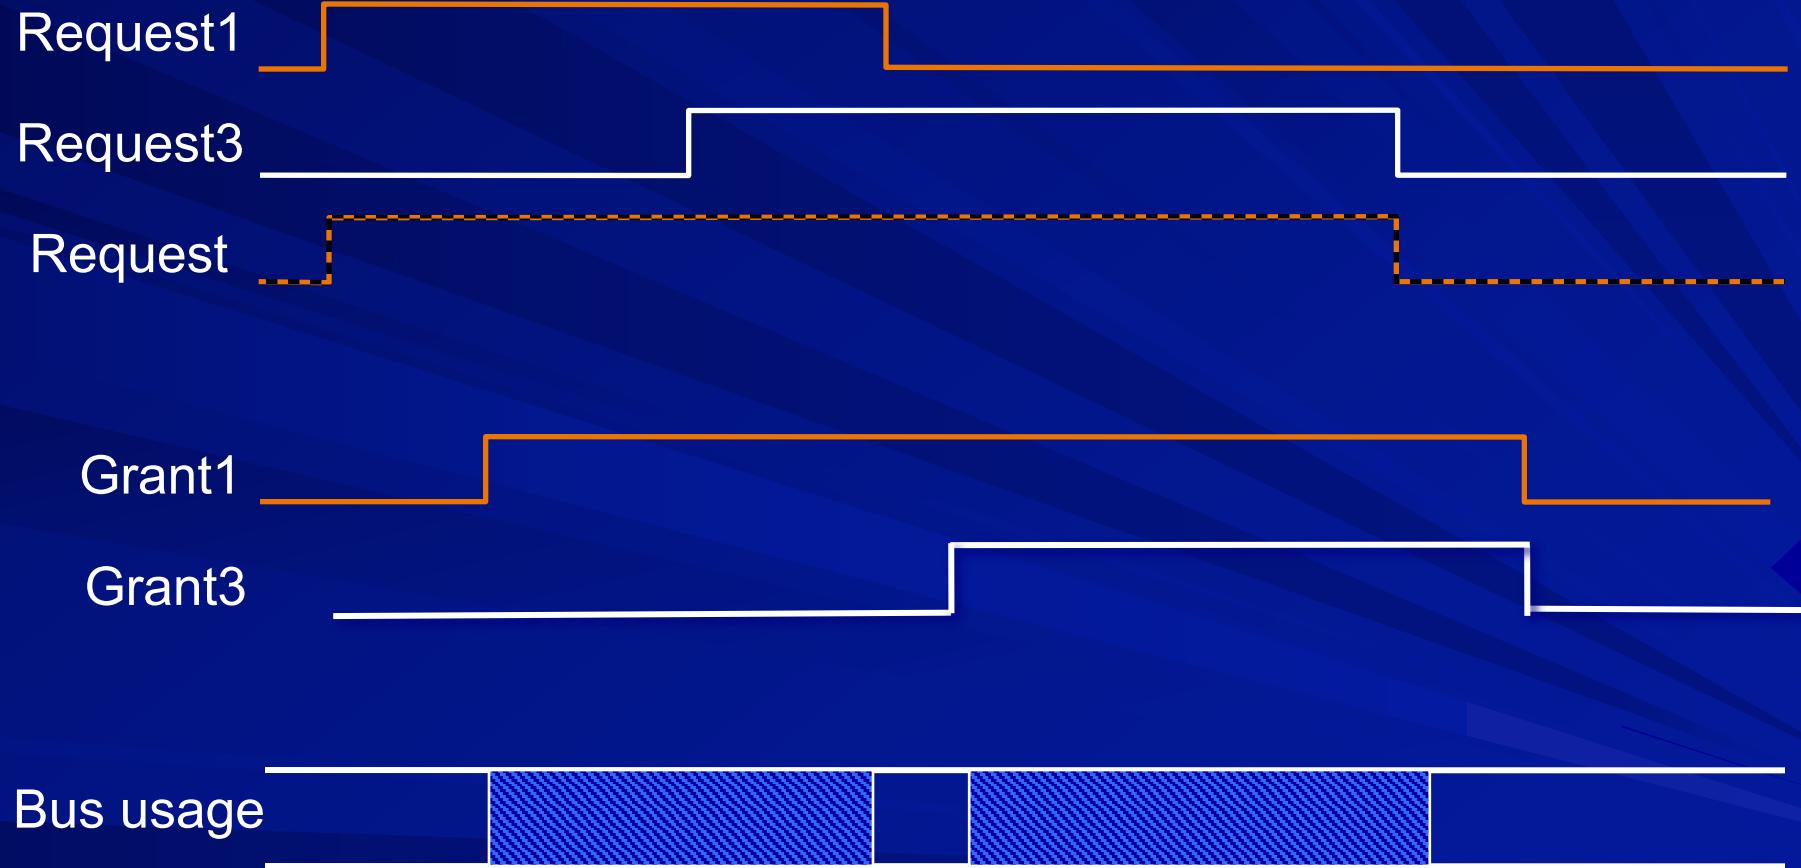
Wired OR



Wired OR

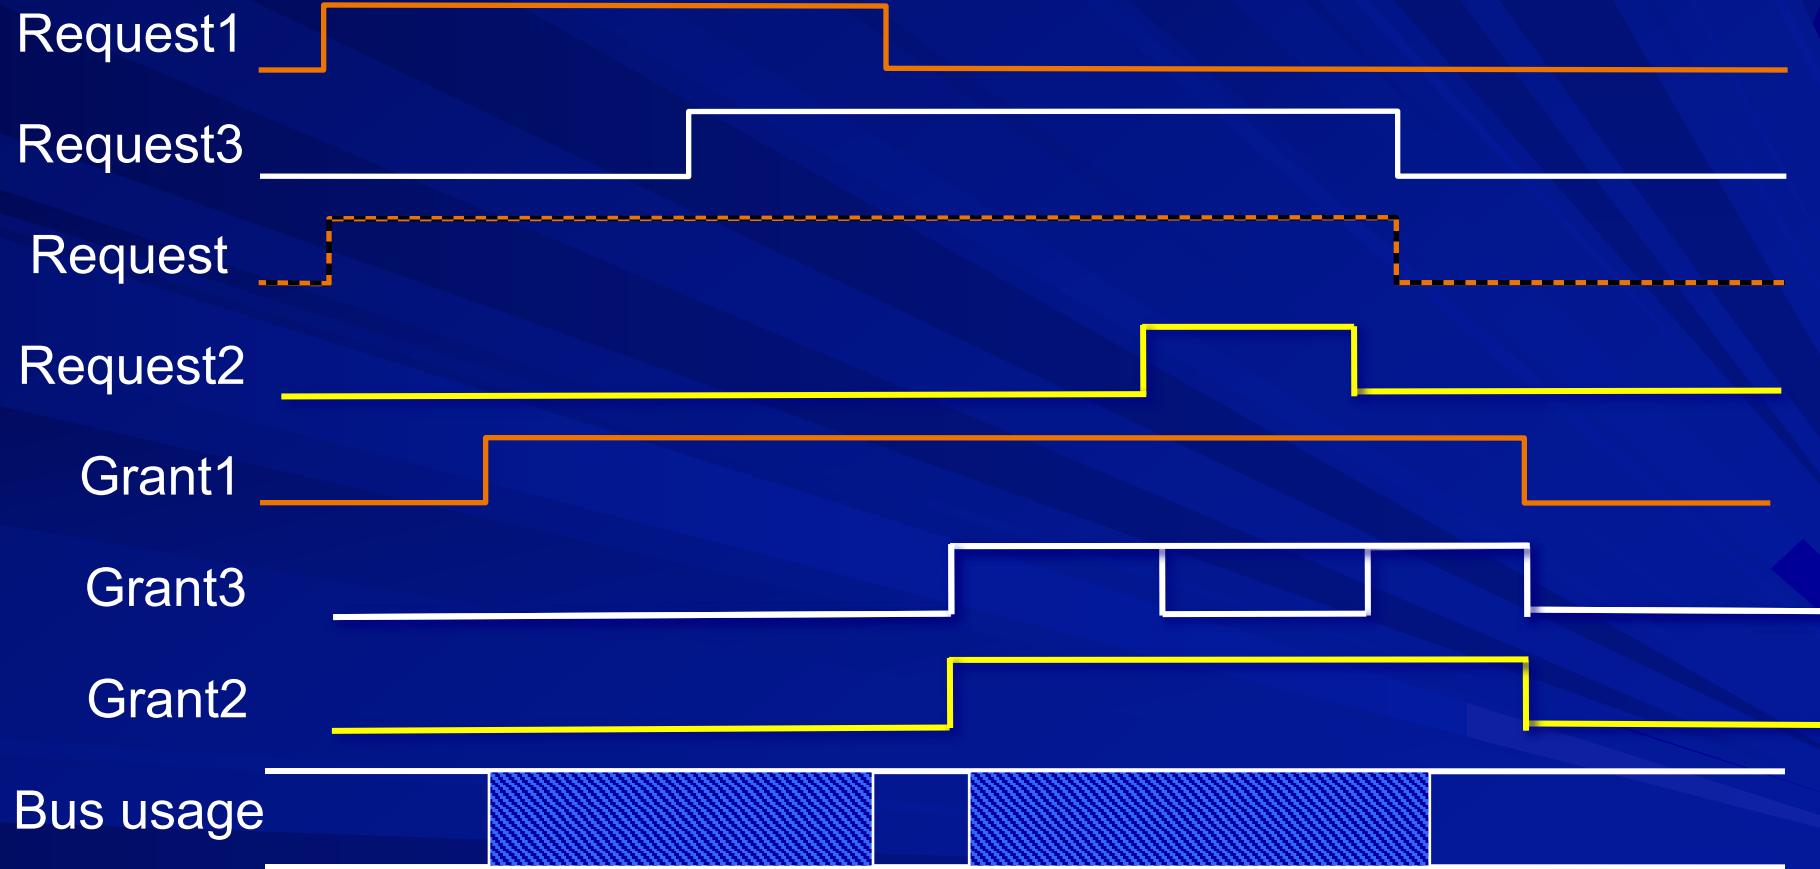


Request and Grant signals



Problem: High priority device can usurp the grant signal

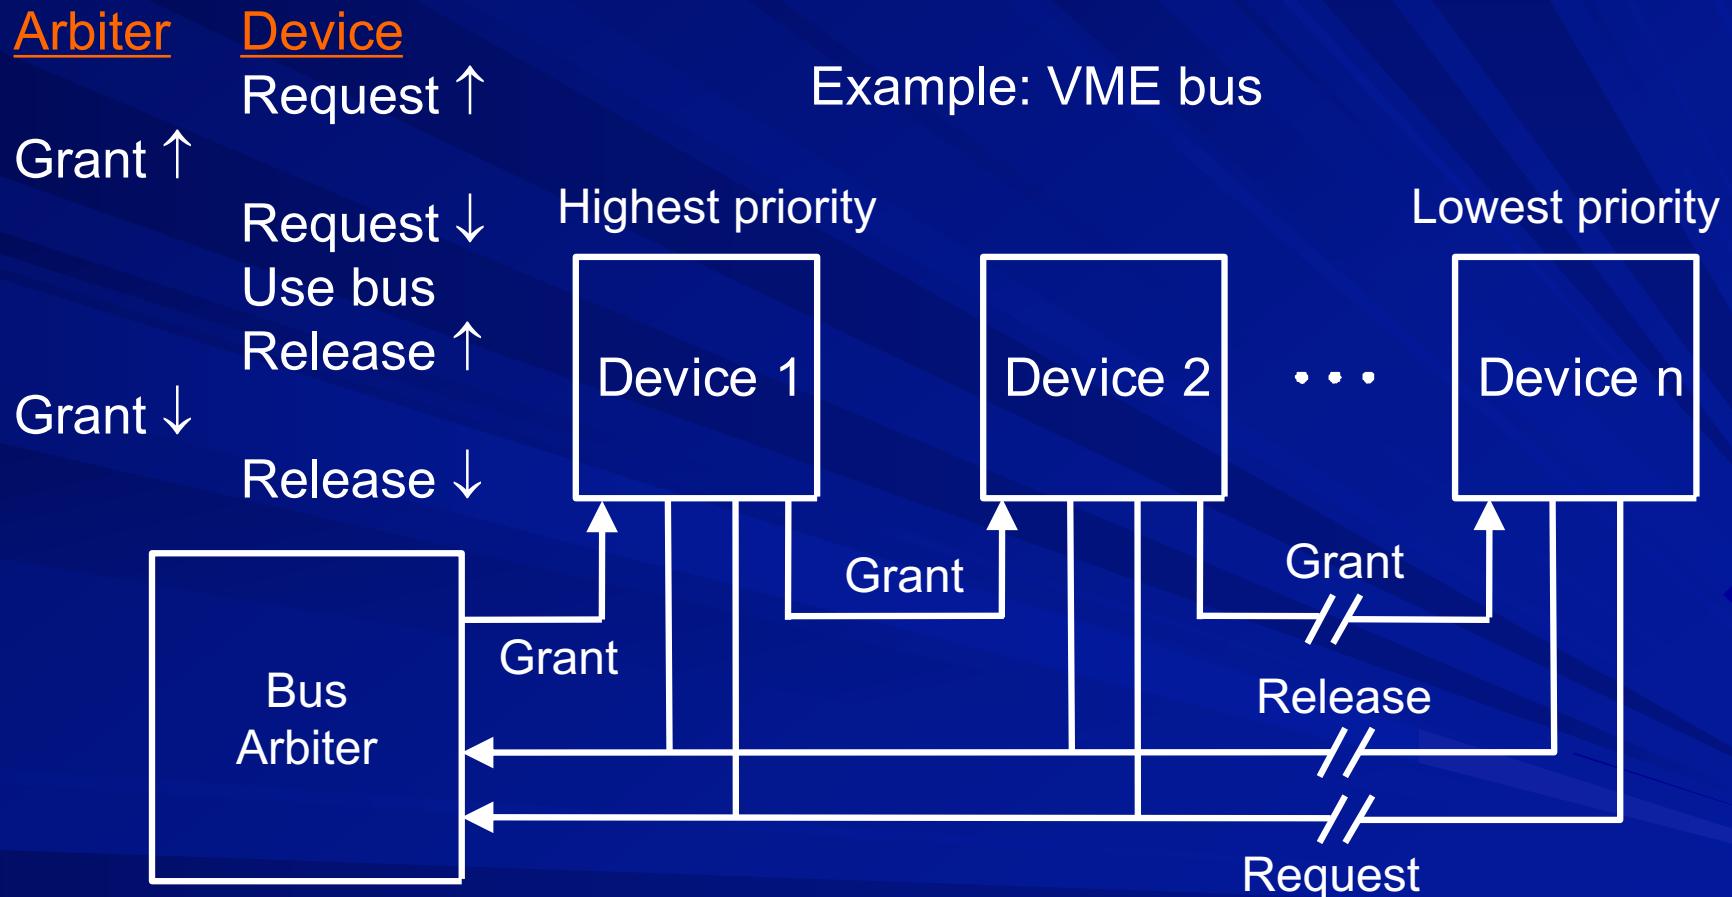
Request and Grant signals



Problem: High priority device can usurp the grant signal
Solution: Daisy chain – request, grant, release

As soon as we see request high, we do grant high and immediately make request low, then use bus and do release high and immediately grant goes down and release goes down after that.

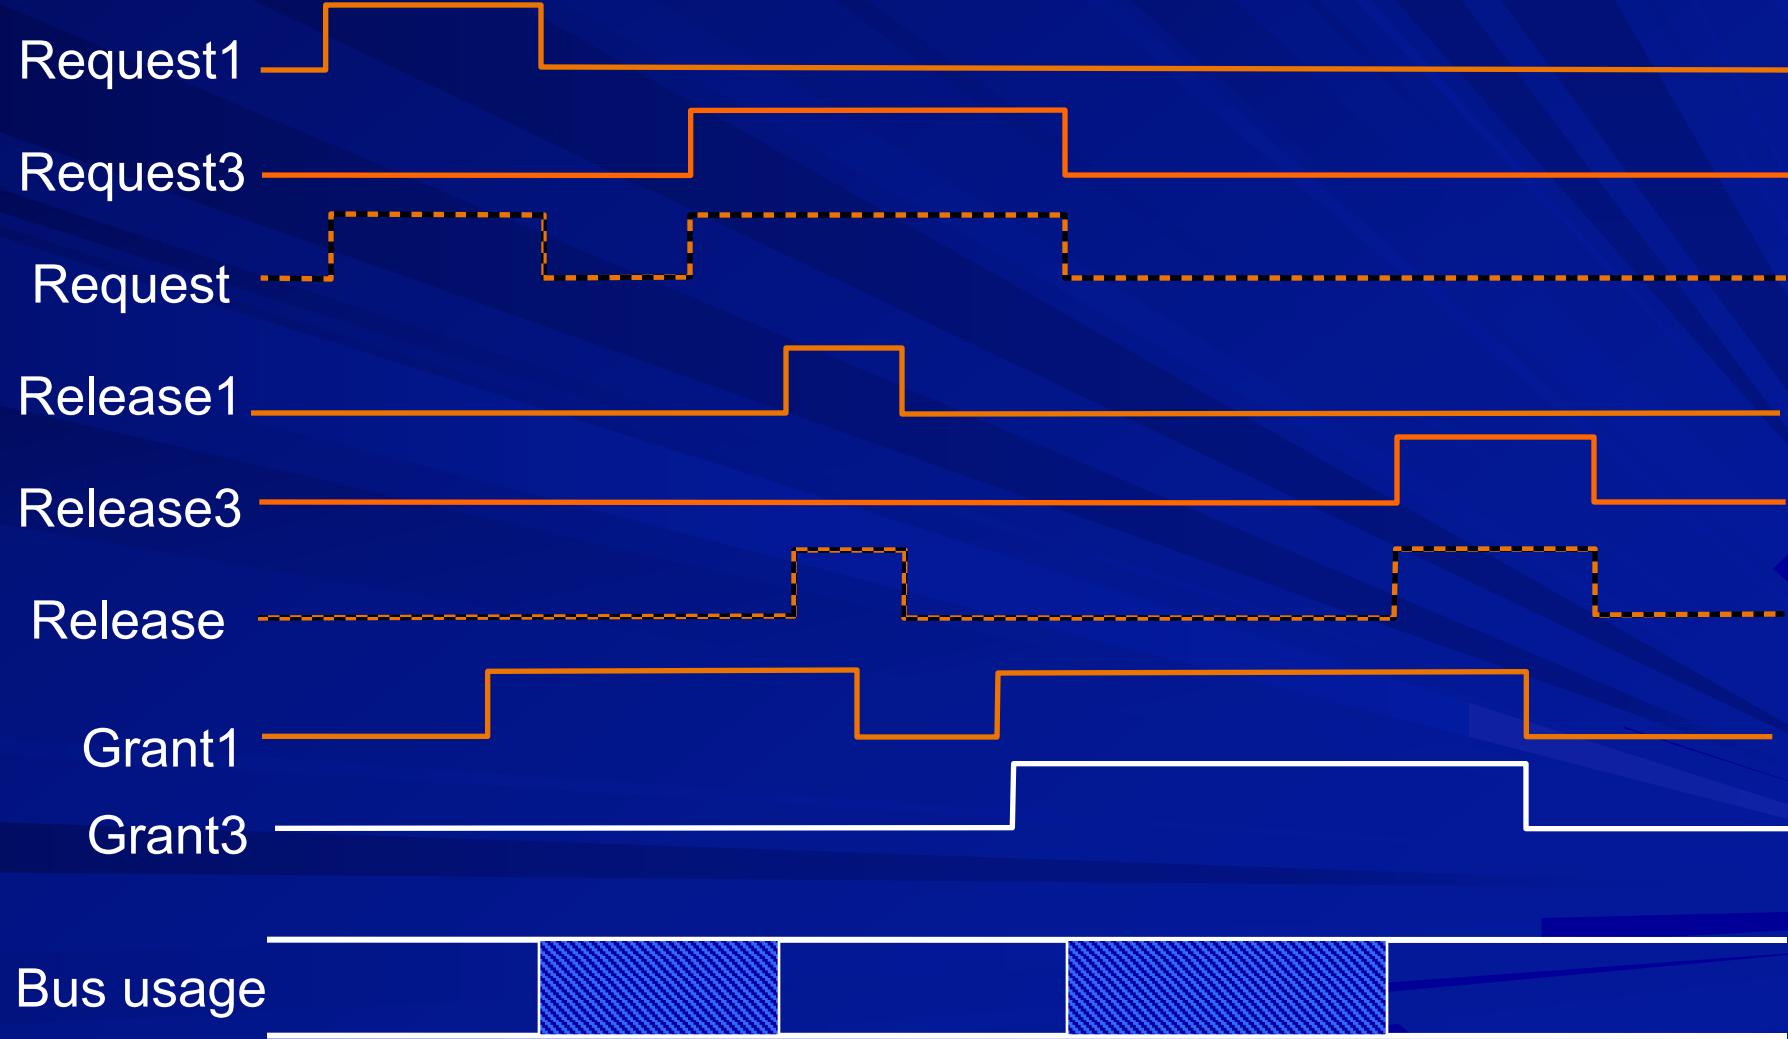
Daisy chain



Note: Rising edge on Grant signal is significant, not level

grant sees request3 is 1 so it immediately goes up and comes down when release3 is 1. bus usage is between grant start and release rising edge

Request, release and grant signals



so if master2 or master1 comes again with a request before master3 had a chance to use it then we will have starving of master3. rising edge of grant can only determine if the master gets to use the bus or not. if master2 comes after master3 has started using the bus, then it has to wait till it sees the grant's rising edge

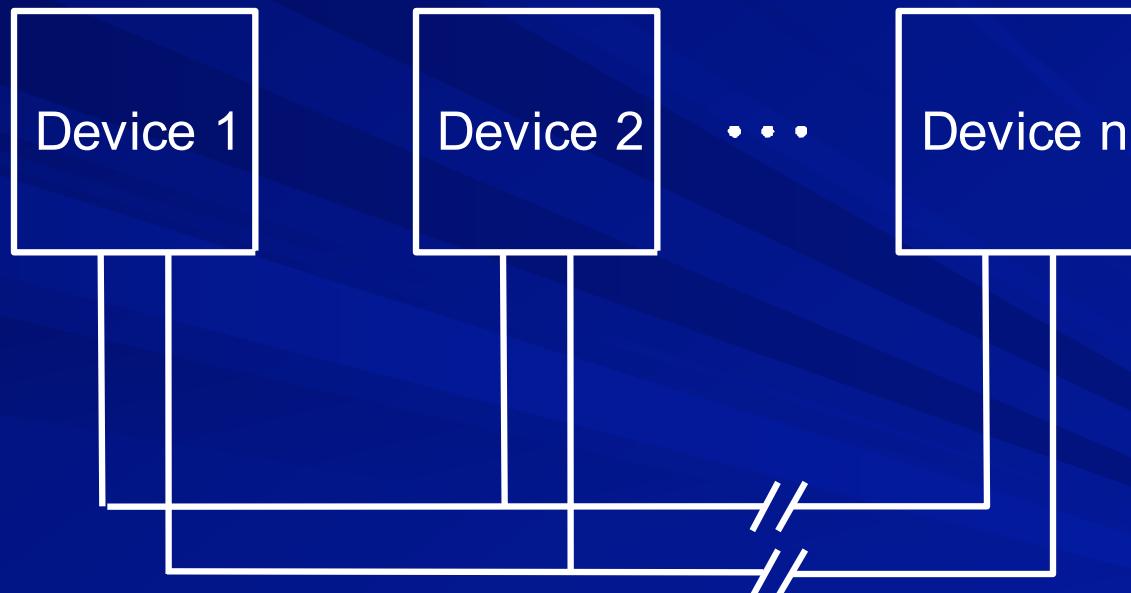
Daisy chain characteristics

- Simple and inexpensive
- Speed is limited due to chaining of grant signal
- Starvation of low priority devices may not be prevented

Rule to improve fairness: A device that has just used the bus cannot reacquire it until it sees the request line go low

Distributed arbitration

Example: NuBus (a backplane bus) in Apple Macintosh



Request lines and identity codes

Distributed arbitration in NuBus

- $\text{req} = \text{req}^0 + \text{req}^1 + \dots + \text{req}^{n-1}$ (wired OR)
- $\text{ID}^i_3 \dots \text{ID}^i_0 = \text{ID}$ of i^{th} requester
- $\text{arb}^i_3 \dots \text{arb}^i_0 = \text{arb}$ output of i^{th} requester
- $\text{arb}_{3..0} = \text{arb}^0_{3..0} + \dots + \text{arb}^{n-1}_{3..0}$ (wired OR)
- $\text{arb}^i_3 = \text{ID}^i_3 \cdot \text{req}^i$ if all are 0, then u can continue, else if u have 1 then u can continue, this logic, basically determines the highest priority device.
- $\text{arb}^i_2 = \text{ID}^i_2 \cdot \text{req}^i \cdot (\text{ID}^i_3 + \text{arb}_3')$
- $\text{arb}^i_1 = \text{ID}^i_1 \cdot \text{req}^i \cdot (\text{ID}^i_3 + \text{arb}_3') \cdot (\text{ID}^i_2 + \text{arb}_2')$
- $\text{arb}^i_0 = \text{ID}^i_0 \cdot \text{req}^i \cdot (\text{ID}^i_3 + \text{arb}_3') \cdot (\text{ID}^i_2 + \text{arb}_2') \cdot (\text{ID}^i_1 + \text{arb}_1')$

we have previous ones since once we have dropped off the race, we should not continue even if we have 1 later in lsb

NuBus arbitration example-1

k	ID^k_3	ID^k_2	ID^k_1	ID^k_0	arb^k_3	arb^k_2	arb^k_1	arb^k_0
1	1	0	0	1				
2	0	1	1	0				
3	1	1	0	0				
4	0	0	1	1				

NuBus arbitration example-1

k	ID^k_3	ID^k_2	ID^k_1	ID^k_0	arb^k_3	arb^k_2	arb^k_1	arb^k_0
1	1	0	0	1	1	0	0	1
2	0	1	1	0	0	1	1	0
3	1	1	0	0	1	1	0	0
4	0	0	1	1	0	0	1	1
<hr/>								
OR					1	1	1	1

NuBus arbitration example-1

k	ID^k_3	ID^k_2	ID^k_1	ID^k_0	arb^k_3	arb^k_2	arb^k_1	arb^k_0
1	1	0	0	1	1			
2	0	1	1	0	0			
3	1	1	0	0	1			
4	0	0	1	1	0			

NuBus arbitration example-1

k	ID^k_3	ID^k_2	ID^k_1	ID^k_0	arb^k_3	arb^k_2	arb^k_1	arb^k_0
1	1	0	0	1	1	0		
2	0	1	1	0	0	x		
3	1	1	0	0	1	1		
4	0	0	1	1	0	x		
					1	1		

NuBus arbitration example-1

k	ID^k_3	ID^k_2	ID^k_1	ID^k_0	arb^k_3	arb^k_2	arb^k_1	arb^k_0
1	1	0	0	1	1	0	x	
2	0	1	1	0	0	x	x	
3	1	1	0	0	1	1	0	
4	0	0	1	1	0	x	x	
					1	1	0	

NuBus arbitration example-1

k	ID^k_3	ID^k_2	ID^k_1	ID^k_0	arb^k_3	arb^k_2	arb^k_1	arb^k_0
1	1	0	0	1	1	0	x	x
2	0	1	1	0	0	x	x	x
3	1	1	0	0	1	1	0	0
4	0	0	1	1	0	x	x	x
<hr/>								
					1	1	0	0

NuBus arbitration example-2

k	ID^k_3	ID^k_2	ID^k_1	ID^k_0	arb^k_3	arb^k_2	arb^k_1	arb^k_0
1	1	0	0	1				
2	0	1	1	0				
3	1	0	1	0				
4	0	0	1	1				

NuBus arbitration example-2

k	ID^k_3	ID^k_2	ID^k_1	ID^k_0	arb^k_3	arb^k_2	arb^k_1	arb^k_0
1	1	0	0	1	1	0	0	1
2	0	1	1	0	0	1	1	0
3	1	0	1	0	1	0	1	0
4	0	0	1	1	0	0	1	1
<hr/>								
OR					1	1	1	1

NuBus arbitration example-2

k	ID^k_3	ID^k_2	ID^k_1	ID^k_0	arb^k_3	arb^k_2	arb^k_1	arb^k_0
1	1	0	0	1	1			
2	0	1	1	0	0			
3	1	0	1	0	1			
4	0	0	1	1	0			

NuBus arbitration example-2

k	ID^k_3	ID^k_2	ID^k_1	ID^k_0	arb^k_3	arb^k_2	arb^k_1	arb^k_0
1	1	0	0	1	1	0		
2	0	1	1	0	0	x		
3	1	0	1	0	1	0		
4	0	0	1	1	0	x		
					1	0		

NuBus arbitration example-2

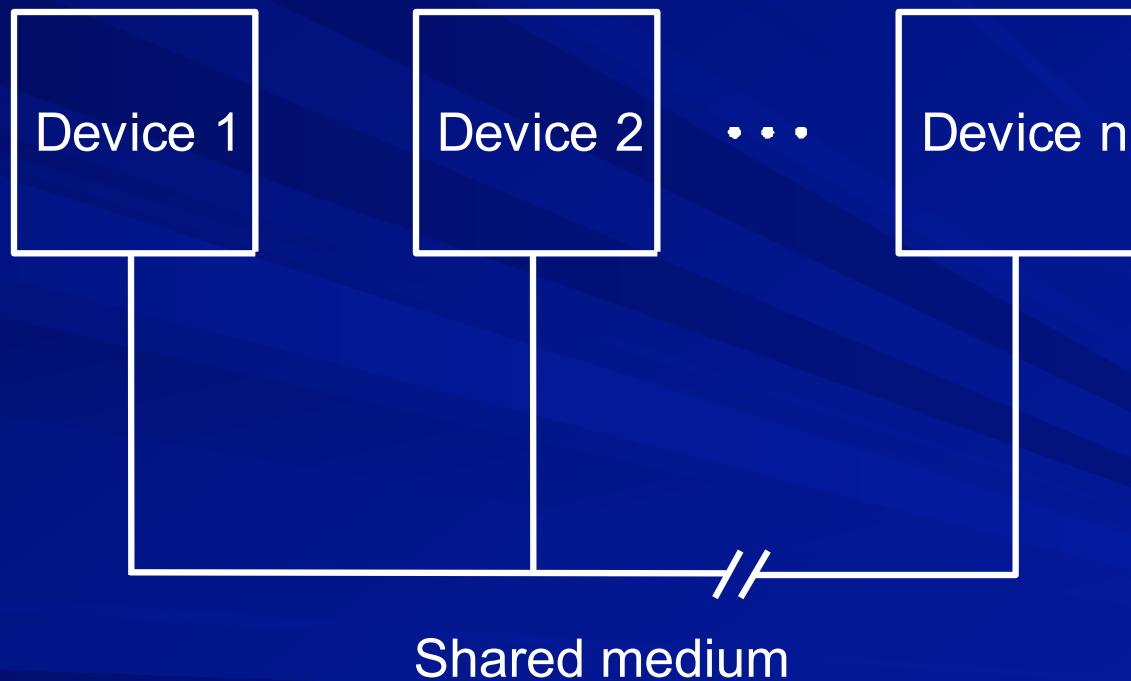
k	ID^k_3	ID^k_2	ID^k_1	ID^k_0	arb^k_3	arb^k_2	arb^k_1	arb^k_0
1	1	0	0	1	1	0	0	
2	0	1	1	0	0	x	x	
3	1	0	1	0	1	0	1	
4	0	0	1	1	0	x	x	
					1	0	1	

NuBus arbitration example-2

k	ID^k_3	ID^k_2	ID^k_1	ID^k_0	arb^k_3	arb^k_2	arb^k_1	arb^k_0
1	1	0	0	1	1	0	0	x
2	0	1	1	0	0	x	x	x
3	1	0	1	0	1	0	1	0
4	0	0	1	1	0	x	x	x
<hr/>								
					1	0	1	0

Arbitration by collision detection

Example: Ethernet

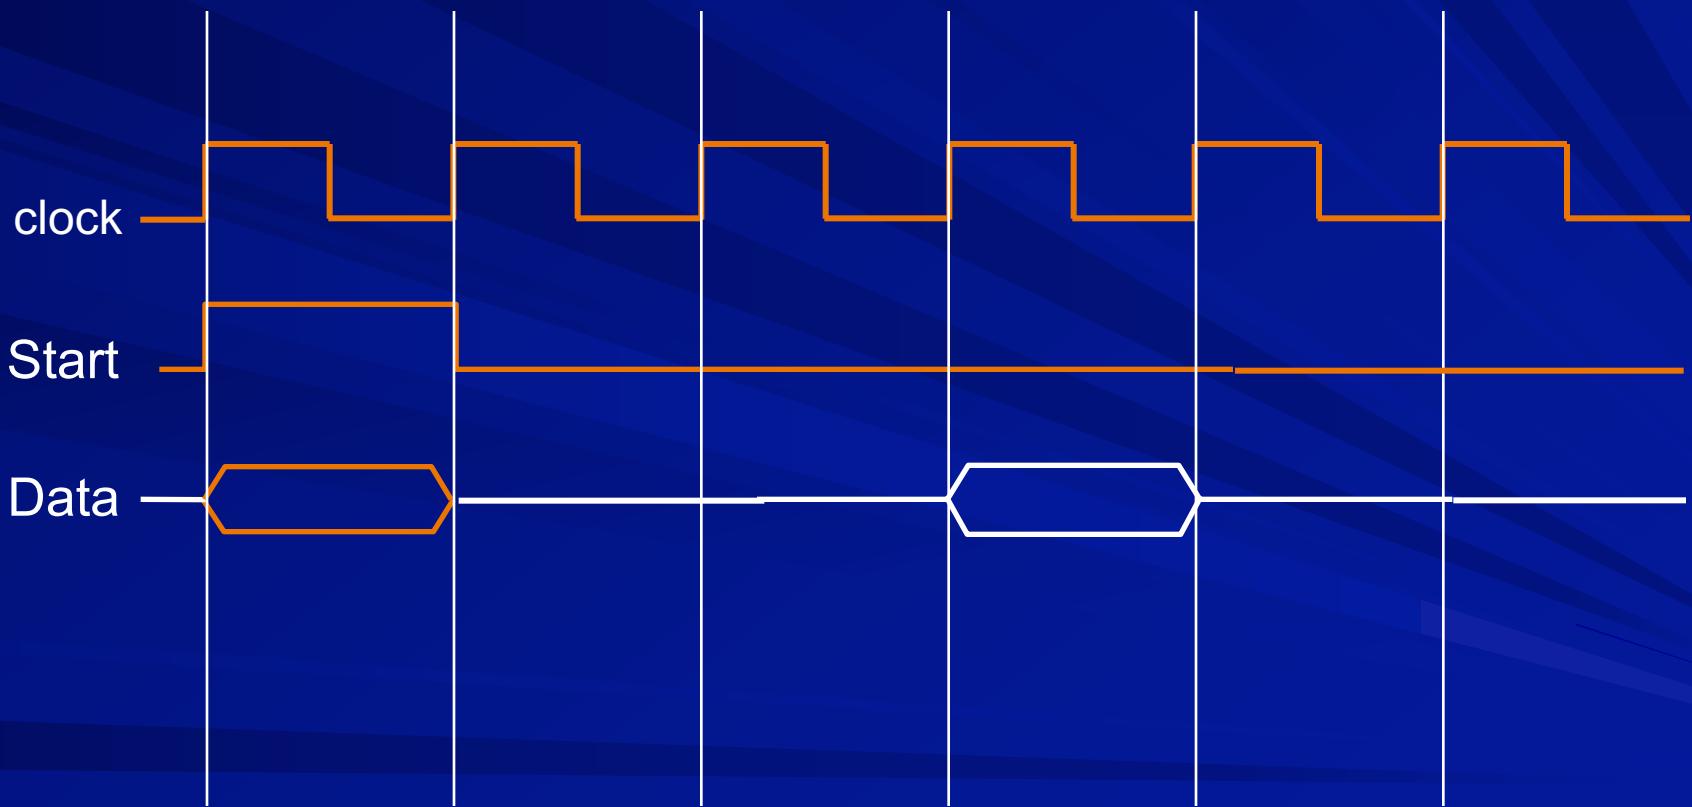


if collision then try again after random interval

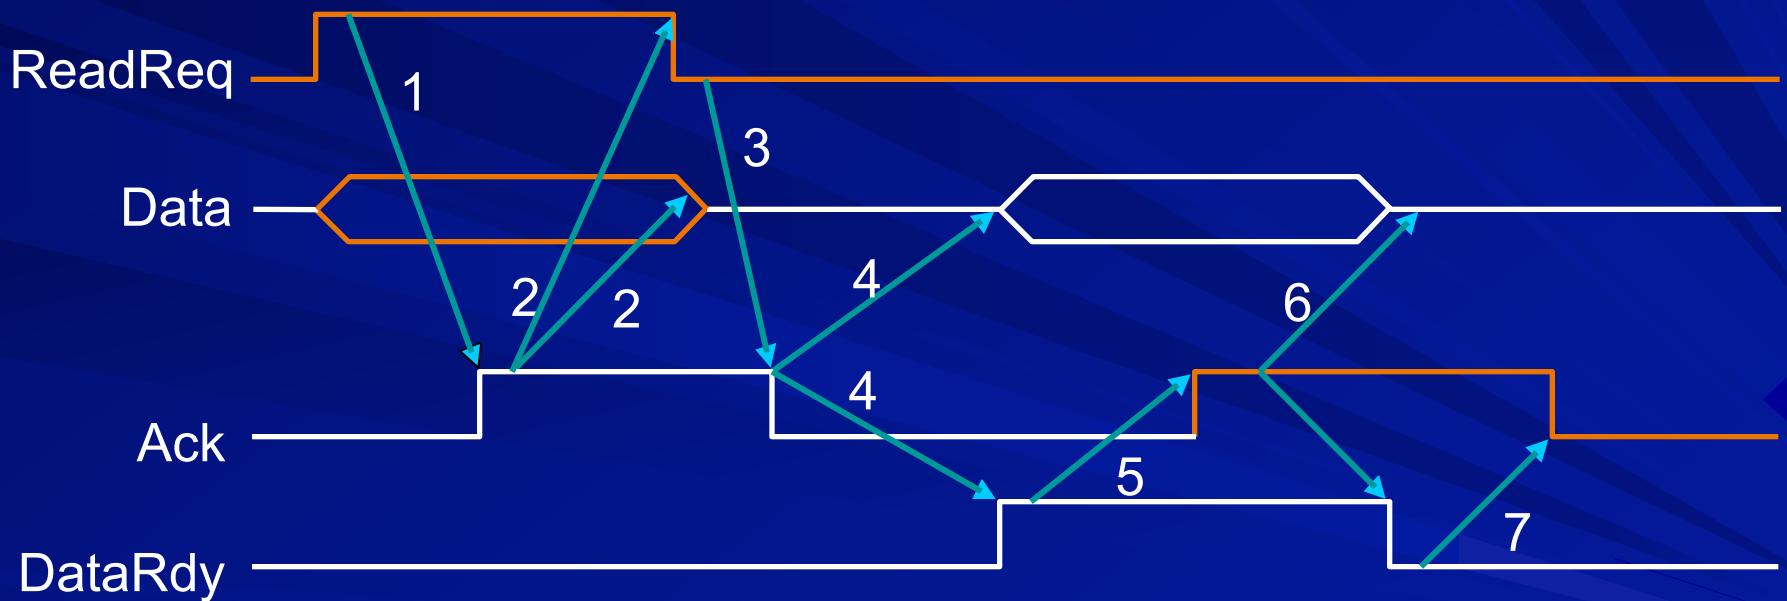
Increasing the bus bandwidth

- increase bus width
- separate data and address lines
- pipelining
- multiple word blocks
- synchronous
- split transaction

Synchronous transfer



Asynchronous handshaking



THANKS

COL216

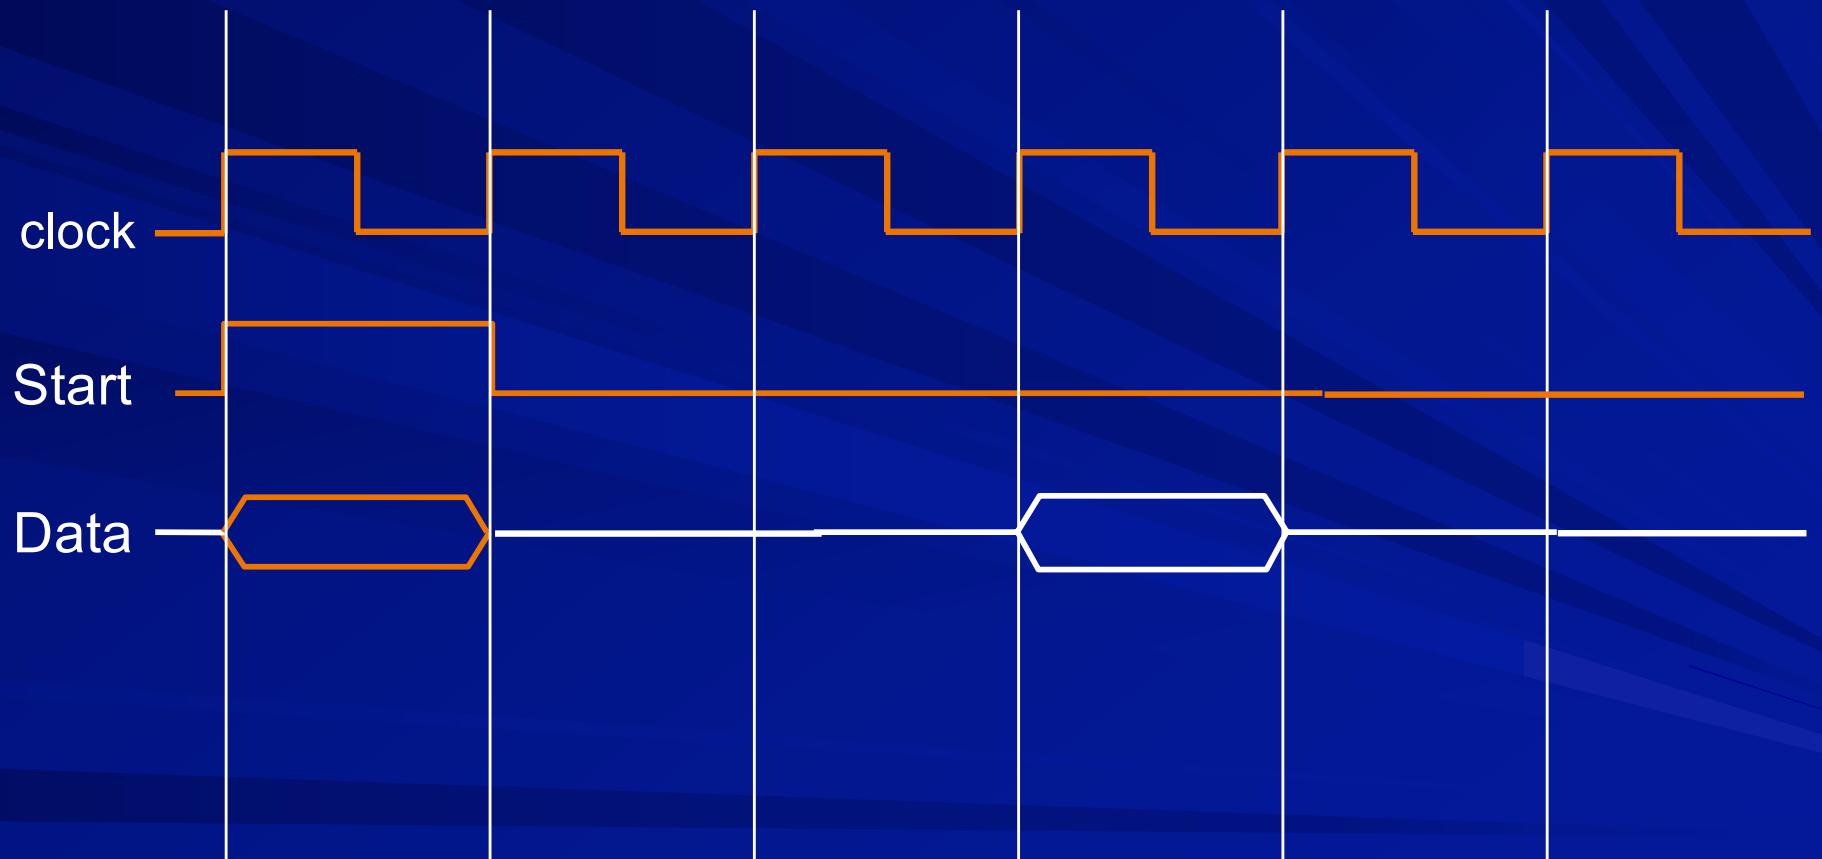
Computer Architecture

Input/Output – 5
Parallel and Serial Buses
28th March 2022

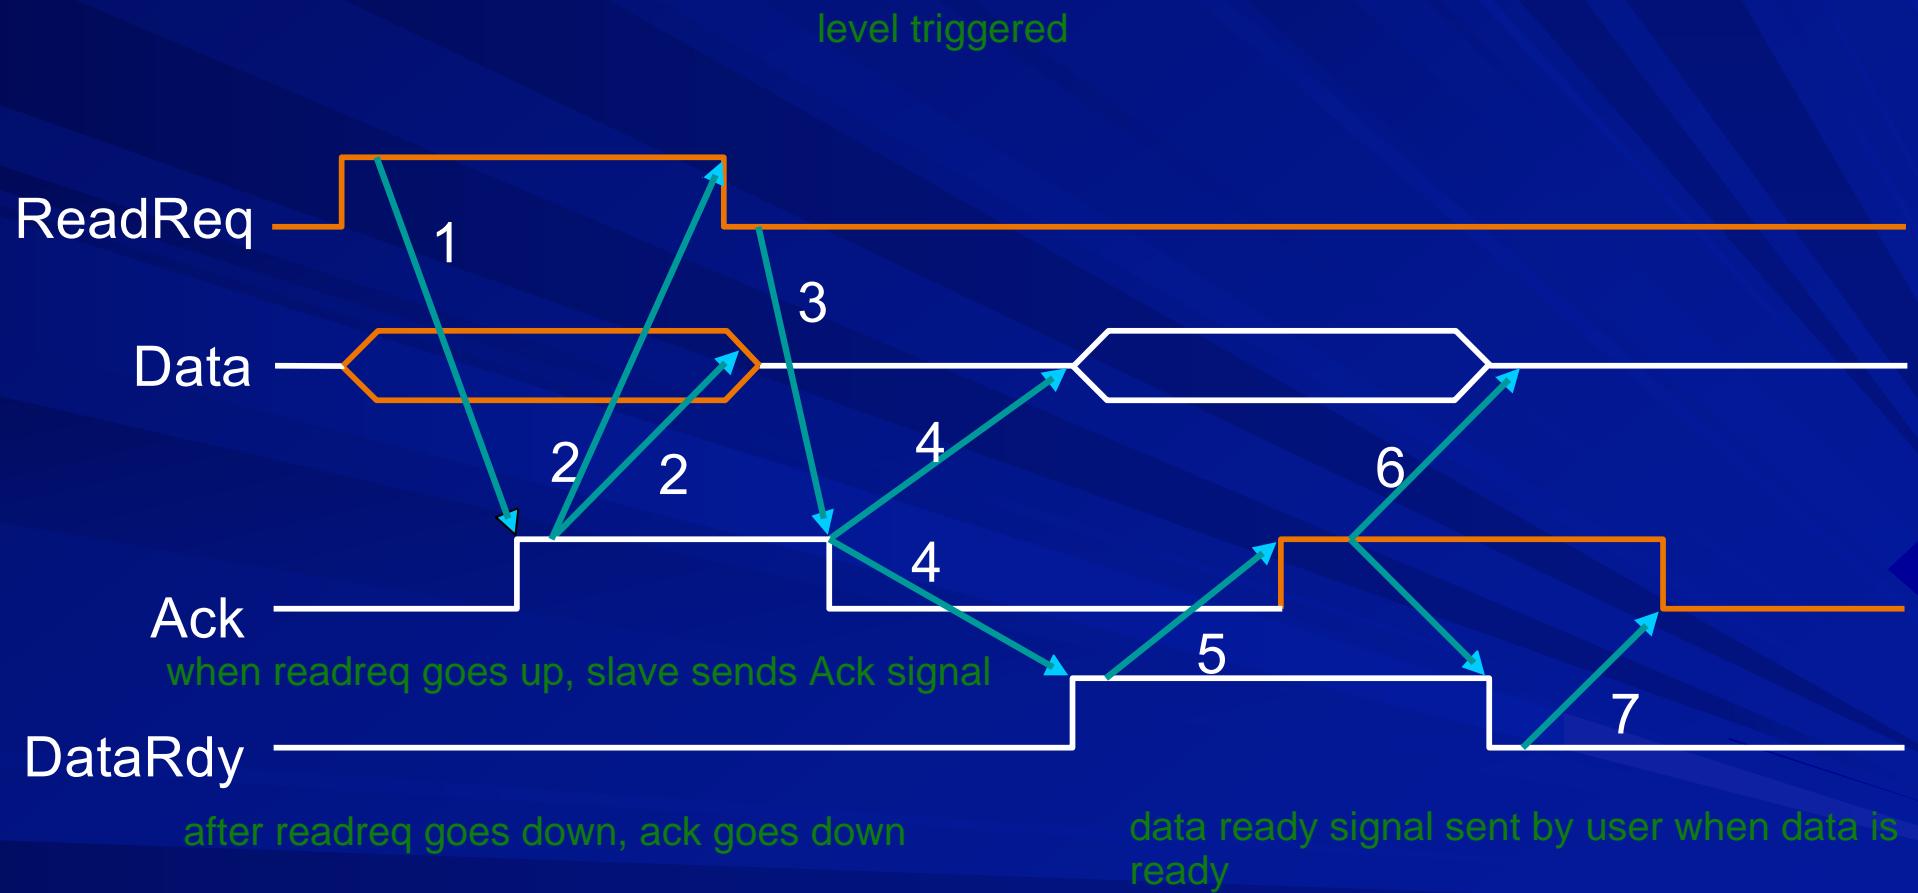
Increasing the bus bandwidth

- increase bus width
- separate data and address lines
- pipelining
- multiple word blocks
- synchronous
- split transaction

Synchronous transfer

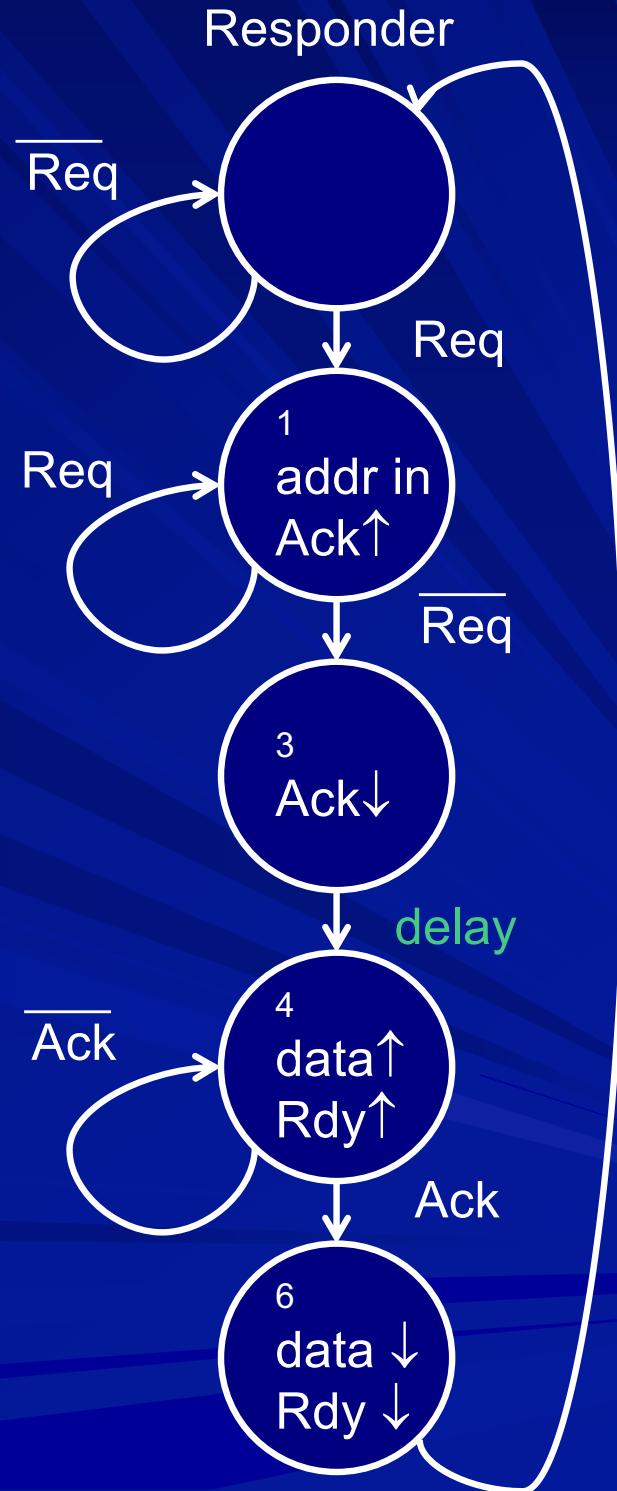
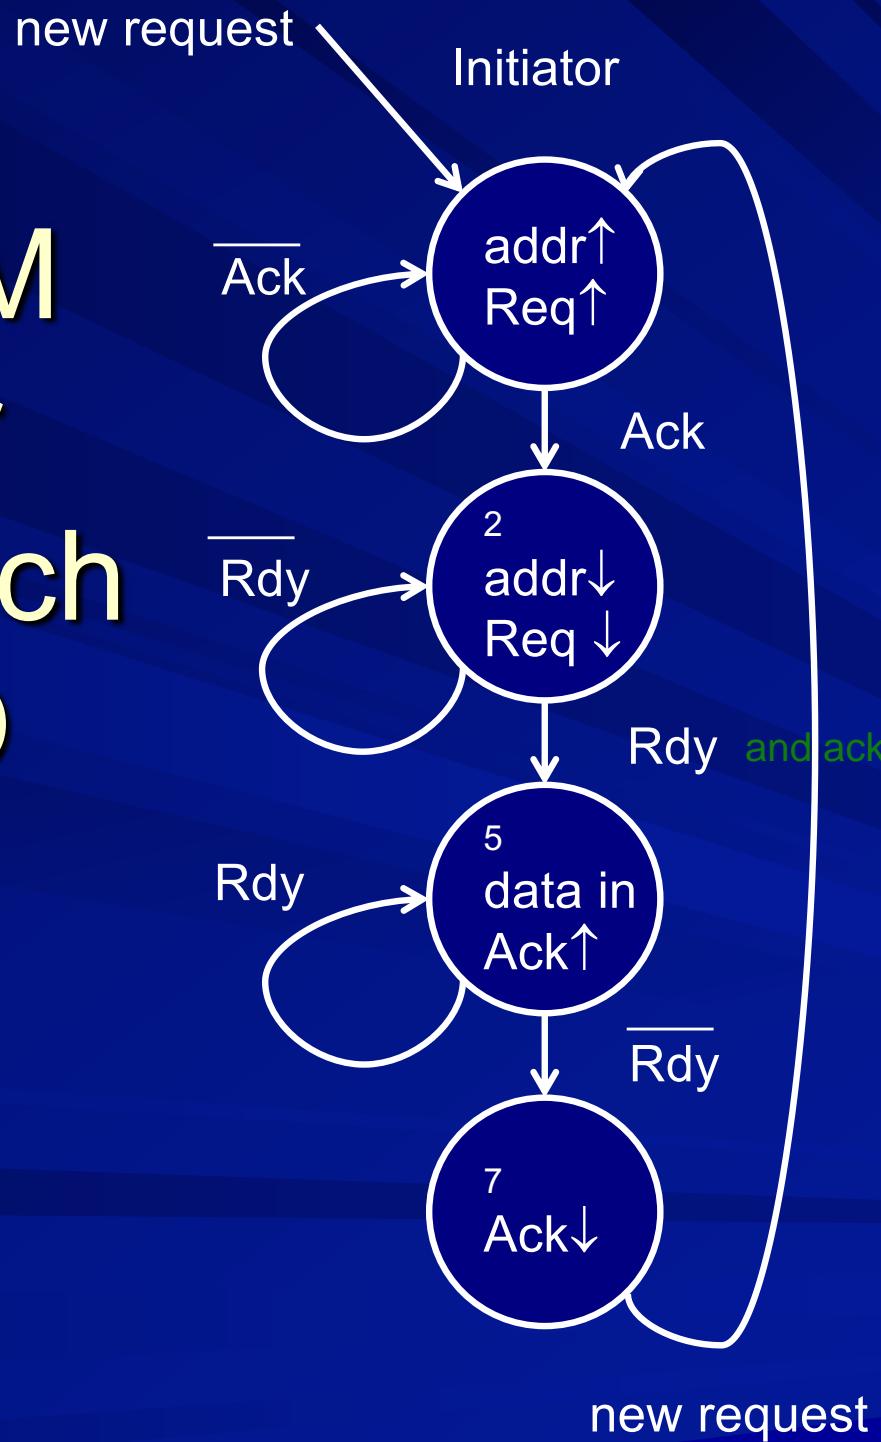


Asynchronous handshaking



FSM for asynch I/O

not clocked



Split transactions



in between, we could release the bus and then a different transaction could begin

Parallel vs Serial

■ Parallel

- multiple wires carry bits in parallel : 8-bit, 16-bit, 32-bit, 64-bit . . .
- address, read data, write data may use same wires or separate wires

■ Serial

- only one bit at a time
- read data, write data may use same wires or separate wires, no separate address wires

Series/parallel bus examples

Parallel

- ISA (Industry Standard Architecture)
- EISA (Extended ISA)
- VLB (VESA (Video Electronics Standards Association) Local Bus)
- PCI (Peripheral Component Interconnect)
- AGP (Accelerated Graphics Port)

Serial

- USB (Universal Serial Bus)
- Fire wire
- Fibre channel
- PCIe (Peripheral Component Interconnect express)
- SATA (Serial Advanced Technology Attachment)

Why serial buses?

- Serial buses are less expensive
- At high speed, keeping all the bits synchronized in parallel buses is very difficult
- Longer the signals travel, more skew is likely among the bits of parallel buses

Move from parallel to serial

PCI Bus

- PCI 1.0: 33MHz, 32-bit, peak BW 133MB/s
- PCI 2.2: 66 MHz, 64-bit, BW 533 MB/s
- PCI-X: 133MHz, 64-bit, BW 1066 MB/s
- PCI-X 2.0: 266MHz, BW 2133 MB/s

PCI-Express (serial bus), 1-16 lanes

- PCIe 1.0: 250 MB/s per lane
- PCIe 2.0: 500 MB/s per lane
- PCIe 3.0: 985 MB/s per lane

Move from parallel to serial

Parallel ATA

33 MB/s =>

66 MB/s =>

100 MB/s =>

133 MB/s

Serial ATA

150 MB/s =>

300 MB/s =>

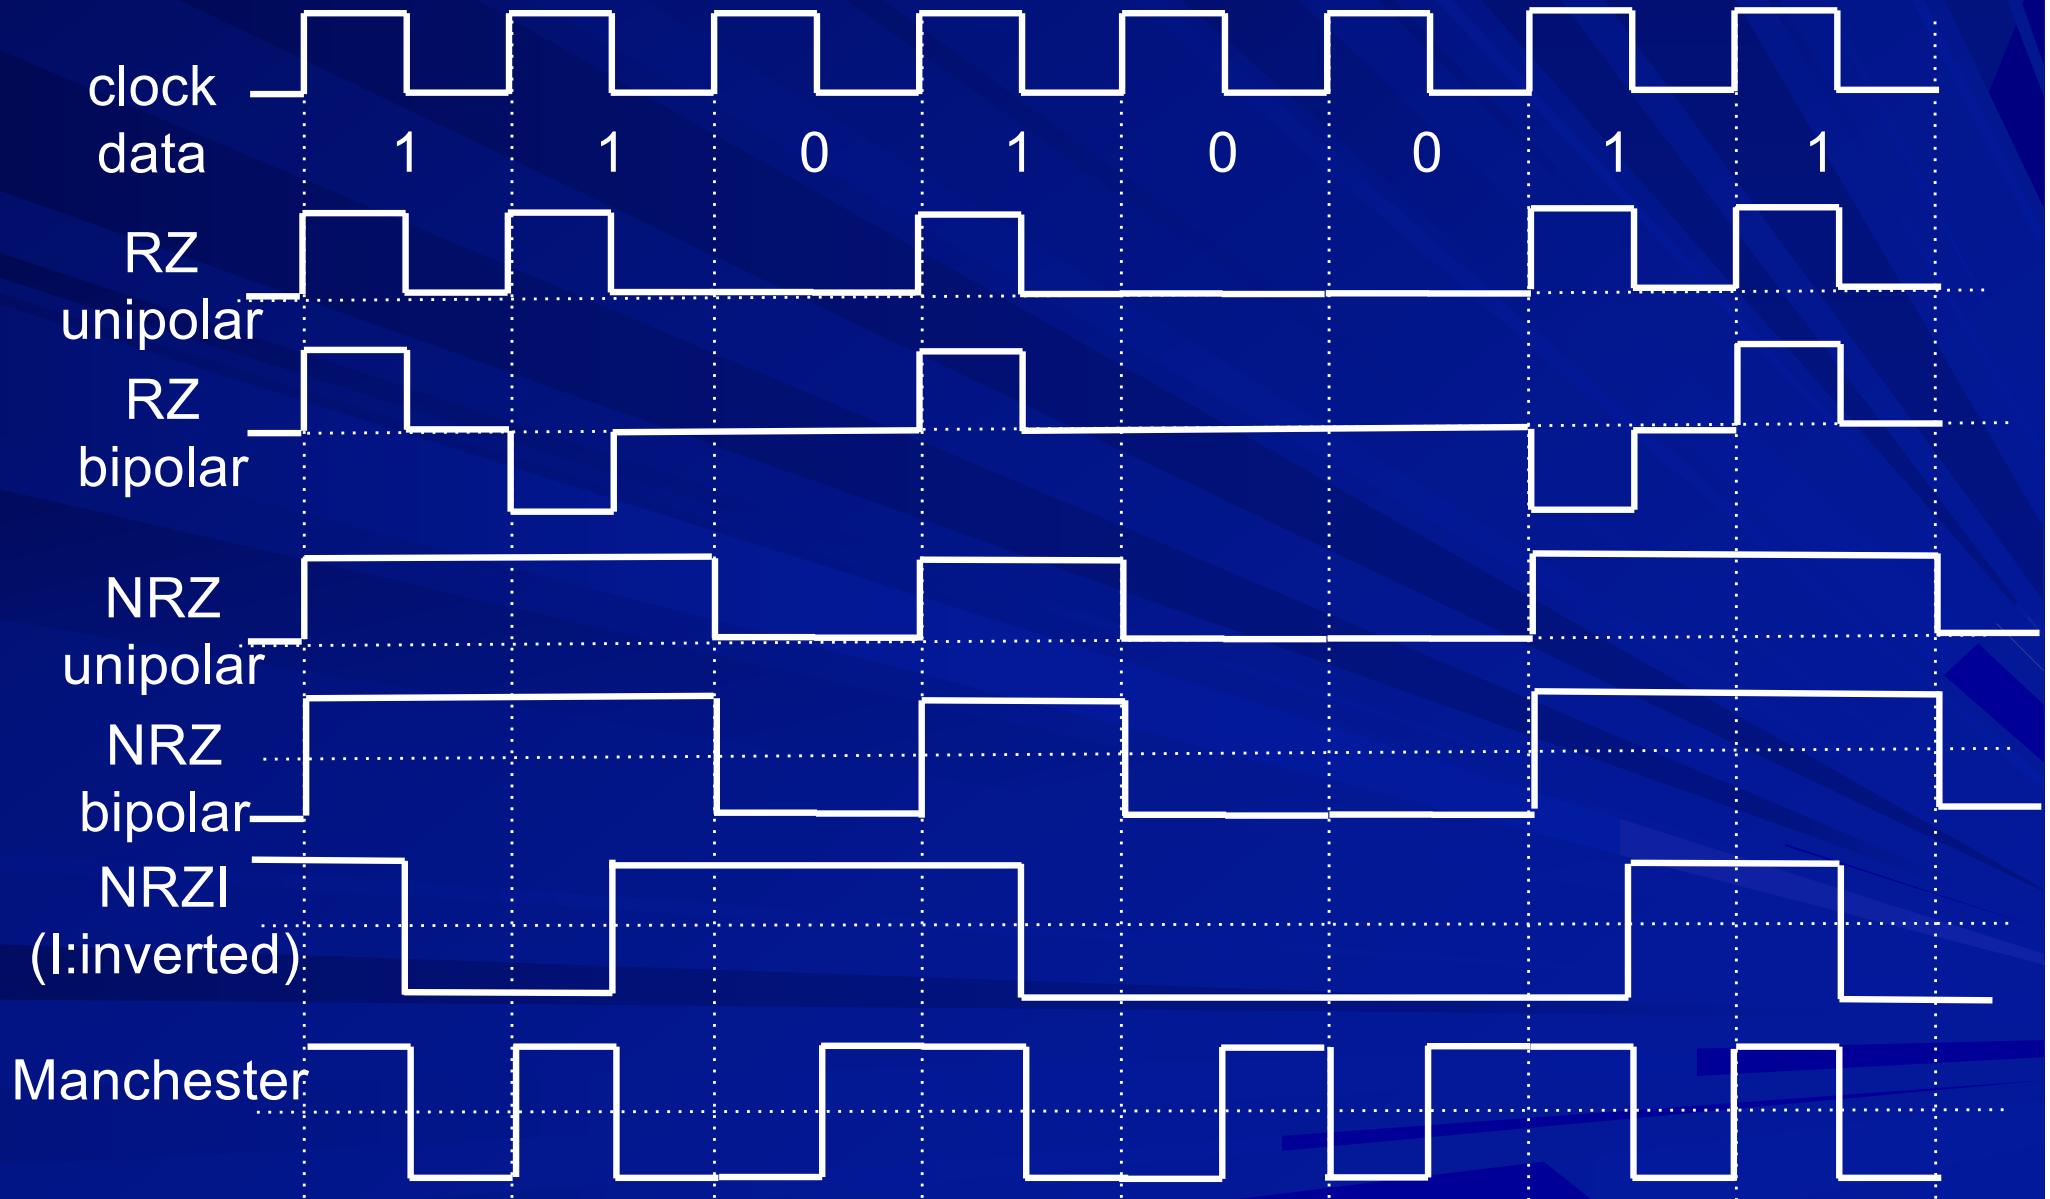
600 MB/s =>

1969 MB/s

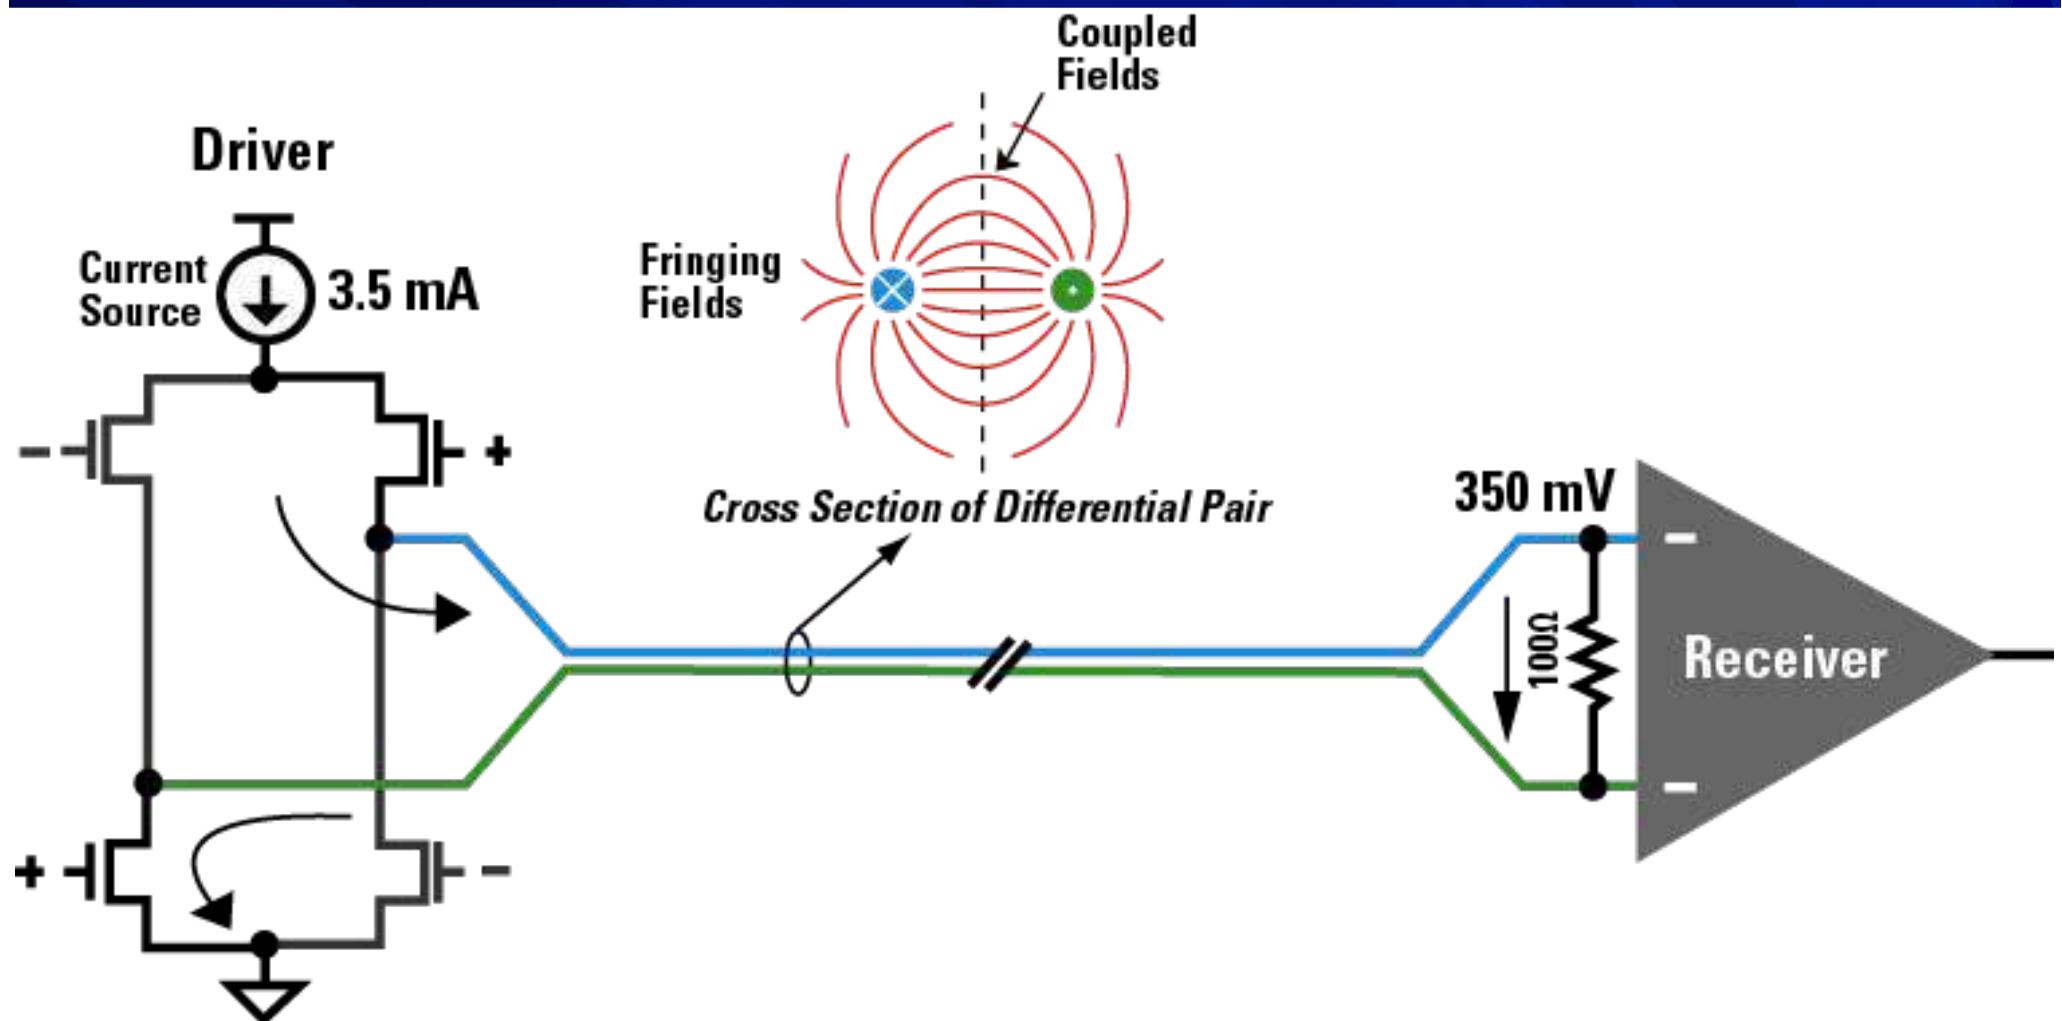
Representing bits

- Unipolar / bipolar
- Active high / active low
- Return to zero (RZ) / Non return to zero (NRZ)
- Encode using levels or transitions
- One wire / two wire (differential)

Representing bits



Low voltage differential signalling (LVDS)

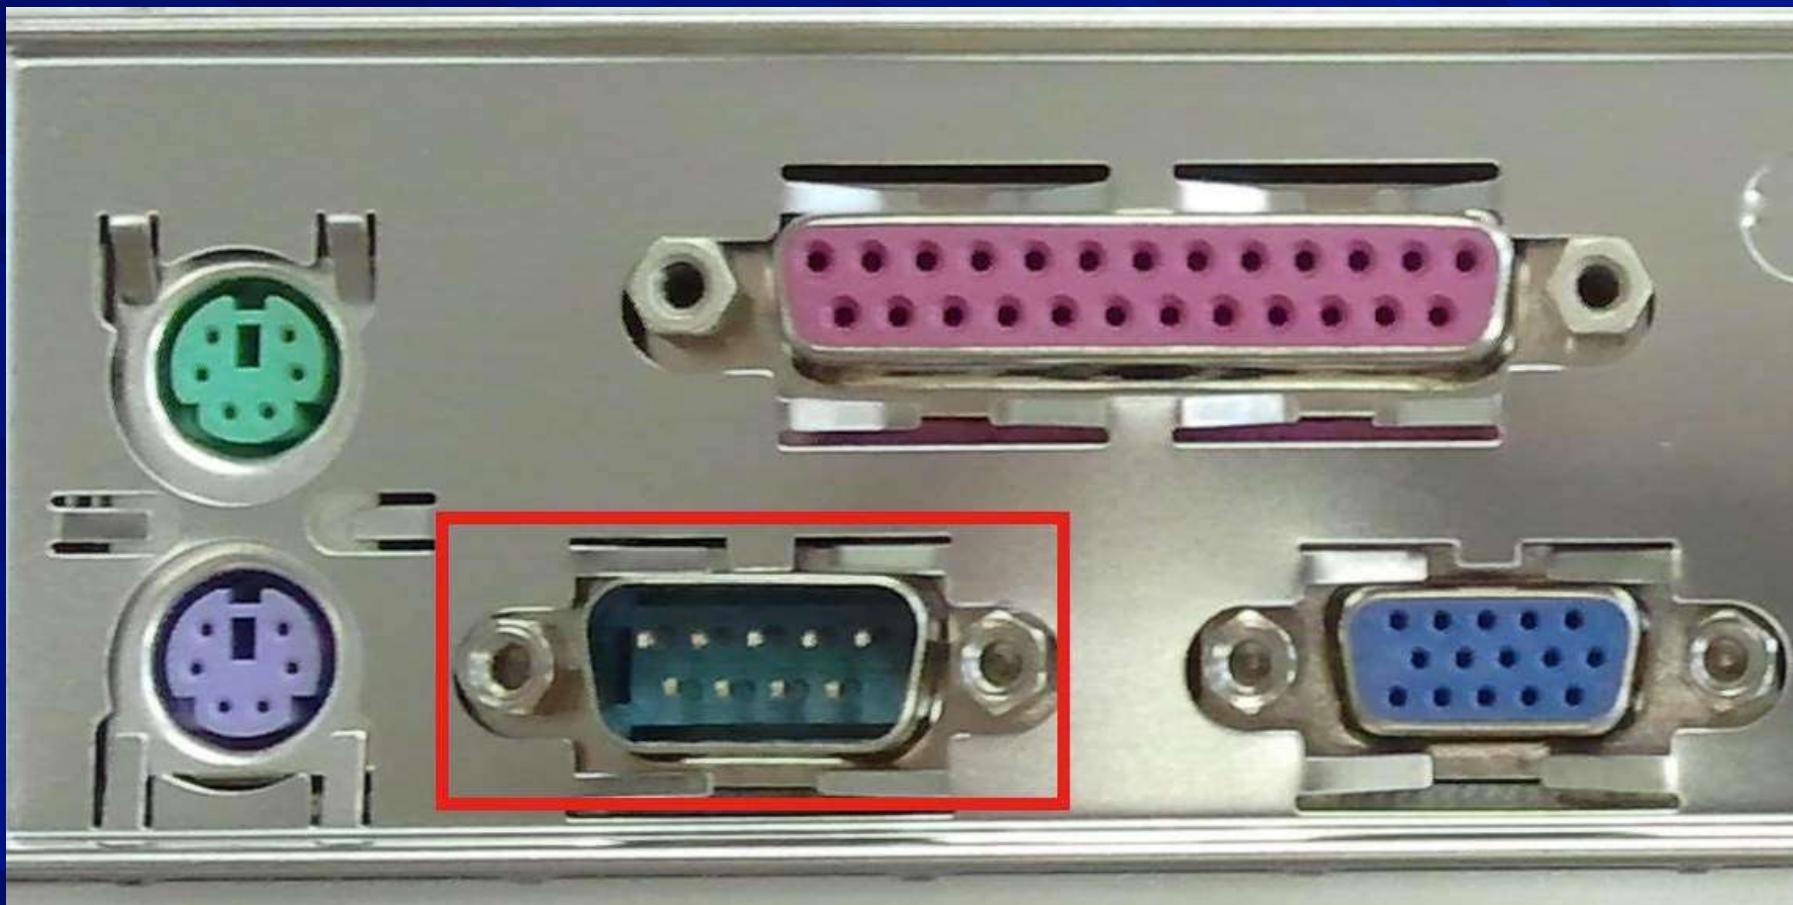


from wikipedia

Timing in serial transfer

- Synchronous
 - shared clock
 - same frequency and phase
- Mesochronous
 - transfer clock signal, separately or with data
 - same frequency but not phase
- Plesiochronous
 - locally generated clock
 - nearly same frequency, changing phase

Serial port / COM port



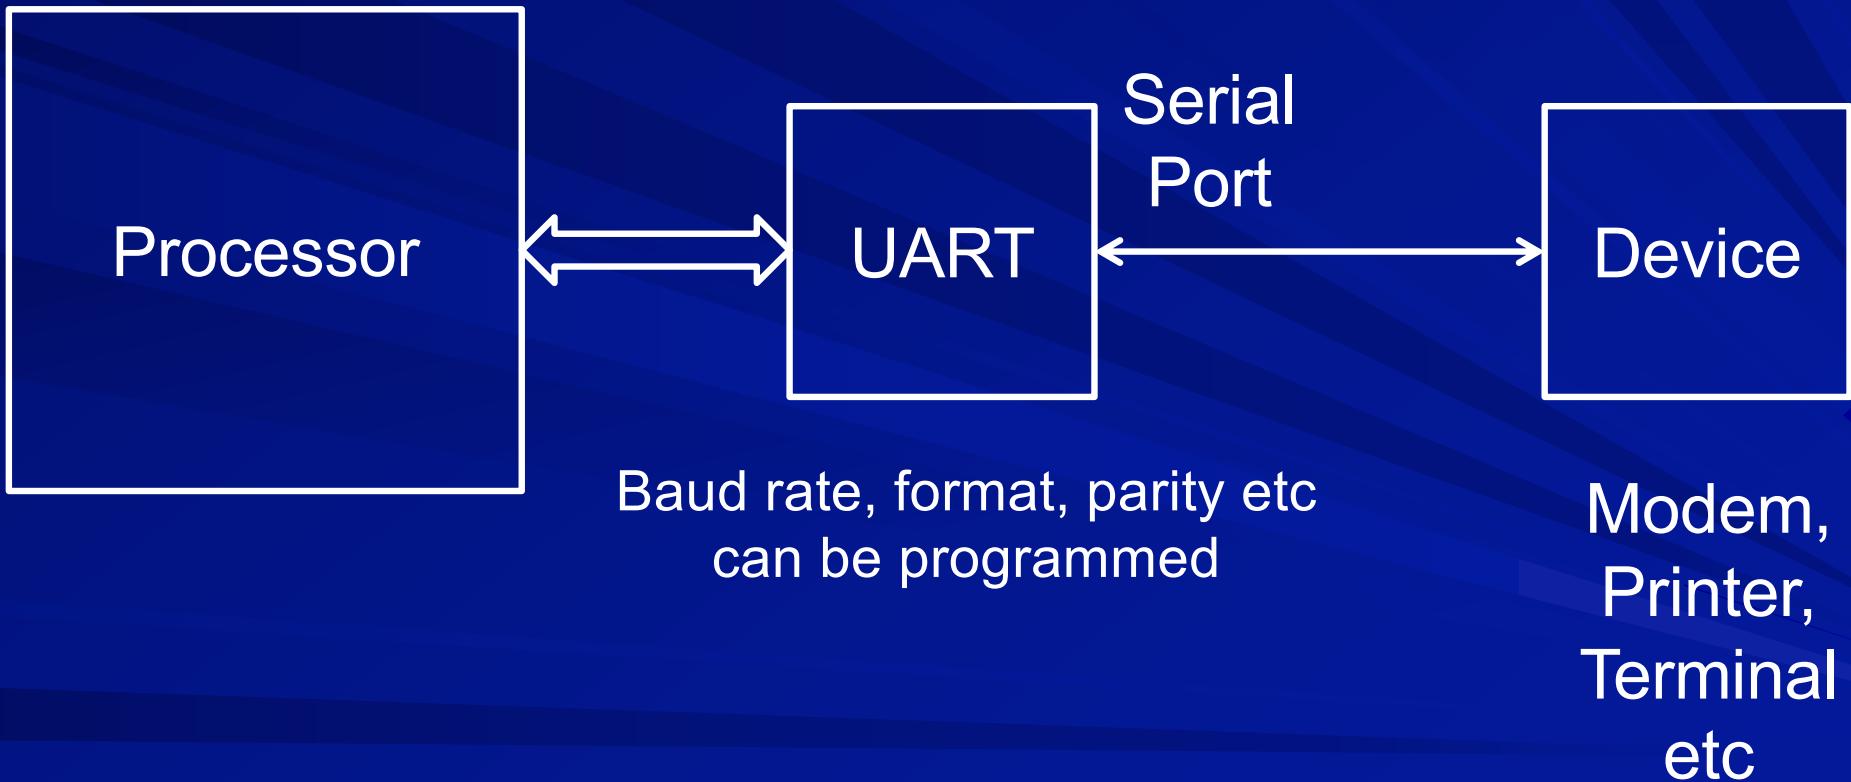
RS232 has been a common standard
75 to 115.2 K bits/s

RS232 Interface Standard



UART

(Universal Asynchronous
Receiver Transmitter)



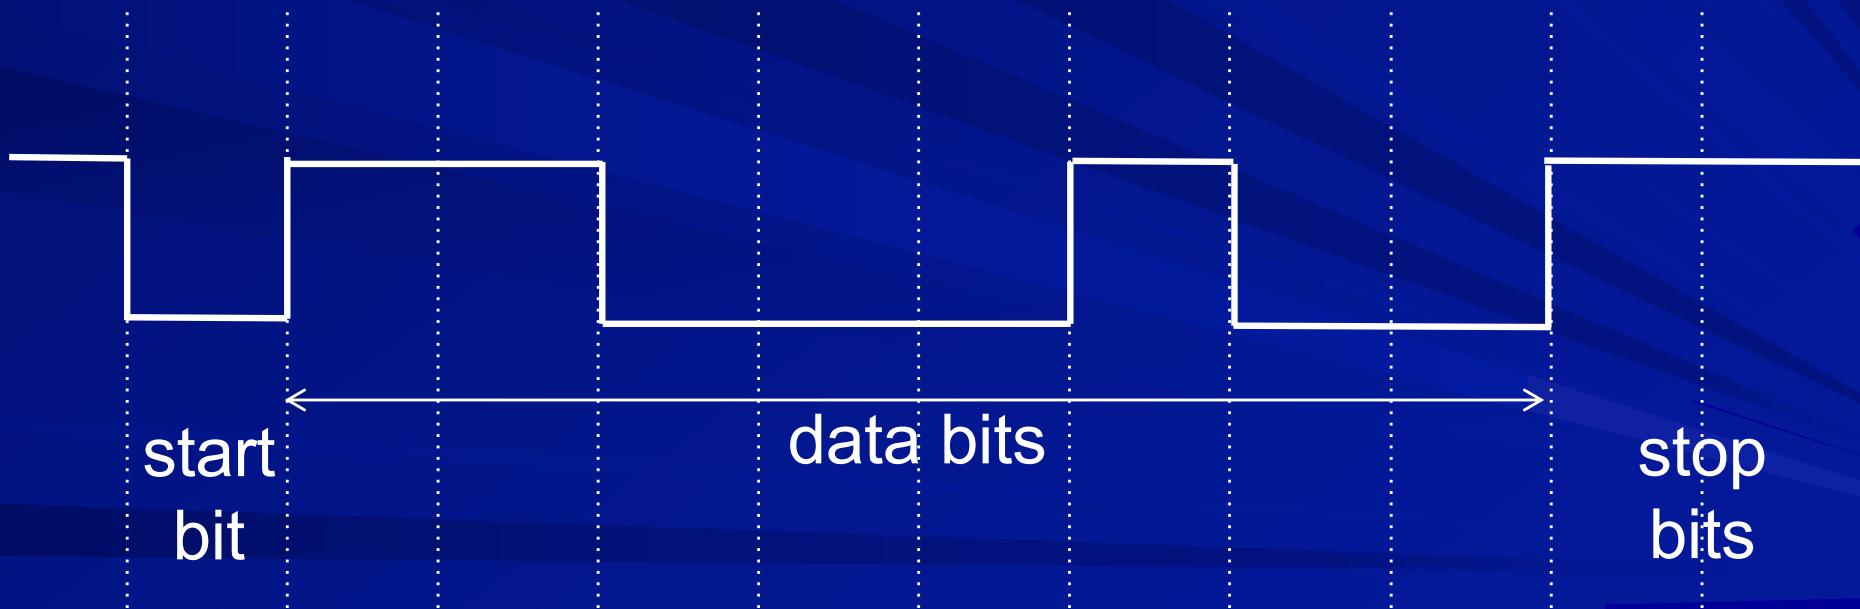
UART



Serial port: data waveform

- Plesiochronous

- Synchronous at bit level
- Asynchronous at frame level



Data Framing



Data bits : 5 to 8

Parity : even / odd / none

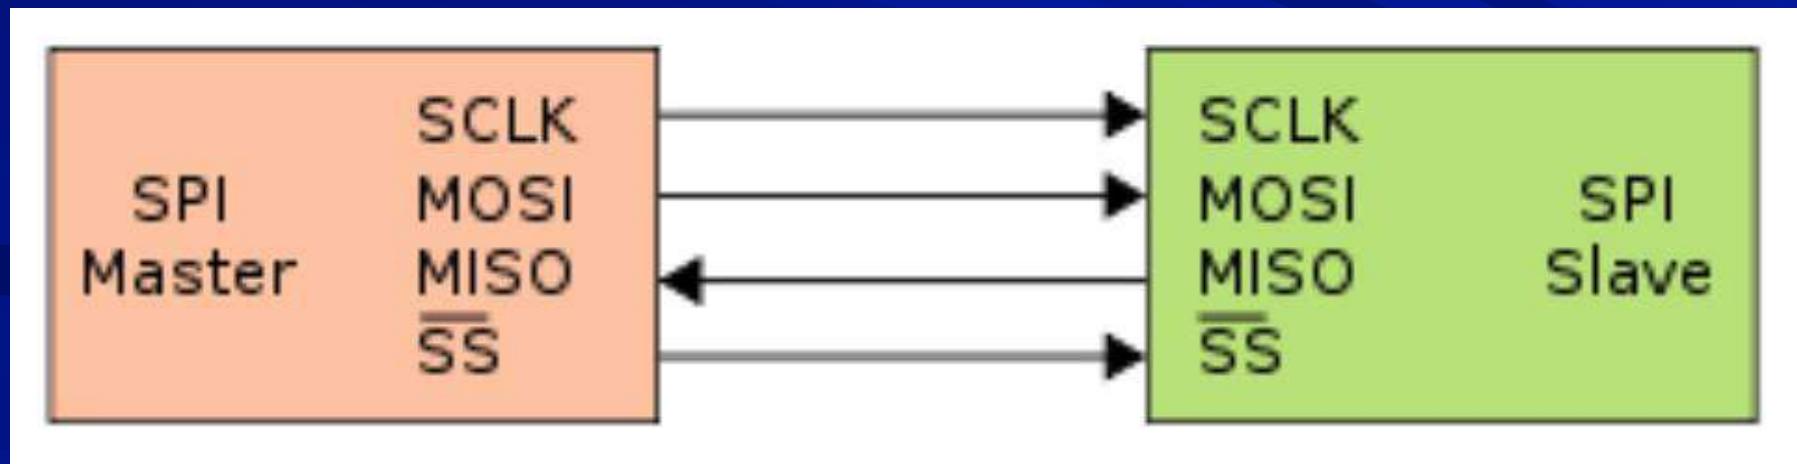
Start bits : 1

Stop bits : 1 or 2

Data rate : 300 / 600 / 1200 / 2400 / 4800 / 9600 . . .
bits/sec (bauds)

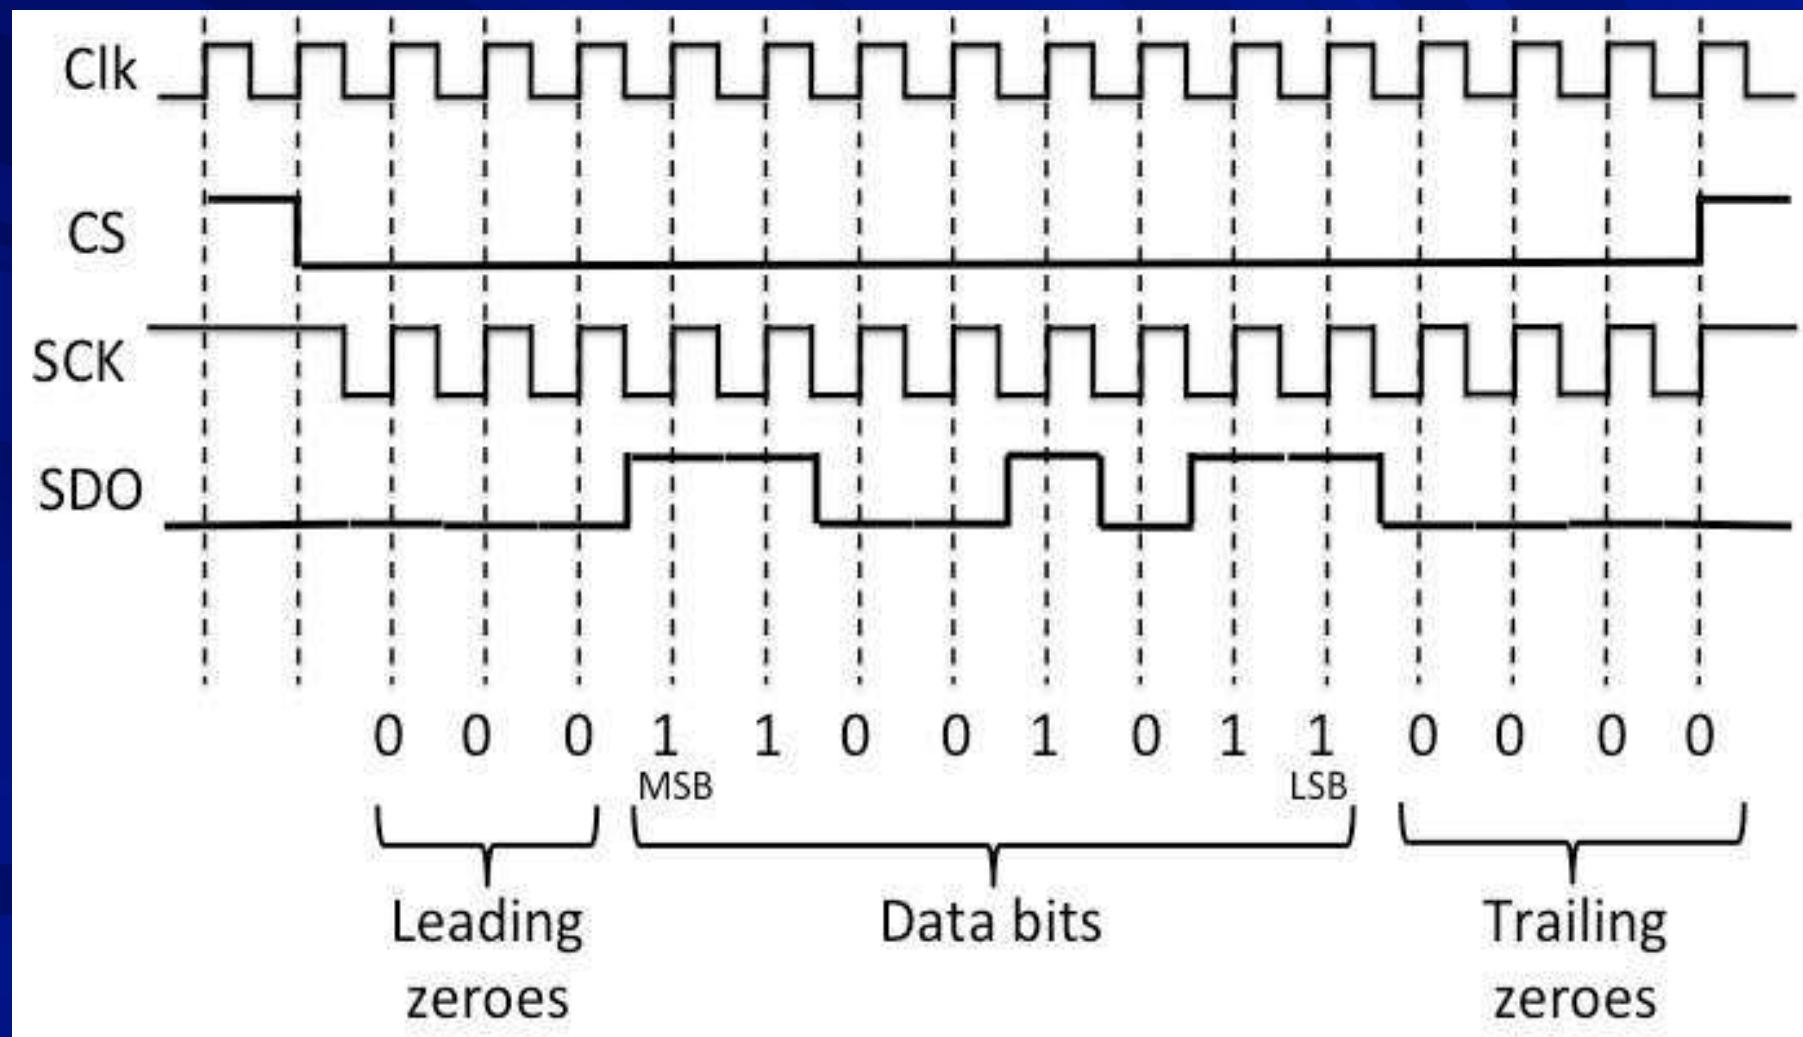
SPI : A simple serial interface

- SPI (Serial Peripheral Interface) is a synchronous serial protocol.
- Used for connecting some Pmods with BASYS-3.



Ref: Wikipedia

SPI Signal Waveforms



COL216

Computer Architecture

Input/Output – 5

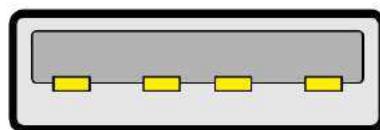
USB

28th March 2022

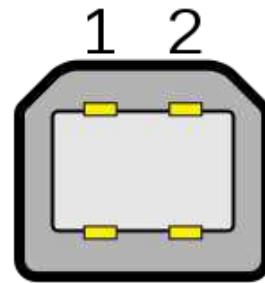
USB : Universal Serial Bus

- Introduced in mid-90's
- Used for a huge variety of devices
- Replaced earlier serial and parallel ports
- Much higher speed than RS232 etc
- Packet oriented, not character oriented
- Clock is recovered from data
- Differential signalling (D+ and D-)
- Up to 127 devices through hubs, each with up to 16 IN and 16 OUT end points

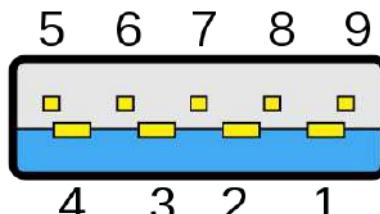
USB connectors



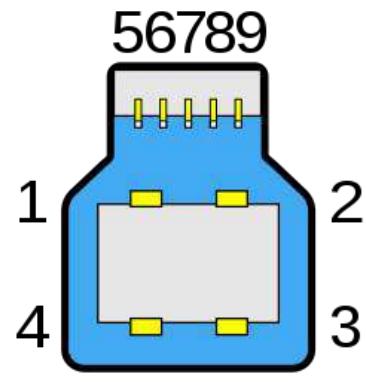
Type A



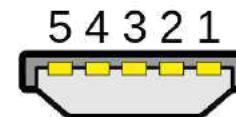
Type B



Type A
SuperSpeed



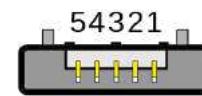
Type B
SuperSpeed



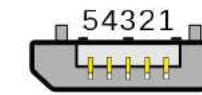
Mini-A



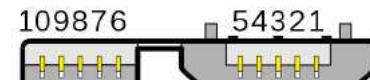
Mini-B



Micro-A



Micro-B



Micro-B SuperSpeed



Type-C

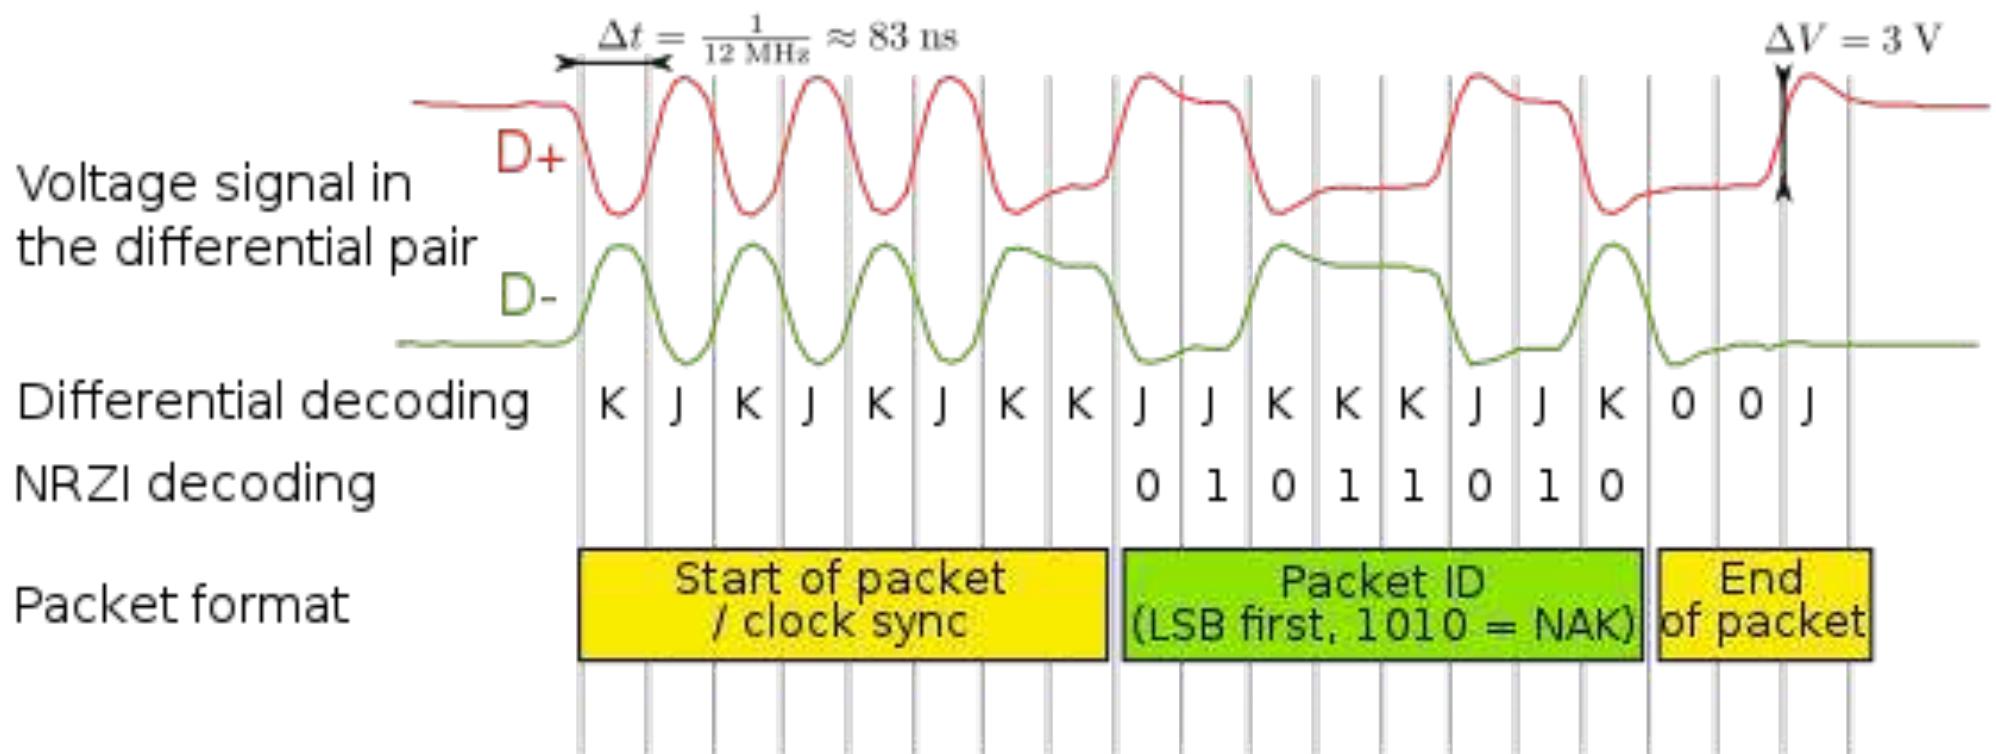
USB speed, levels

■ USB 1.0	Low Speed	1.5 Mb/s
■ USB 1.0	Full Speed	12 Mb/s
■ USB 2.0	High Speed	480 Mb/s
■ USB 3.0	Super Speed	5 Gb/s
■ USB 3.1	Super Speed+	10 Gb/s
■ USB 3.2	Super Speed+	20 Gb/s
■ USB 4	forthcoming	
LS, FS :	0.0 to 0.3 V	2.8 to 3.6 V
HS, SS :	-10 to 10 mV	-360 to 440mV

What D+,D- patterns can indicate

- Device found detached
- Device wakes up host
- Host and device idling
- Start of packet
- End of Packet
- Reset device to a known state
- Power down device
- Host wakes up device
- Host wants device to stay awake
- Device wants to wake up

USB NAK packet



source : wikipedia

Packet IDs

■ Token packets

- SPLIT: Split transaction
- PING
- OUT, IN: Address for data transfer
- SOF: start of frame marker
- SETUP: Device management

■ Handshake packets

- ACK: Packet accepted
- NAK: Packet not accepted, retransmit
- NYET: Data not ready yet
- STALL: Transfer impossible, do recovery

■ DATA0, DATA1 etc: data packet

COL216

Computer Architecture

Input/Output – 5
Secondary Storage
28th March 2022

Hard Disk Drive Example

Capacity: 14TB

Size: 3.5"

RPM: 7200

Data rate: 250 MB/s

Disks: 8 (16 heads)

Sector: 512B (logical)

4096B (physical)

Density: 2426 Kb/in

436 Ktracks/in

Cache: 256 MB



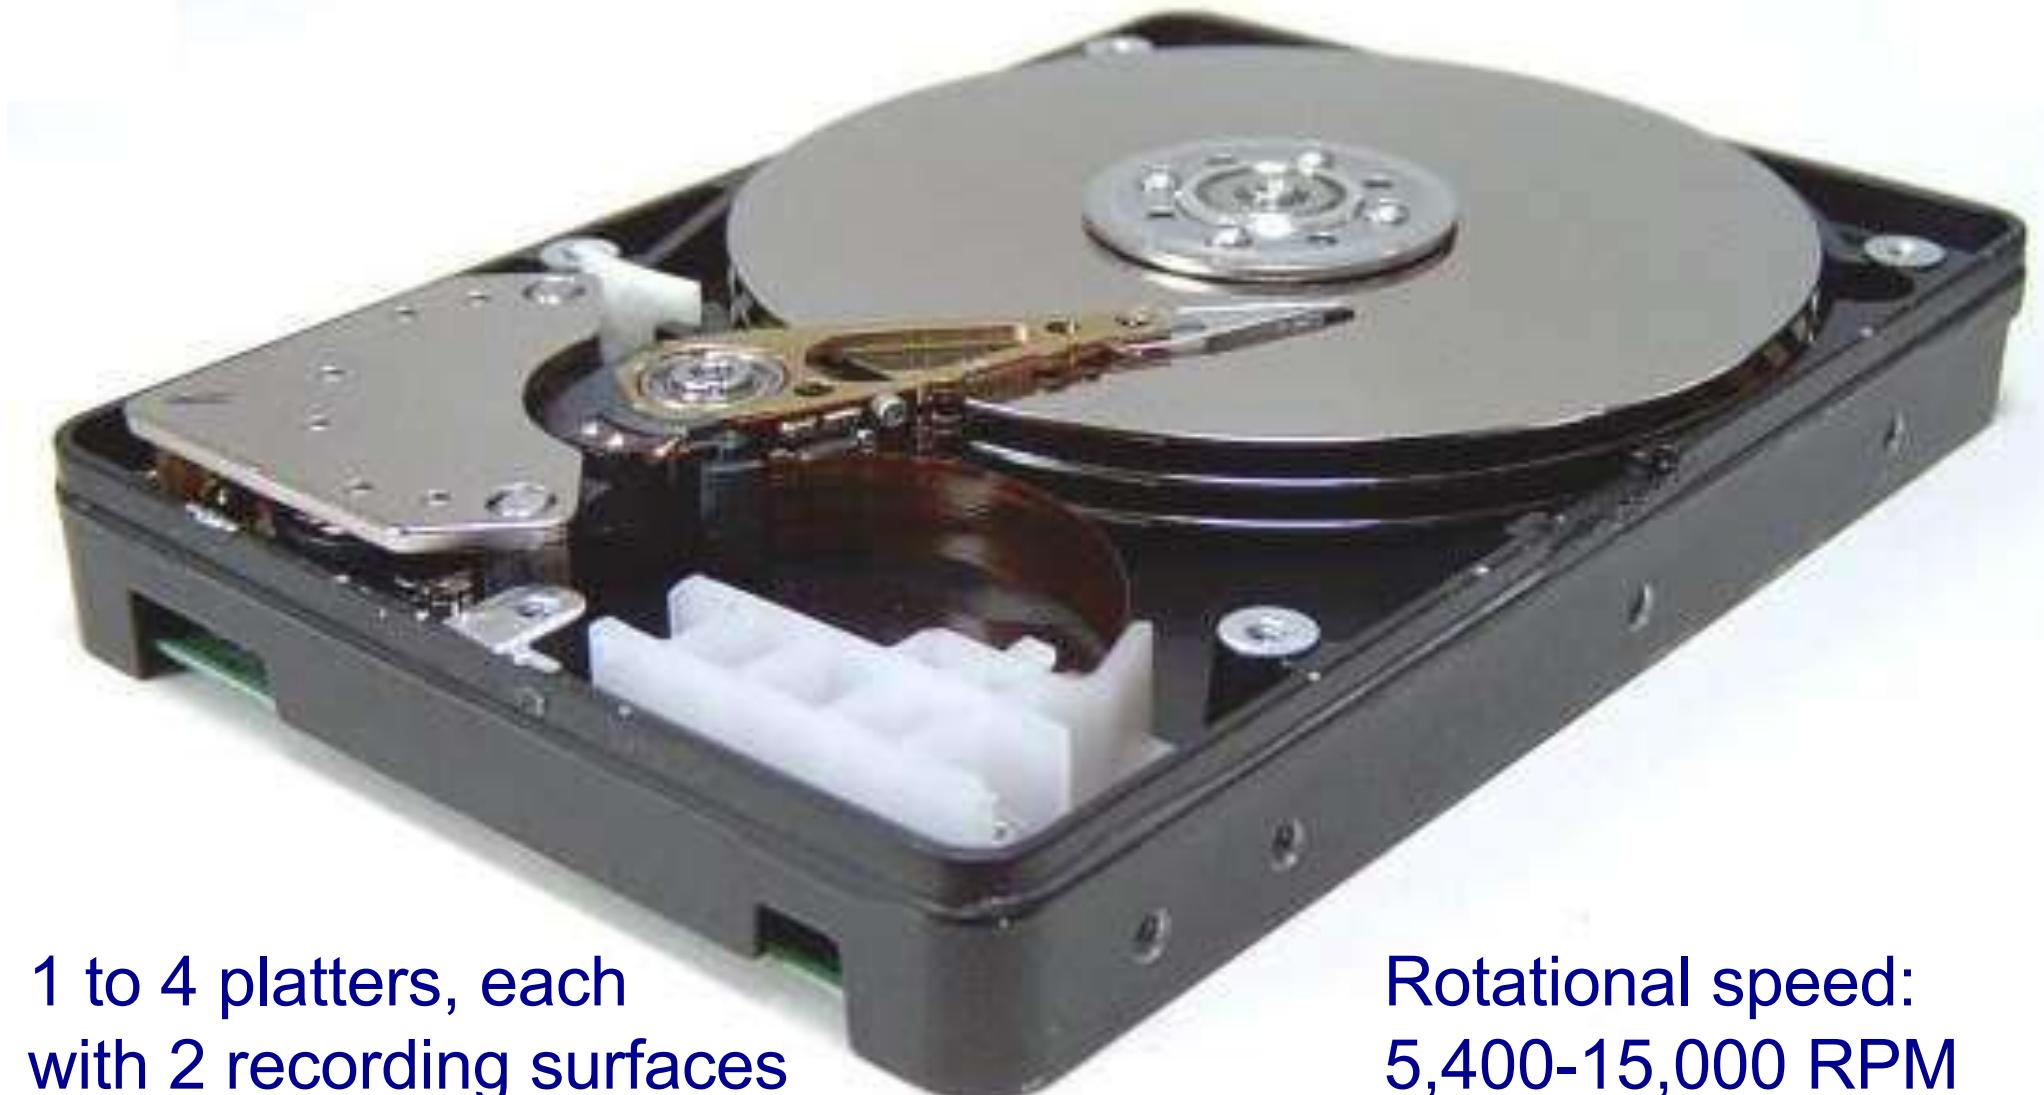
HDD comparison – then and now

Parameter	1957	2018-20	Improvement
Capacity	3.75 MB	18 TB	4.8×10^6
Volume	1.9 m ³	34 cm ³	5.6×10^4
Weight	910 kg	62 g	1.5×10^4
Access Time	600 ms	2.5 – 10 ms	2×10^2
Price	USD 9.2 / KB	USD 24 / TB	3×10^6
Density	2K b / in ²	1.3 TB / in ²	6.5×10^8
MTBF	2000 hrs	285 yrs	1.25×10^3

mean time between failures

Source: Wikipedia

Hard Disk Drive



1 to 4 platters, each
with 2 recording surfaces

Rotational speed:
5,400-15,000 RPM

Hard disk drive close-up



Tracks:
10-50 K
Sector/track
100-500
Bytes/sector
512
Seek time:

- Track to track
.5-2 ms
- Average
~10 ms

Disk Latency

Disk Latency =

Seek time + Rotational Latency + data
transfer time + controller overhead

Average rotational latency =

time for $\frac{1}{2}$ rotation =

15000 rpm

$(.5 \times 60 \times 1000 / \text{rpm}) \text{ msec}$

Transfer rate = $10^1 - 10^2 \text{ MB/sec}$

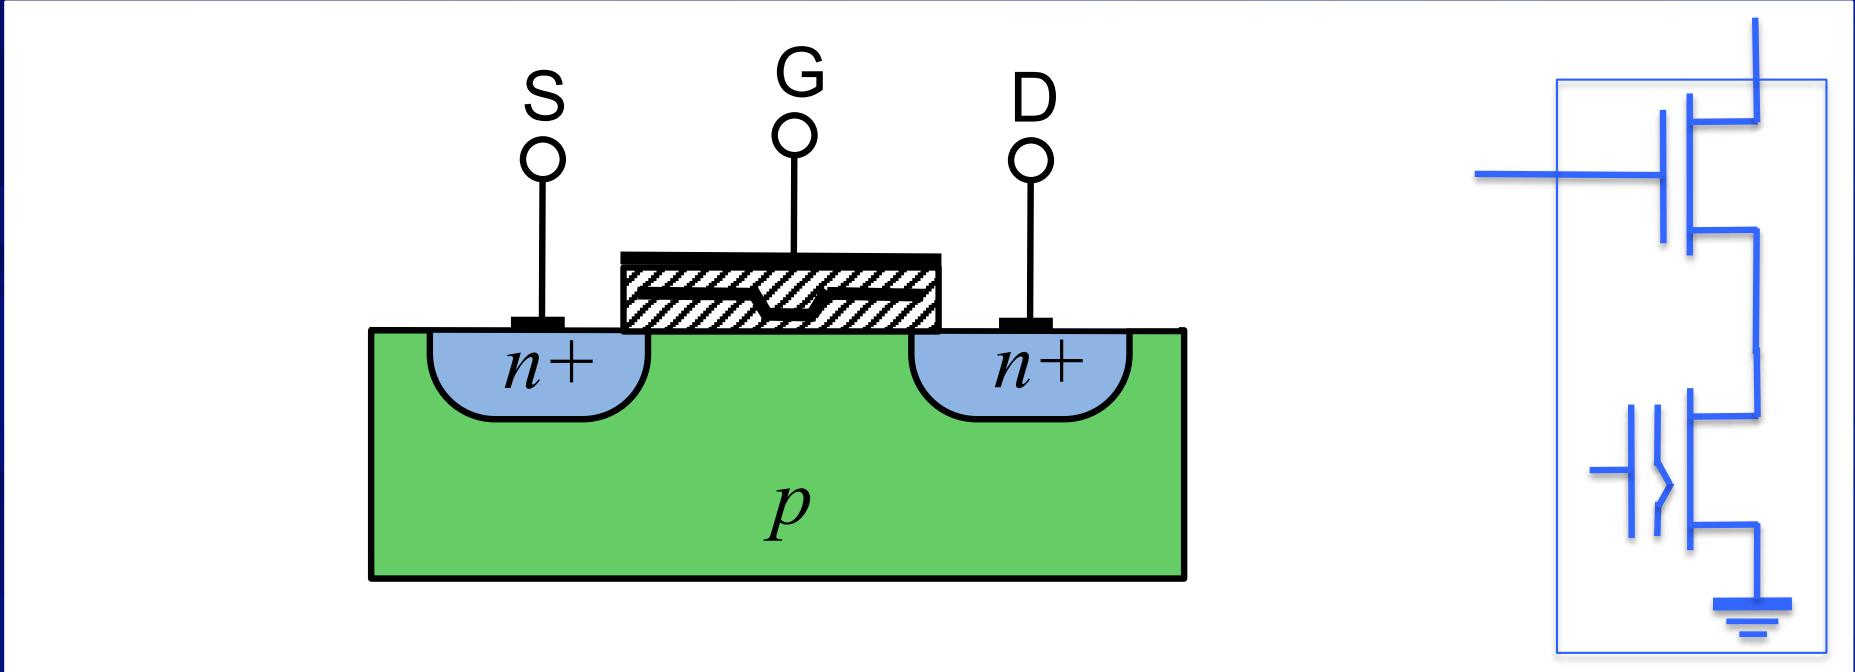
Bulk storage without moving parts

- Flash: semiconductor memory like SRAM, DRAM, but non-volatile like magnetic disks
- Latency is 100–1000 times less than disk
- Smaller, lighter, more power efficient, and more shock resistant than disk
- Has caught up with HDD in capacity, but more expensive
- SSD wears out faster than HDD

Flash Memory

- Semiconductor memory made up of “Floating gate” transistors – these store information in a non-volatile manner
- Reading is fast, but writing is slow as it requires large current
- Writing involves one way state change (e.g. 0 to 1), therefore needs erasing first (make all bits 0 first)
- Block erase, unlike byte erase of EEPROM

Floating gate transistor



Programming: pass high current through channel
Electrons are trapped on the floating gate

Before trapping: like a normal NMOS transistor
After trapping: transistor remains OFF

NOR and NAND flash

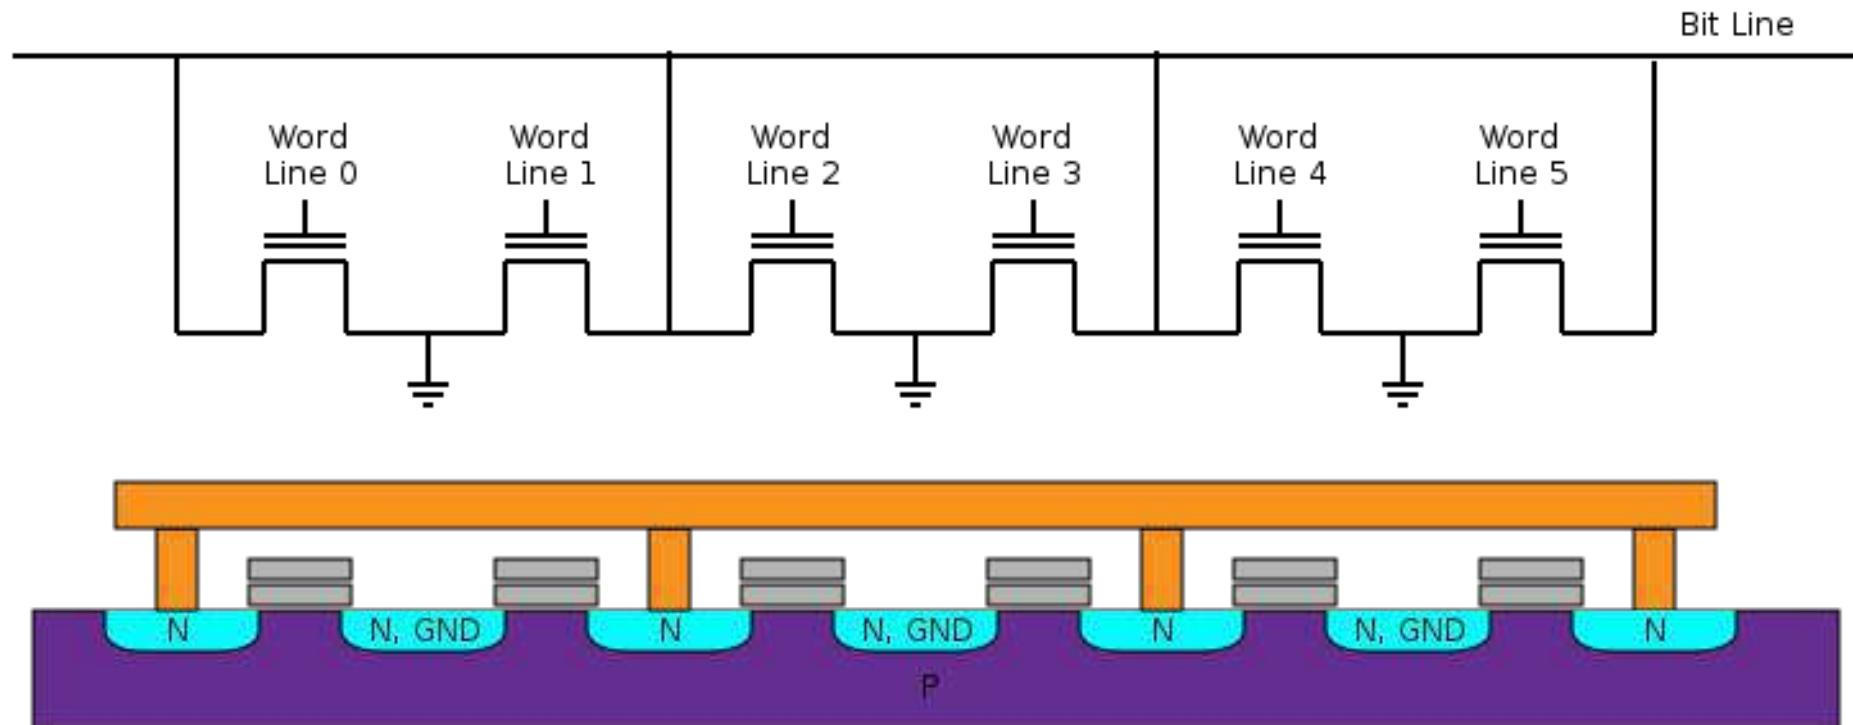
NOR flash

- Structure like NOR gate
- Random access
- Lower density
- More expensive
- Used for BIOS memory

NAND flash

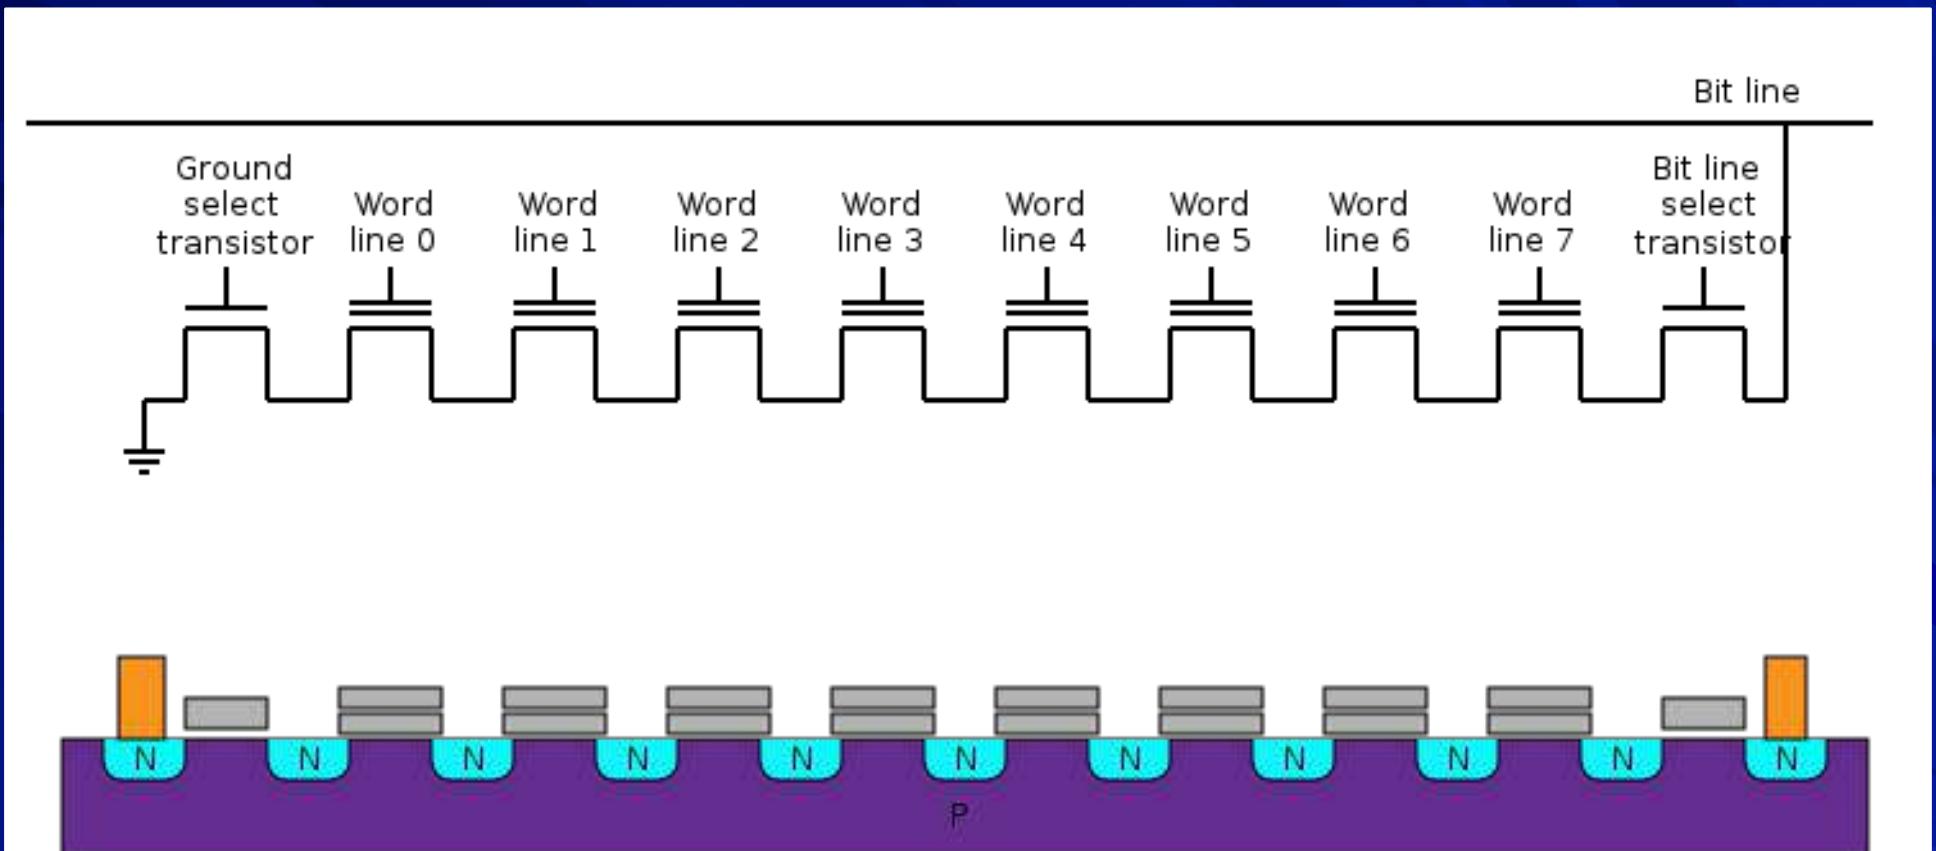
- Structure like NAND gate
- Random plus sequential access
- Higher density
- Less expensive
- Used for pen drive, SD card, SSD
- 32 Mb on BASYS 3

NOR Flash



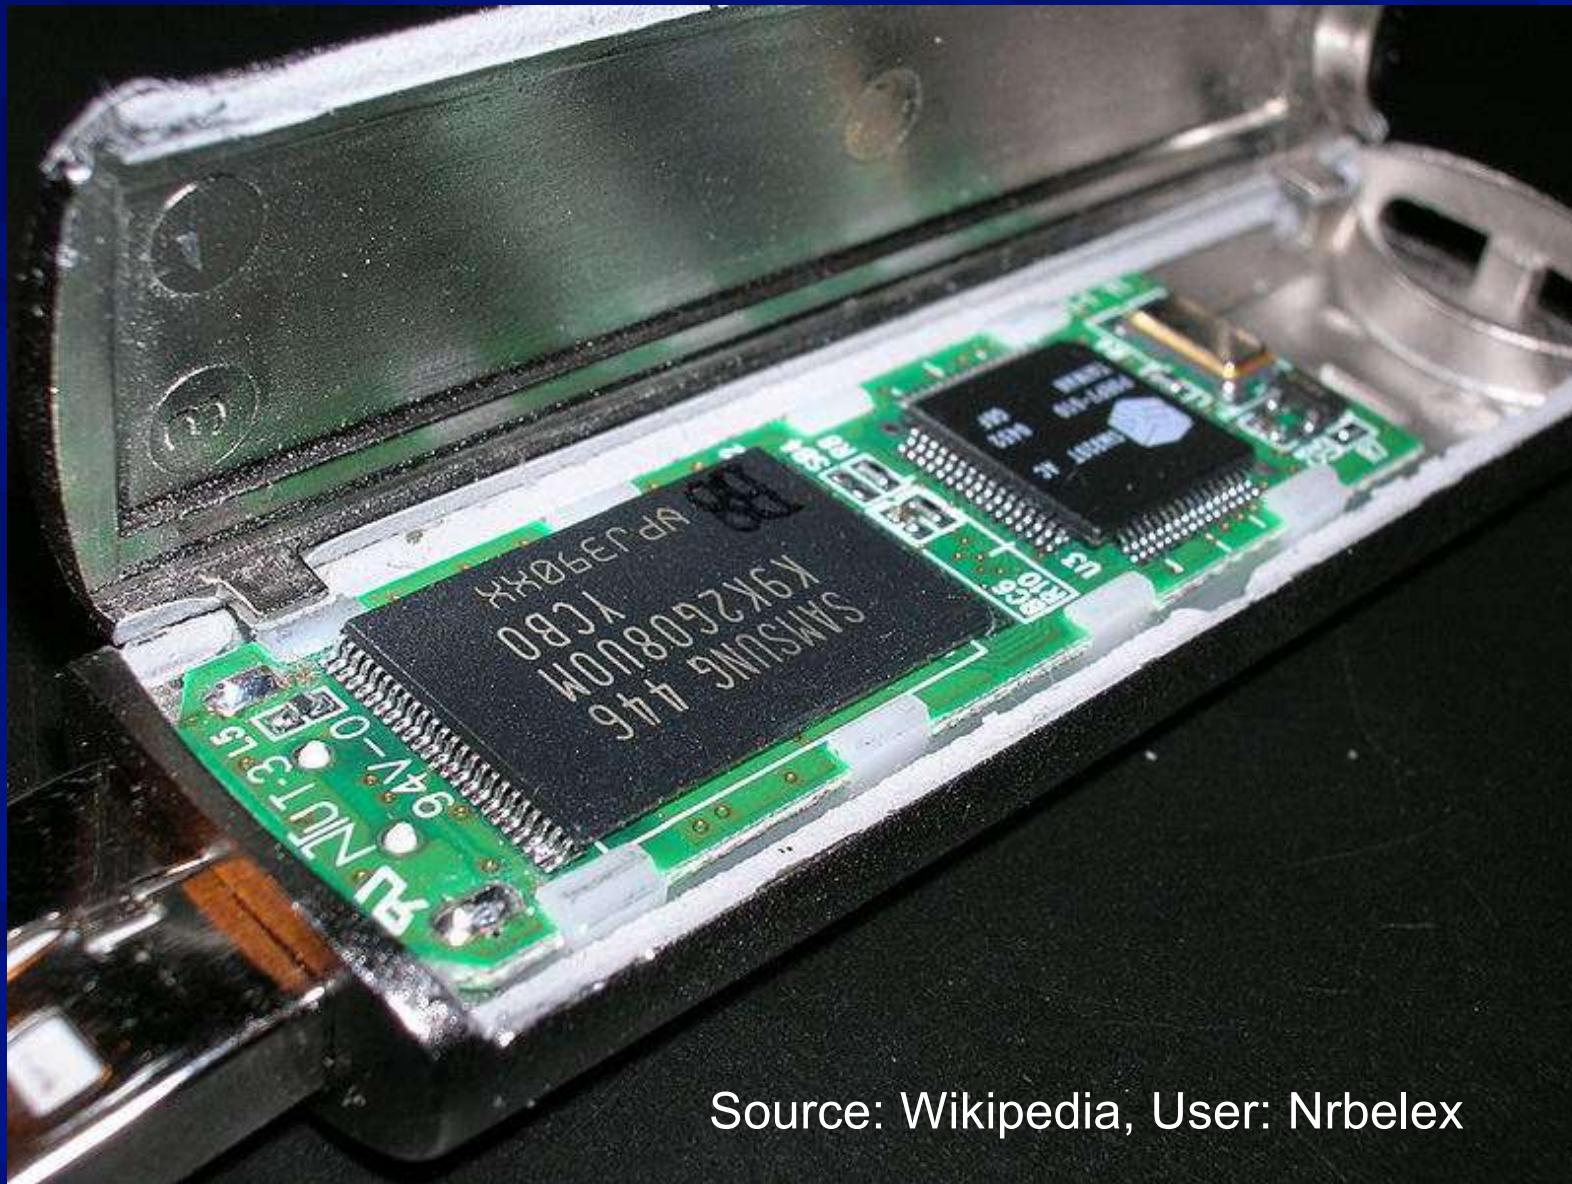
Source: Wikipedia, Author: Cyferz

NAND Flash



Source: Wikipedia, Author: Cyferz

A USB Flash Drive



Source: Wikipedia, User: Nrbelex

COL216

Computer Architecture

Input/Output – 5
I/O Performance
28th March 2022

I/O Performance

- It often gets neglected
- Suppose a benchmark executes in 100 sec
 - CPU 90%
 - I/O 10%
- CPU performance improves by 50% every year for 5 years
 - $90 \rightarrow 60 \rightarrow 40 \rightarrow 27 \rightarrow 18 \rightarrow 12$ (total = 7.5 times)
- I/O performance unchanged
- Overall performance
 - $100 \rightarrow 70 \rightarrow 50 \rightarrow 37 \rightarrow 28 \rightarrow 22$ (total = 4.5 times)
- If program is I/O bound, improvement is even smaller

I/O performance definitions

- Throughput

- Amount of data transfer in unit time (KB/s, MB/s)
 - Number of I/O operations in unit time

- Response time

- Relevant in interactive environment, embedded systems

- Both important in some cases

Examples

- Supercomputers
 - data throughput (KB or MB/s), more output than input
- Transaction processing
 - transactions per sec, response time
- File server
 - file operations per sec

Discrepancy in units

- Memory size

- $1 \text{ KB} = 1024 \text{ B} (2^{10})$

- $1 \text{ MB} = 1024 \text{ KB} (2^{10}) = 1048576 \text{ B} (2^{20})$

- I/O rate

- $1 \text{ KB/s} = 1000 \text{ B/s} (10^3)$

- $1 \text{ MB/s} = 1000 \text{ KB/s} (10^3) = 1000000 \text{ B/s} (10^6)$

I/O system performance example

Given :

- Server configuration
- I/O intensive application characteristics

Required :

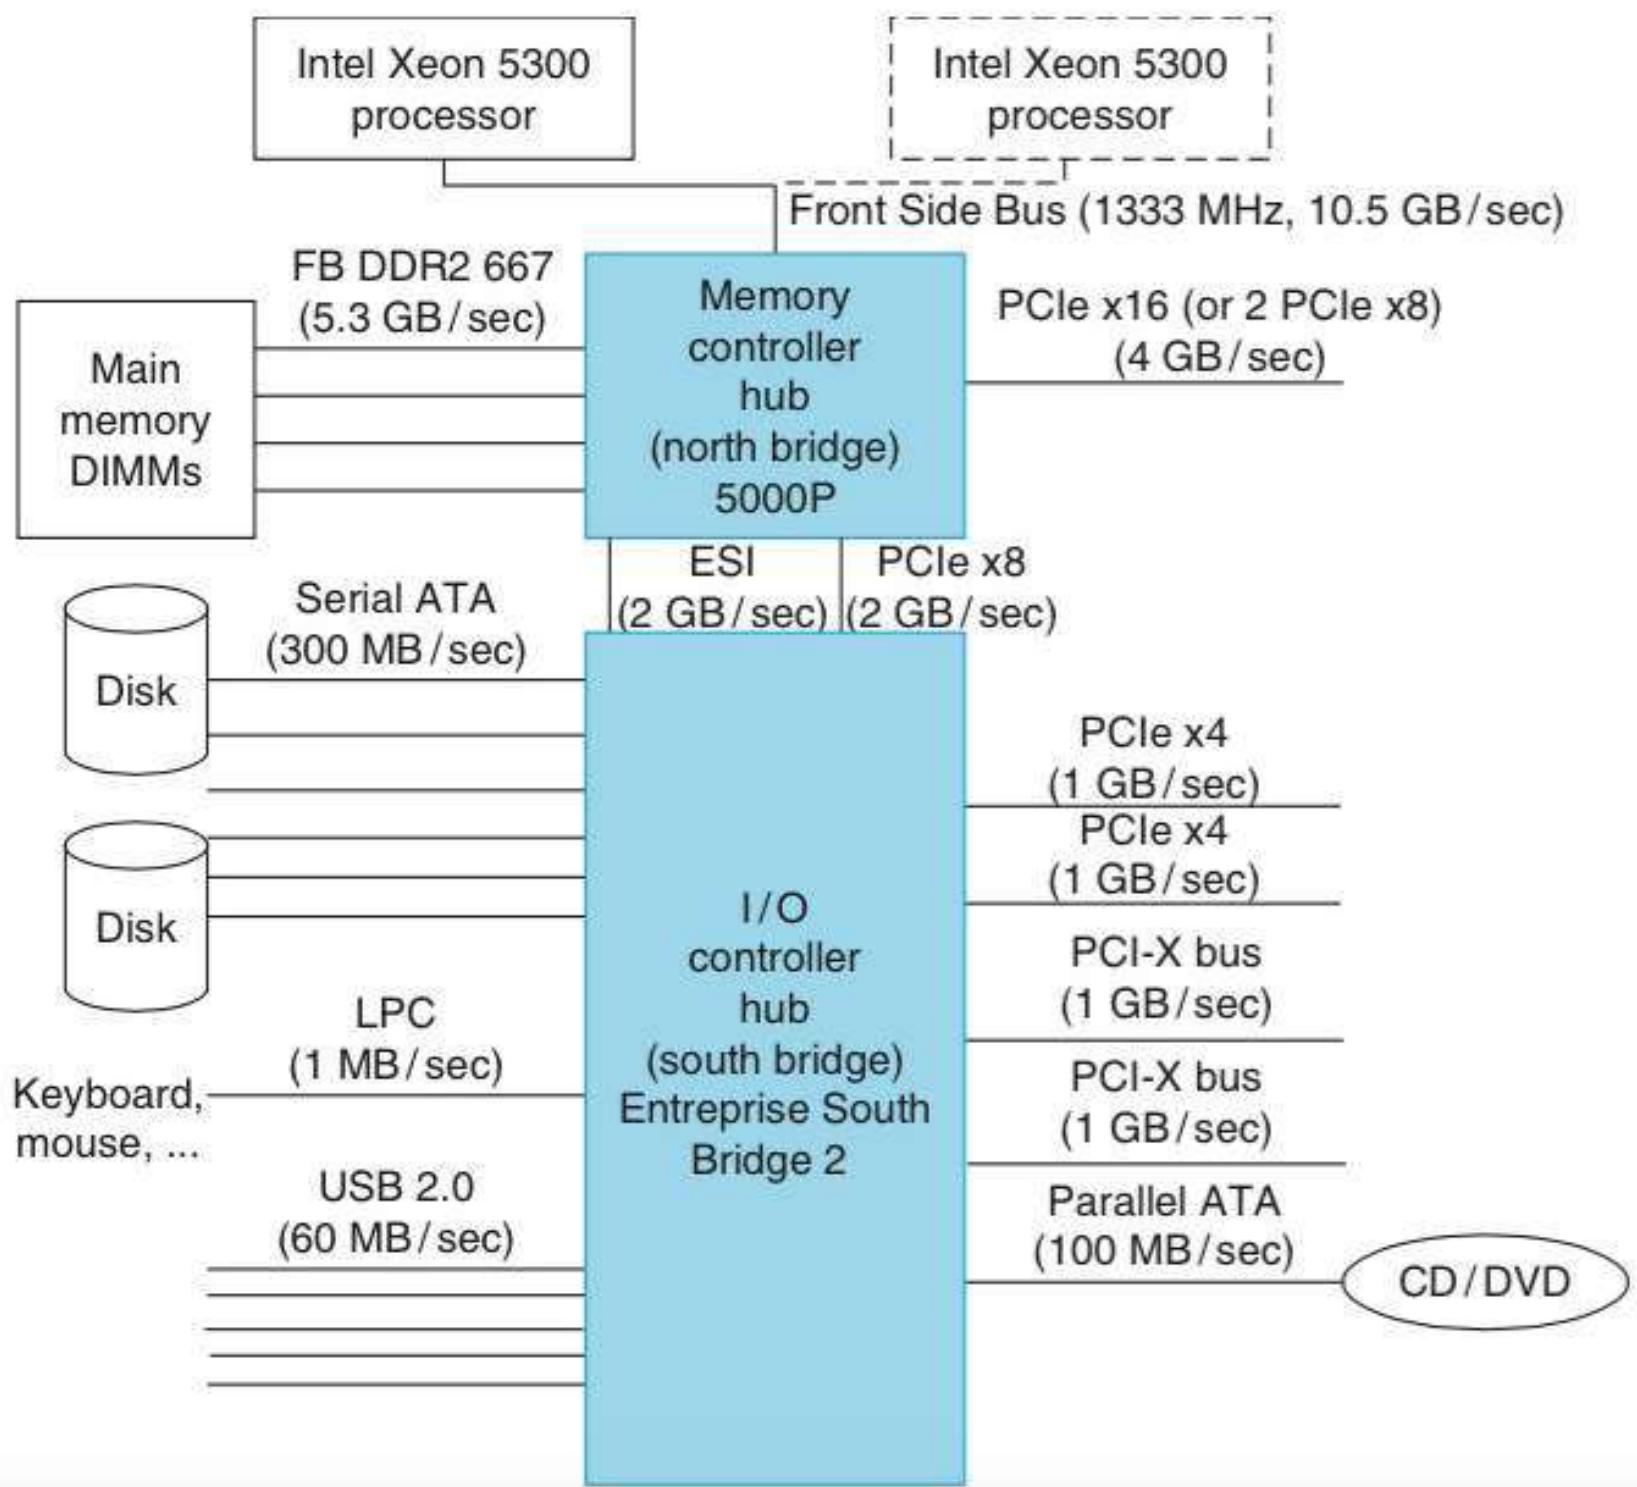
- Maximum sustainable performance
(in terms of rate of I/O operations)

I/O example: server configuration

Sun Fire x4150 server

- 8 processors, across two sockets (Xeon 5345)
- 64 GB DRAM (DDR2 667), across 16 FB DIMMs
- 8 HDDs, 2.5 inch, 15,000 RPM 73 GB SAS
- 4 Ethernet ports 10/100/1000 Gbps
- 3 PCI Express x8 ports
- 4 external and 1 internal USB 2.0 ports

Server Architecture



I/O example: application characteristics

- User program uses 200,000 instructions per I/O operation
- Operating system averages 100,000 instructions per I/O operation
- Workload consists of 64 KB reads
- Each processor sustains 1 billion instructions per second

I/O example: requirement

- Find the maximum sustainable I/O rate for a fully loaded Sun Fire x4150
 - for random reads
 - sequential reads
- Assume that the reads can always be done on an idle disk if one exists (i.e., ignore disk conflicts)

200 for user program, 100 for OS

IOPS for processors

Maximum I/O rate of 1 processor =

Instruction execution rate / Instructions per I/O

$$= 1 \times 10^9 / (200 + 100) \times 10^3 = 3,333 \text{ IOPS}$$

One socket has four processors

$$\Rightarrow 4 \times 3,333 = 13,333 \text{ IOPS}$$

Two sockets with eight processors

$$\Rightarrow 2 \times 13,333 = 26,666 \text{ IOPS.}$$

IOPS for disks

Time per I/O at disk (for random access)

$$= \text{seek} + \text{rotational time} + \text{transfer time}$$

$$= .725 \text{ ms} + 2 \text{ ms} + 64 \text{ KB} / 112 \text{ MB/s}$$

$$= 3.3 \text{ ms}$$

Each disk can do $1000/3.3 = 303$ IOPS

8 disks can do $8 \times 303 = 2,424$ IOPS

For sequential access,

in sequential, we don't bother about seek and rotational time, so we have only transfer time

$$\text{time} = 64 \text{ KB} / 112 \text{ MB/s} \Rightarrow 1,750 \text{ IOPS}$$

i.e., 14,000 IOPS for 8 disks

IOPS for PCI express

Bandwidth of a PCIe lane = 250 MB/s

Bandwidth of 8 PCIe lanes = 8×250 MB/s

= 2,000 MB/s throughput in terms of I/O blocks

IOPS for 8 PCIe lanes = $2,000$ MB/s / 64 KB

= 31,250

each I/O operation requires 64KB to be transferred

IOPS for DRAM

memory latency is 667 MHz

Bandwidth of 667 MHz FBDIMM

$$= 8 \times 667 = 5,336 \text{ MB/s}$$

$$\begin{aligned} \text{IOPS for one FBDIMM} &= (5,336 / 64) \times 1000 \\ &= 83,375 \end{aligned}$$

Max no. of DIMMs = 16

finally we compare all the IOPS to find the bottleneck

FSB limit

FSB peak bandwidth = 10.6 GB/s

Sustained bandwidth = 5.3 GB/s

IOPS for each FSB = $5.3 \times 10^6 / 64 = 81,540$

For 2 FSBs = $> 2 \times 81,540 = 163,080$ IOPS

IOPS at a glance

■ 8 Processors :	26,667
■ 8 Disks : <small>disk read is slowest</small>	2,424 (random) 14,000 (sequential)
■ PCIe 8x :	31,250
■ 1 FBDIMM :	83,375
■ 1 FSB :	81,540

THANKS